



UNIVERSITY OF TECHNOLOGY  
IN THE EUROPEAN CAPITAL OF CULTURE  
CHEMNITZ

# Deep Reinforcement Learning

Bandits

Julien Vitay

Professur für Künstliche Intelligenz - Fakultät für Informatik

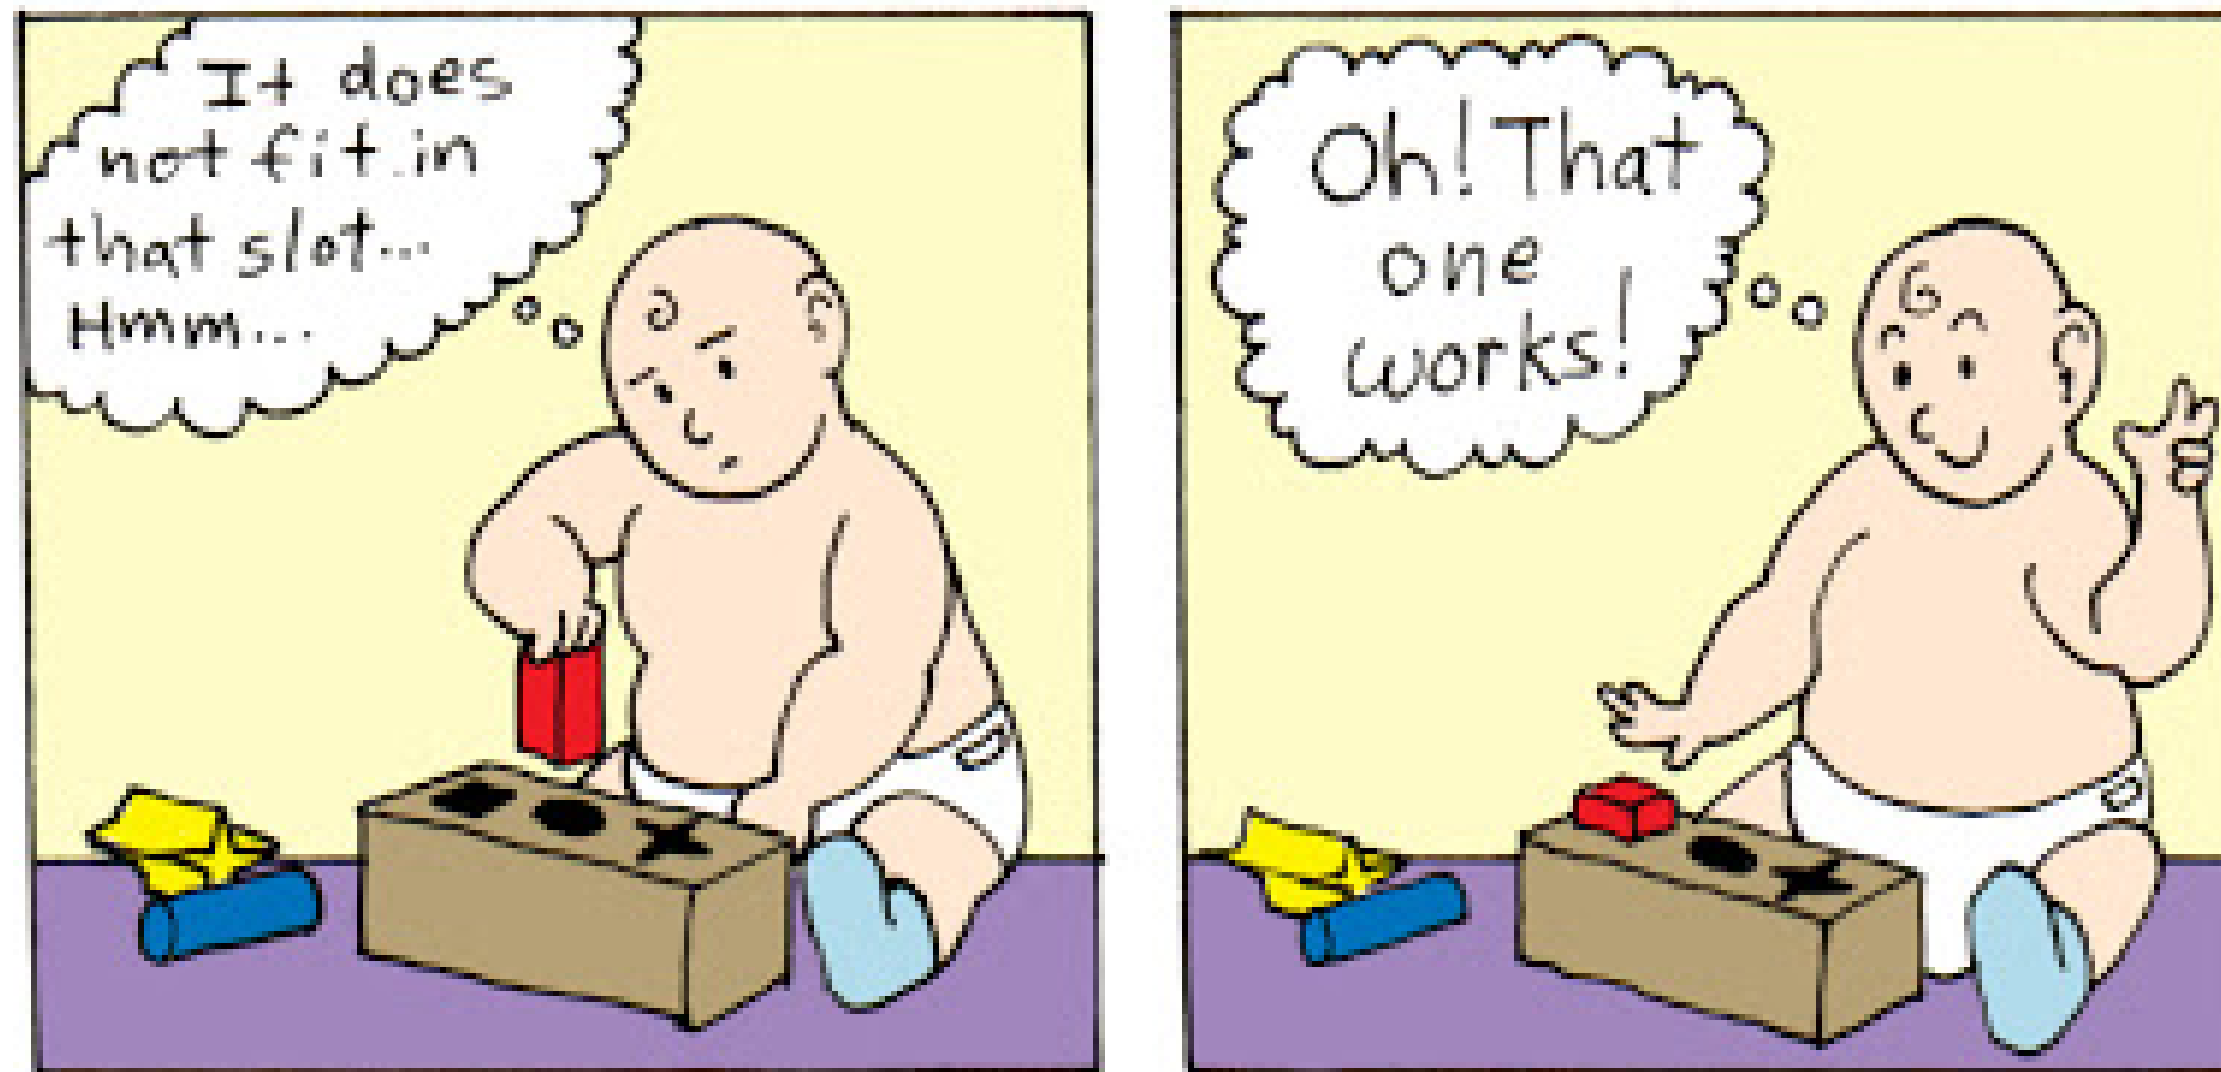
<https://tu-chemnitz.de/informatik/KI/edu/deepri>

# Outline

1. n-armed bandits
2. Sampling
3. Action selection
  1. Greedy action selection
  2.  $\epsilon$ -greedy action selection
  3. Softmax action selection
  4. Optimistic initialization
  5. Reinforcement comparison
  6. Gradient bandit algorithm
  7. Upper-Confidence-Bound action selection
4. Contextual bandits

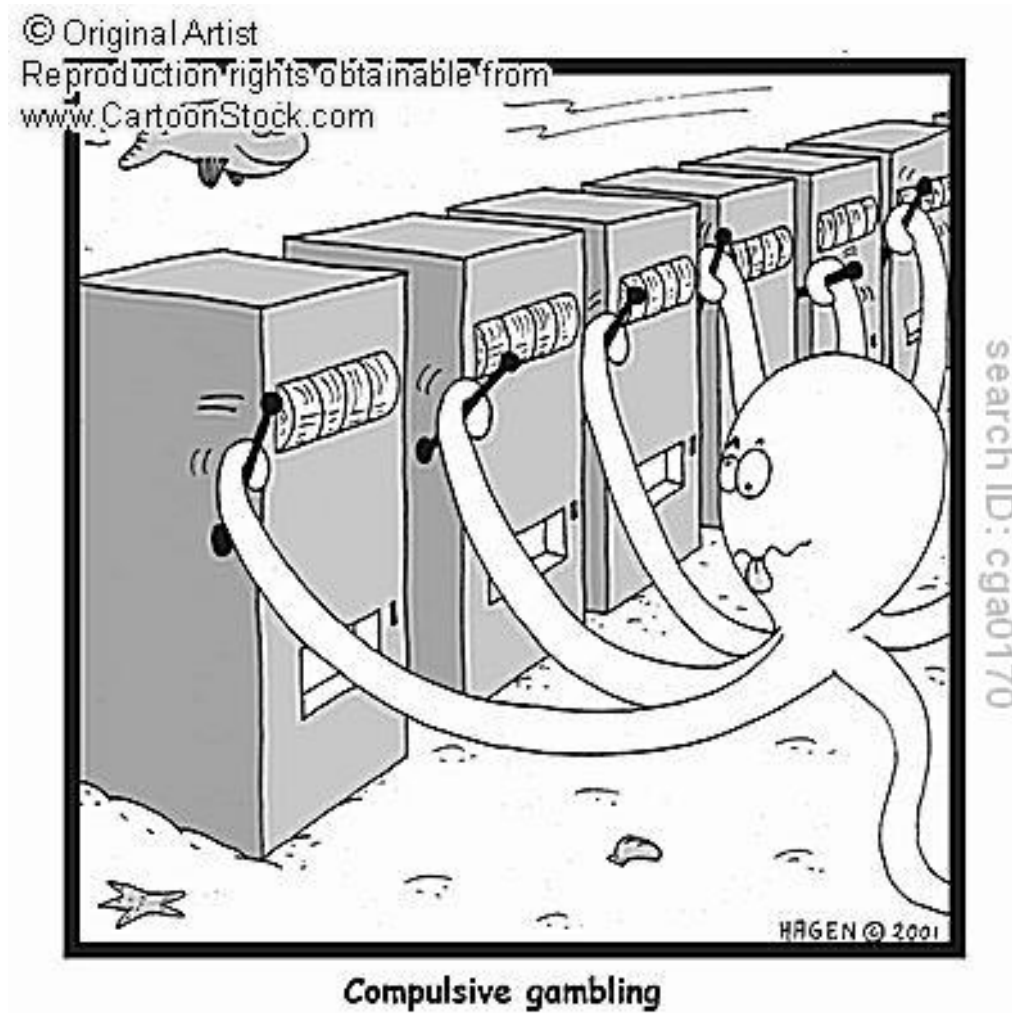
# 1 - n-armed bandits

# Evaluative Feedback




- RL evaluates actions through **trial-and-error** rather than comparing its predictions to the correct actions.
  - RL: **evaluative feedback** depends completely on the action taken.
  - SL: **instructive feedback** depends not at all on the action taken.
- Evaluative feedback indicates how good the action is, but not whether it is the best or worst action possible.
  - **Associative learning**: inputs are mapped to the best possible outputs (general RL).
  - **Non-associative learning**: finds one best output, regardless of the current state or past history (bandits).

# n-armed bandits



- The **n-armed bandit** (or multi-armed bandit) is a non-associative evaluative feedback procedure.
- Learning and action selection take place in the same single state.
- The  $n$  actions have different reward distributions.
- The goal is to find out through trial and error which action provides the most reward on average.

# n-armed bandits



- We have the choice between  $N$  different actions  $(a_1, \dots, a_N)$ .
- Each action  $a$  taken at time  $t$  provides a **reward**  $r_t$  drawn from the action-specific probability distribution  $r(a)$ .
- The mathematical expectation of that distribution is the **expected reward**, called the **true value** of the action  $Q^*(a)$ .

$$Q^*(a) = \mathbb{E}[r(a)]$$


- The reward distribution also has a **variance**: we usually ignore it in RL, as all we care about is the **optimal action**  $a^*$  (but see distributional RL later).

$$a^* = \operatorname{argmax}_a Q^*(a)$$

- If we take the optimal action an infinity of times, we maximize the reward intake **on average**.

## n-armed bandits

- The question is how to find out the optimal action through **trial and error**, i.e. without knowing the exact reward distribution  $r(a)$ .



- We only have access to **samples** of  $r(a)$  by taking the action  $a$  at time  $t$  (a **trial**, **play** or **step**).

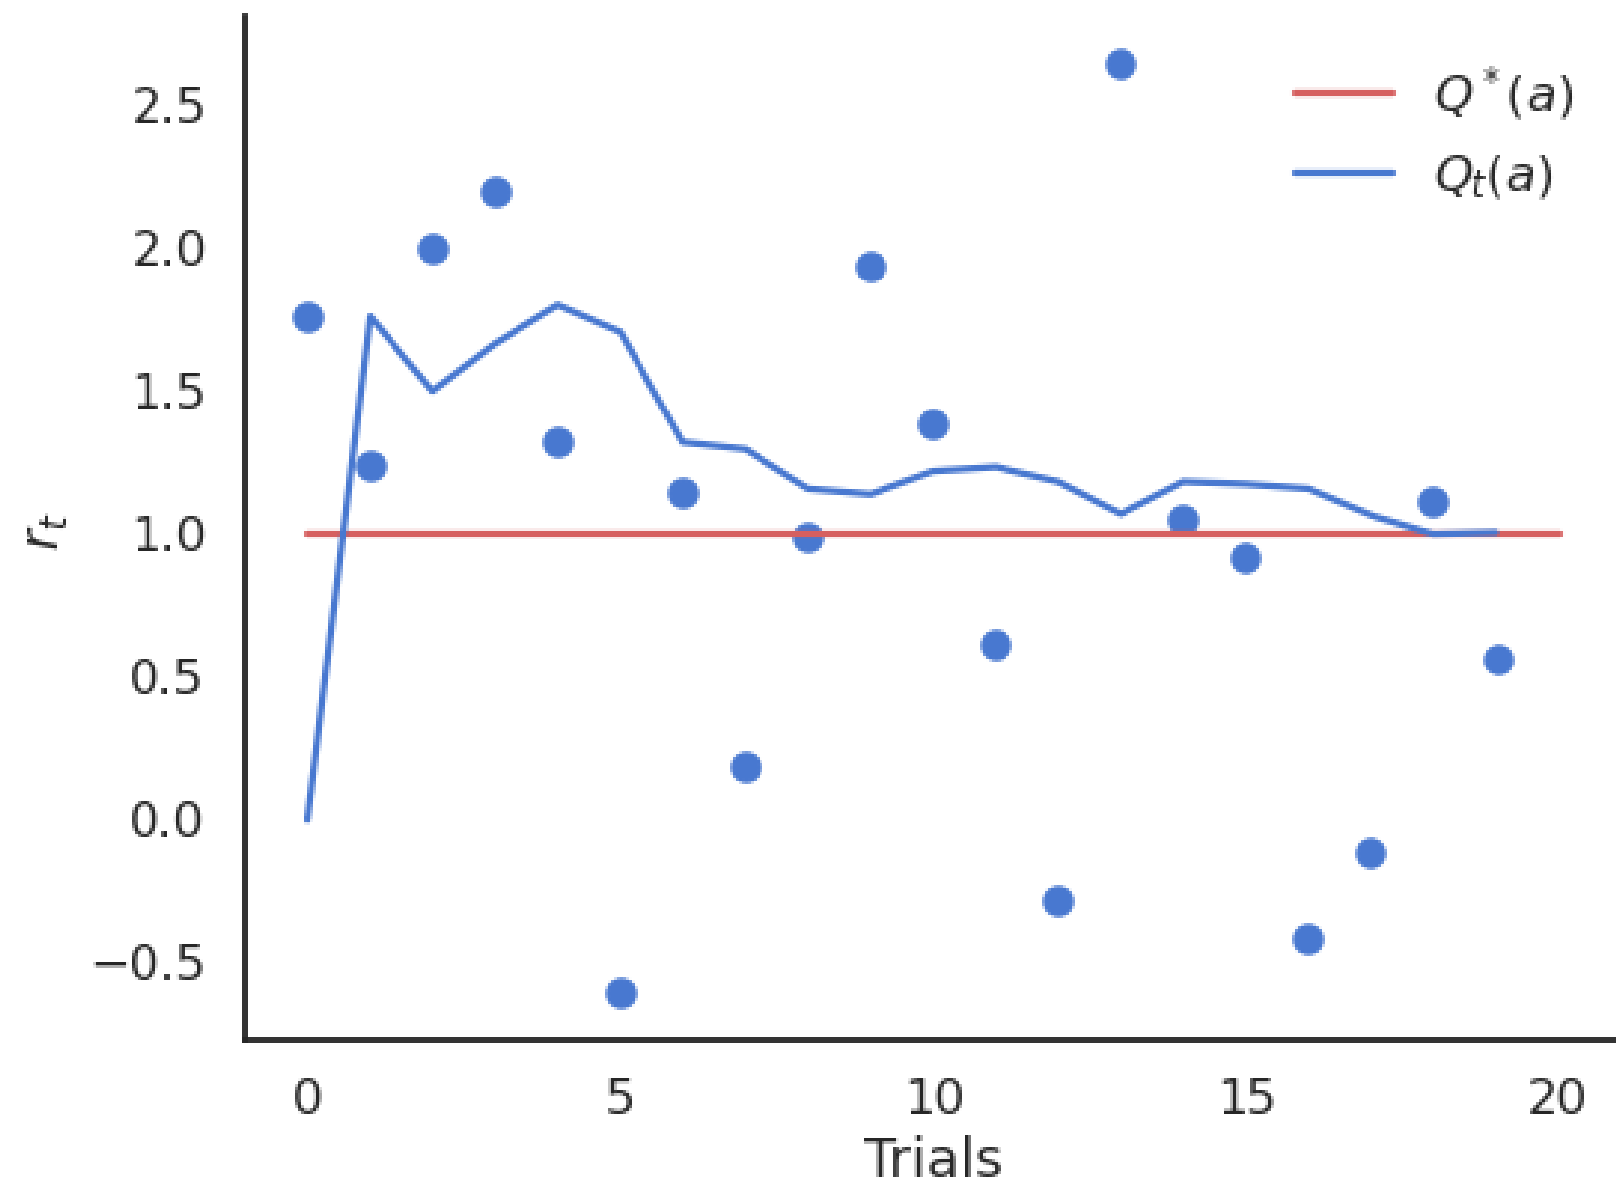
$$r_t \sim r(a)$$

- The received rewards  $r_t$  vary around the true value over time.
- We need to build **estimates**  $Q_t(a)$  of the value of each action based on the samples.
- These estimates will be very wrong at the beginning, but should get better over time.

## 2 - Sampling-based evaluation



# Sampling-based evaluation



- The expectation of the reward distribution can be approximated by the **mean** of its samples:

$$\mathbb{E}[r(a)] \approx \frac{1}{N} \sum_{t=1}^N r_t |_{a_t=a}$$

- Suppose that the action  $a$  had been selected  $t$  times, producing rewards

$$(r_1, r_2, \dots, r_t)$$

- The estimated value of action  $a$  at play  $t$  is then:

$$Q_t(a) = \frac{r_1 + r_2 + \dots + r_t}{t}$$

- Over time, the estimated action-value converges to the true action-value:

$$\lim_{t \rightarrow \infty} Q_t(a) = Q^*(a)$$

# Online evaluation

- The drawback of maintaining the mean of the received rewards is that it consumes a lot of memory:

$$Q_t(a) = \frac{r_1 + r_2 + \dots + r_t}{t} = \frac{1}{t} \sum_{i=1}^t r_i$$

- It is possible to update an estimate of the mean in an **online** or incremental manner:

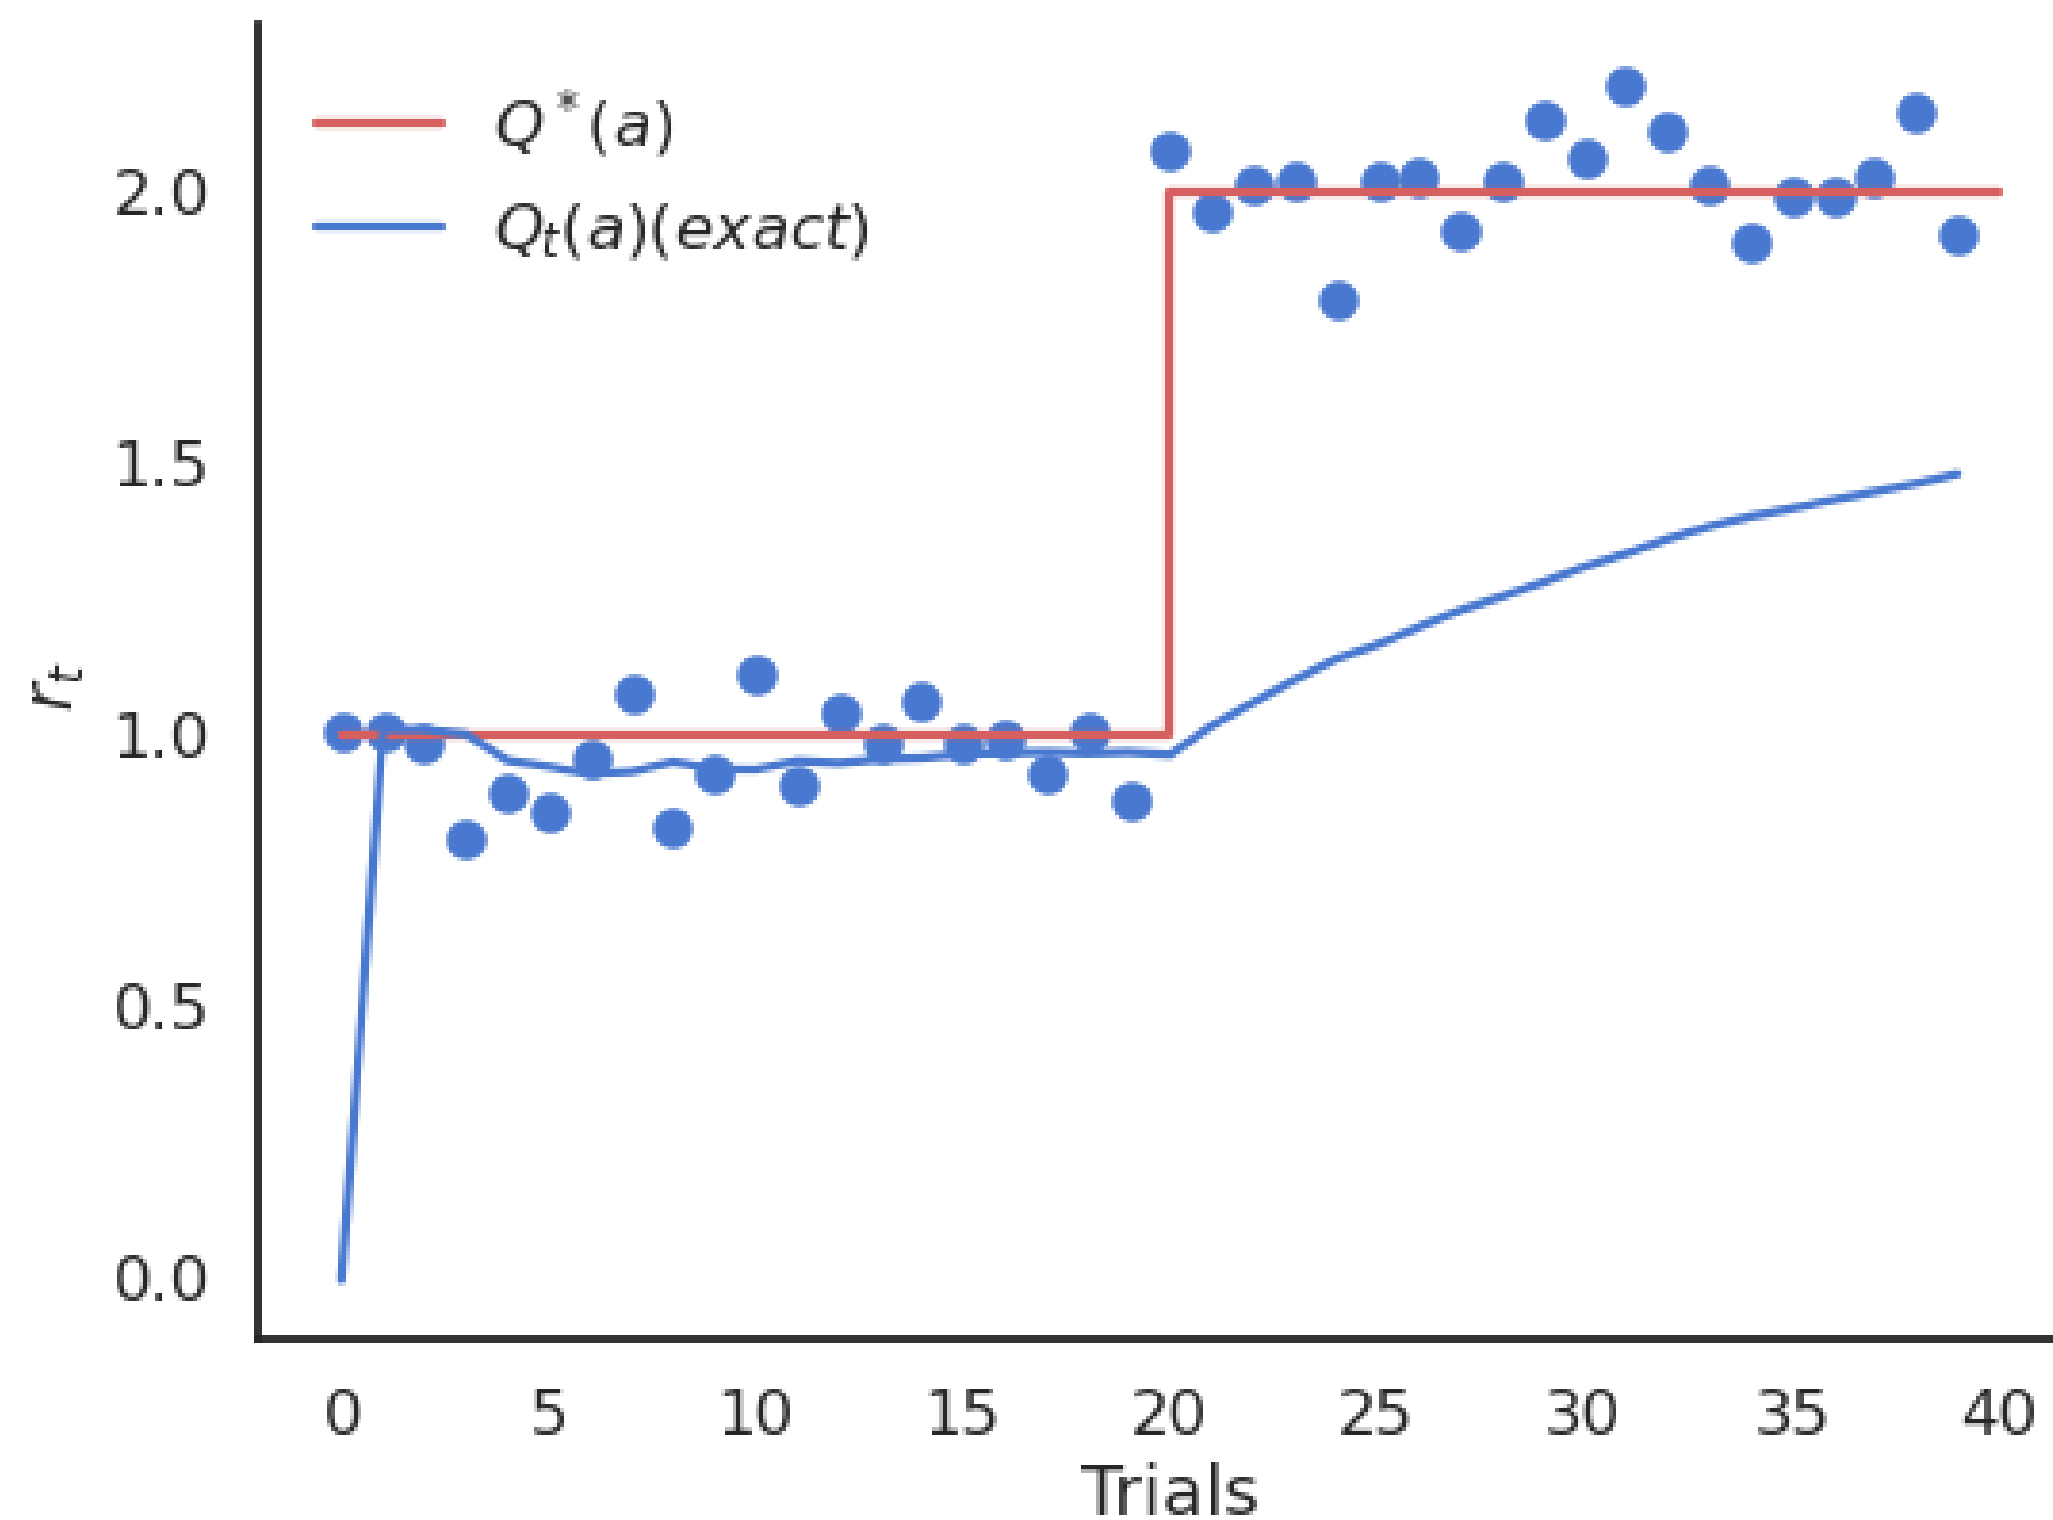
$$\begin{aligned} Q_{t+1}(a) &= \frac{1}{t+1} \sum_{i=1}^{t+1} r_i = \frac{1}{t+1} \left( r_{t+1} + \sum_{i=1}^t r_i \right) \\ &= \frac{1}{t+1} (r_{t+1} + t Q_t(a)) \\ &= \frac{1}{t+1} (r_{t+1} + (t+1) Q_t(a) - Q_t(a)) \end{aligned}$$

- The estimate at time  $t+1$  depends on the previous estimate at time  $t$  and the last reward  $r_{t+1}$ :

$$Q_{t+1}(a) = Q_t(a) + \frac{1}{t+1} (r_{t+1} - Q_t(a))$$

# Online evaluation

- The problem with the exact mean is that it is only exact when the reward distribution is **stationary**, i.e. when the probability distribution does not change over time.
- If the reward distribution is **non-stationary**, the  $\frac{1}{t+1}$  term will become very small and prevent rapid updates of the mean.



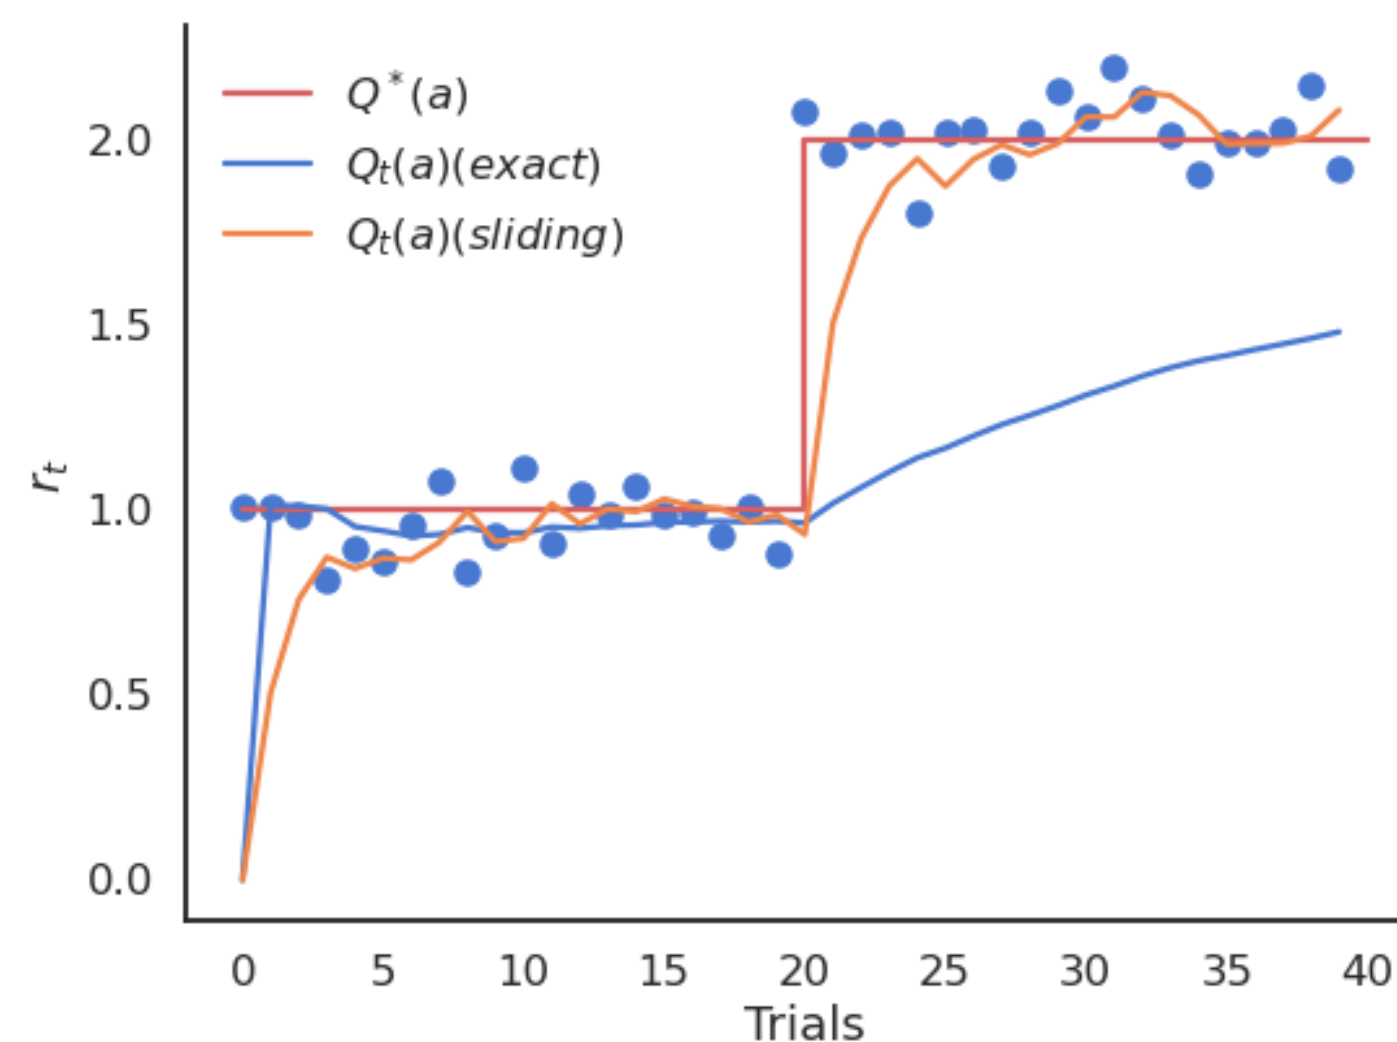
# Online evaluation

- The solution is to replace  $\frac{1}{t+1}$  with a fixed parameter called the **learning rate** (or **step size**)  $\alpha$ :


$$Q_{t+1}(a) = Q_t(a) + \alpha (r_{t+1} - Q_t(a))$$

$$= (1 - \alpha) Q_t(a) + \alpha r_{t+1}$$

- The computed value is called a **moving average** (or sliding average), as if one used only a small window of the past history.



# Online evaluation



- Moving average:

$$Q_{t+1}(a) = Q_t(a) + \alpha (r_{t+1} - Q_t(a))$$

or:

$$\Delta Q(a) = \alpha (r_{t+1} - Q_t(a))$$

- The moving average adapts very fast to changes in the reward distribution and should be used in **non-stationary problems**.
- It is however not exact and sensible to noise.
- Choosing the right value for  $\alpha$  can be difficult.

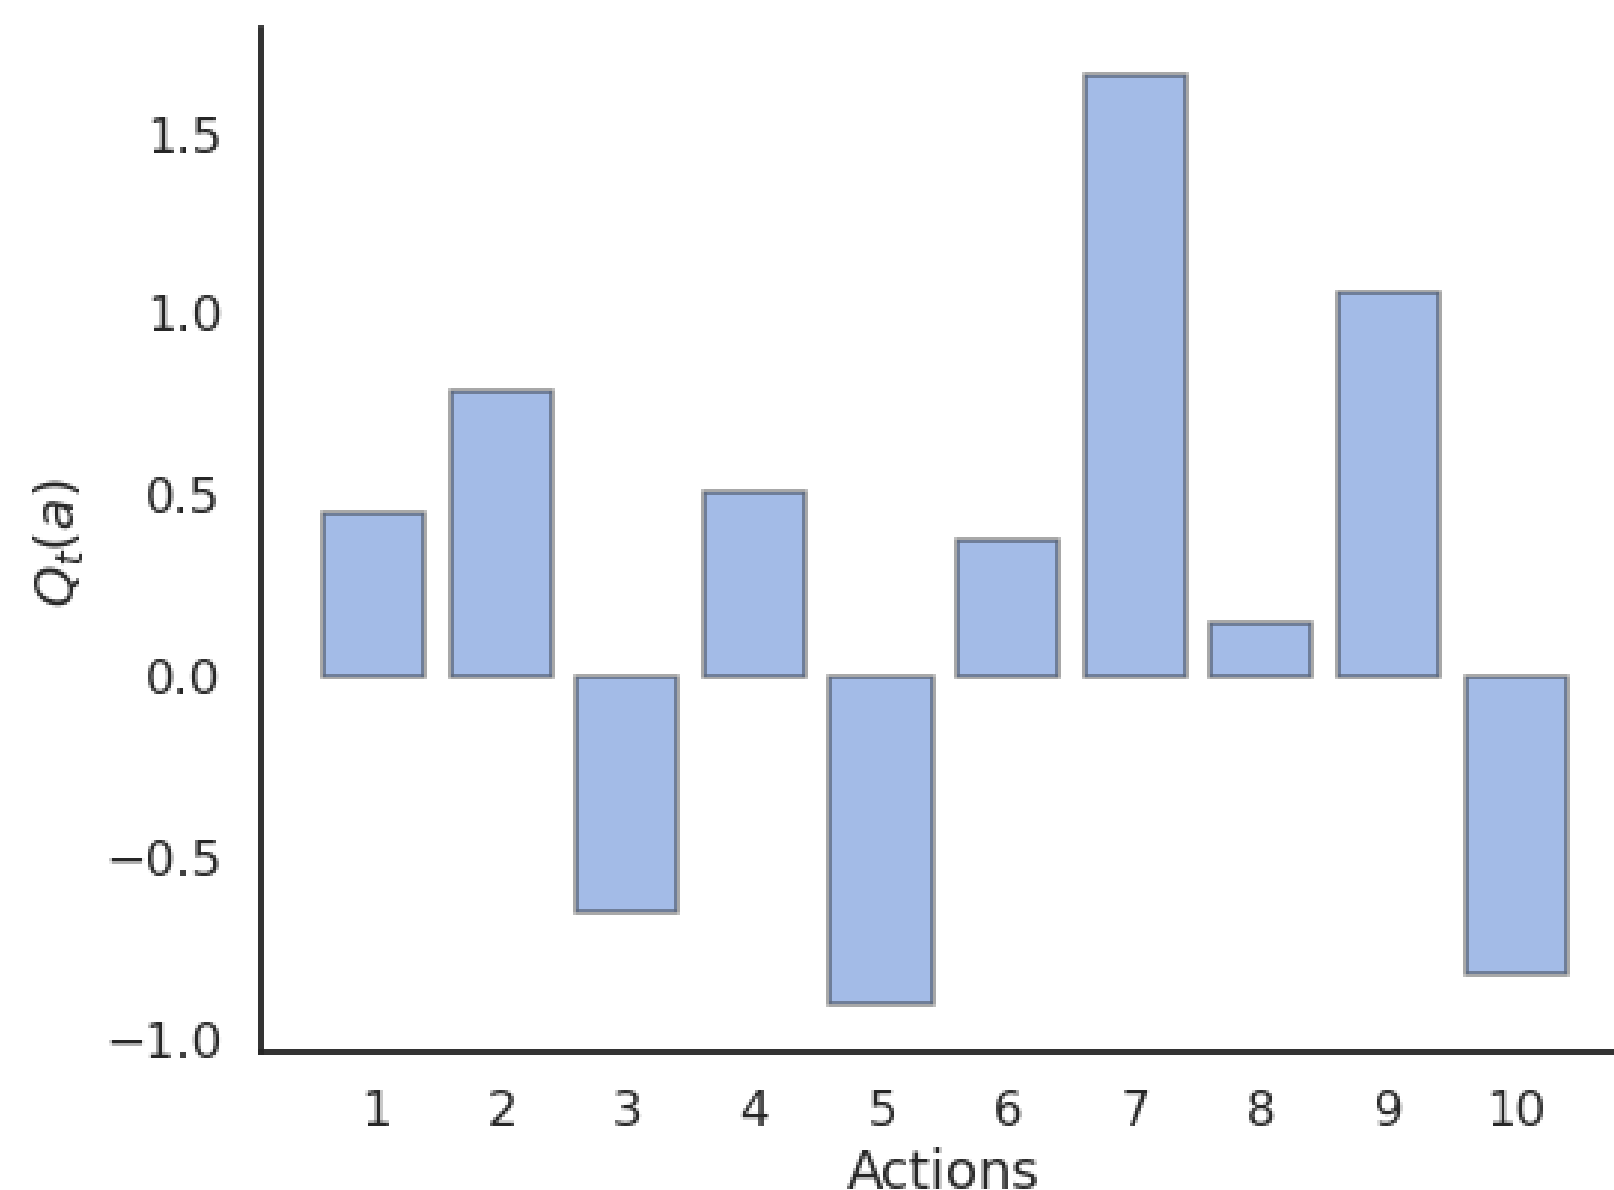
- The form of this **update rule** is very important to remember:

$$\text{new estimate} = \text{current estimate} + \alpha (\text{target} - \text{current estimate})$$

- Estimates following this update rule track the mean of their sampled target values.
- $\text{target} - \text{current estimate}$  is the **prediction error** between the target and the estimate.

## 3 - Action selection

# Action selection



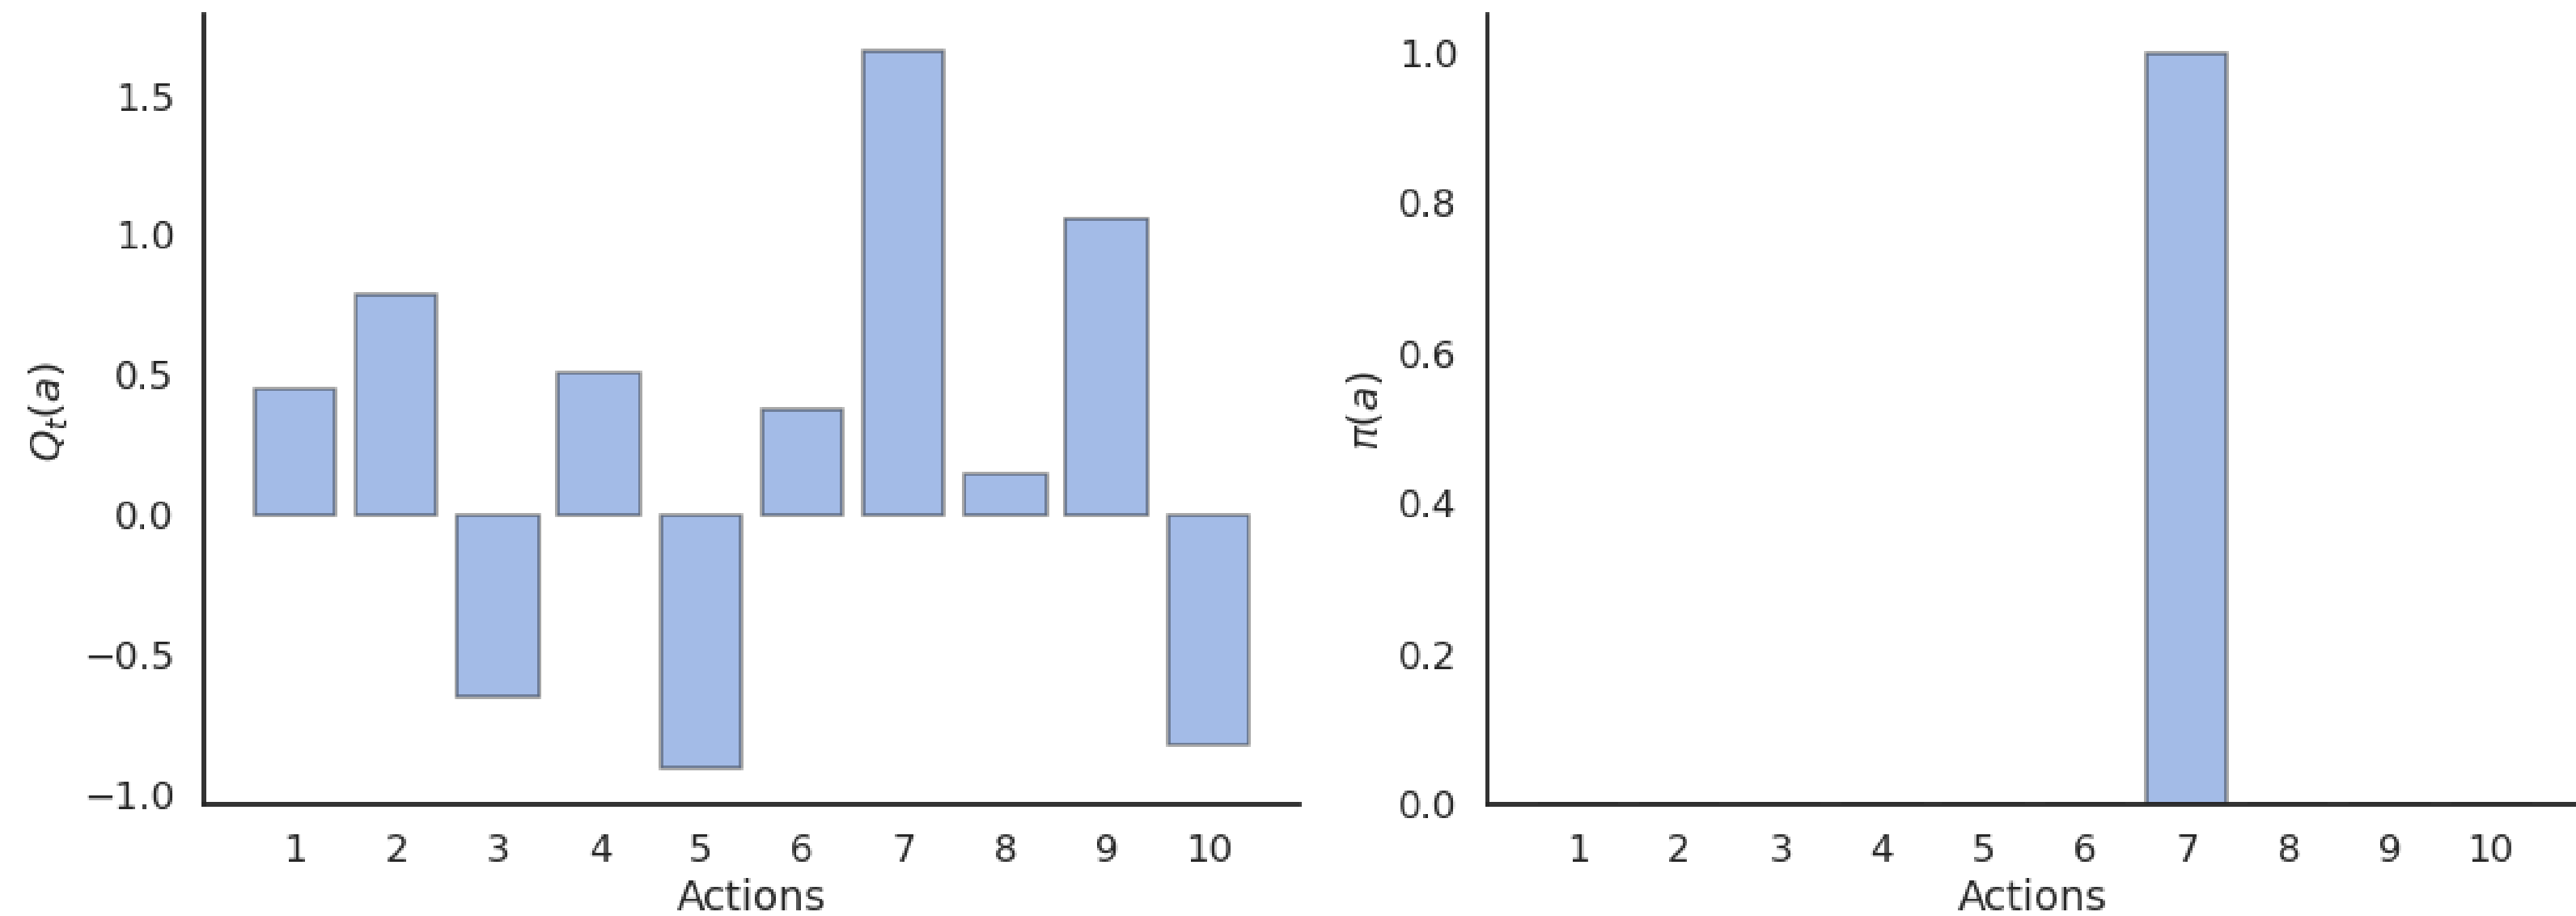
- Let's suppose we have formed reasonable estimates of the Q-values  $Q_t(a)$  at time  $t$ .
- Which action should we do next?
- If we select the next action  $a_{t+1}$  randomly (**random agent**), we do not maximize the rewards we receive, but we can continue learning the Q-values.
- Choosing the action to perform next is called **action selection** and several schemes are possible.

# Action selection

1. Greedy action selection
2.  $\epsilon$ -greedy action selection
3. Softmax action selection
4. Optimistic initialization
5. Reinforcement comparison
6. Gradient bandit algorithm
7. Upper-Confidence-Bound action selection



# Greedy action selection

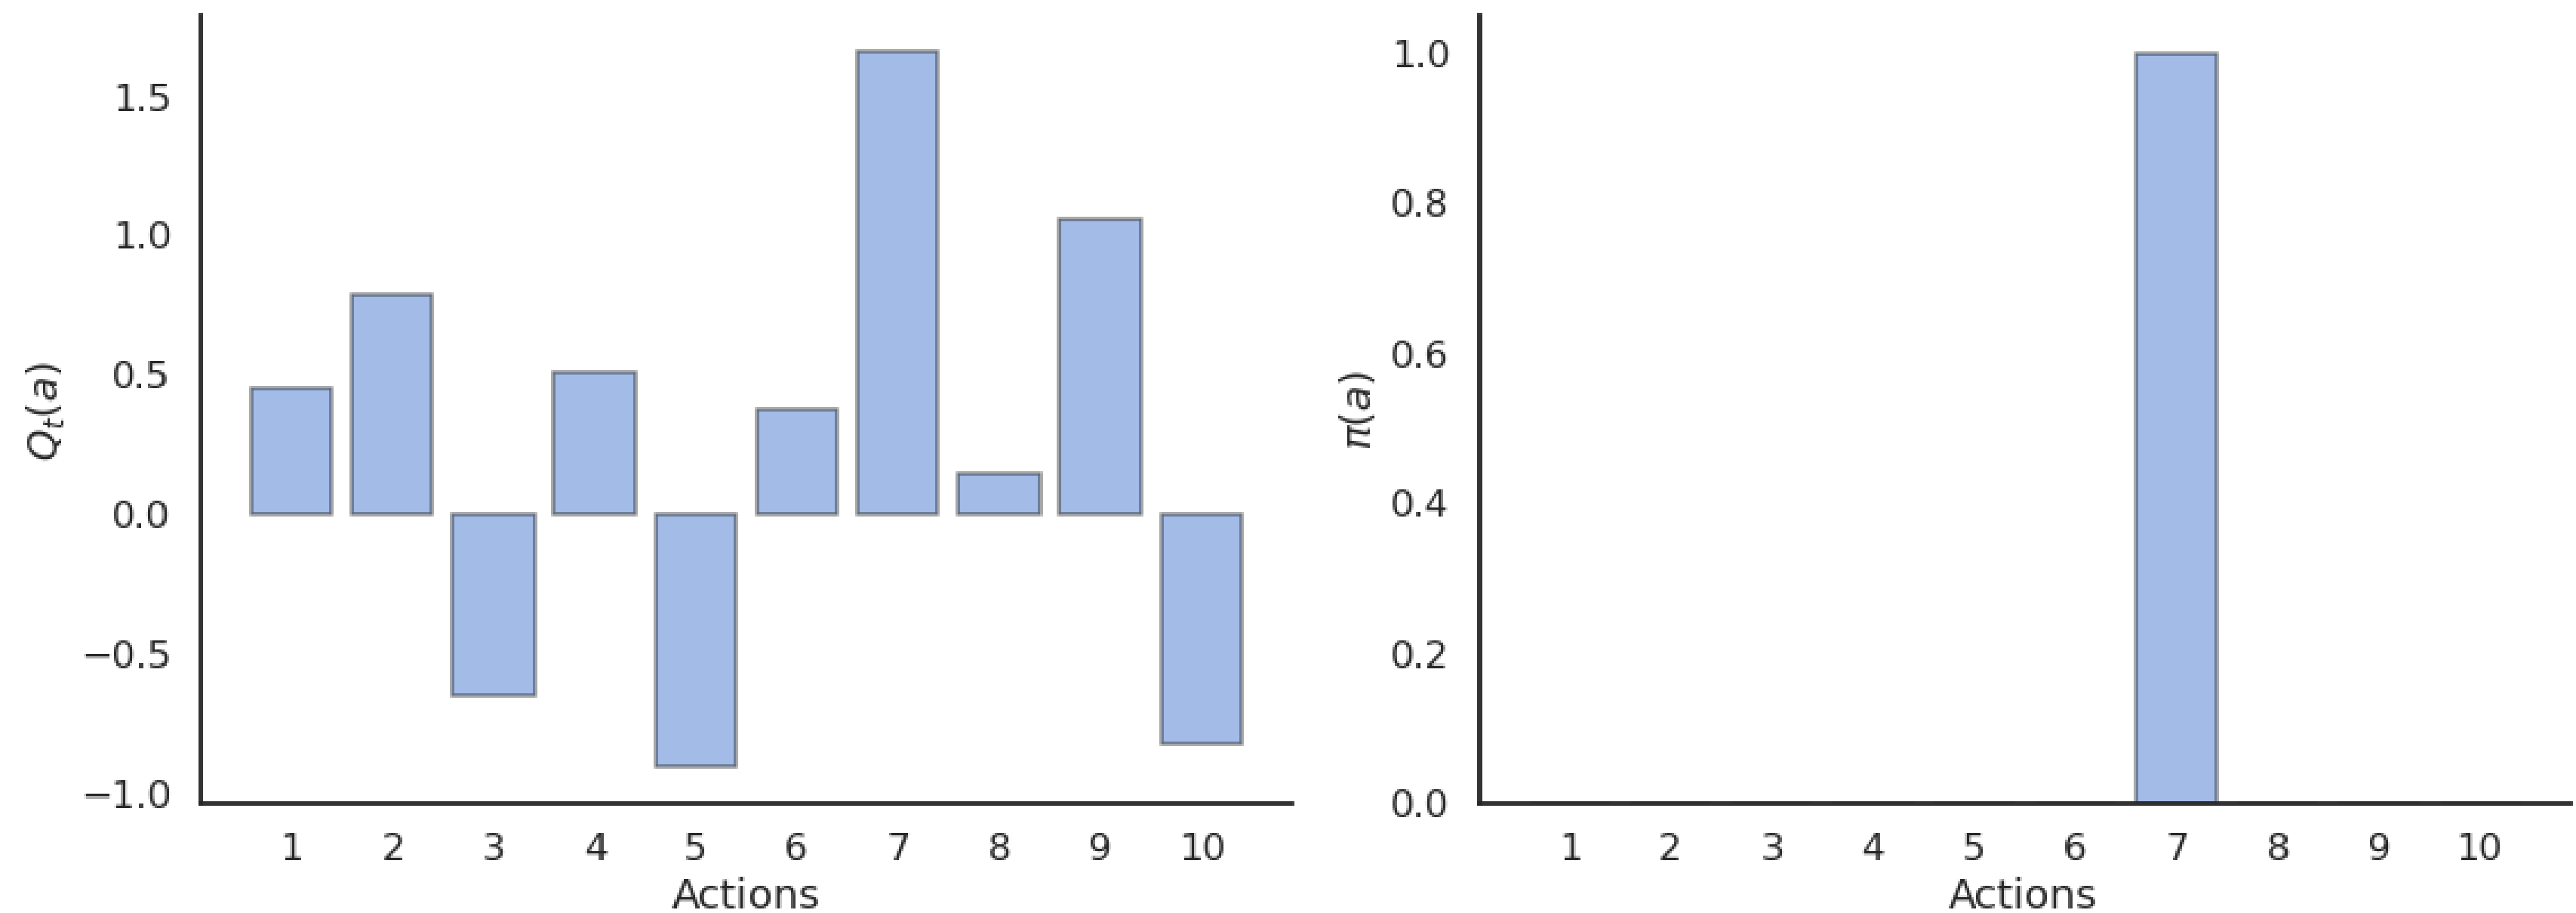


- The **greedy action** is the action whose estimated value is **maximal** at time  $t$  based on our current estimates:

$$a_t^* = \operatorname{argmax}_a Q_t(a)$$

- If our estimates  $Q_t$  are correct (i.e. close from  $Q^*$ ), the greedy action is the **optimal action** and we maximize the rewards on average.
- If our estimates are wrong, the agent will perform **sub-optimally**.

# Greedy action selection



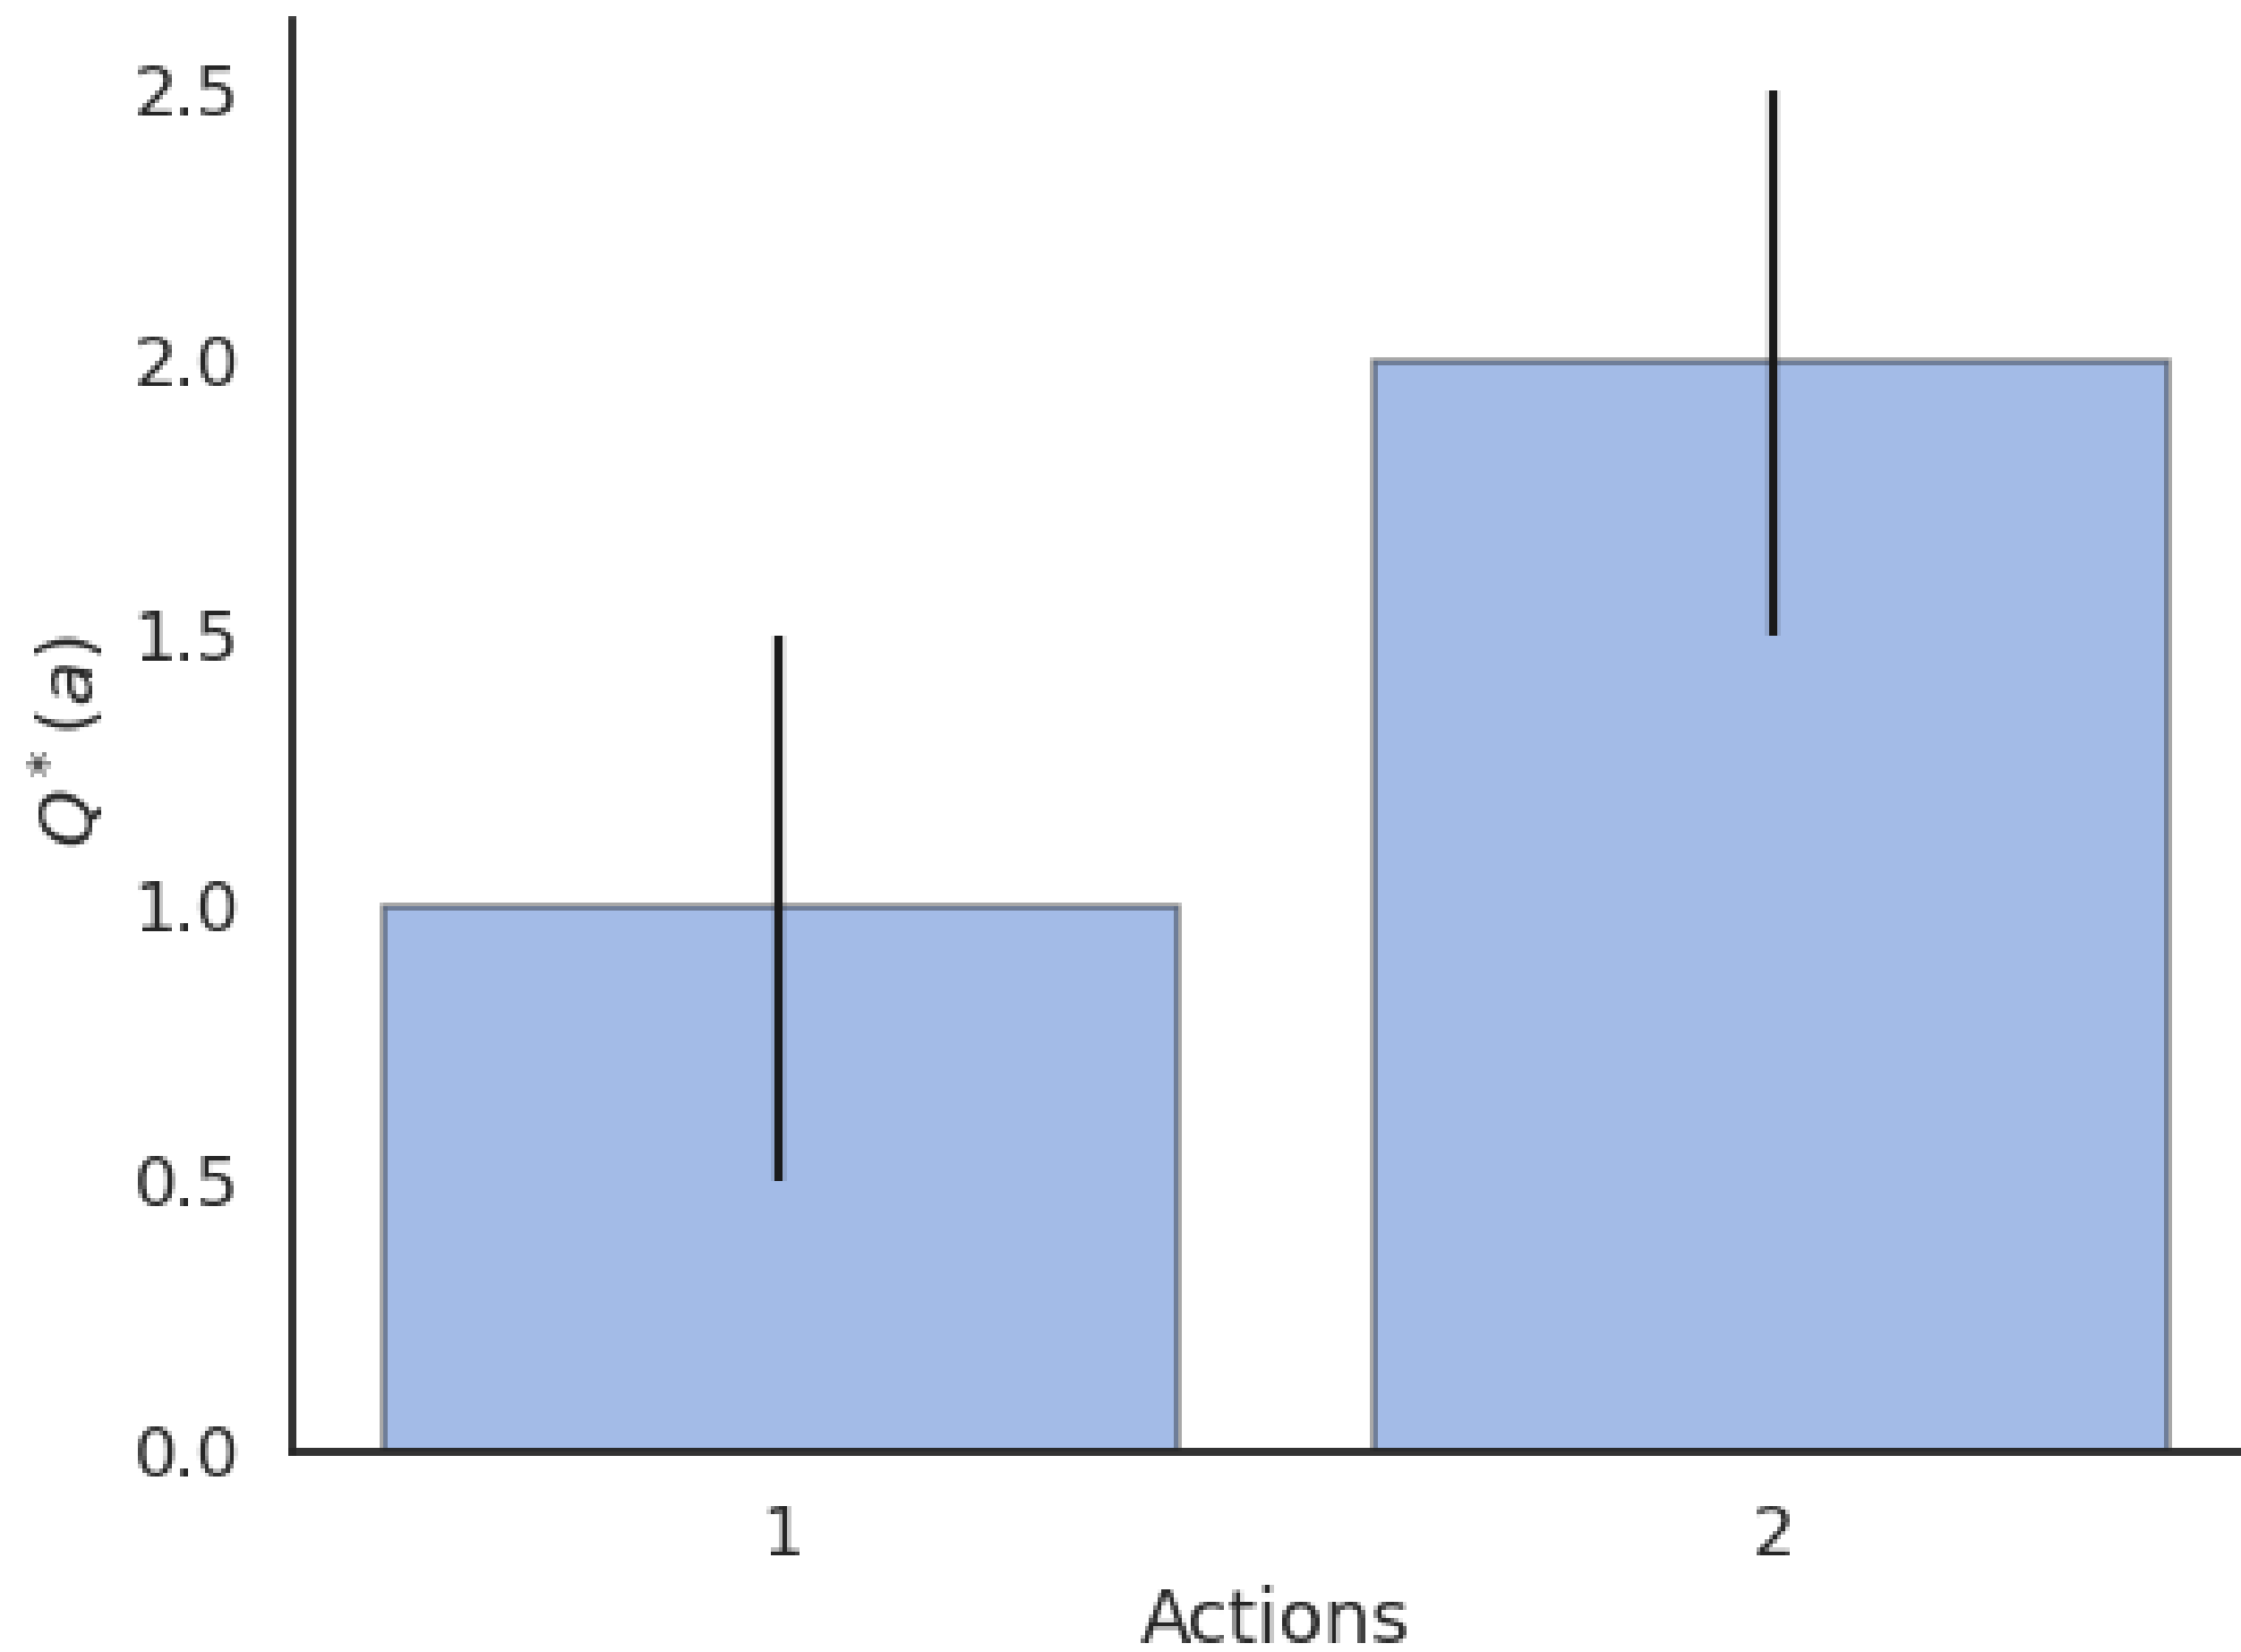
- This defines the **greedy policy**, where the probability of taking the greedy action is 1 and the probability of selecting another action is 0:

$$\pi(a) = \begin{cases} 1 & \text{if } a = a_t^* \\ 0 & \text{otherwise.} \end{cases}$$

- The greedy policy is **deterministic**: the action taken is always the same for a fixed  $Q_t$ .

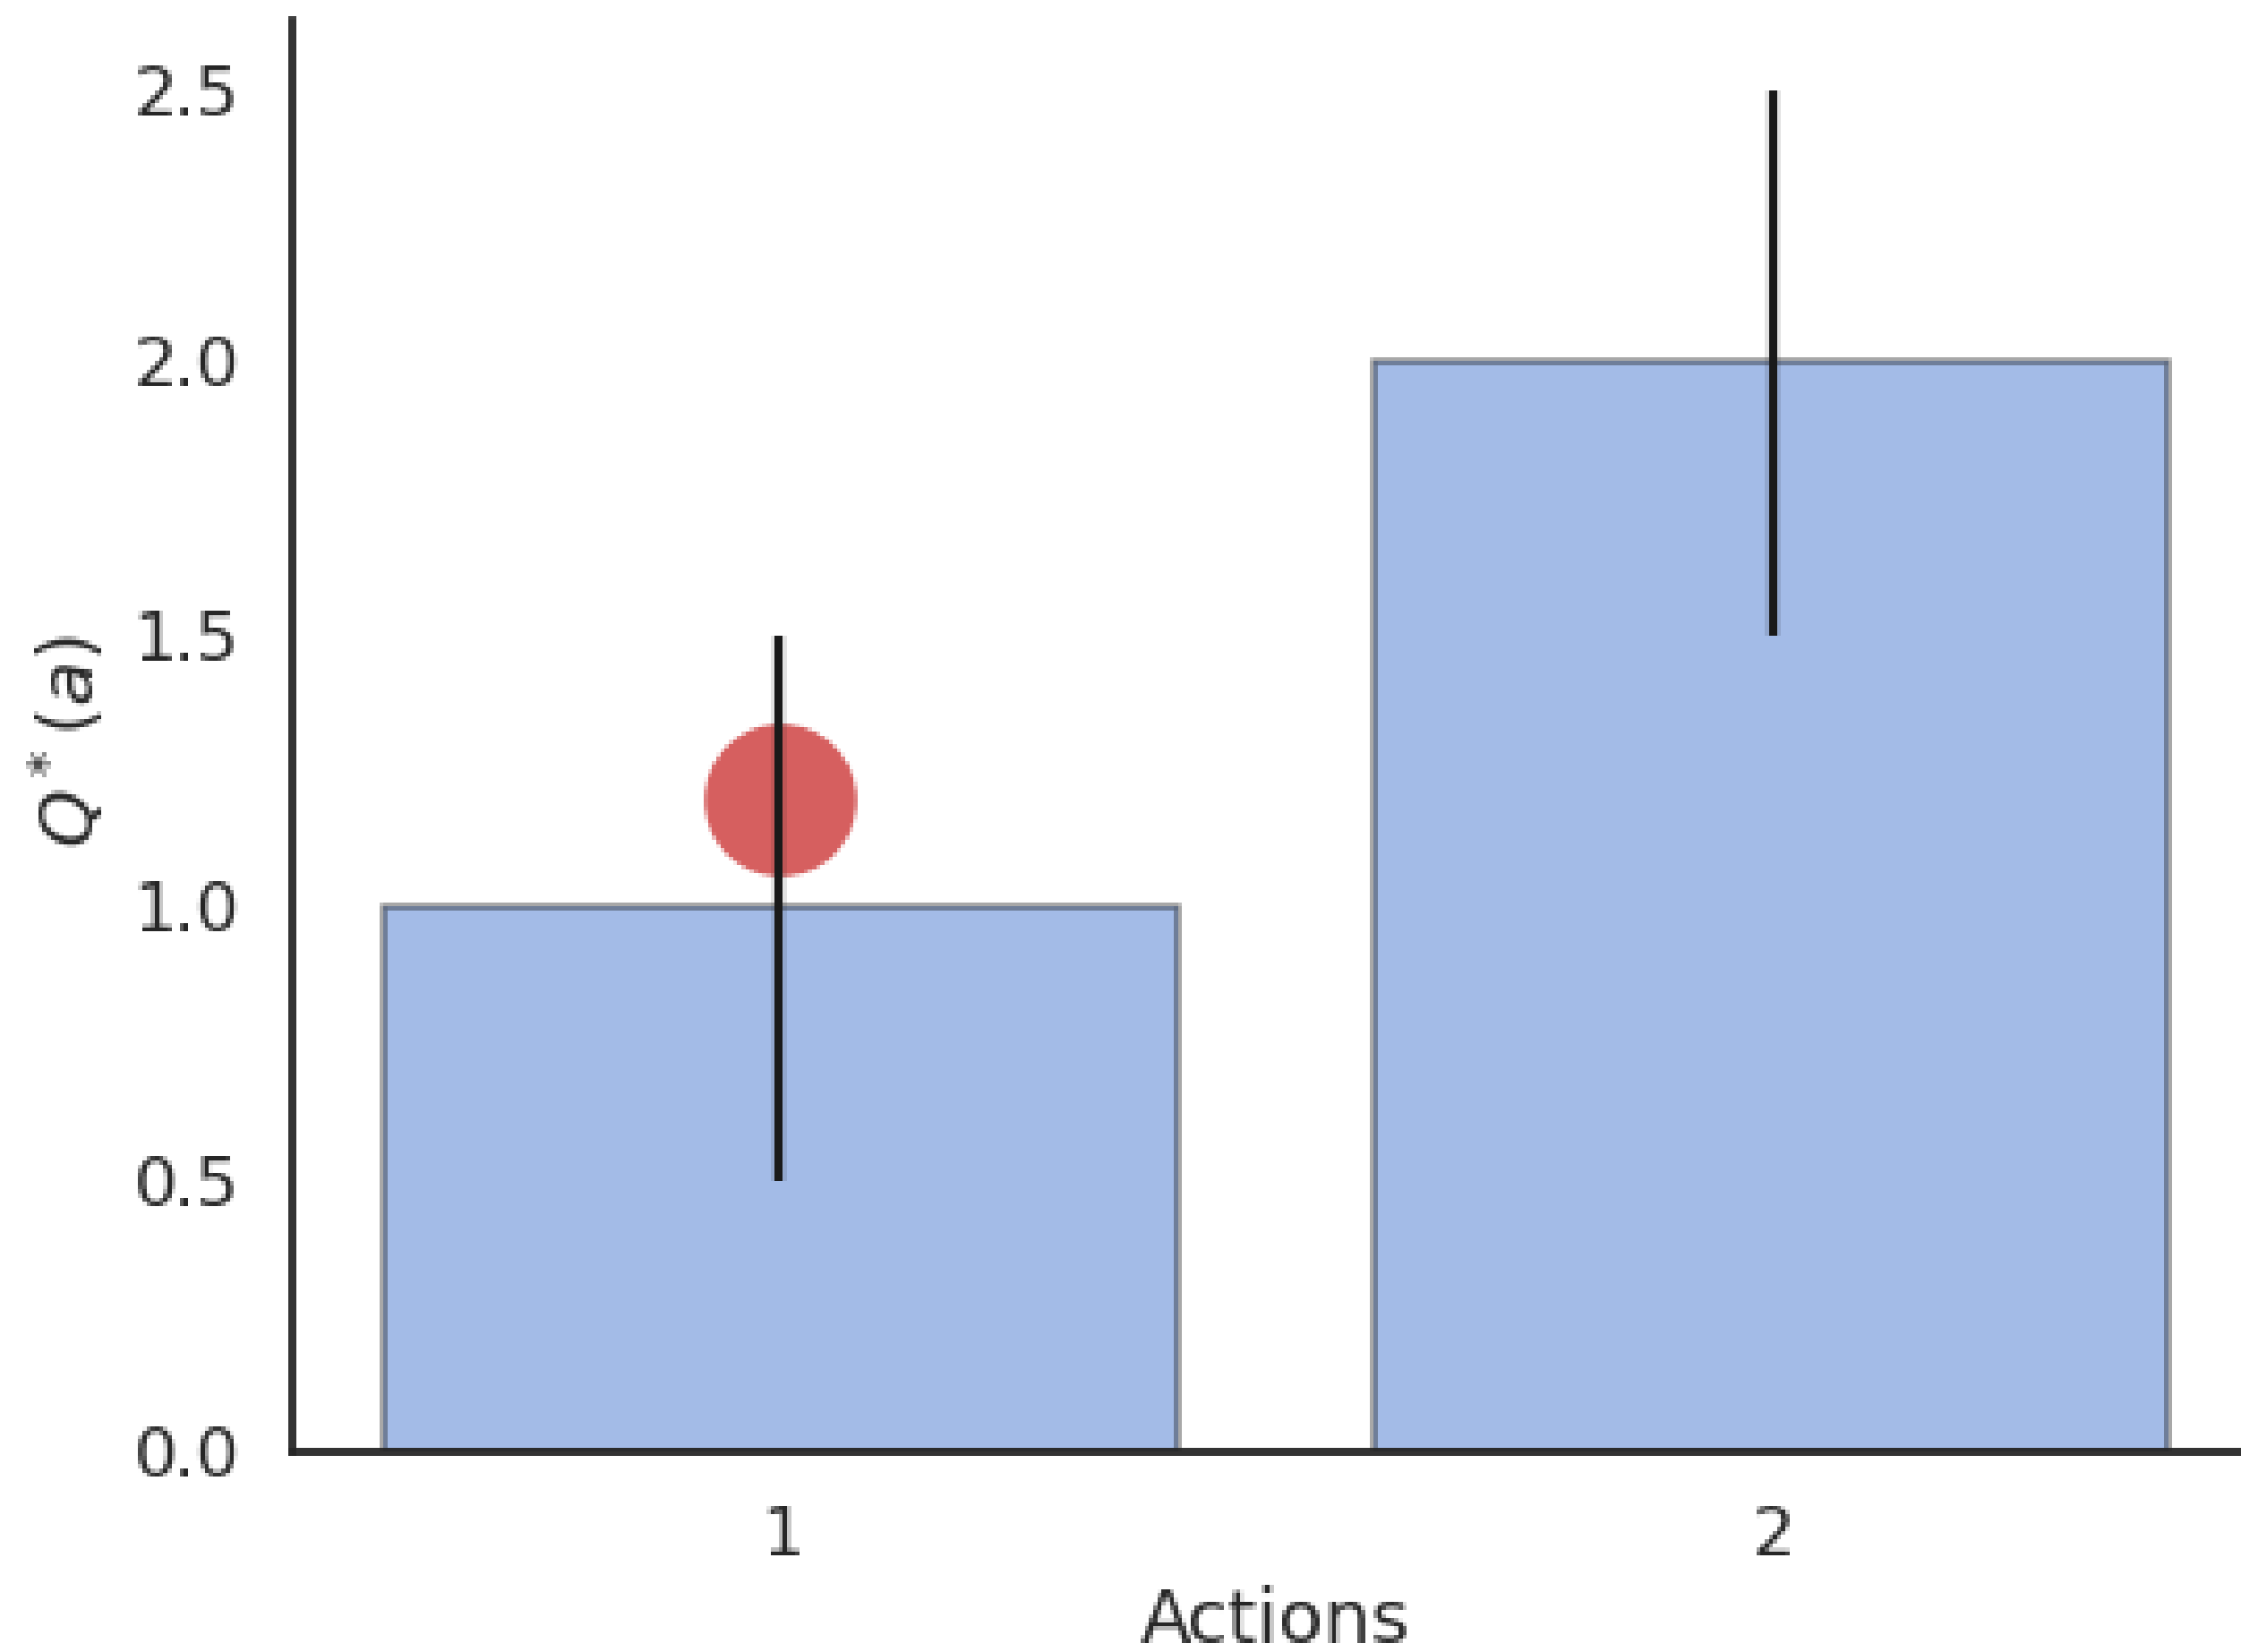
# Problem with greedy action selection

- Greedy action selection only works when the estimates are good enough.



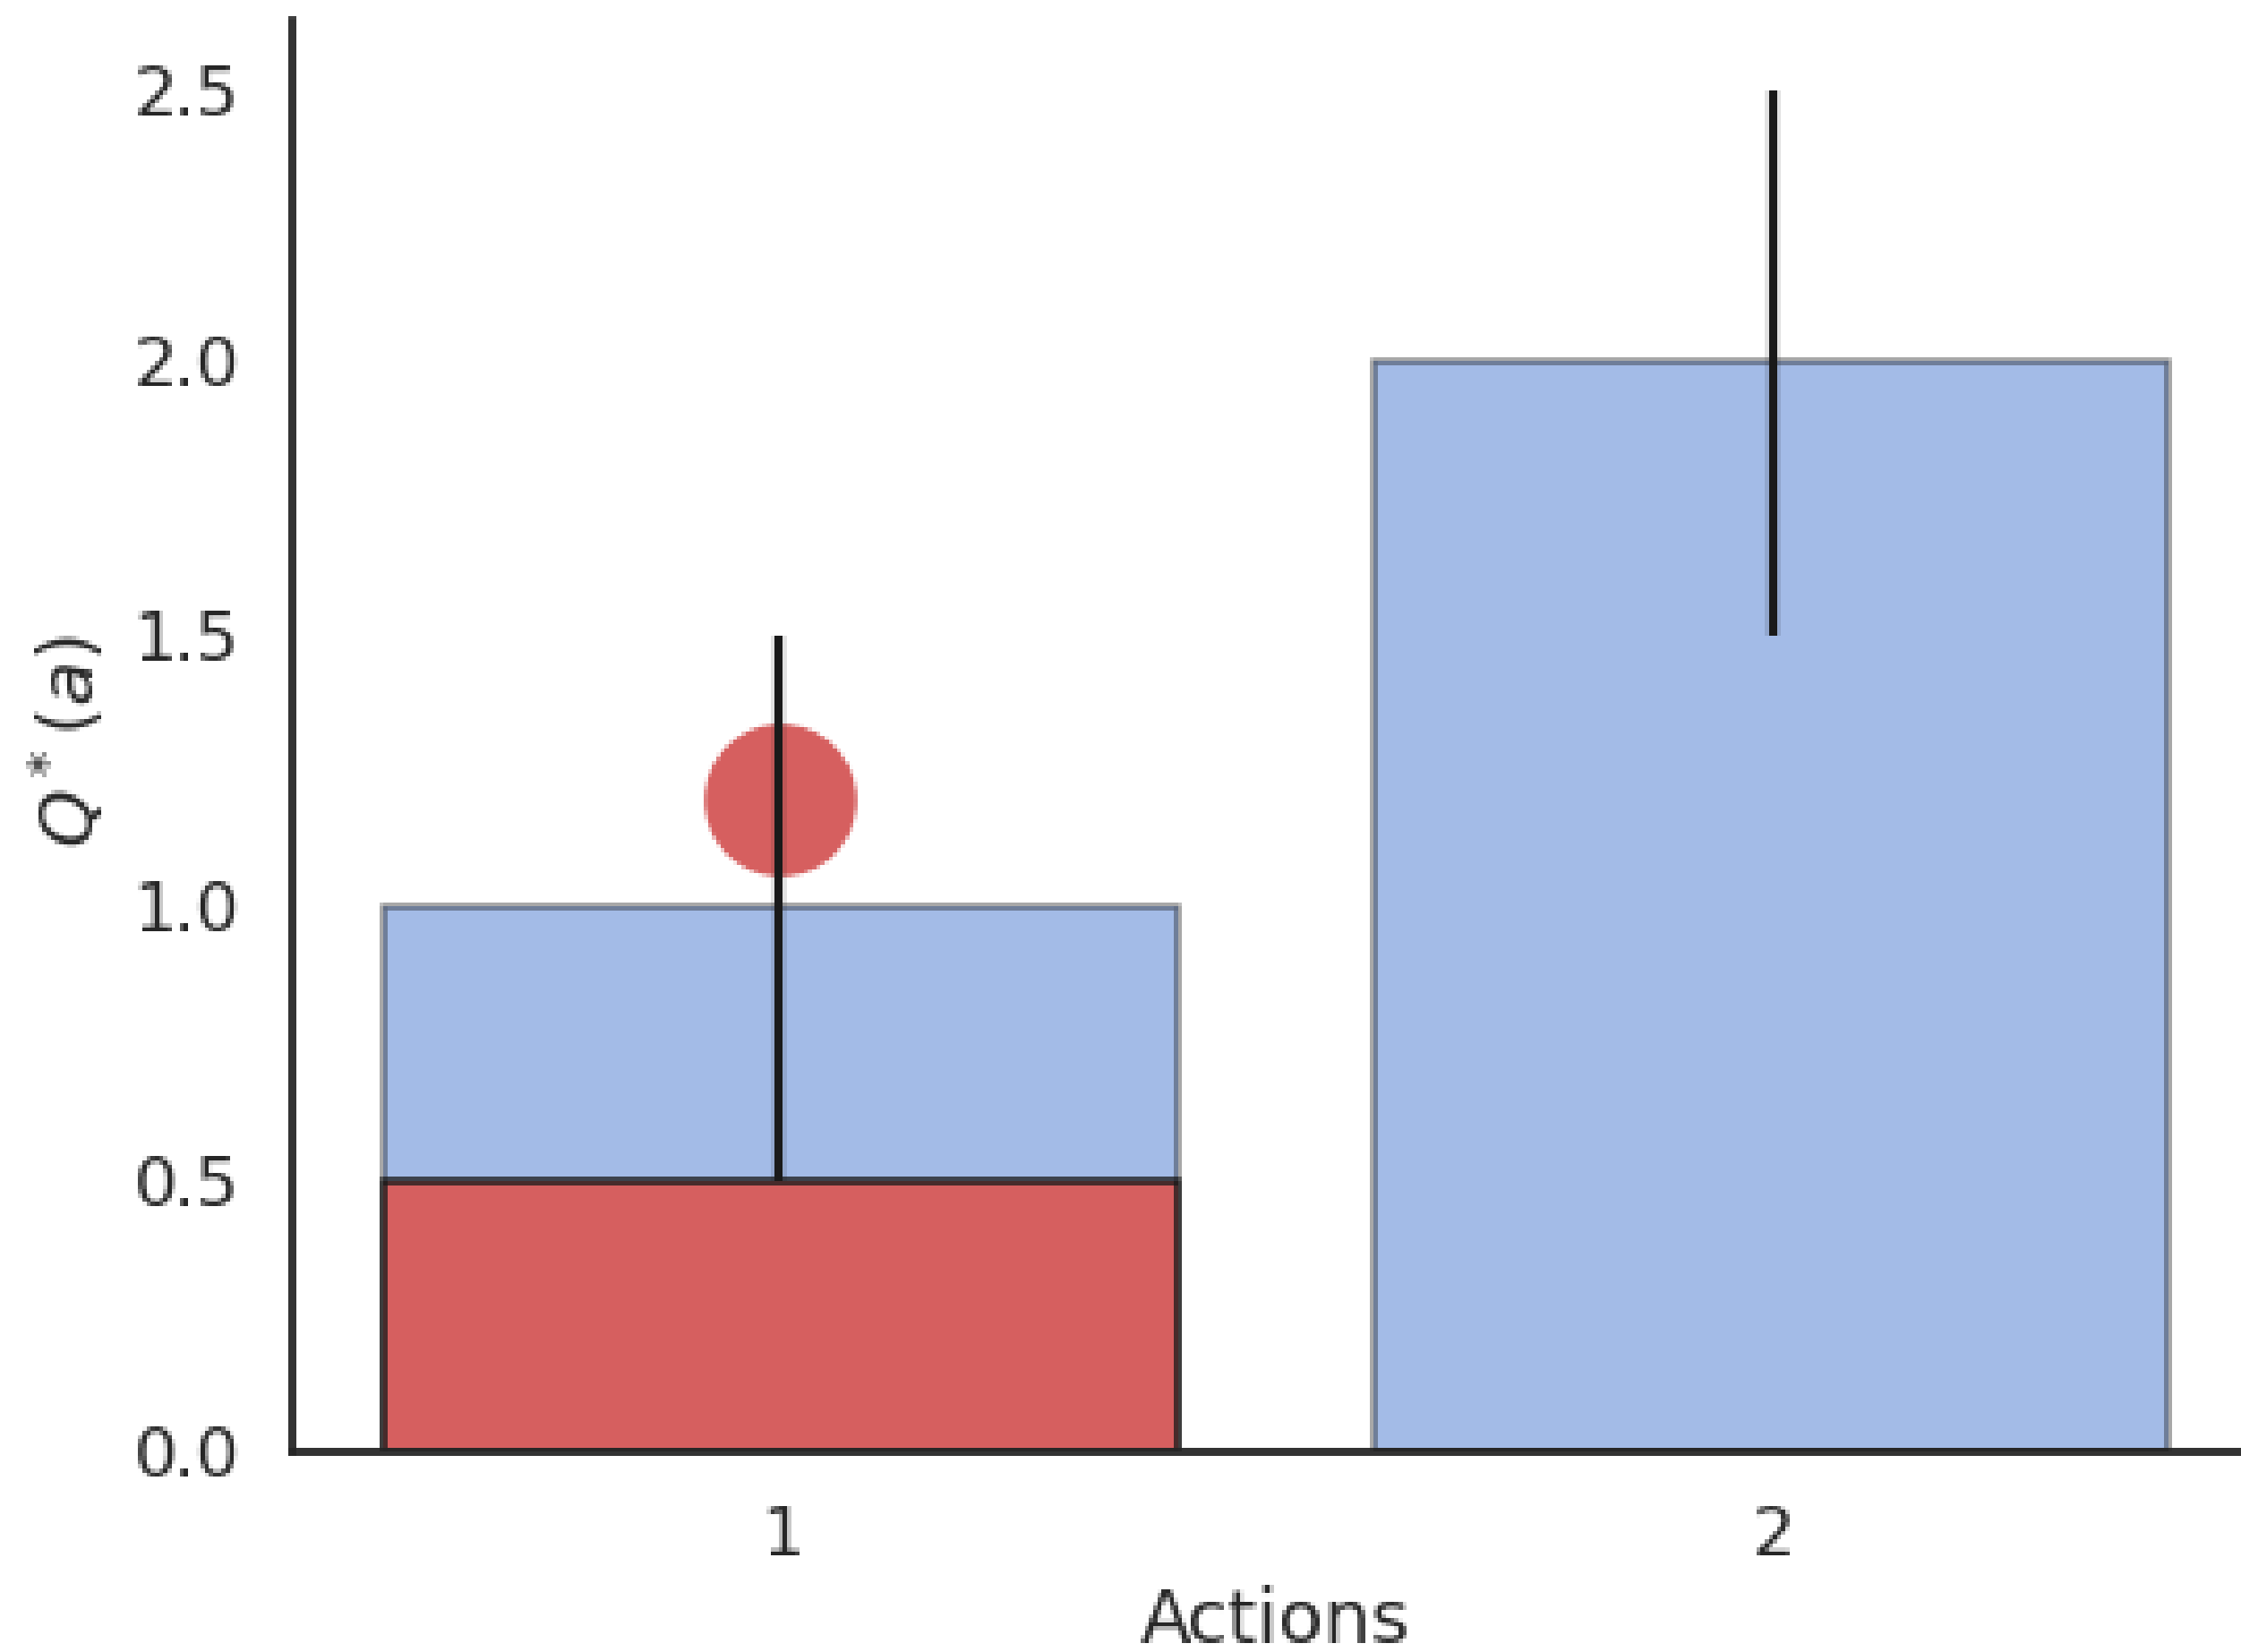
# Problem with greedy action selection

- Estimates are initially bad (e.g. 0 here), so an action is sampled randomly and a reward is received.



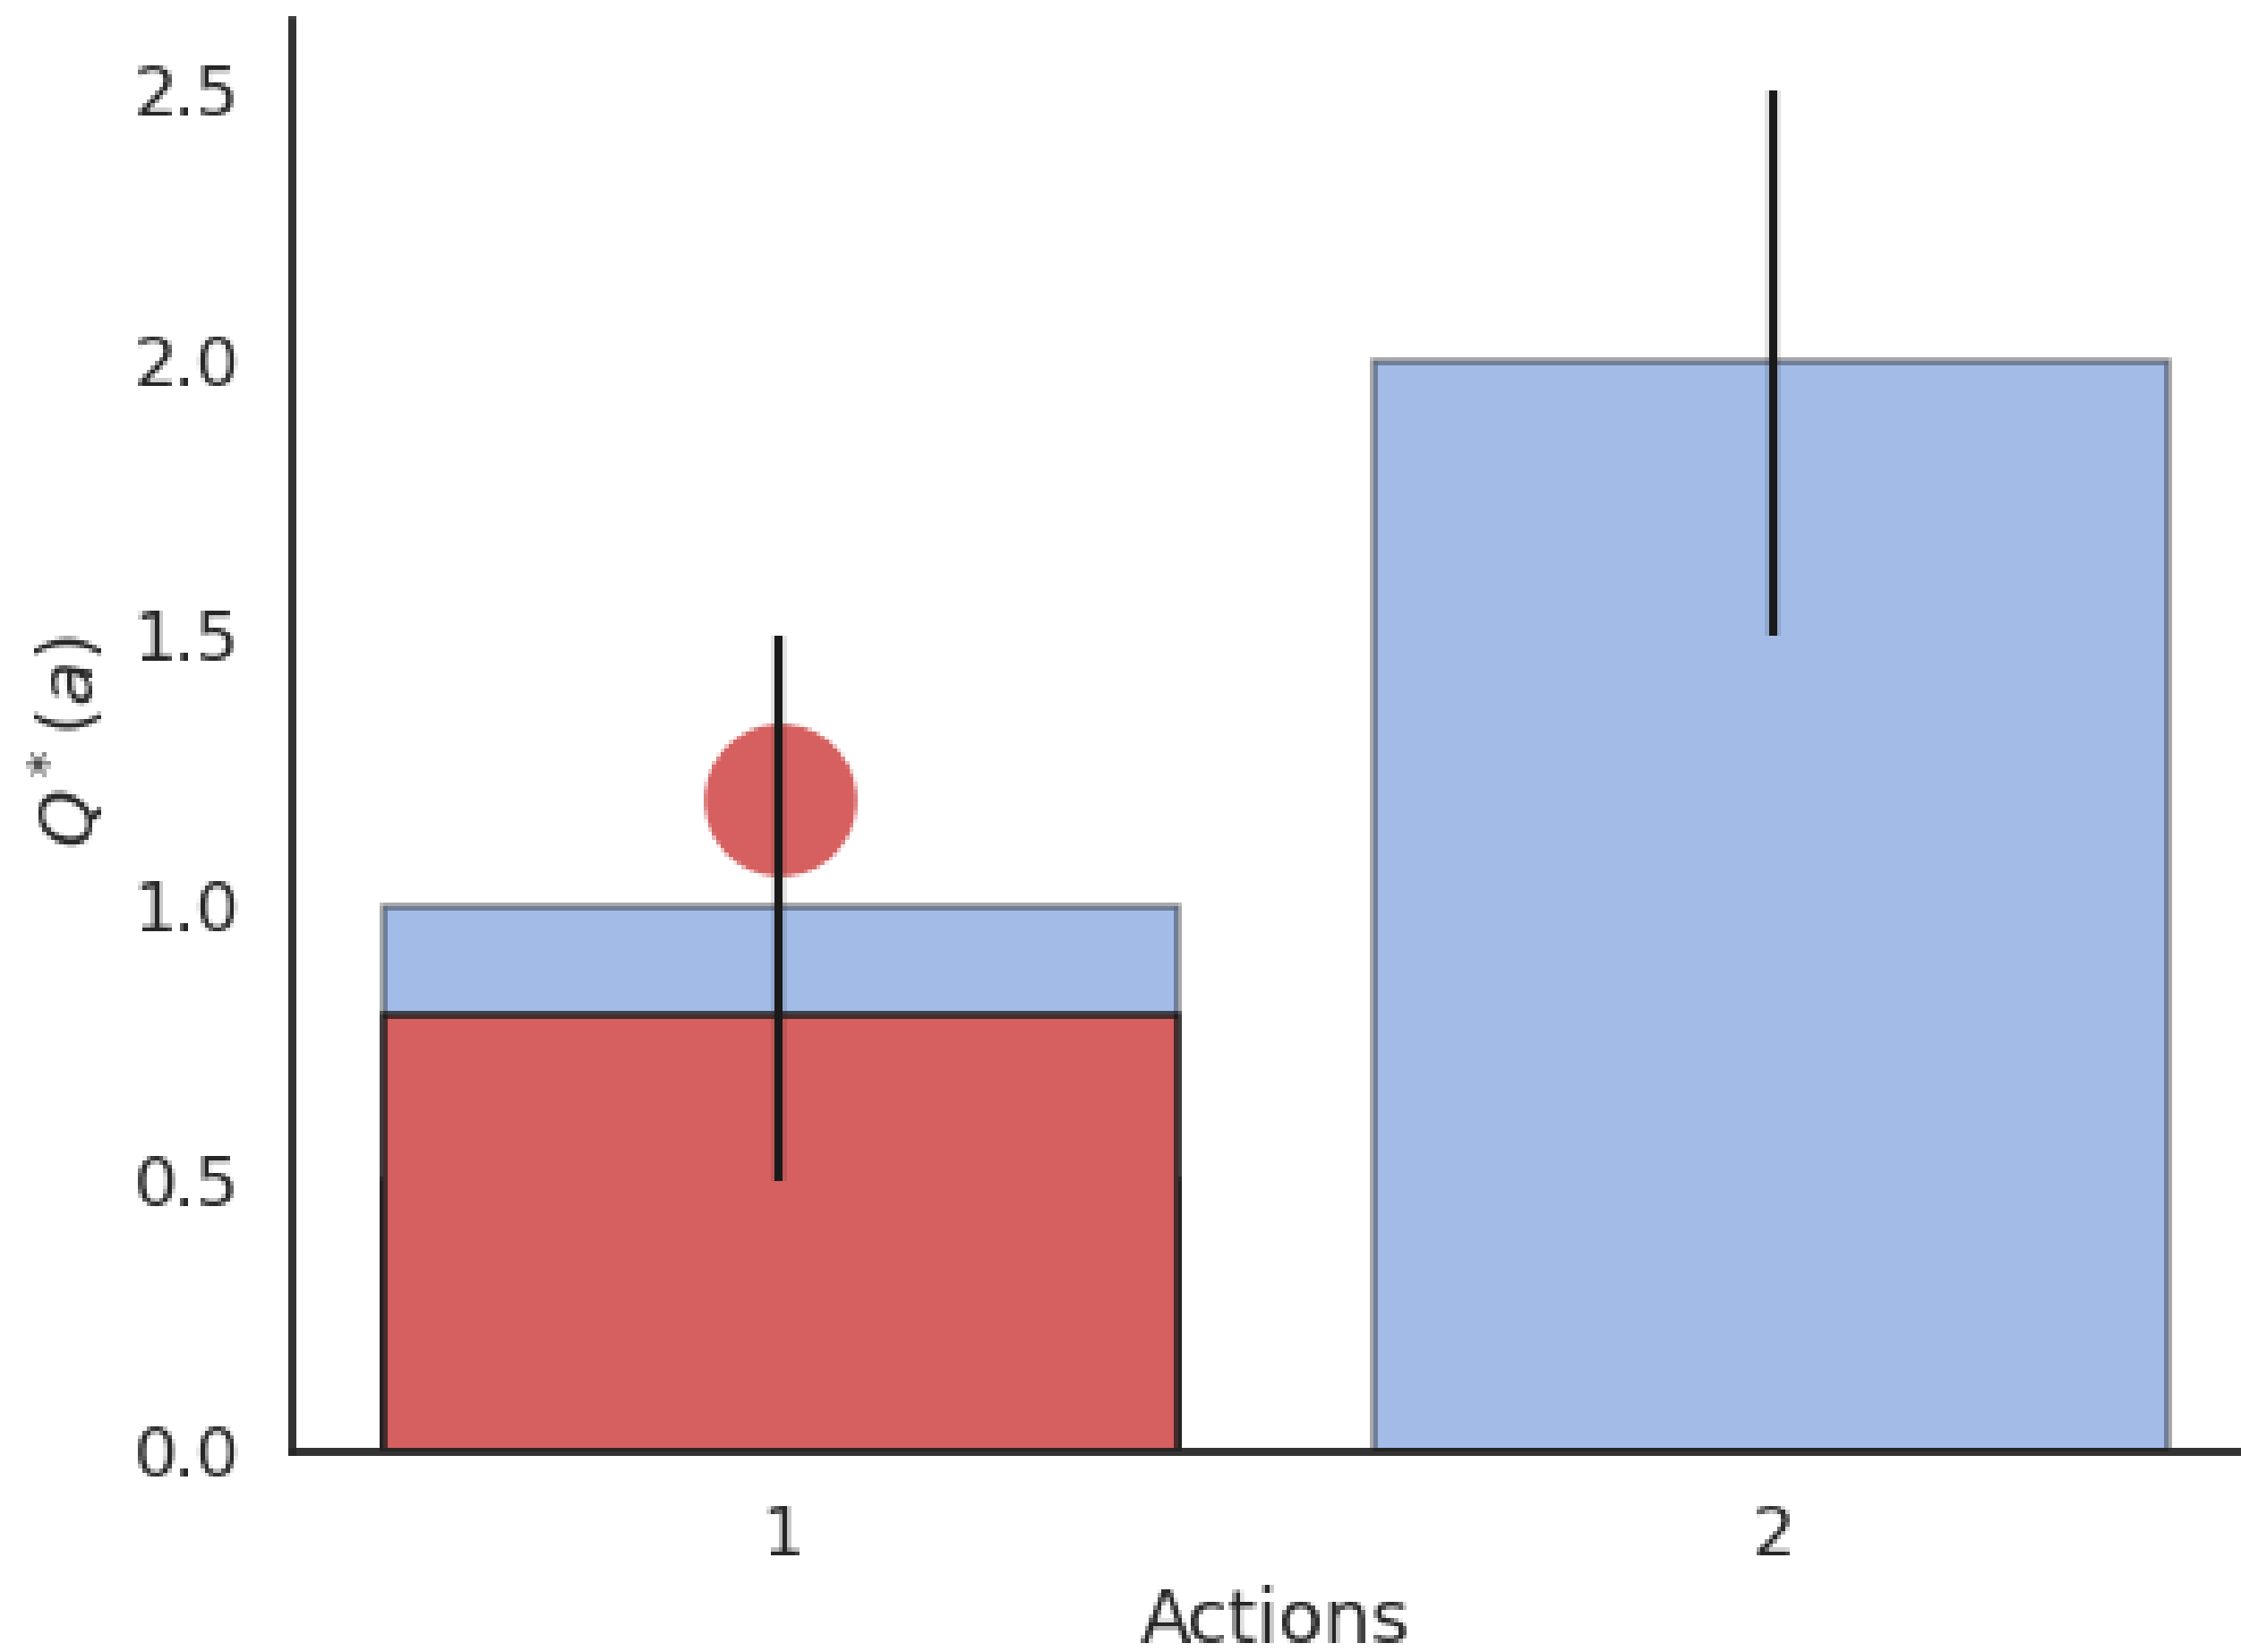
# Problem with greedy action selection

- The Q-value of that action becomes positive, so it becomes the greedy action.

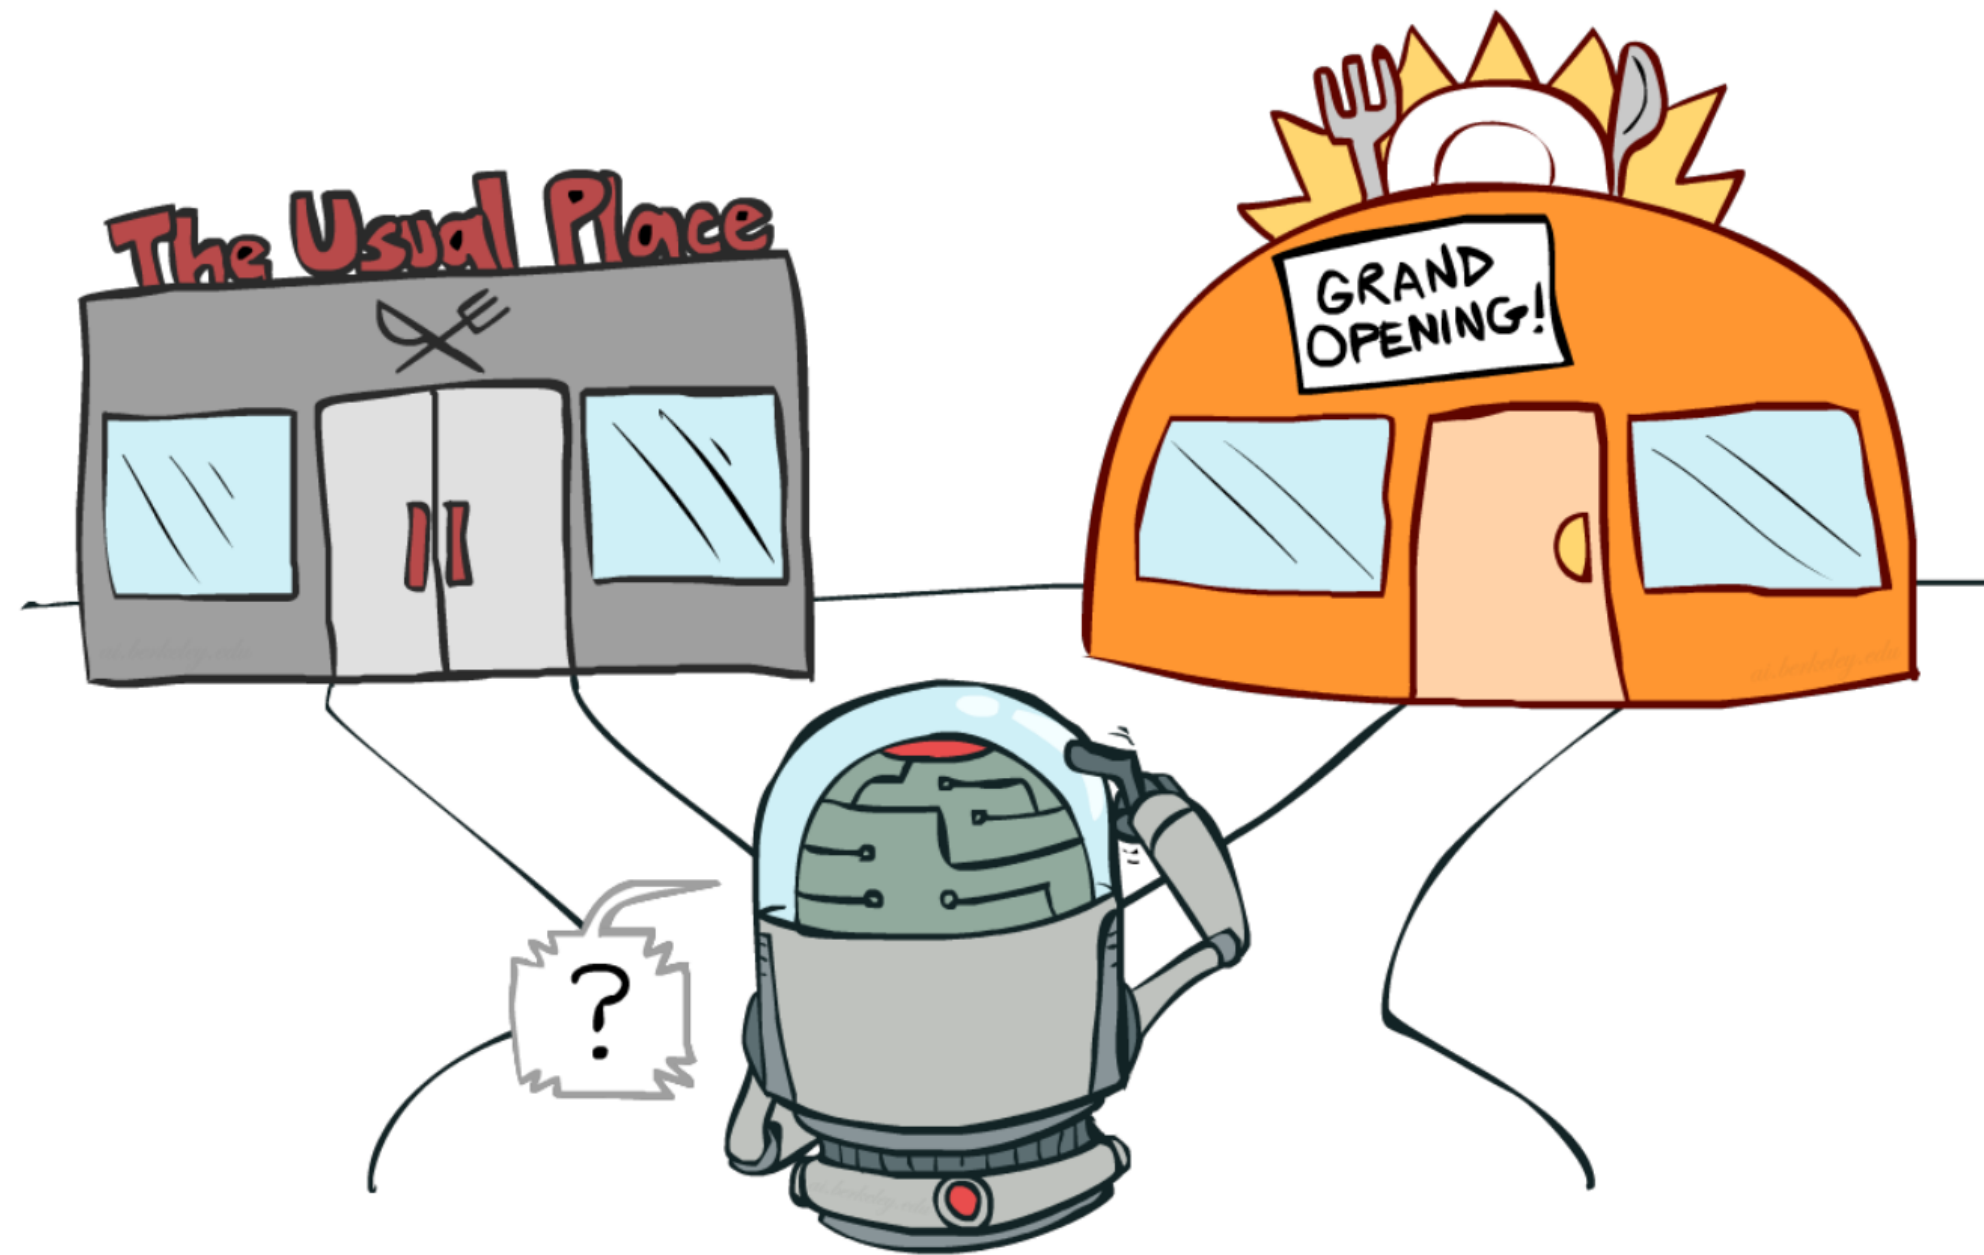


# Problem with greedy action selection

- Greedy action selection will always select that action, although the second one would have been better.

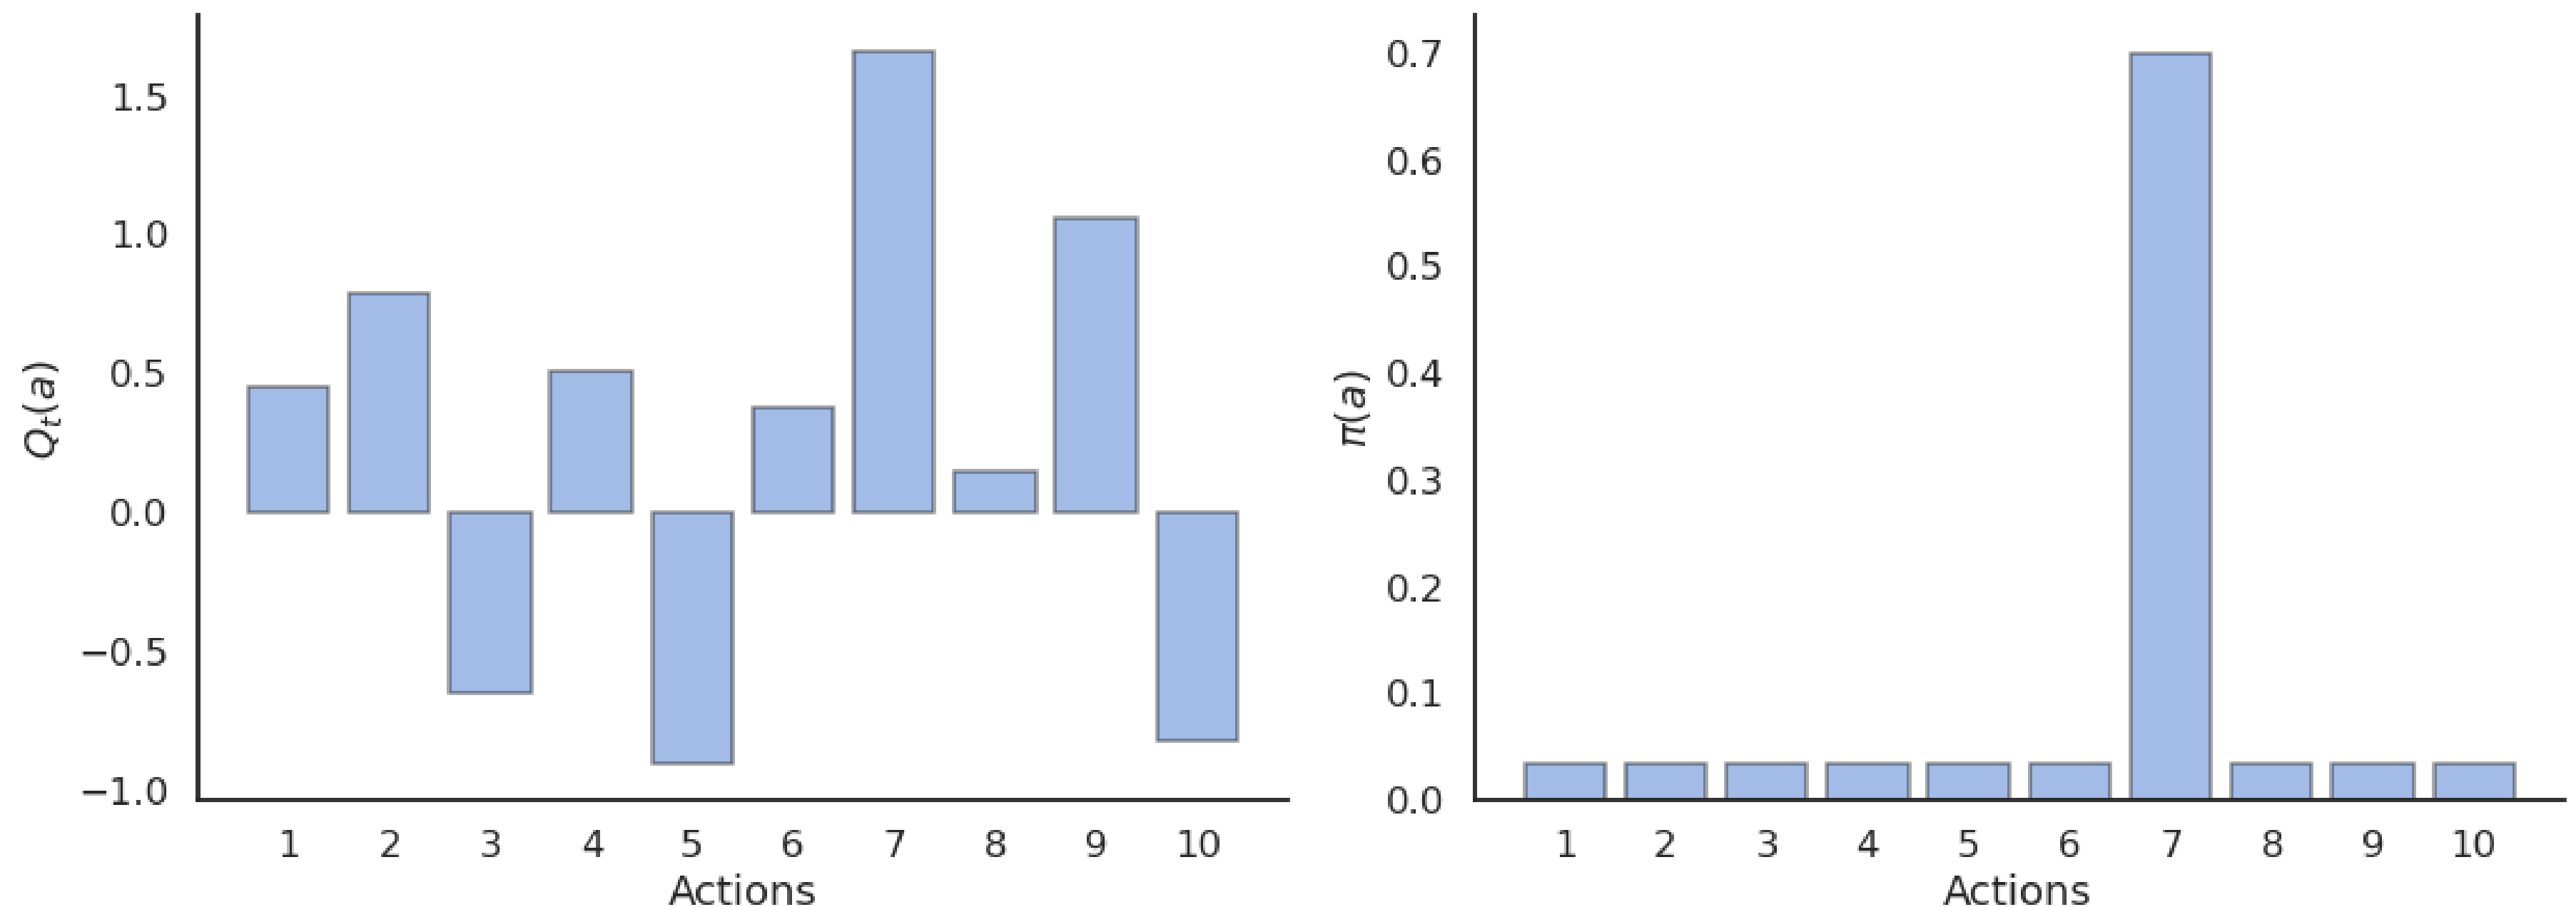


# Exploration-exploitation dilemma



- This **exploration-exploitation** dilemma is the hardest problem in RL:
  - **Exploitation** is using the current estimates to select an action: they might be wrong!
  - **Exploration** is selecting non-greedy actions in order to improve their estimates: they might not be optimal!
- One has to balance exploration and exploitation over the course of learning:
  - More exploration at the beginning of learning, as the estimates are initially wrong.
  - More exploitation at the end of learning, as the estimates get better.

## $\epsilon$ -greedy action selection

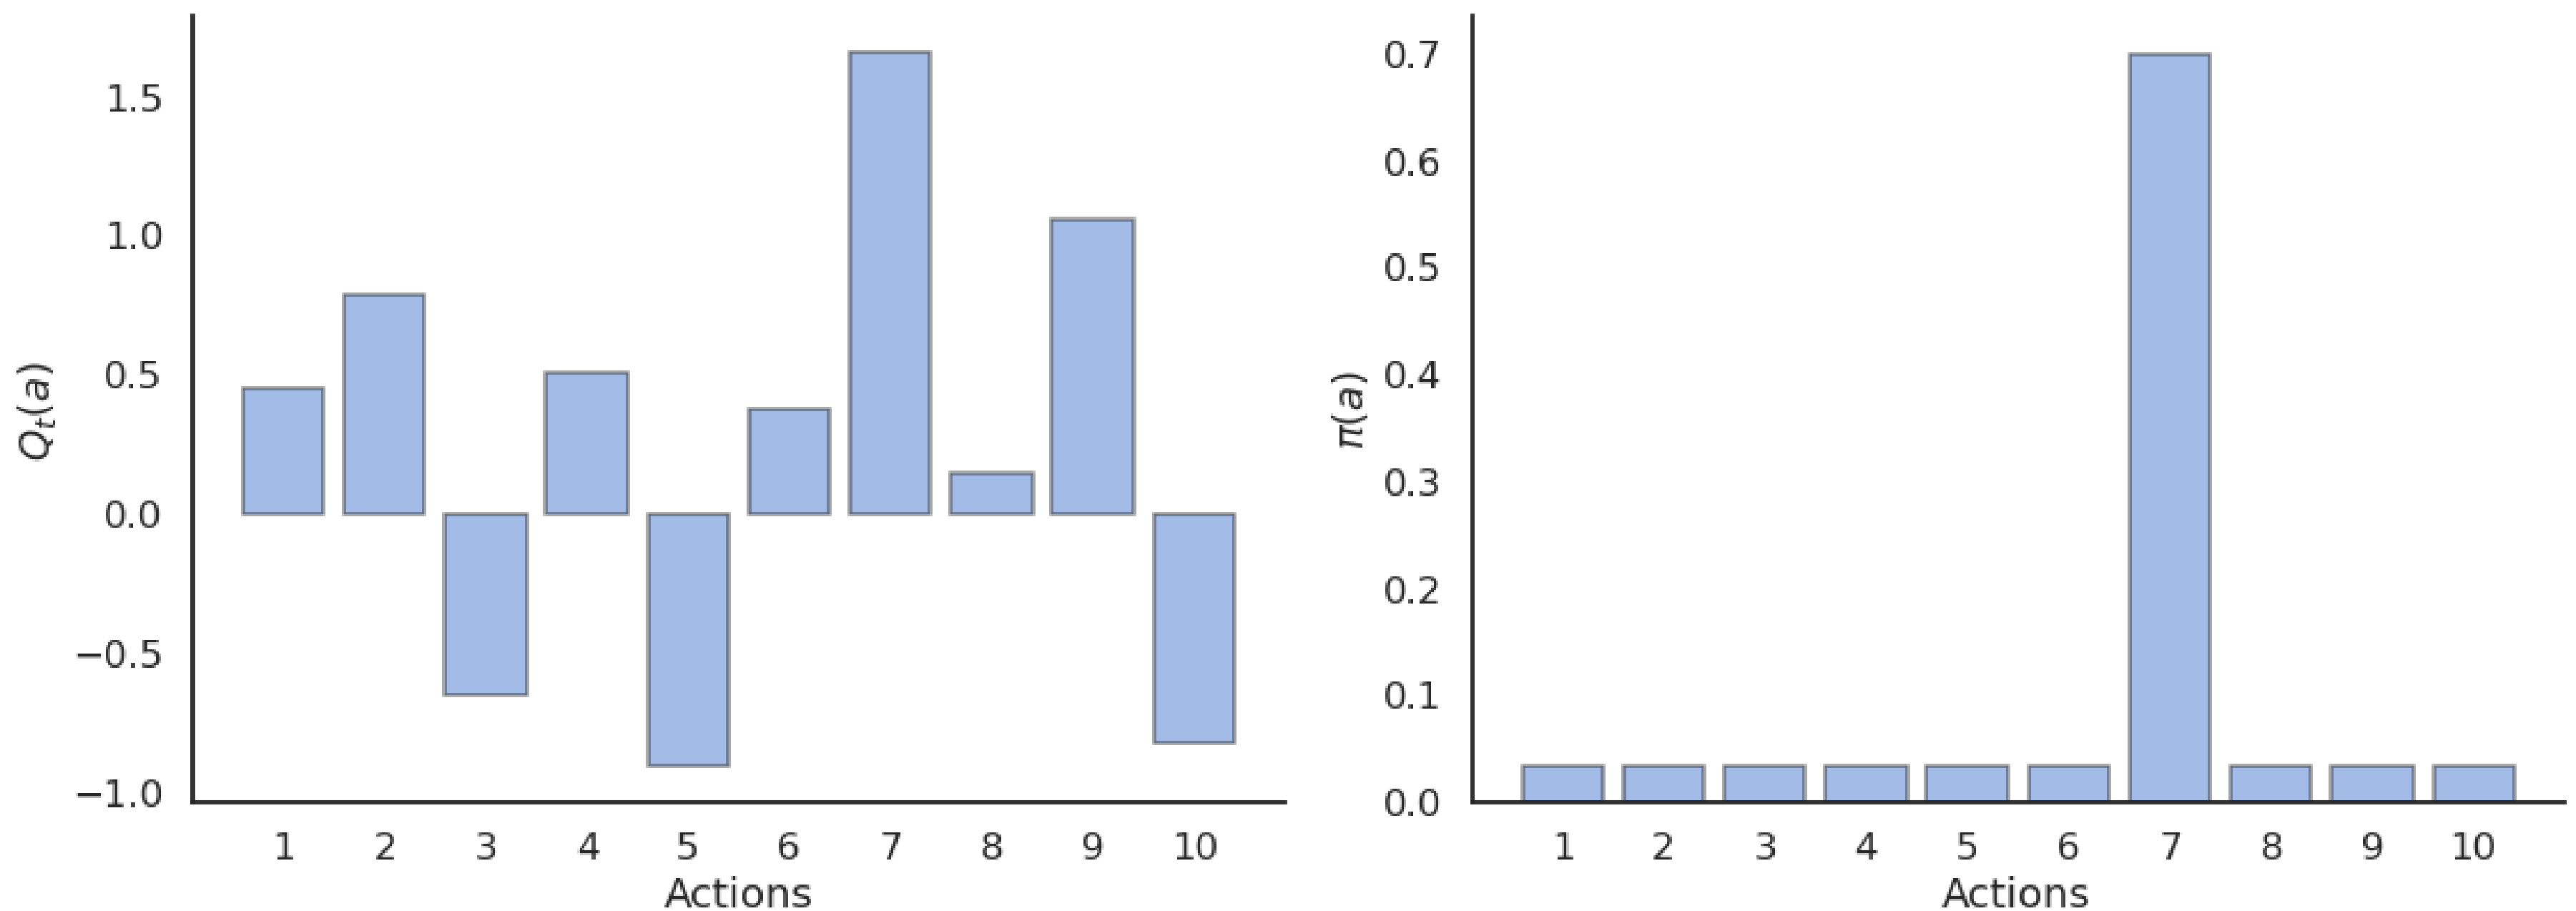


- **$\epsilon$ -greedy action selection** ensures a trade-off between exploitation and exploration.
- The greedy action is selected with probability  $1 - \epsilon$  (with  $0 < \epsilon < 1$ ), the others with probability  $\epsilon$ :

$$\pi(a) = \begin{cases} 1 - \epsilon & \text{if } a = a_t^* \\ \frac{\epsilon}{|\mathcal{A}|-1} & \text{otherwise.} \end{cases}$$




## $\epsilon$ -greedy action selection



- The parameter  $\epsilon$  controls the level of exploration: the higher  $\epsilon$ , the more exploration.
- One can set  $\epsilon$  high at the beginning of learning and progressively reduce it to exploit more.
- However, it chooses equally among all actions: the worst action is as likely to be selected as the next-to-best action.

# Softmax action selection

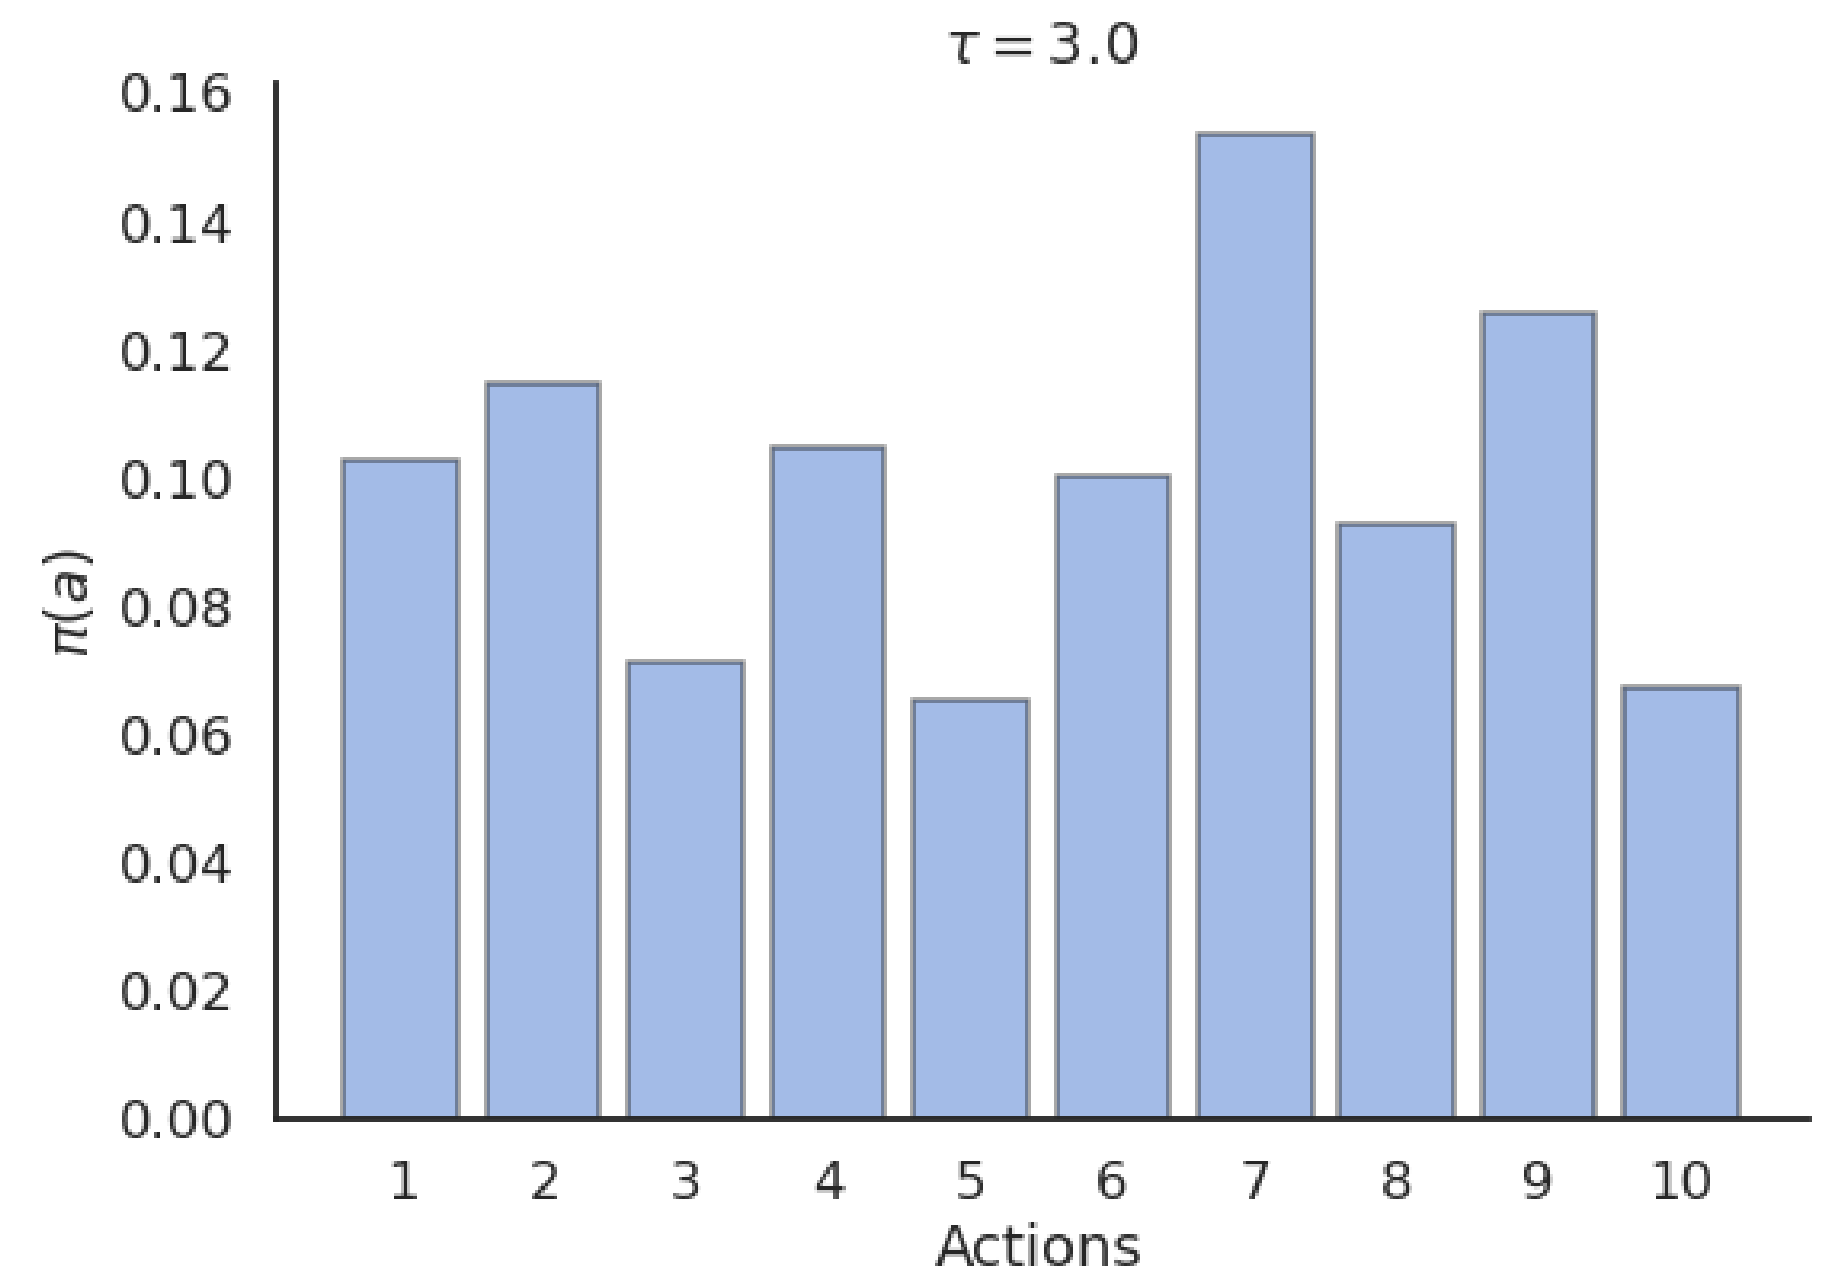
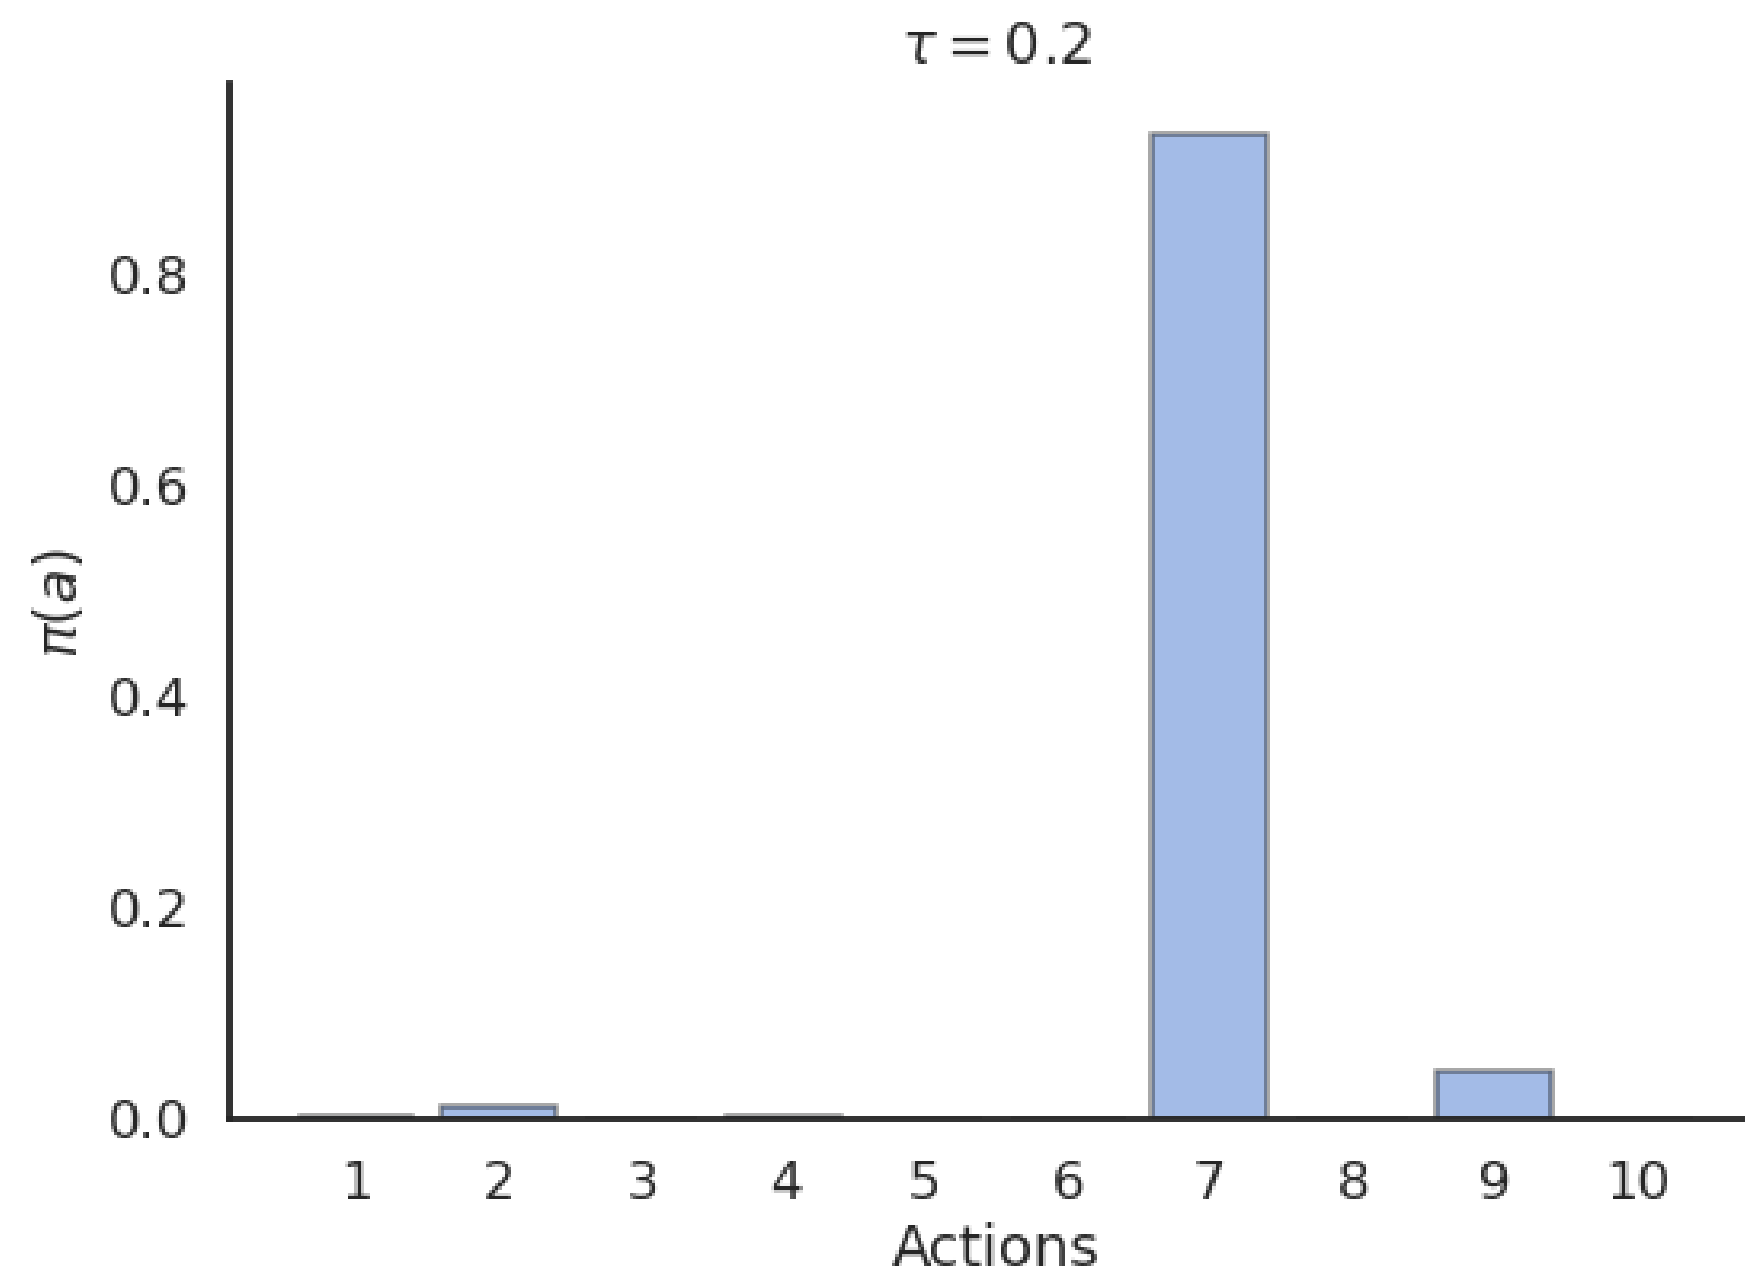


- **Softmax action selection** defines the probability of choosing an action using all estimated value.
- It represents the policy using a Gibbs (or Boltzmann) distribution:

$$\pi(a) = \frac{\exp \frac{Q_t(a)}{\tau}}{\sum_{a'} \exp \frac{Q_t(a')}{\tau}}$$


where  $\tau$  is a positive parameter called the **temperature**.

# Softmax action selection



- Just as  $\epsilon$ , the temperature  $\tau$  controls the level of exploration:
  - High temperature causes the actions to be nearly equiprobable (**random agent**).
  - Low temperature causes the greediest actions only to be selected (**greedy agent**).

## Example of action selection for the 10-armed bandit

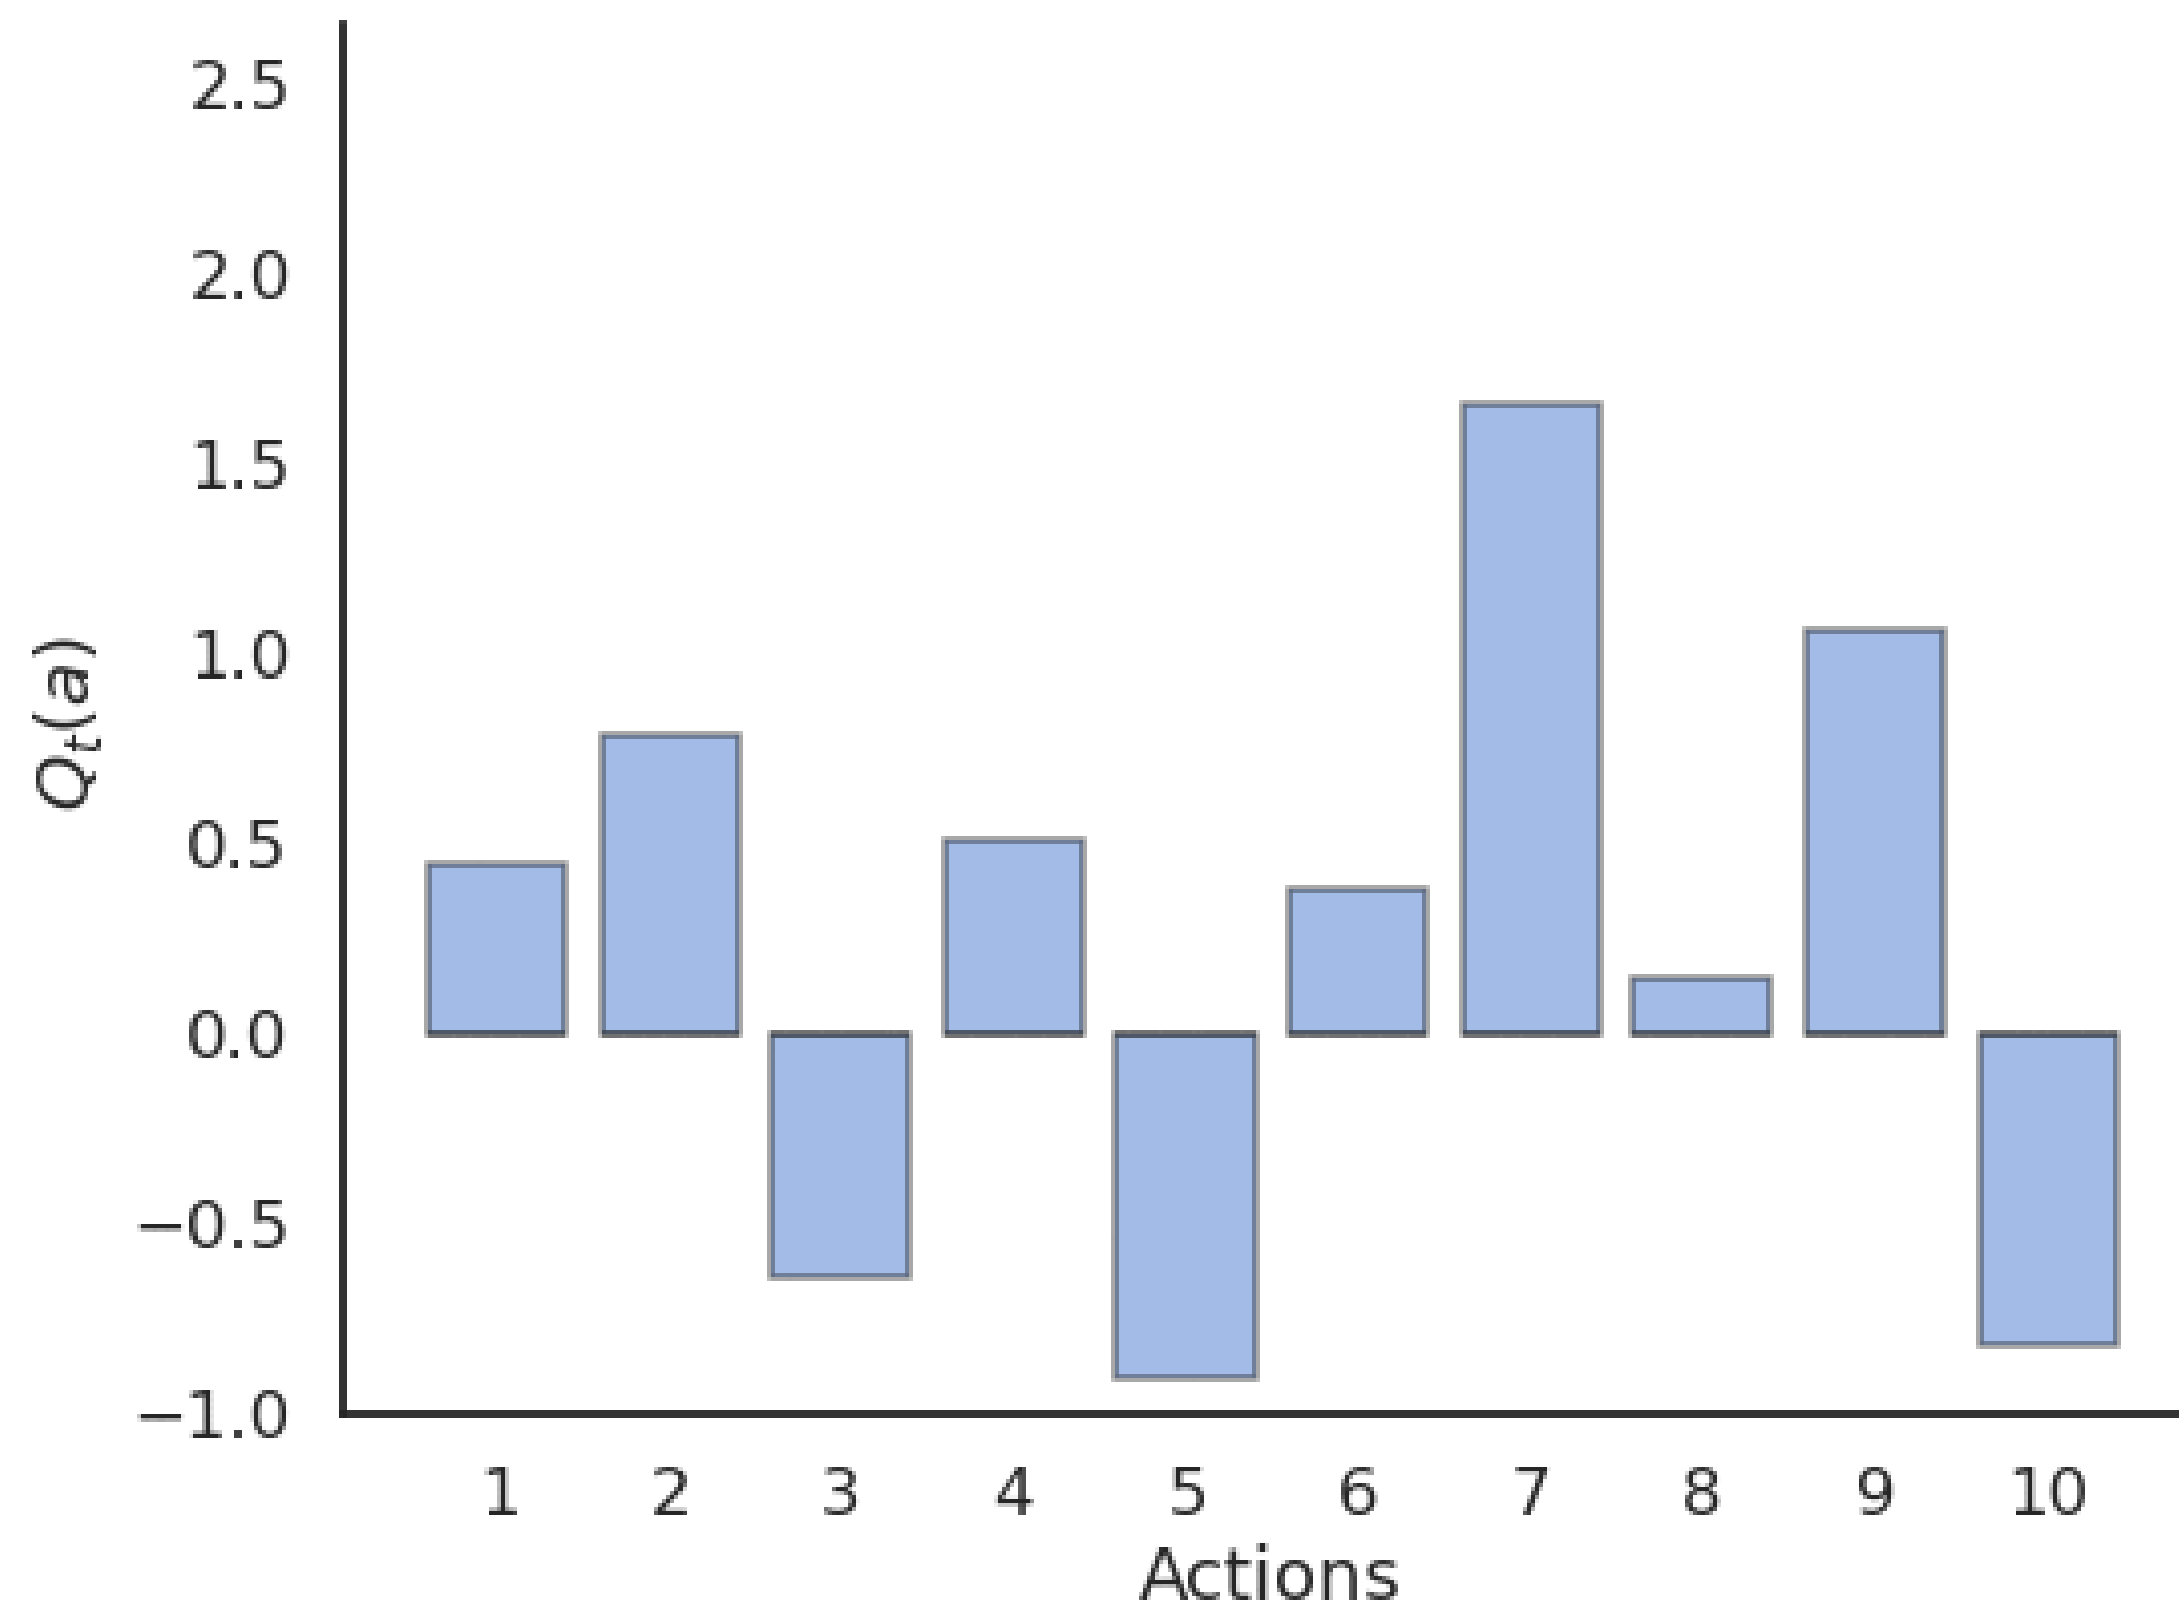


### Procedure as in (Sutton and Barto, 2017):

- $N = 10$  possible actions with Q-values  $Q^*(a_1), \dots, Q^*(a_{10})$  randomly chosen in  $\mathcal{N}(0, 1)$ .
- Each reward  $r_t$  is drawn from a normal distribution  $\mathcal{N}(Q^*(a), 1)$  depending on the selected action.
- Estimates  $Q_t(a)$  are initialized to 0.
- The algorithms run for 1000 plays, and the results are averaged 200 times.

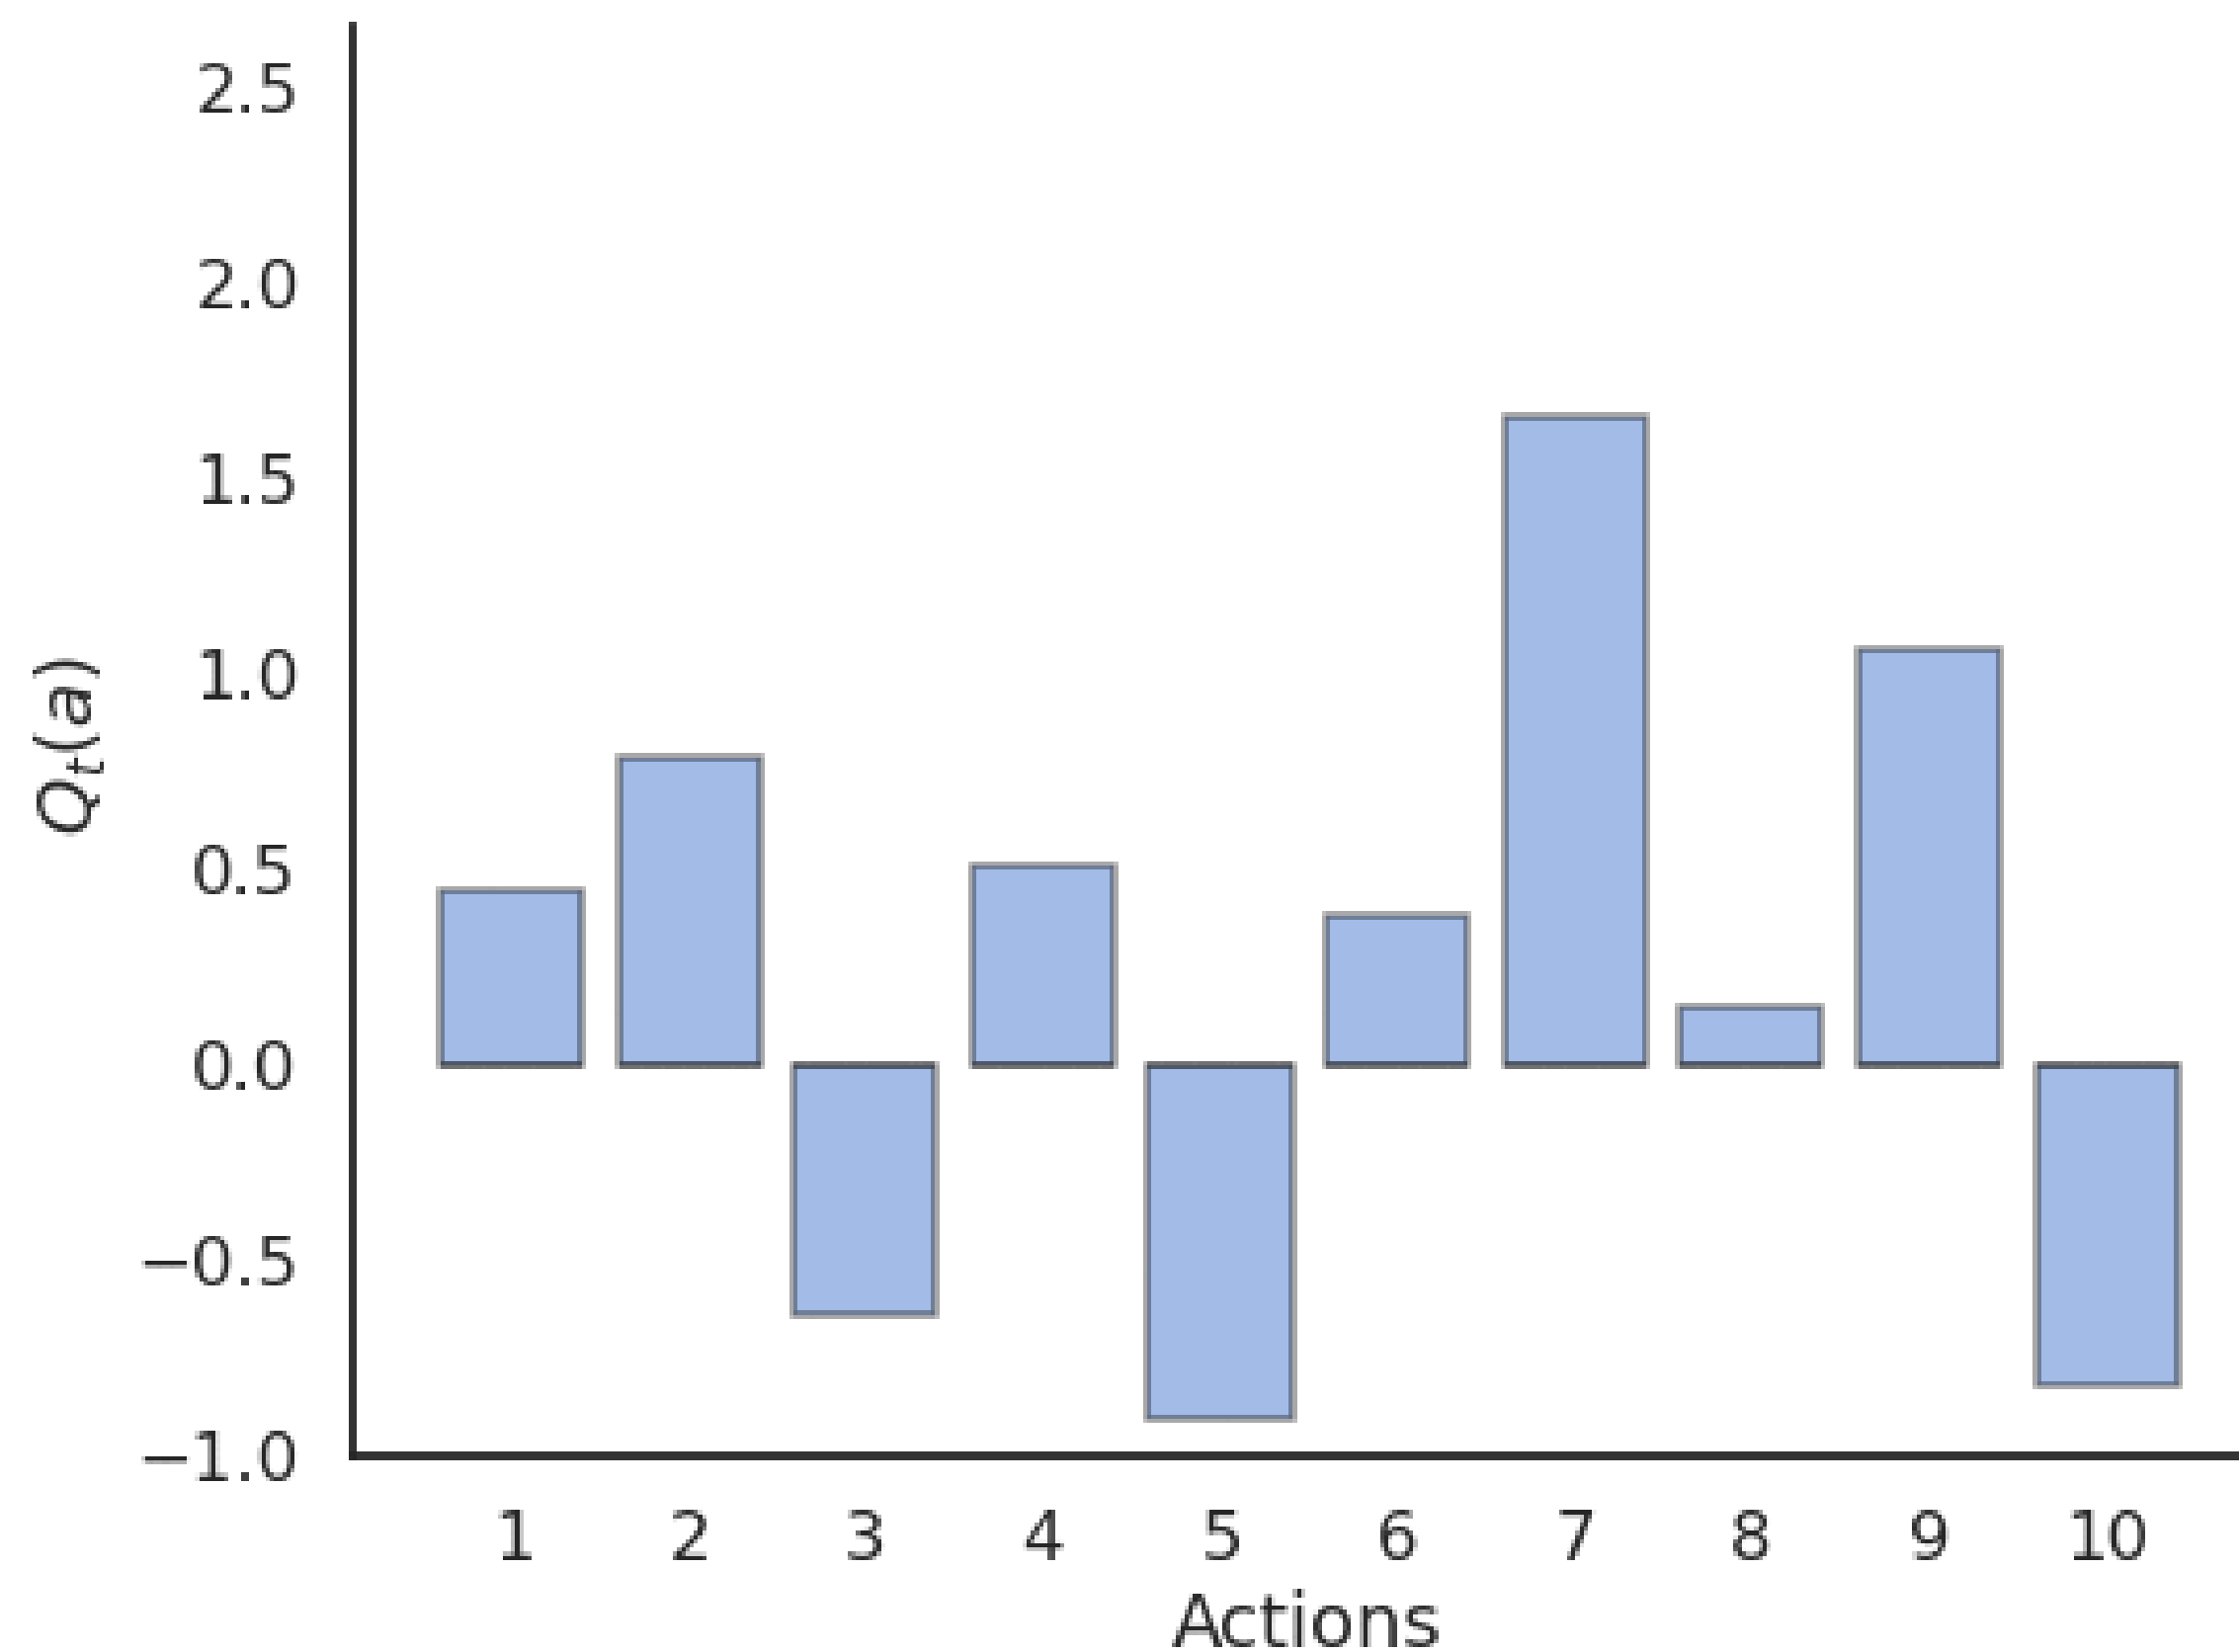
# Greedy action selection

- Greedy action selection allows to get rid quite early of the actions with negative rewards.
- However, it may stick with the first positive action it finds, probably not the optimal one.
- The more actions you have, the more likely you will get stuck in a **suboptimal policy**.



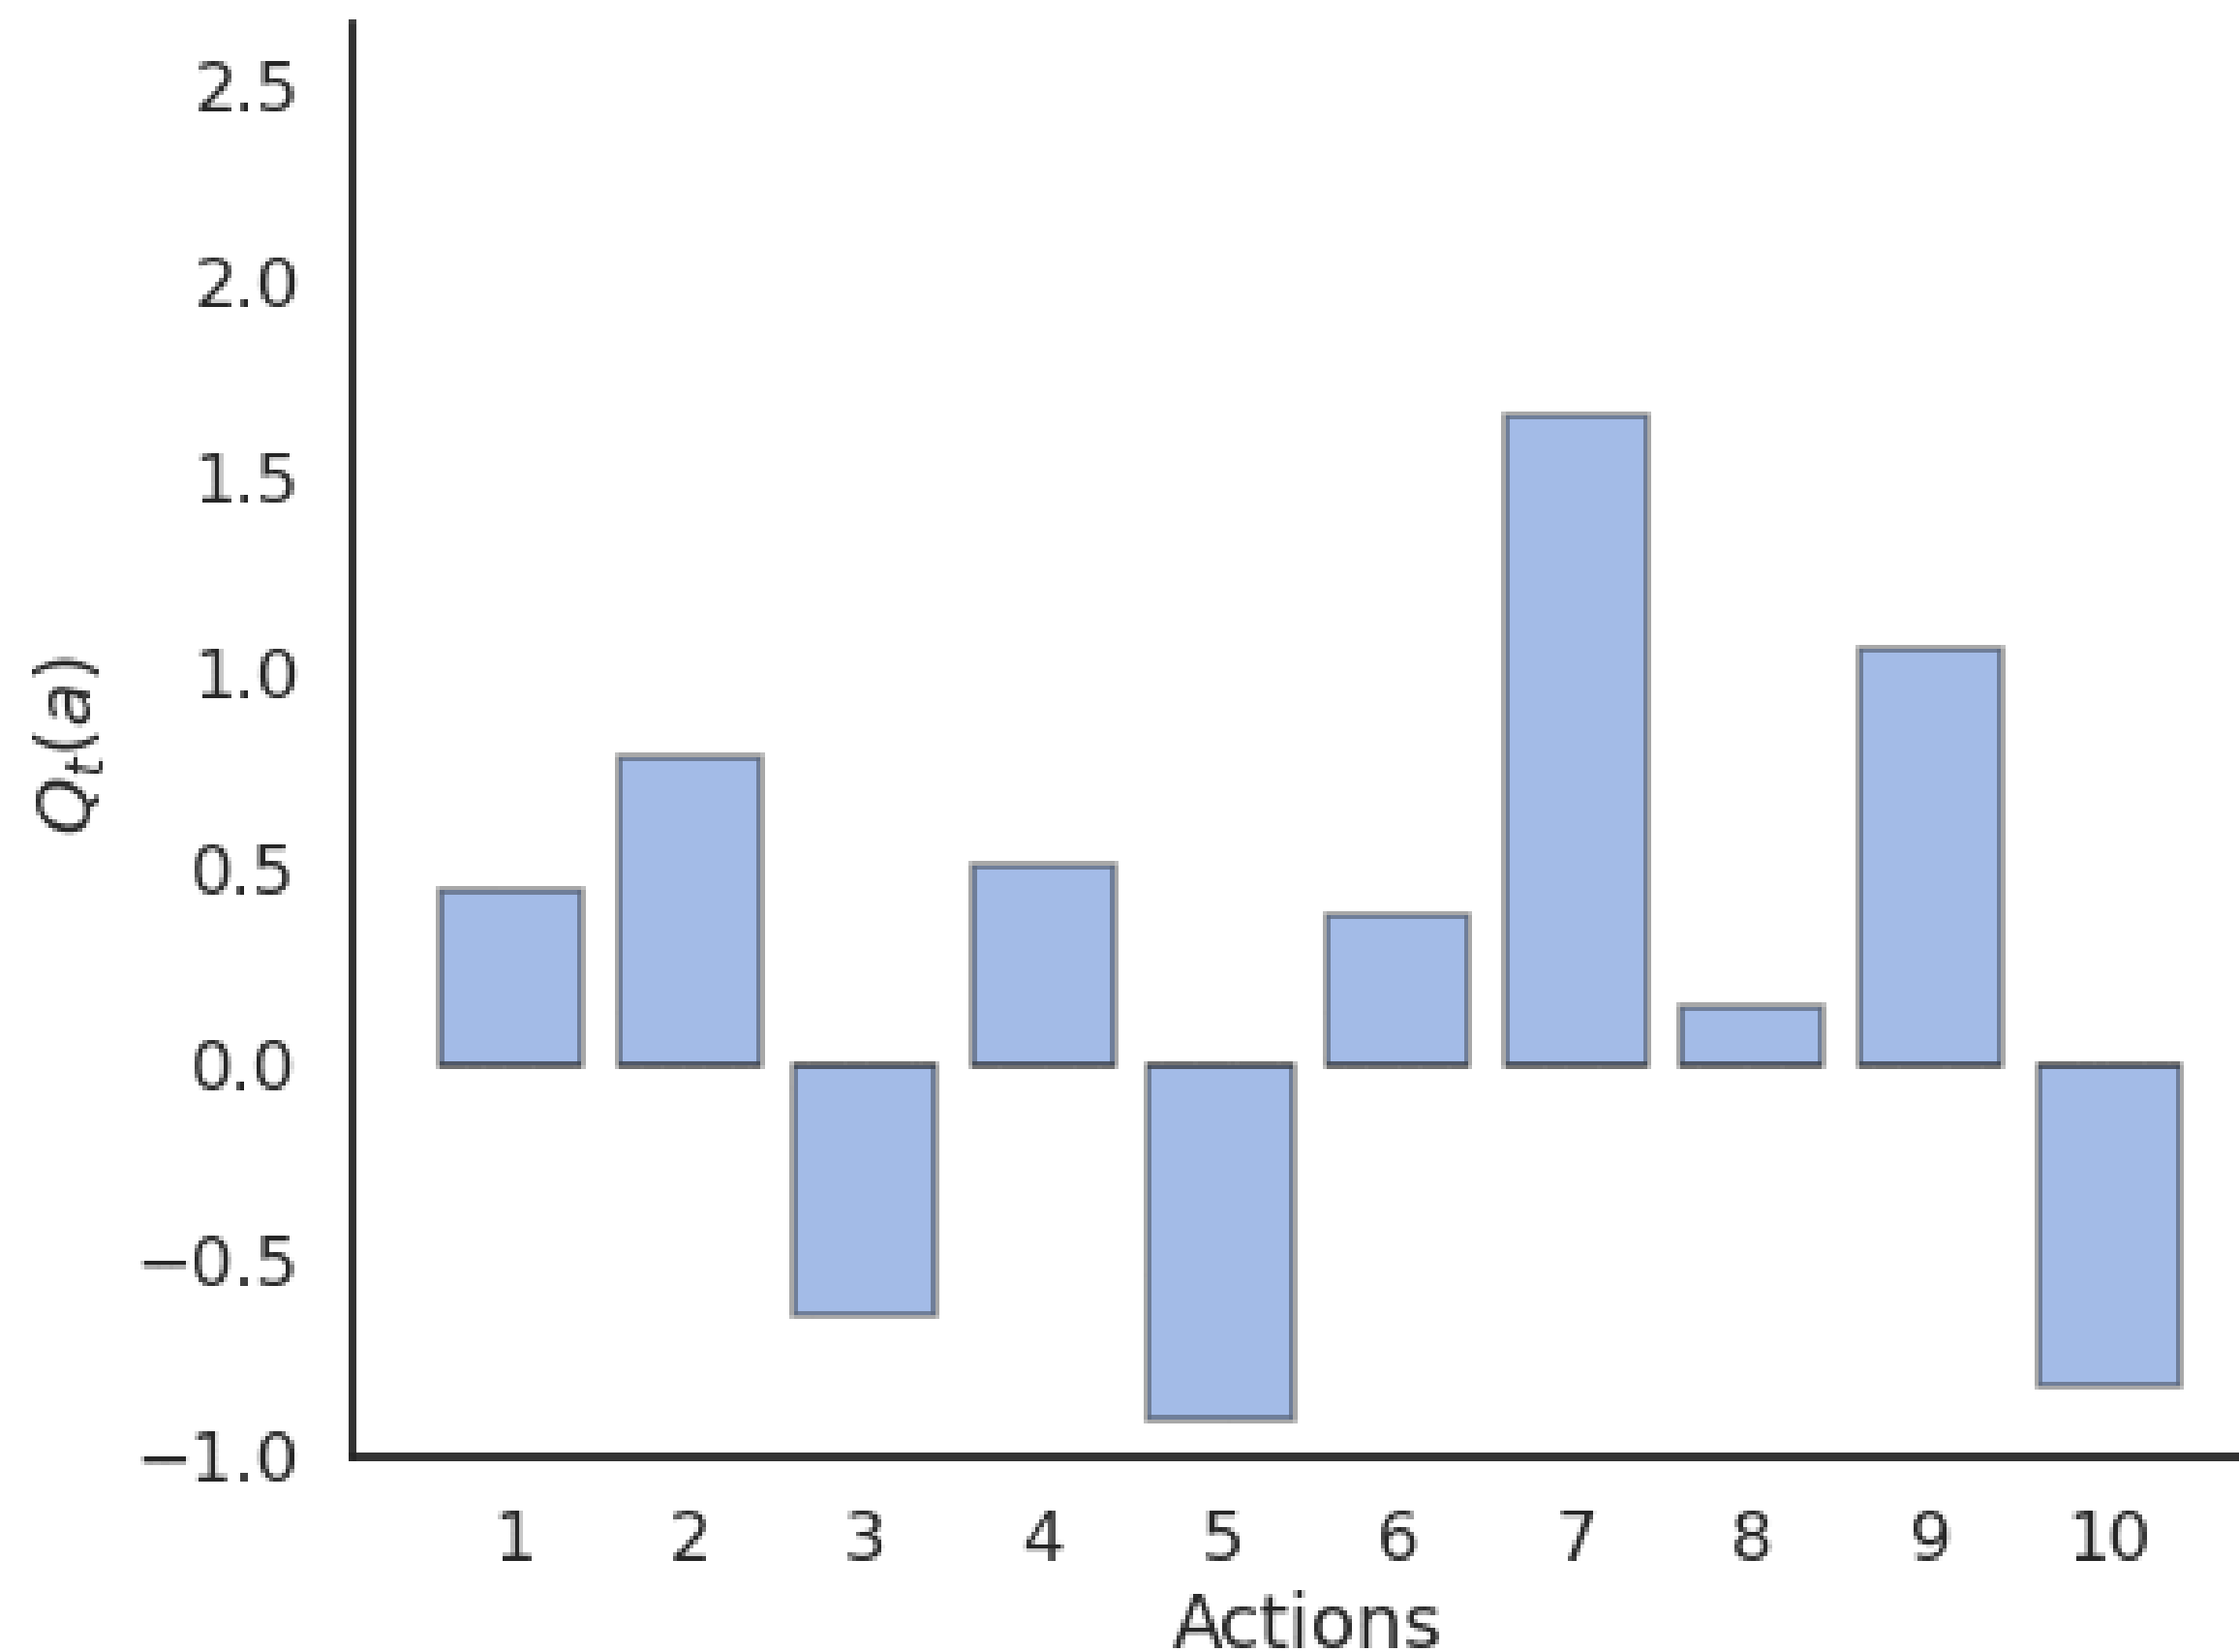
## $\epsilon$ -greedy action selection

- $\epsilon$ -greedy action selection continues to explore after finding a good (but often suboptimal) action.
- It is not always able to recognize the optimal action (it depends on the variance of the rewards).

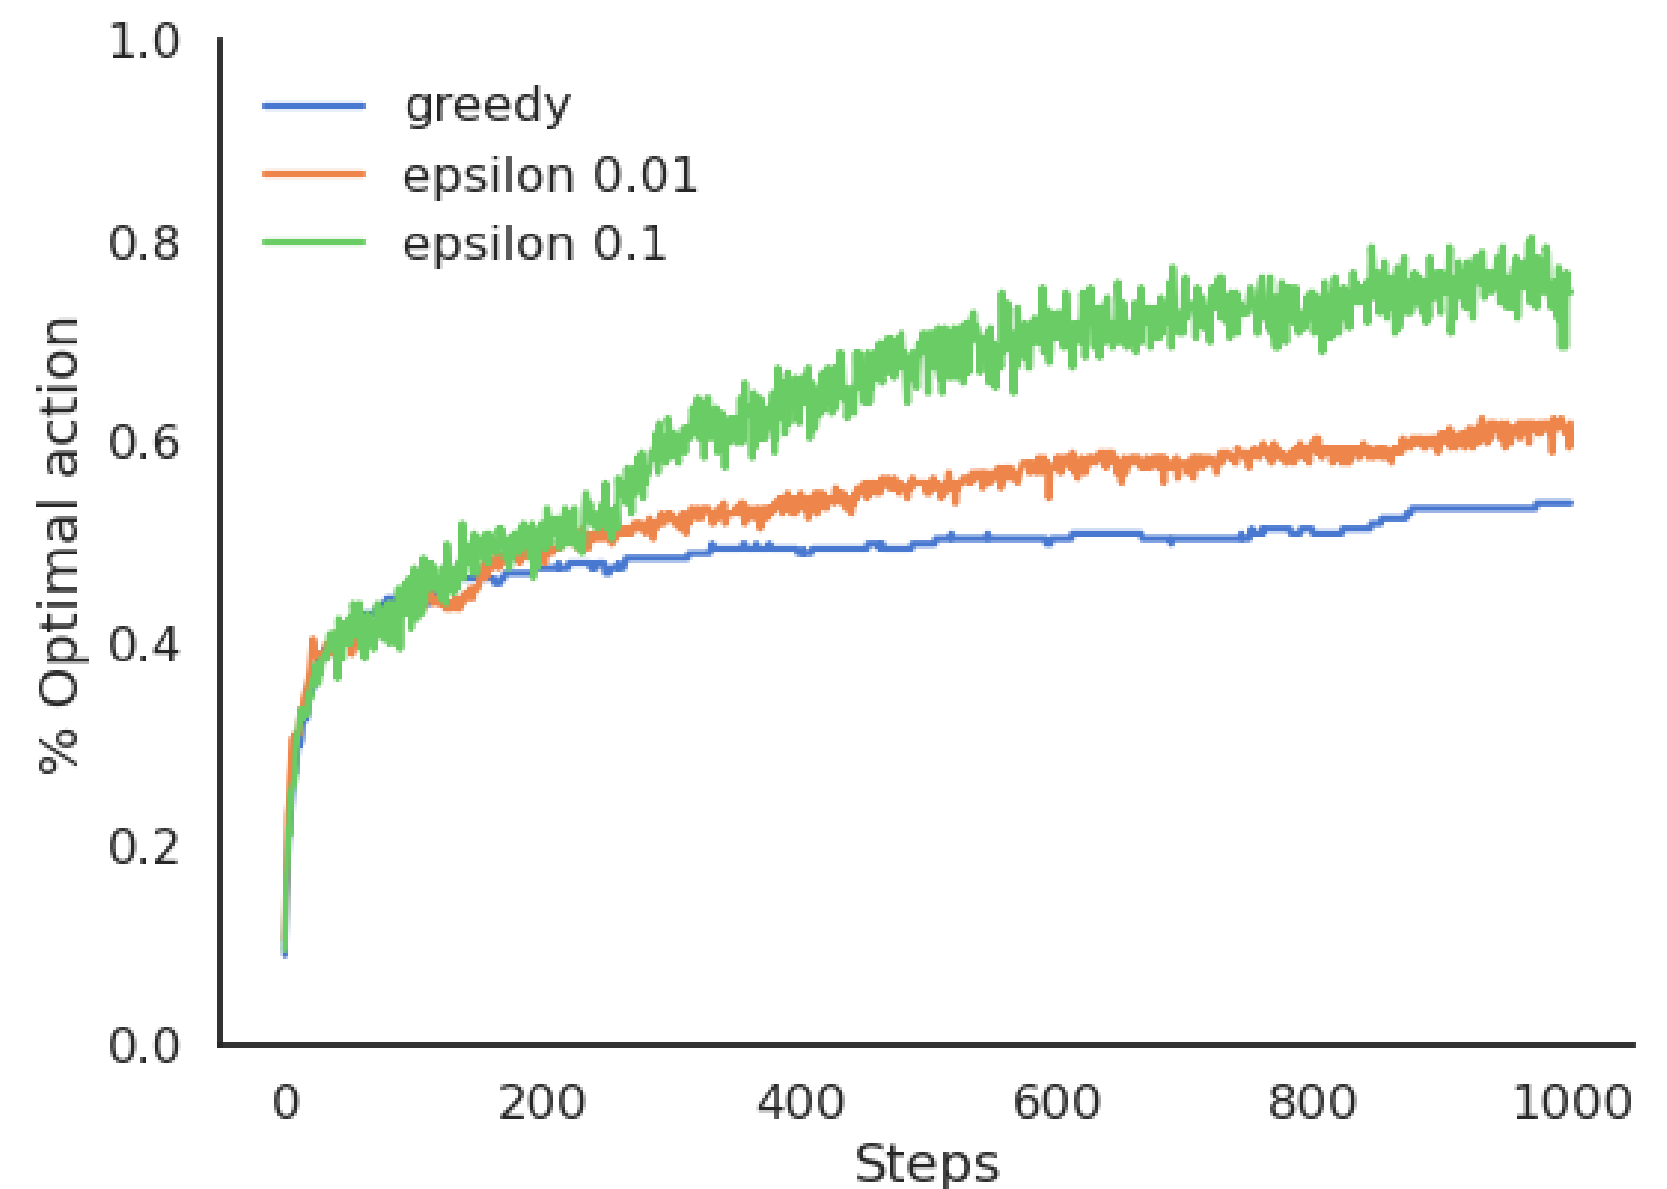
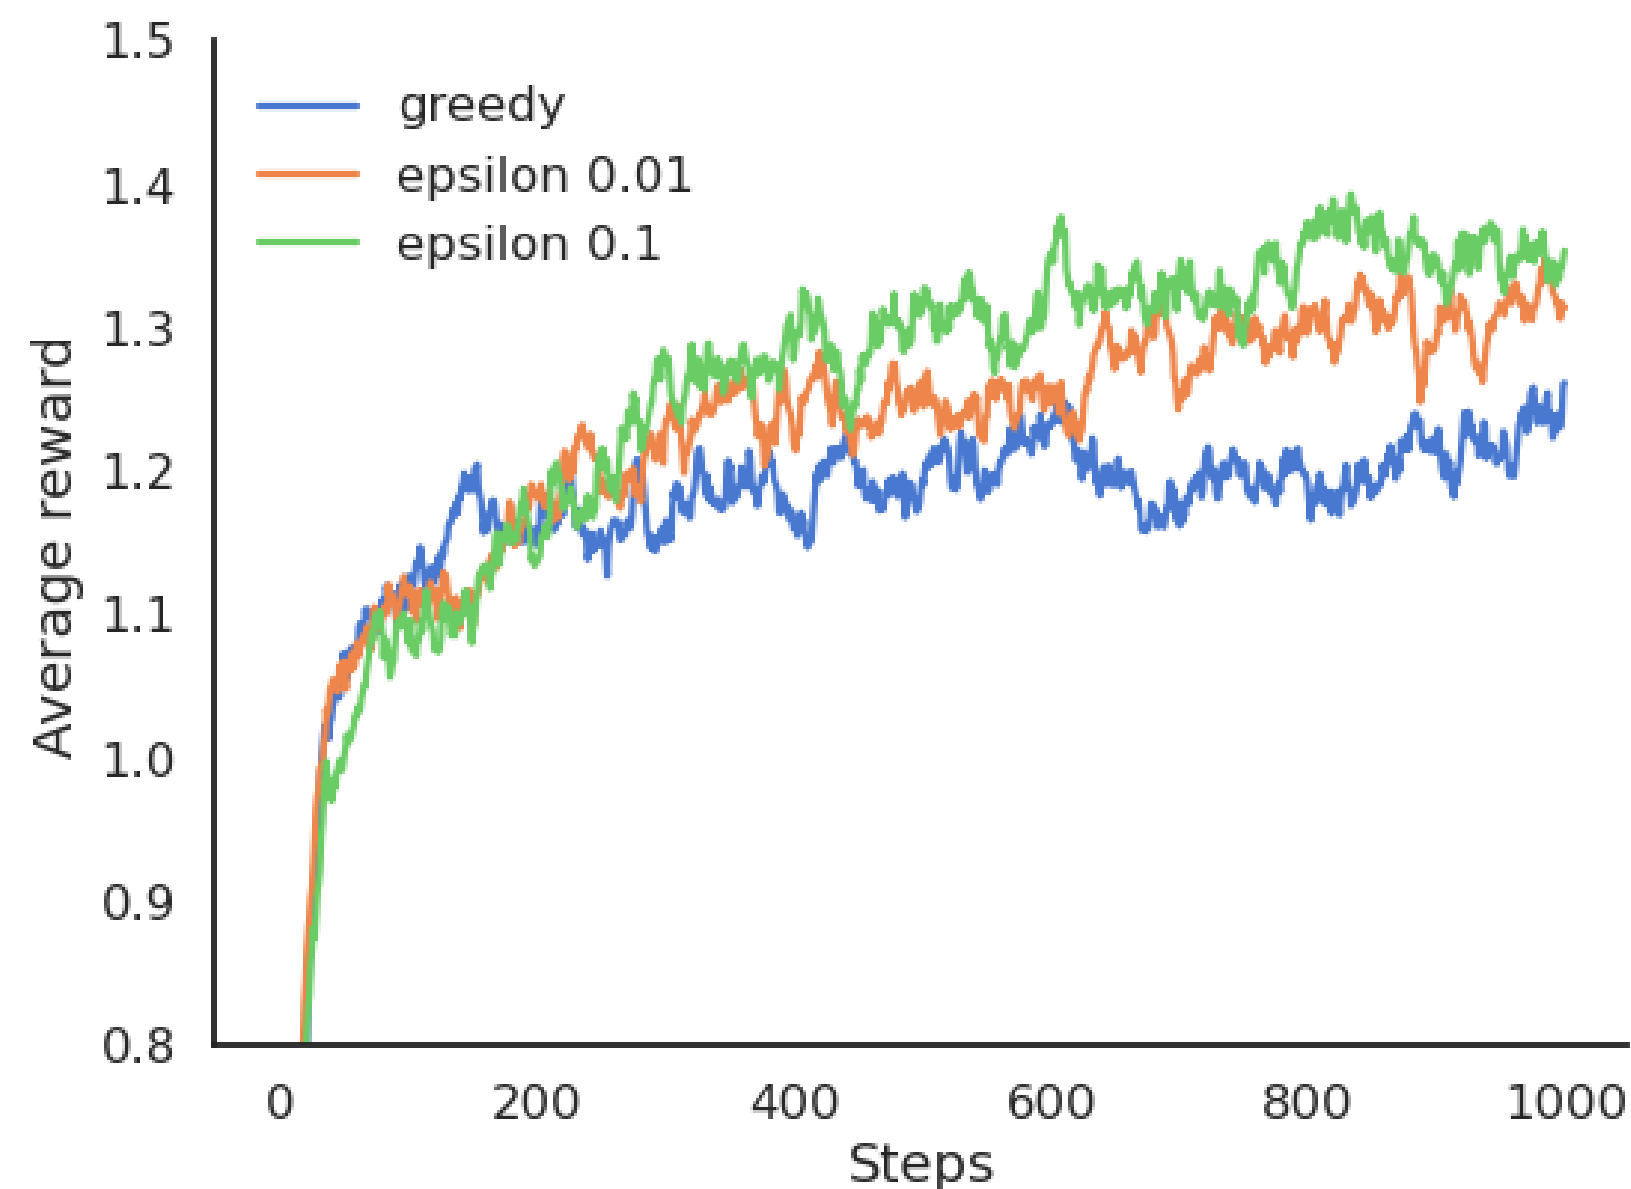


# Softmax action selection

- Softmax action selection explores more consistently the available actions.
- The estimated Q-values are much closer to the true values than with ( $\epsilon$ -)greedy methods.



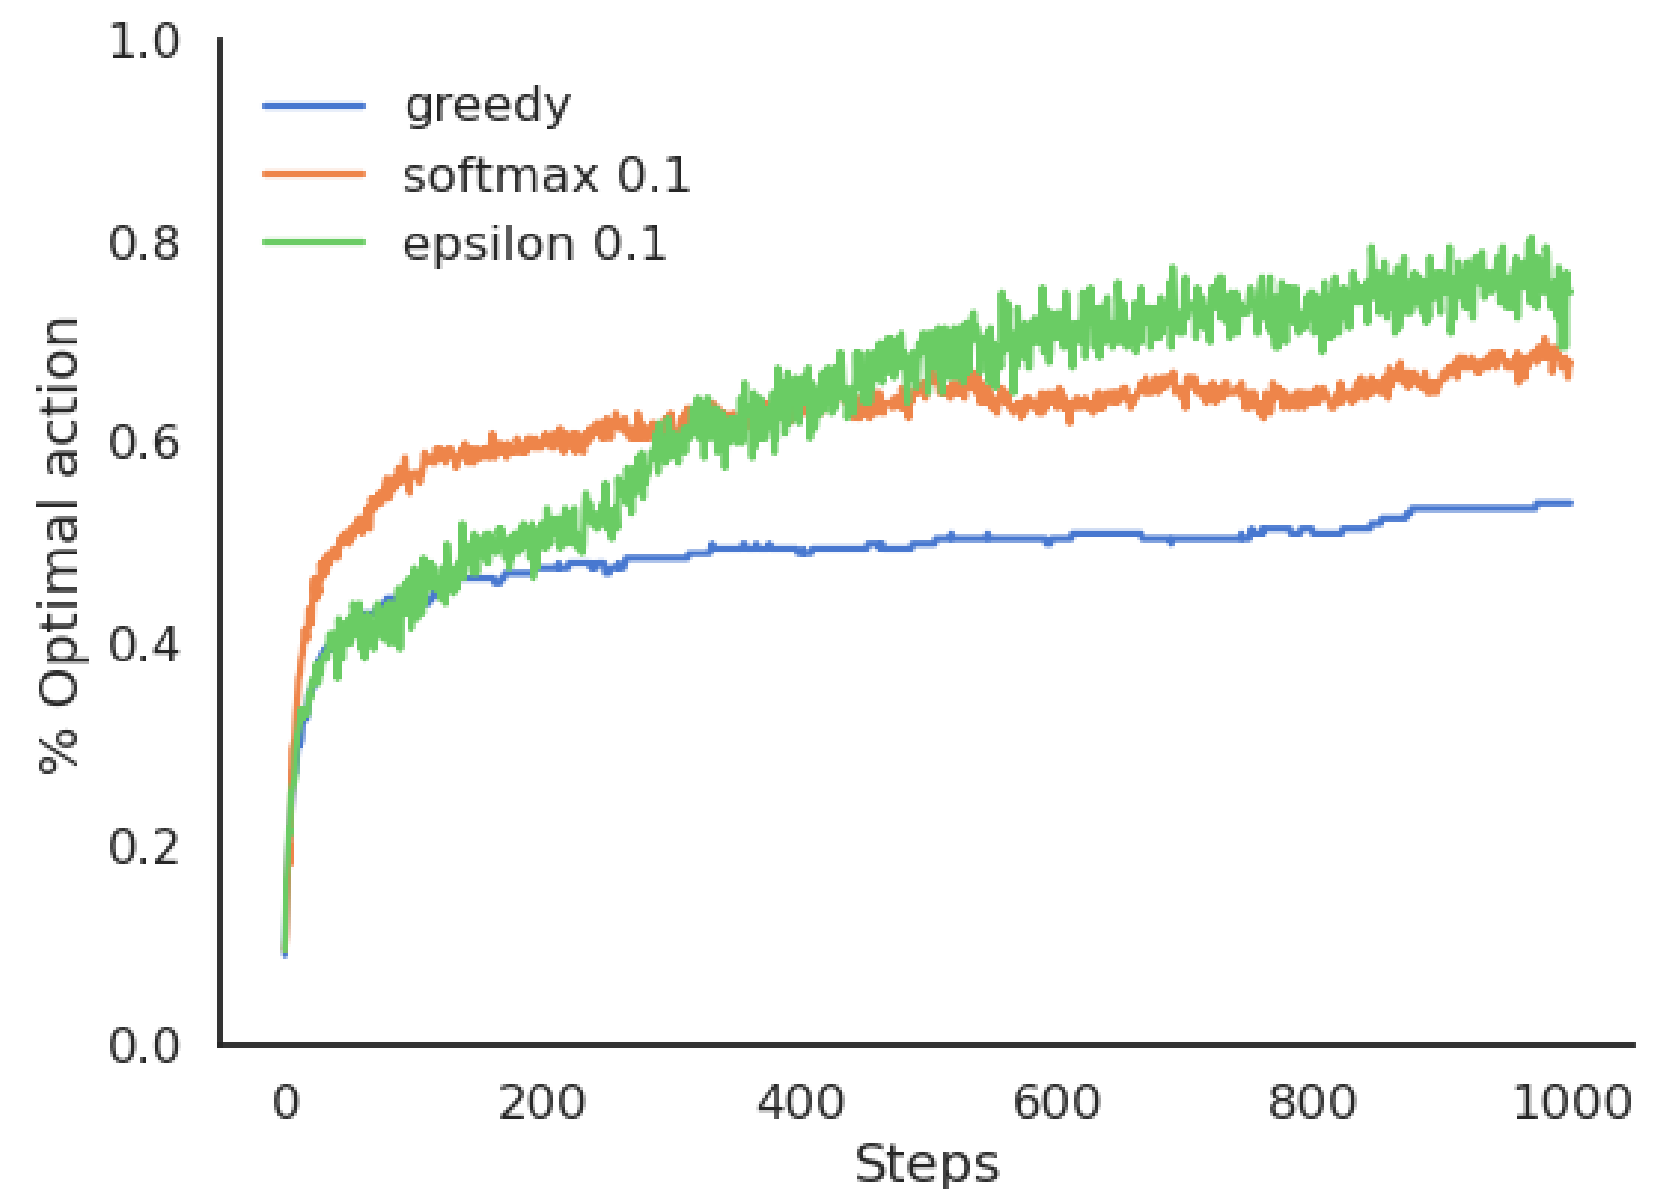
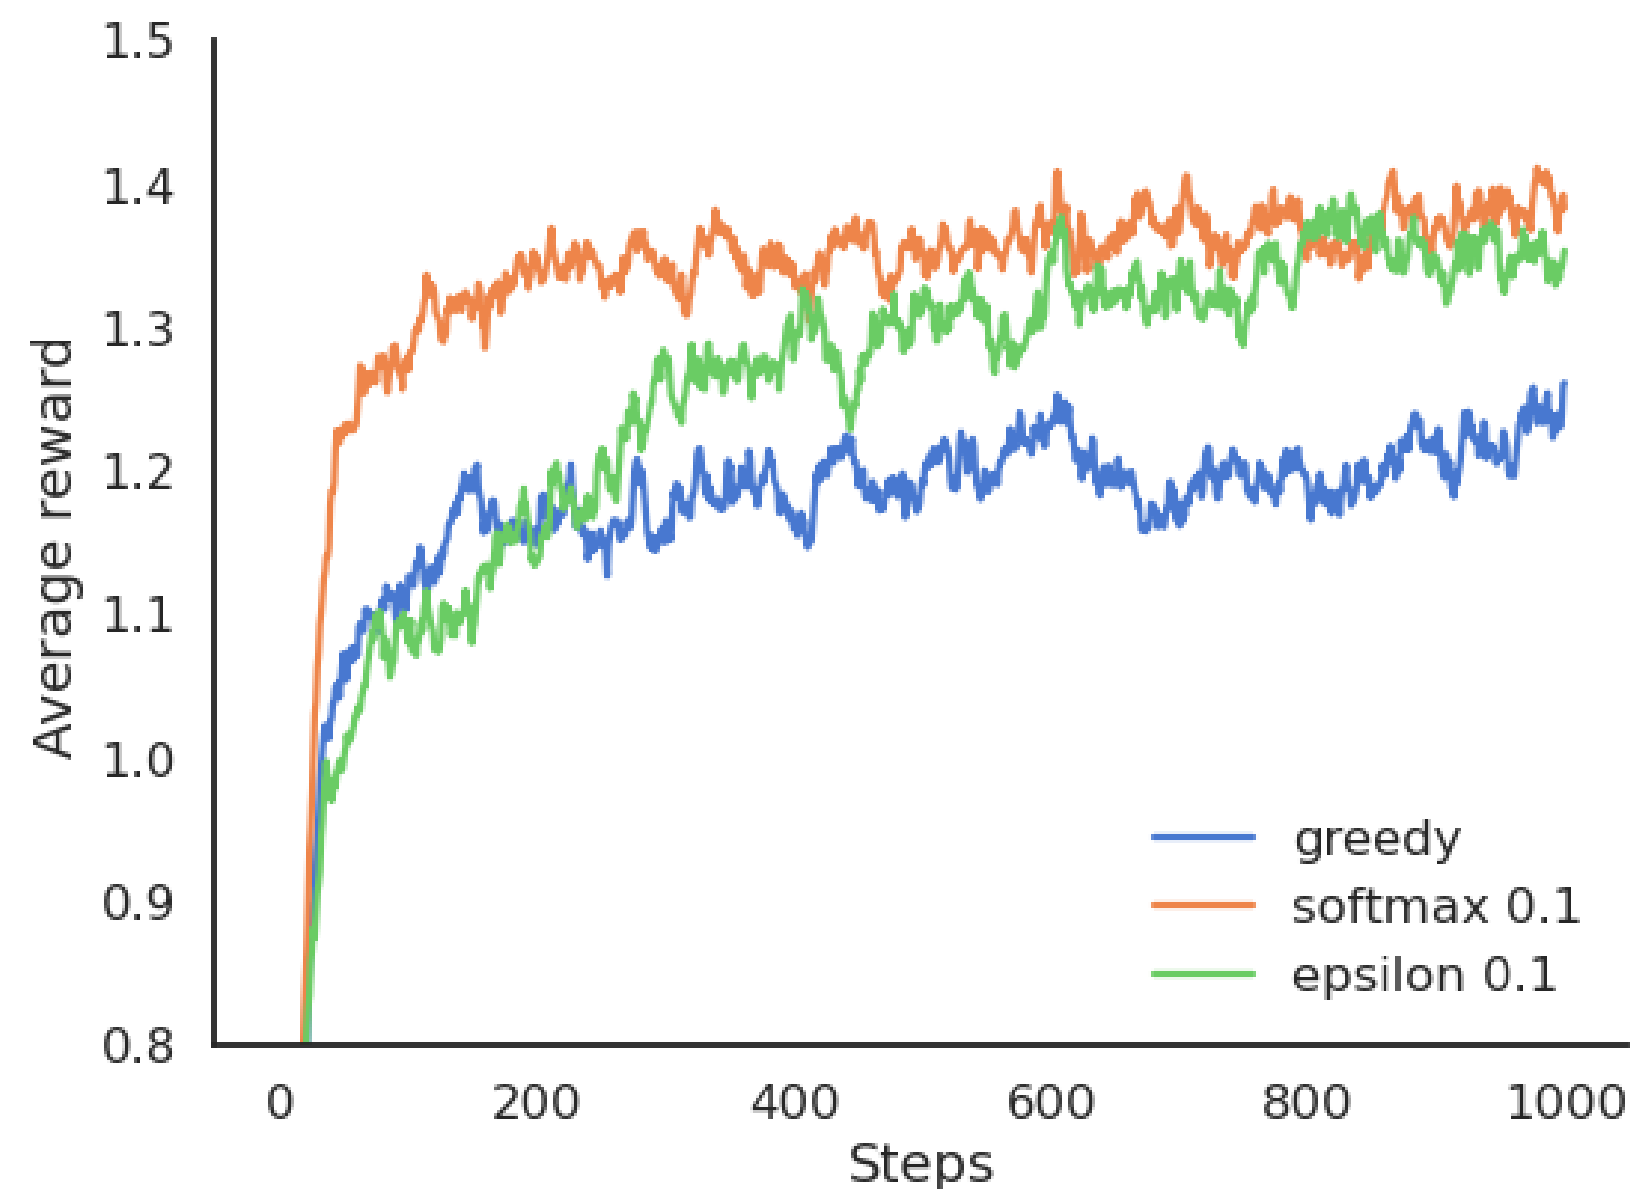
# Greedy vs. $\epsilon$ -greedy



- The **greedy** method learns faster at the beginning, but get stuck in the long-term by choosing **suboptimal** actions (50% of trials).
- $\epsilon$ -greedy methods perform better on the long term, because they continue to explore.
- High values for  $\epsilon$  provide more exploration, hence find the optimal action earlier, but also tend to deselect it more often: with a limited number of plays, it may collect less reward than smaller values of  $\epsilon$ .




## Softmax vs. $\epsilon$ -greedy



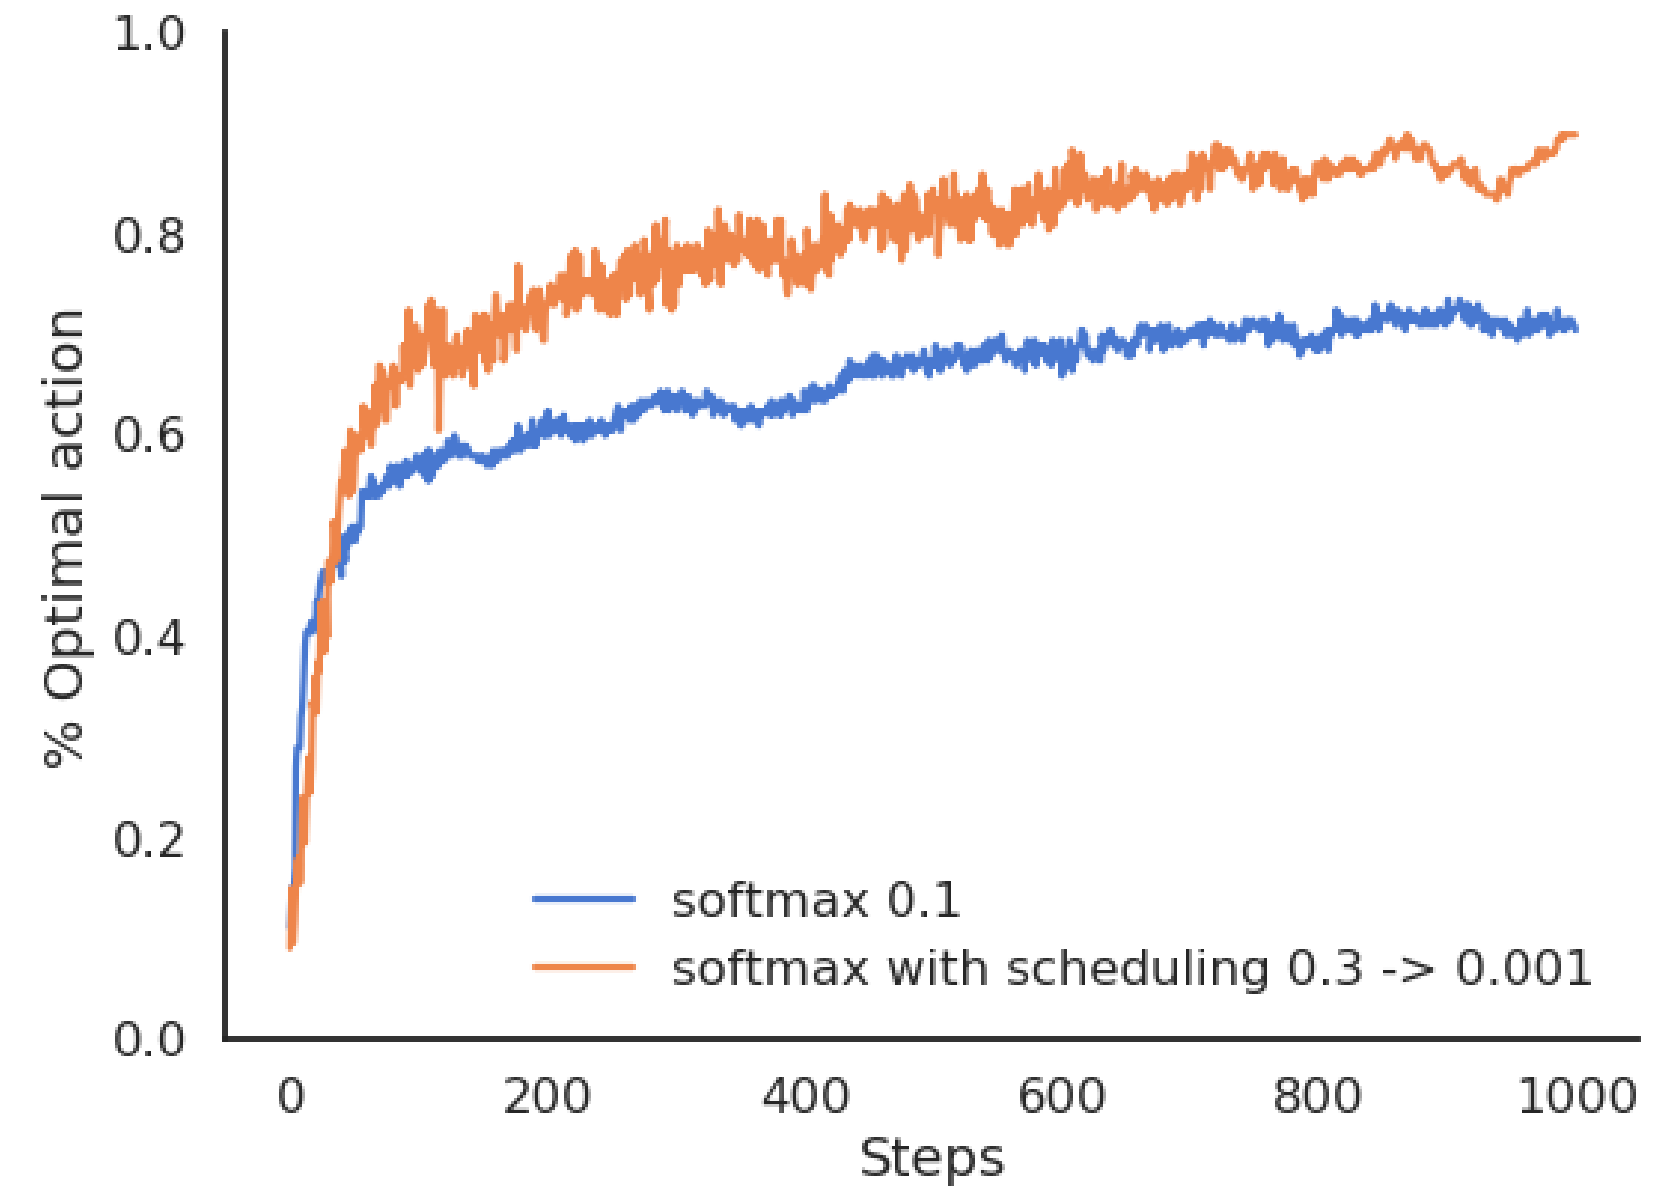

- The softmax does not necessarily find a better solution than  $\epsilon$ -greedy, but it tends to find it **faster** (depending on  $\epsilon$  or  $\tau$ ), as it does not lose time exploring obviously bad solutions.
- $\epsilon$ -greedy or softmax methods work best when the variance of rewards is high.
- If the variance is zero (always the same reward value), the greedy method would find the optimal action more rapidly: the agent only needs to try each action once.

# Exploration schedule

- A useful technique to cope with the **exploration-exploitation dilemma** is to slowly decrease the value of  $\epsilon$  or  $\tau$  with the number of plays.
- This allows for more exploration at the beginning of learning and more exploitation towards the end.
- It is however hard to find the right **decay rate** for the exploration parameters.



# Exploration schedule



- The performance is worse at the beginning, as the agent explores with a high temperature.
- But as the agent becomes greedier and greedier, the performance become more **optimal** than with a fixed temperature.

## Optimistic initial values

- The problem with online evaluation is that it depends a lot on the initial estimates  $Q_0$ .
  - If the initial estimates are already quite good (expert knowledge), the Q-values will converge very fast.
  - If the initial estimates are very wrong, we will need a lot of updates to correctly estimate the true values.

$$Q_{t+1}(a) = (1 - \alpha) Q_t(a) + \alpha r_{t+1}$$

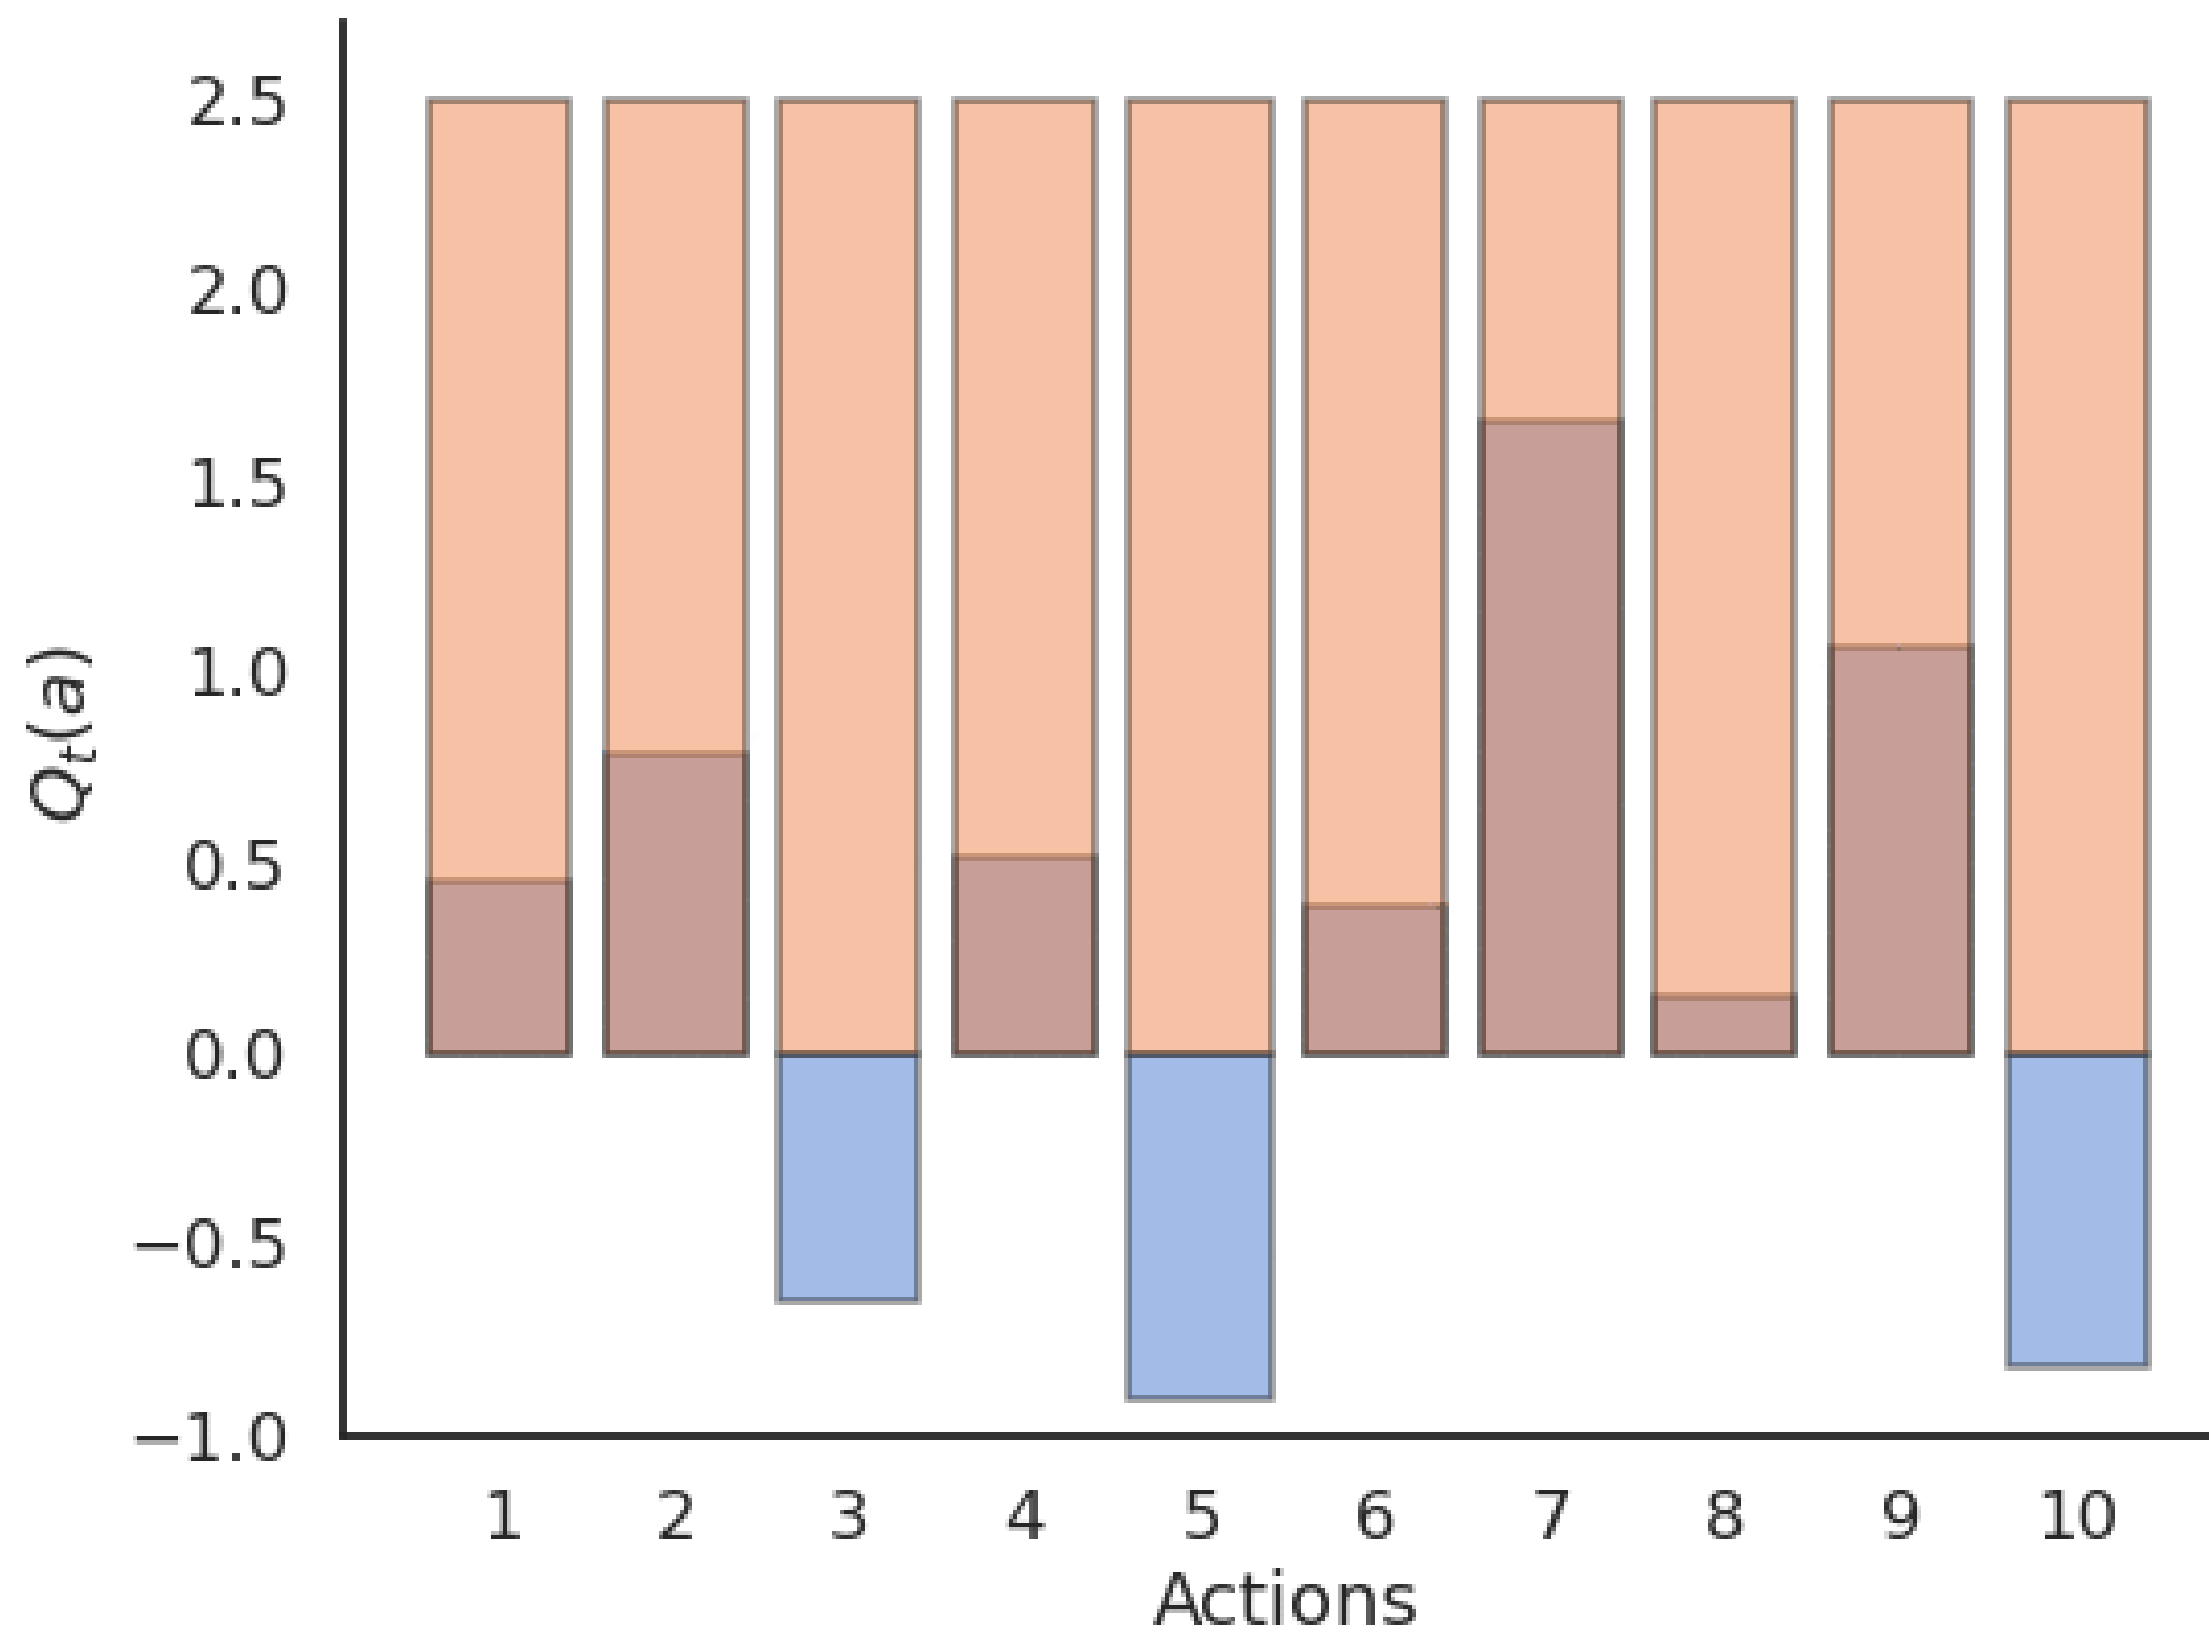
$$\rightarrow Q_1(a) = (1 - \alpha) Q_0(a) + \alpha r_1$$

$$\rightarrow Q_2(a) = (1 - \alpha) Q_1(a) + \alpha r_2 = (1 - \alpha)^2 Q_0(a) + (1 - \alpha)\alpha r_1 + \alpha r_2$$

- The influence of  $Q_0$  on  $Q_t$  **fades** quickly with  $(1 - \alpha)^t$ , but that can be lost time or lead to a suboptimal policy.
- However, we can use this at our advantage with **optimistic initialization**.

# Optimistic initial values

- By choosing very high initial values for the estimates (they can only decrease), one can ensure that all possible actions will be selected during learning by the greedy method, solving the **exploration problem**.
- This leads however to an **overestimation** of the value of other actions.



# Reinforcement comparison

- Actions followed by large rewards should be made more likely to recur, whereas actions followed by small rewards should be made less likely to recur.
- But what is a large/small reward? Is a reward of 5 large or small?
- **Reinforcement comparison** methods only maintain a **preference**  $p_t(a)$  for each action, which is not exactly its Q-value.
- The preference for an action is updated after each play, according to the update rule:

$$p_{t+1}(a_t) = p_t(a_t) + \beta (r_t - \tilde{r}_t)$$

where  $\tilde{r}_t$  is the moving average of the recently received rewards (regardless the action):

$$\tilde{r}_{t+1} = \tilde{r}_t + \alpha (r_t - \tilde{r}_t)$$

- If an action brings more reward than usual (**good surprise**), we increase the preference for that action.
- If an action brings less reward than usual (**bad surprise**), we decrease the preference for that action.
- $\beta > 0$  and  $0 < \alpha < 1$  are two constant parameters.


# Reinforcement comparison

- Preferences are updated by replacing the action-dependent Q-values by a baseline  $\tilde{r}_t$ :

$$p_{t+1}(a_t) = p_t(a_t) + \beta (r_t - \tilde{r}_t)$$

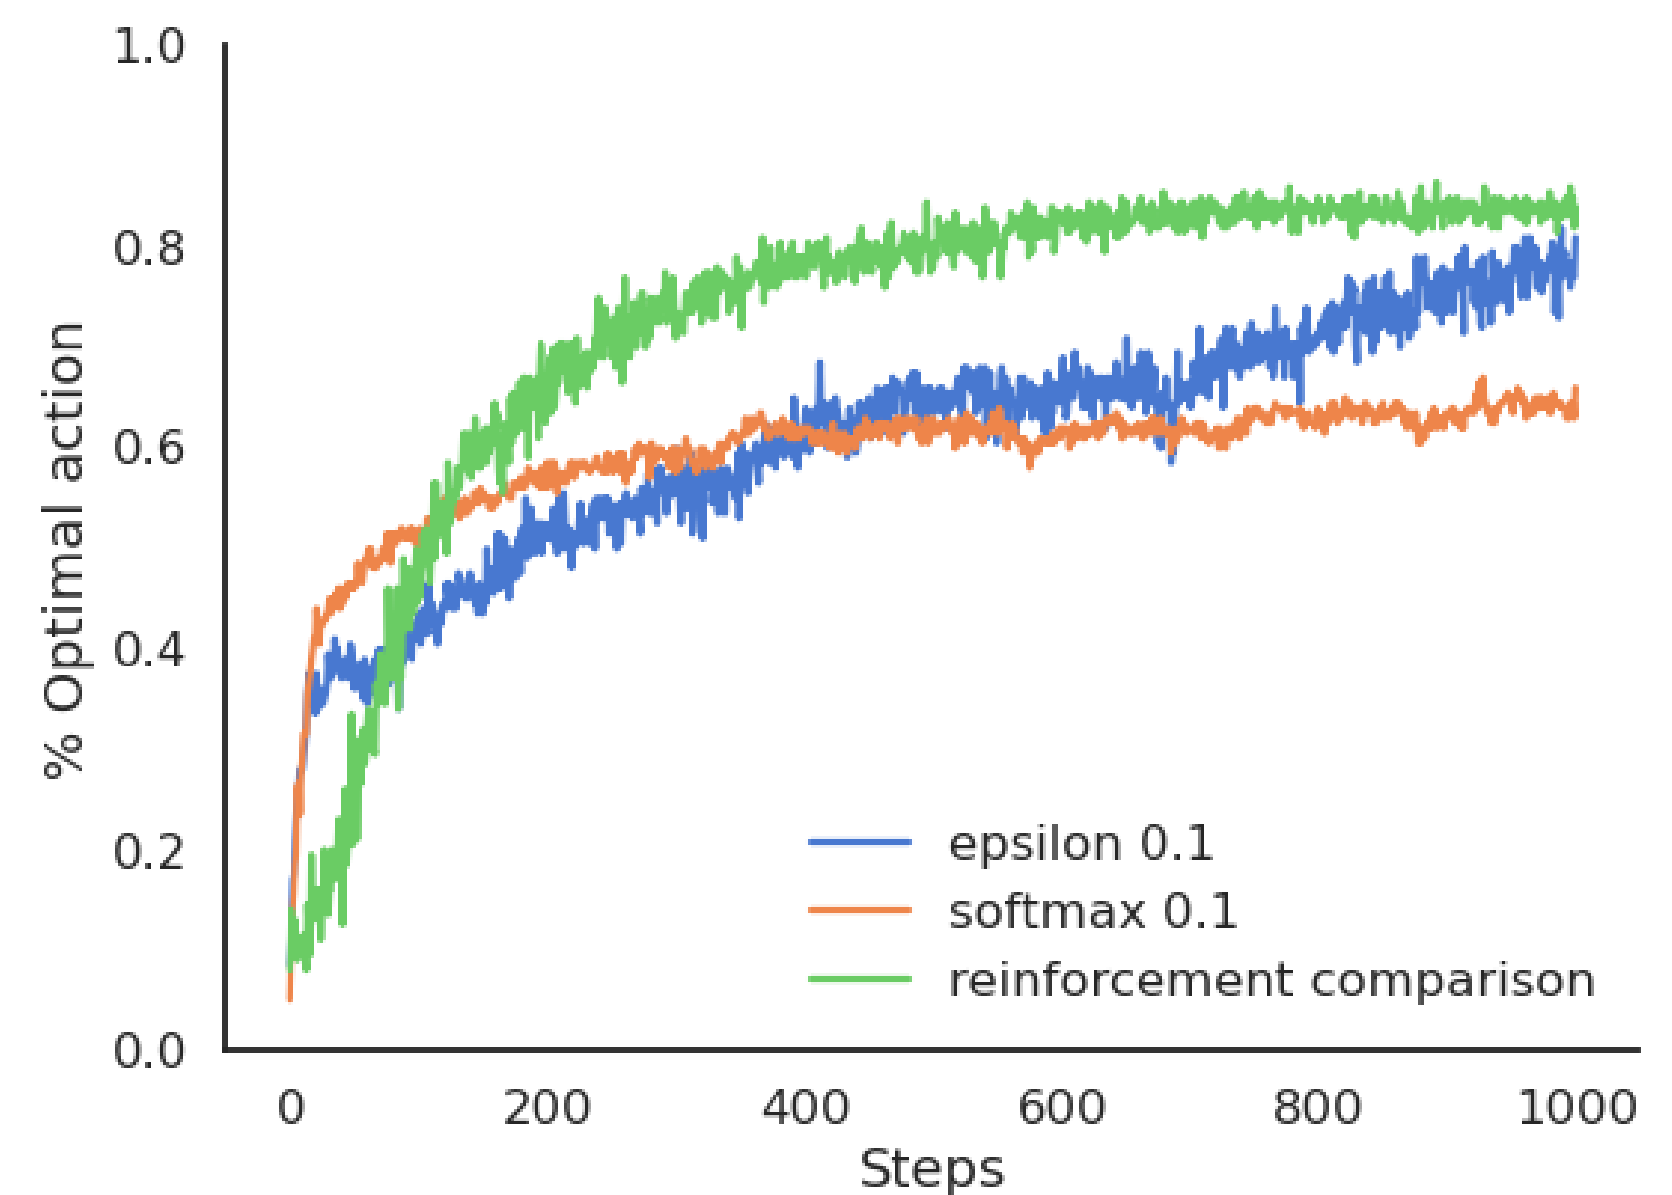
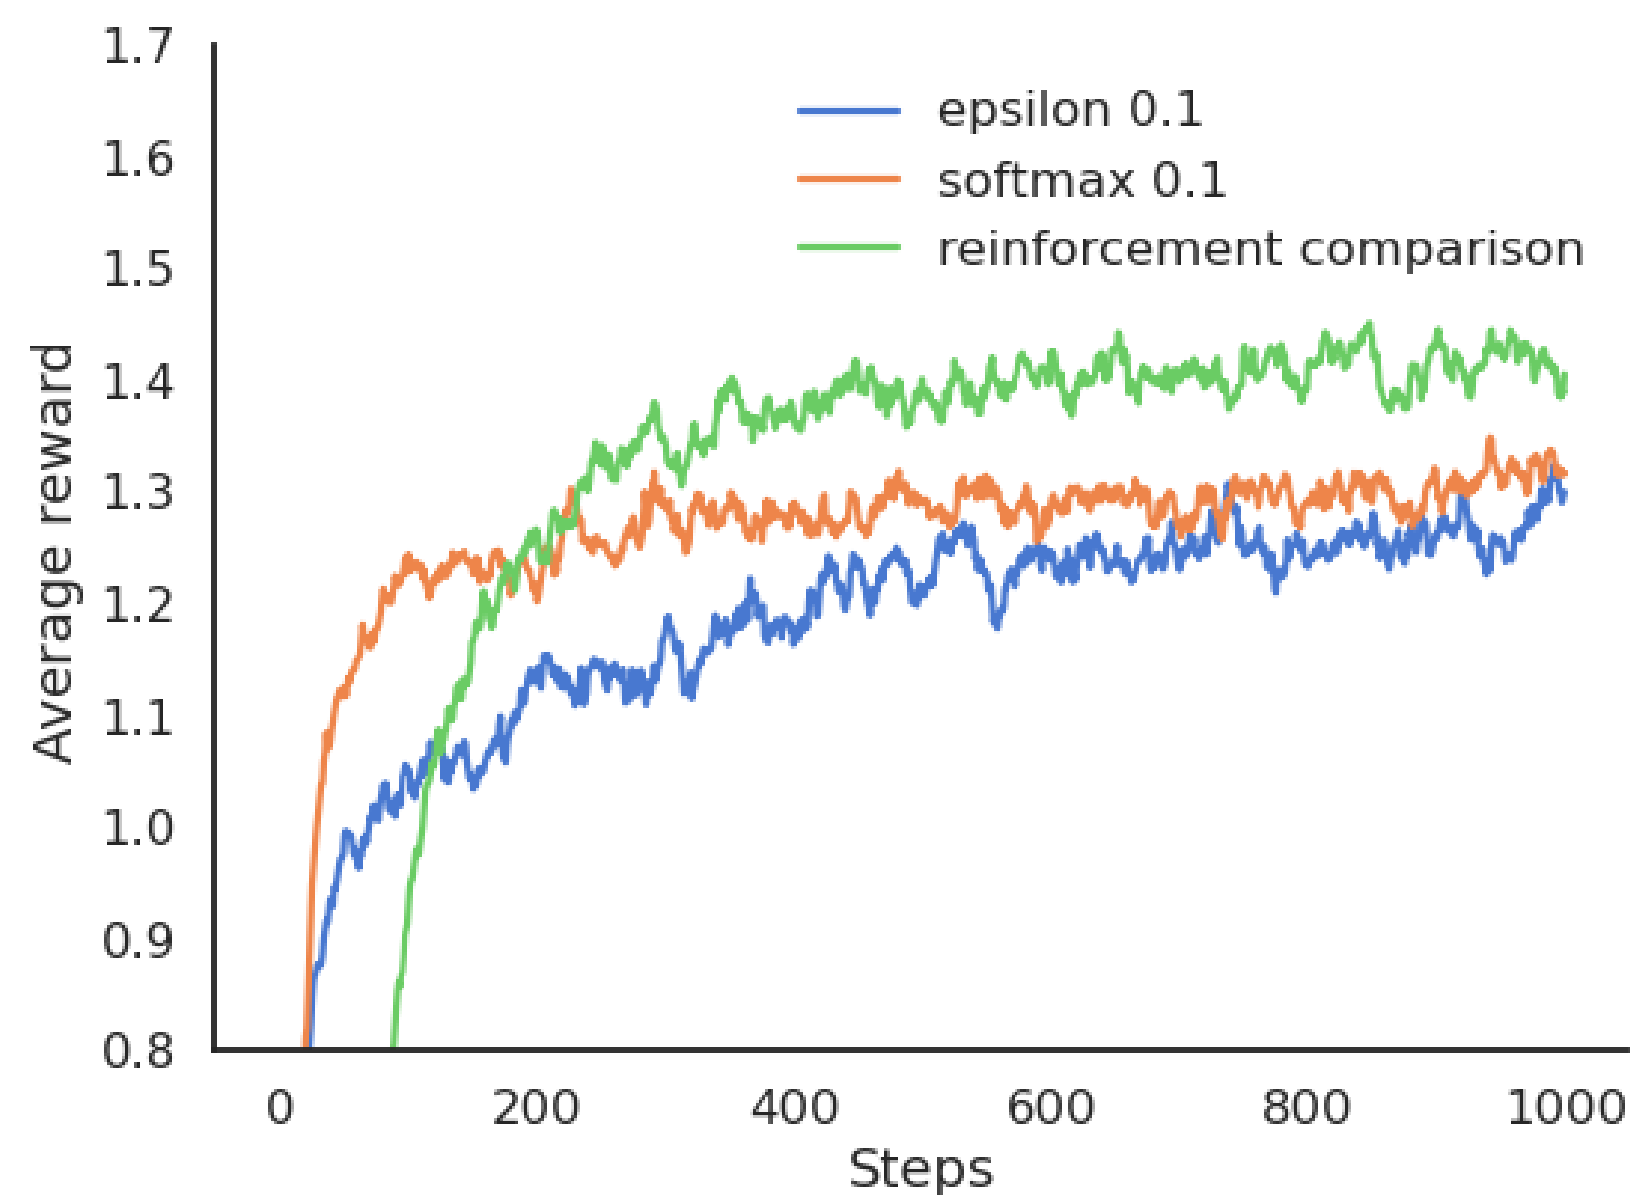
- The preferences can be used to select the action using the softmax method just as the Q-values (without temperature):

$$\pi_t(a) = \frac{\exp p_t(a)}{\sum_{a'} \exp p_t(a')}$$



# Reinforcement comparison

- Reinforcement comparison can be very effective, as it does not rely only on the rewards received, but also on their comparison with a **baseline**, the average reward.
- This idea is at the core of **actor-critic** architectures which we will see later.
- The initial average reward  $\tilde{r}_0$  can be set optimistically to encourage exploration.





# Gradient bandit algorithm

- Instead of only increasing the preference for the executed action if it brings more reward than usual, we could also decrease the preference for the other actions.
- The preferences are used to select an action  $a_t$  via softmax:

$$\pi_t(a) = \frac{\exp p_t(a)}{\sum_{a'} \exp p_t(a')}$$

- Update rule for the **action taken**  $a_t$ :

$$p_{t+1}(a_t) = p_t(a_t) + \beta (r_t - \tilde{r}_t) (1 - \pi_t(a_t))$$

- Update rule for the **other actions**  $a \neq a_t$ :

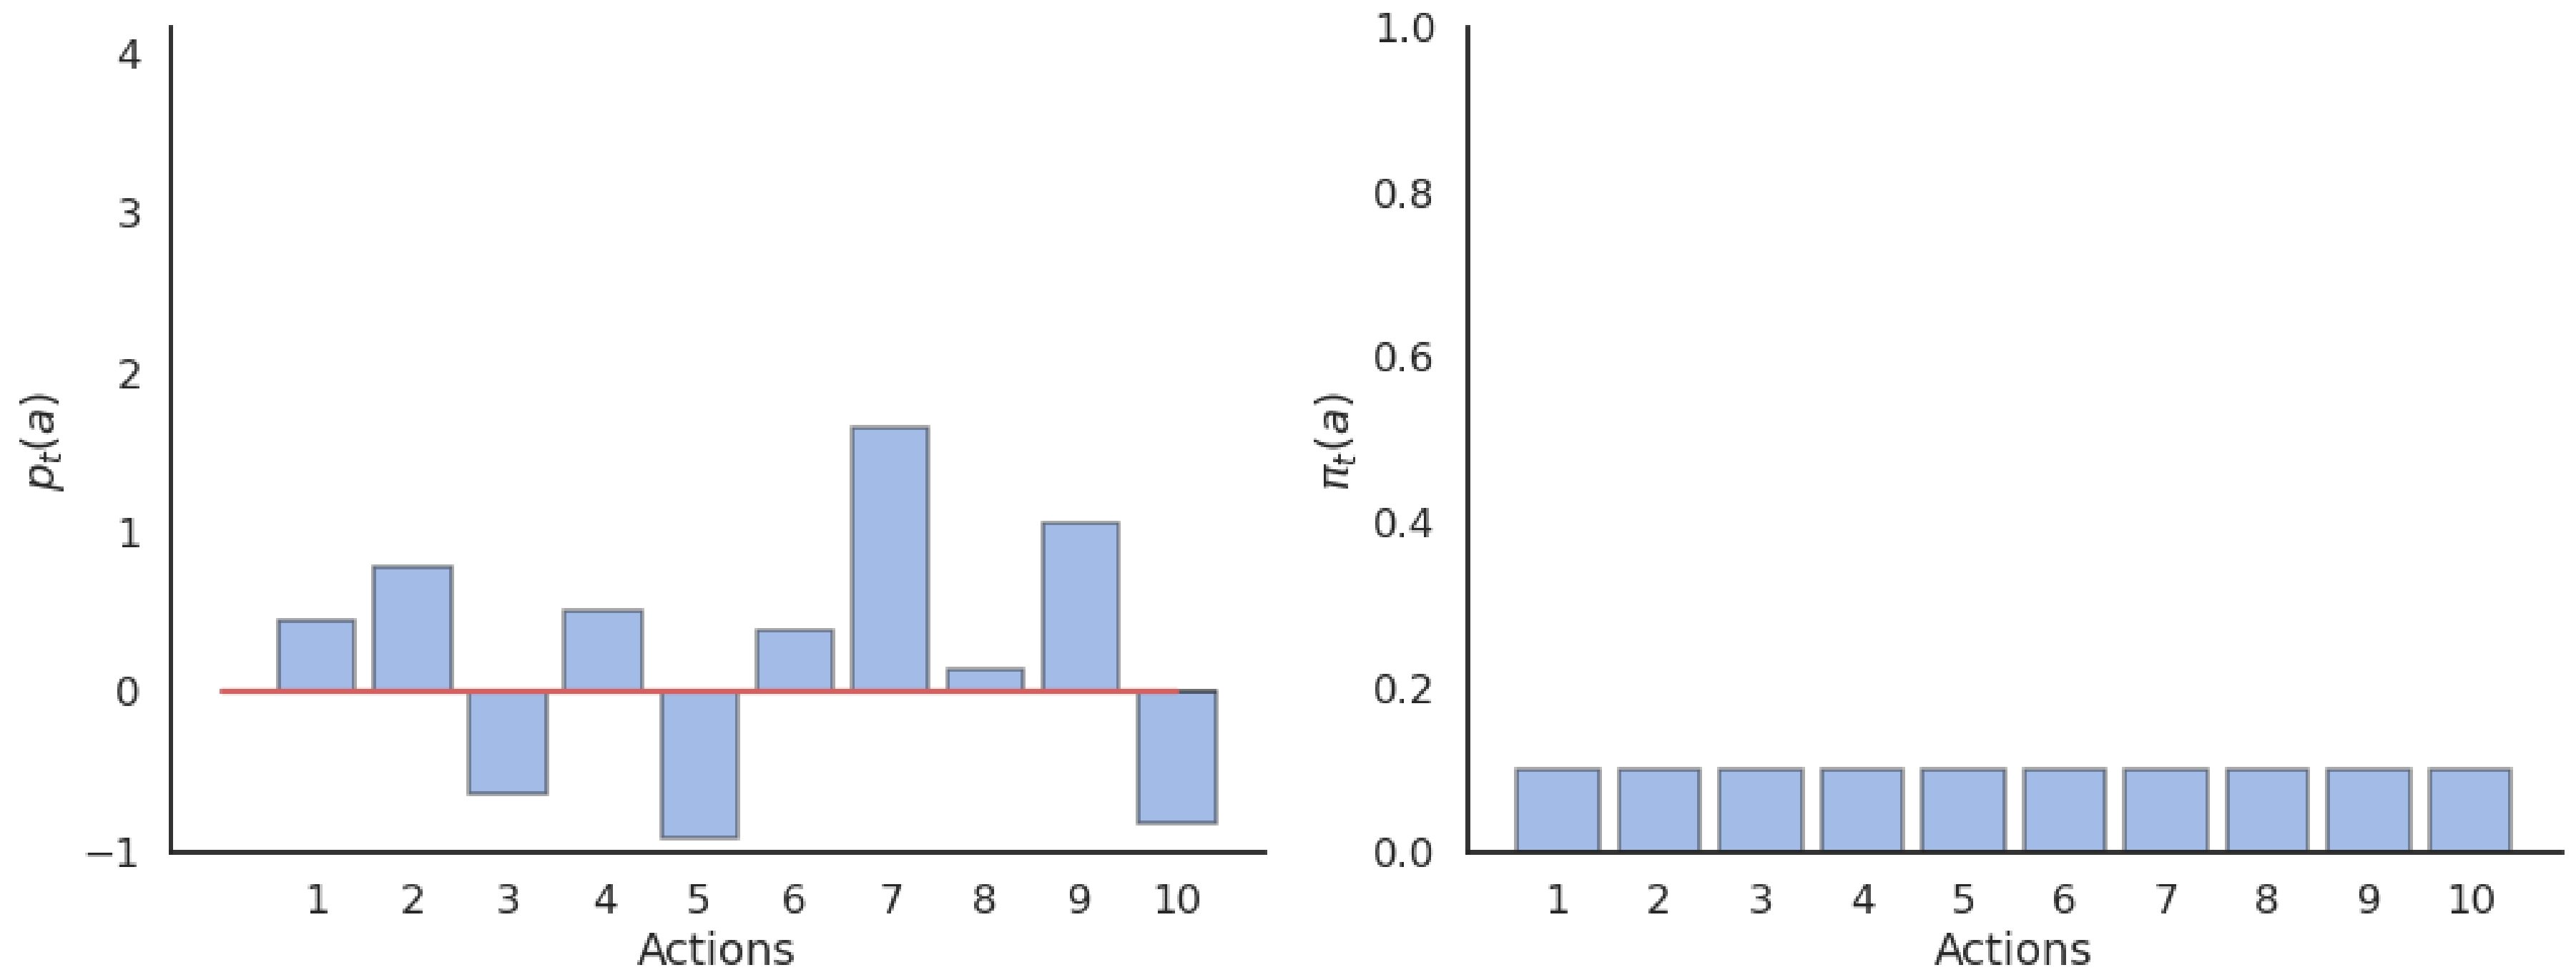
$$p_{t+1}(a) = p_t(a) - \beta (r_t - \tilde{r}_t) \pi_t(a)$$

- Update of the reward **baseline**:

$$\tilde{r}_{t+1} = \tilde{r}_t + \alpha (r_t - \tilde{r}_t)$$

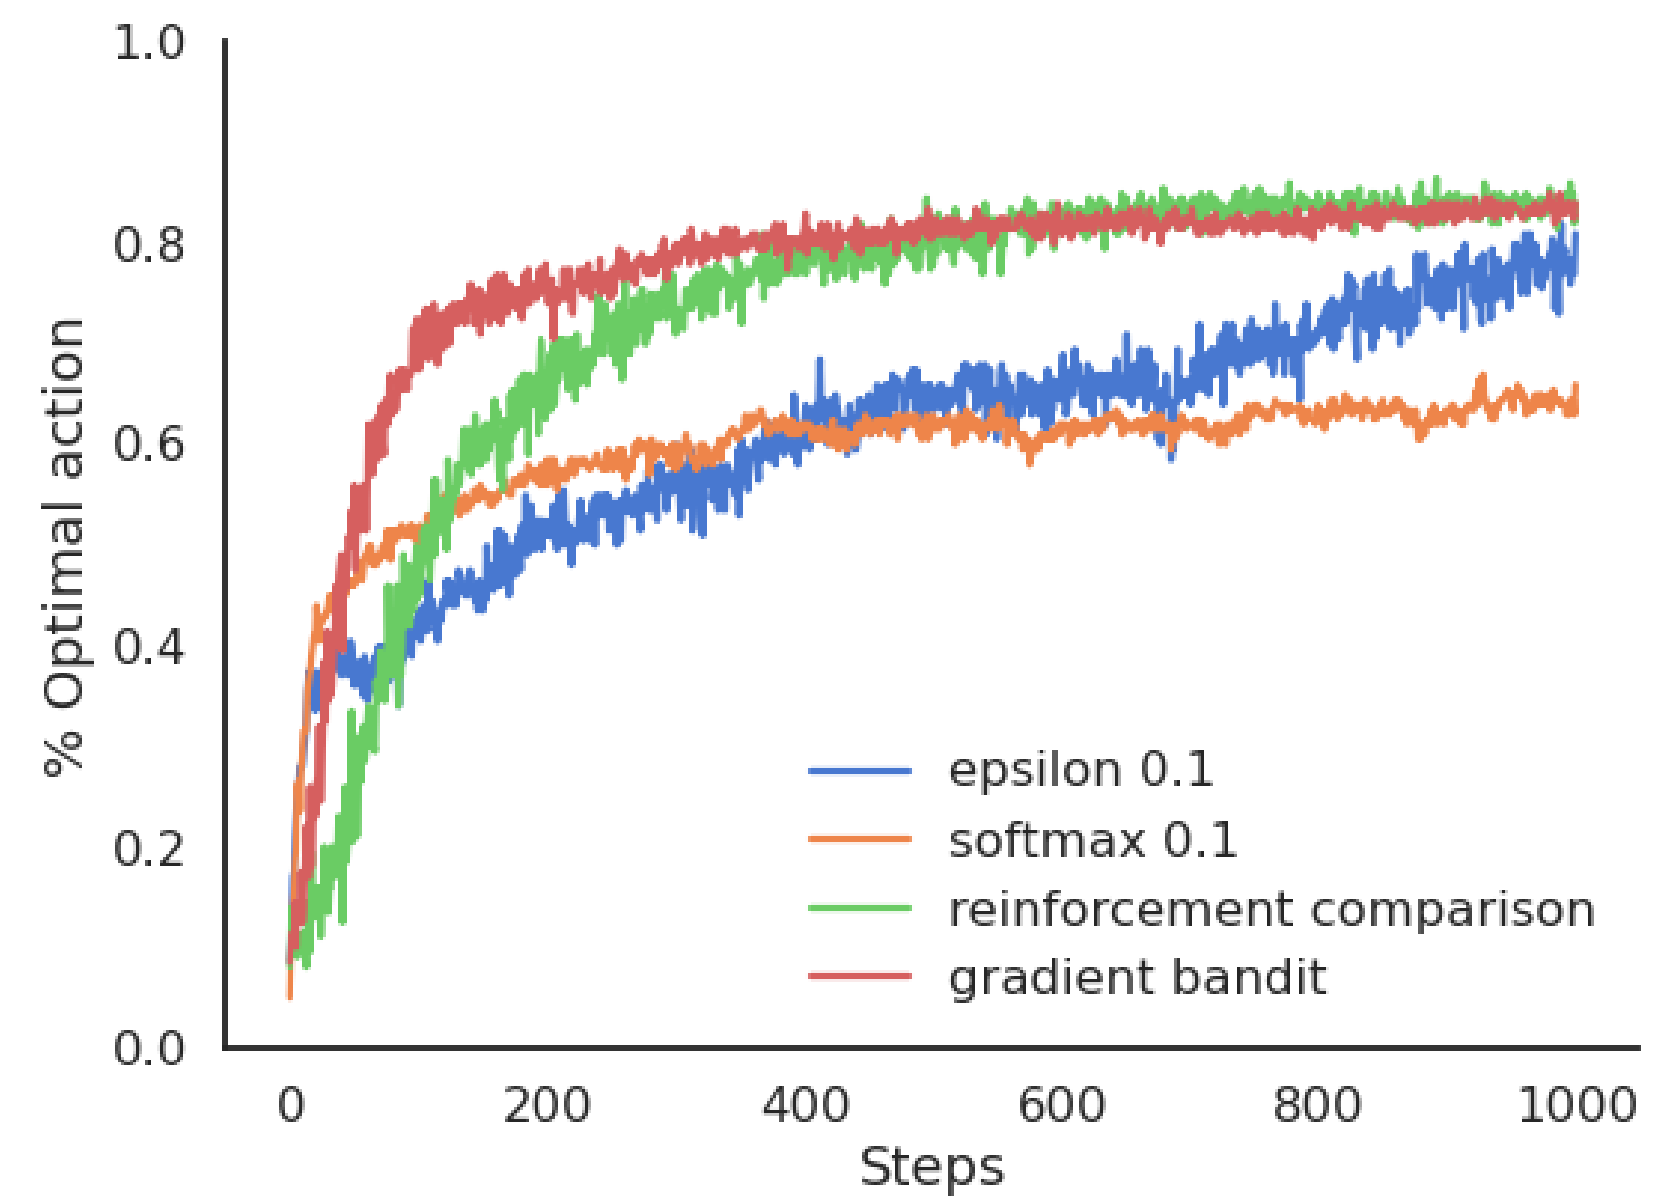
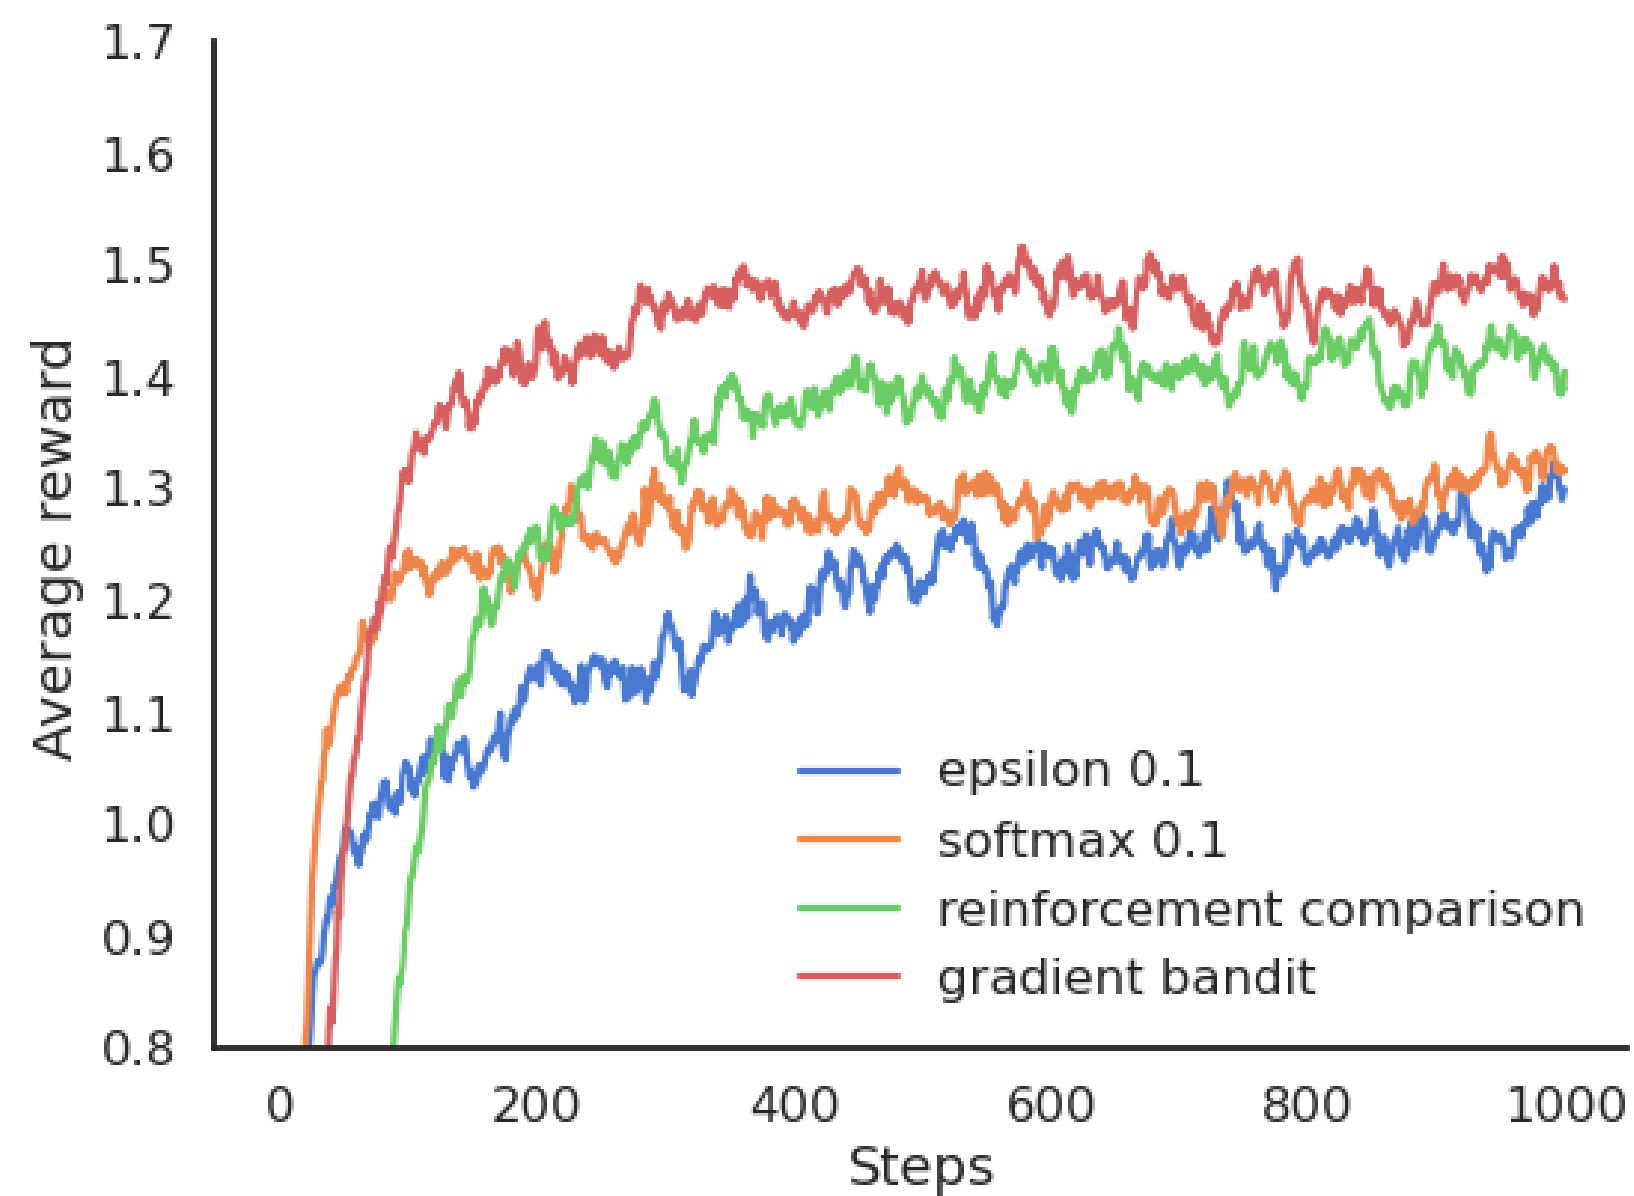
# Gradient bandit algorithm

- The preference can increase become quite high, making the policy greedy towards the end.
- No need for a temperature parameter!



# Gradient bandit algorithm

- Gradient bandit is not always better than reinforcement comparison, but learns initially faster (depending on the parameters  $\alpha$  and  $\beta$ ).



## Upper-Confidence-Bound action selection

- In the previous methods, **exploration** is controlled by an external parameter ( $\epsilon$  or  $\tau$ ) which is **global** to each action and must be scheduled.
- A much better approach would be to decide whether to explore an action based on the **uncertainty** about its Q-value:
  - If we are certain about the value of an action, there is no need to explore it further, we only have to exploit it if it is good.
- The **central limit theorem** tells us that the variance of a sampling estimator decreases with the number of samples:
  - The distribution of sample averages is normally distributed with mean  $\mu$  and variance  $\frac{\sigma^2}{N}$ .

$$S_N \sim \mathcal{N}(\mu, \frac{\sigma}{\sqrt{N}})$$

- The more you explore an action  $a$ , the smaller the variance of  $Q_t(a)$ , the more certain you are about the estimation, the less you need to explore it.

# Upper-Confidence-Bound action selection

- **Upper-Confidence-Bound** (UCB) action selection is a **greedy** action selection method that uses an **exploration** bonus:

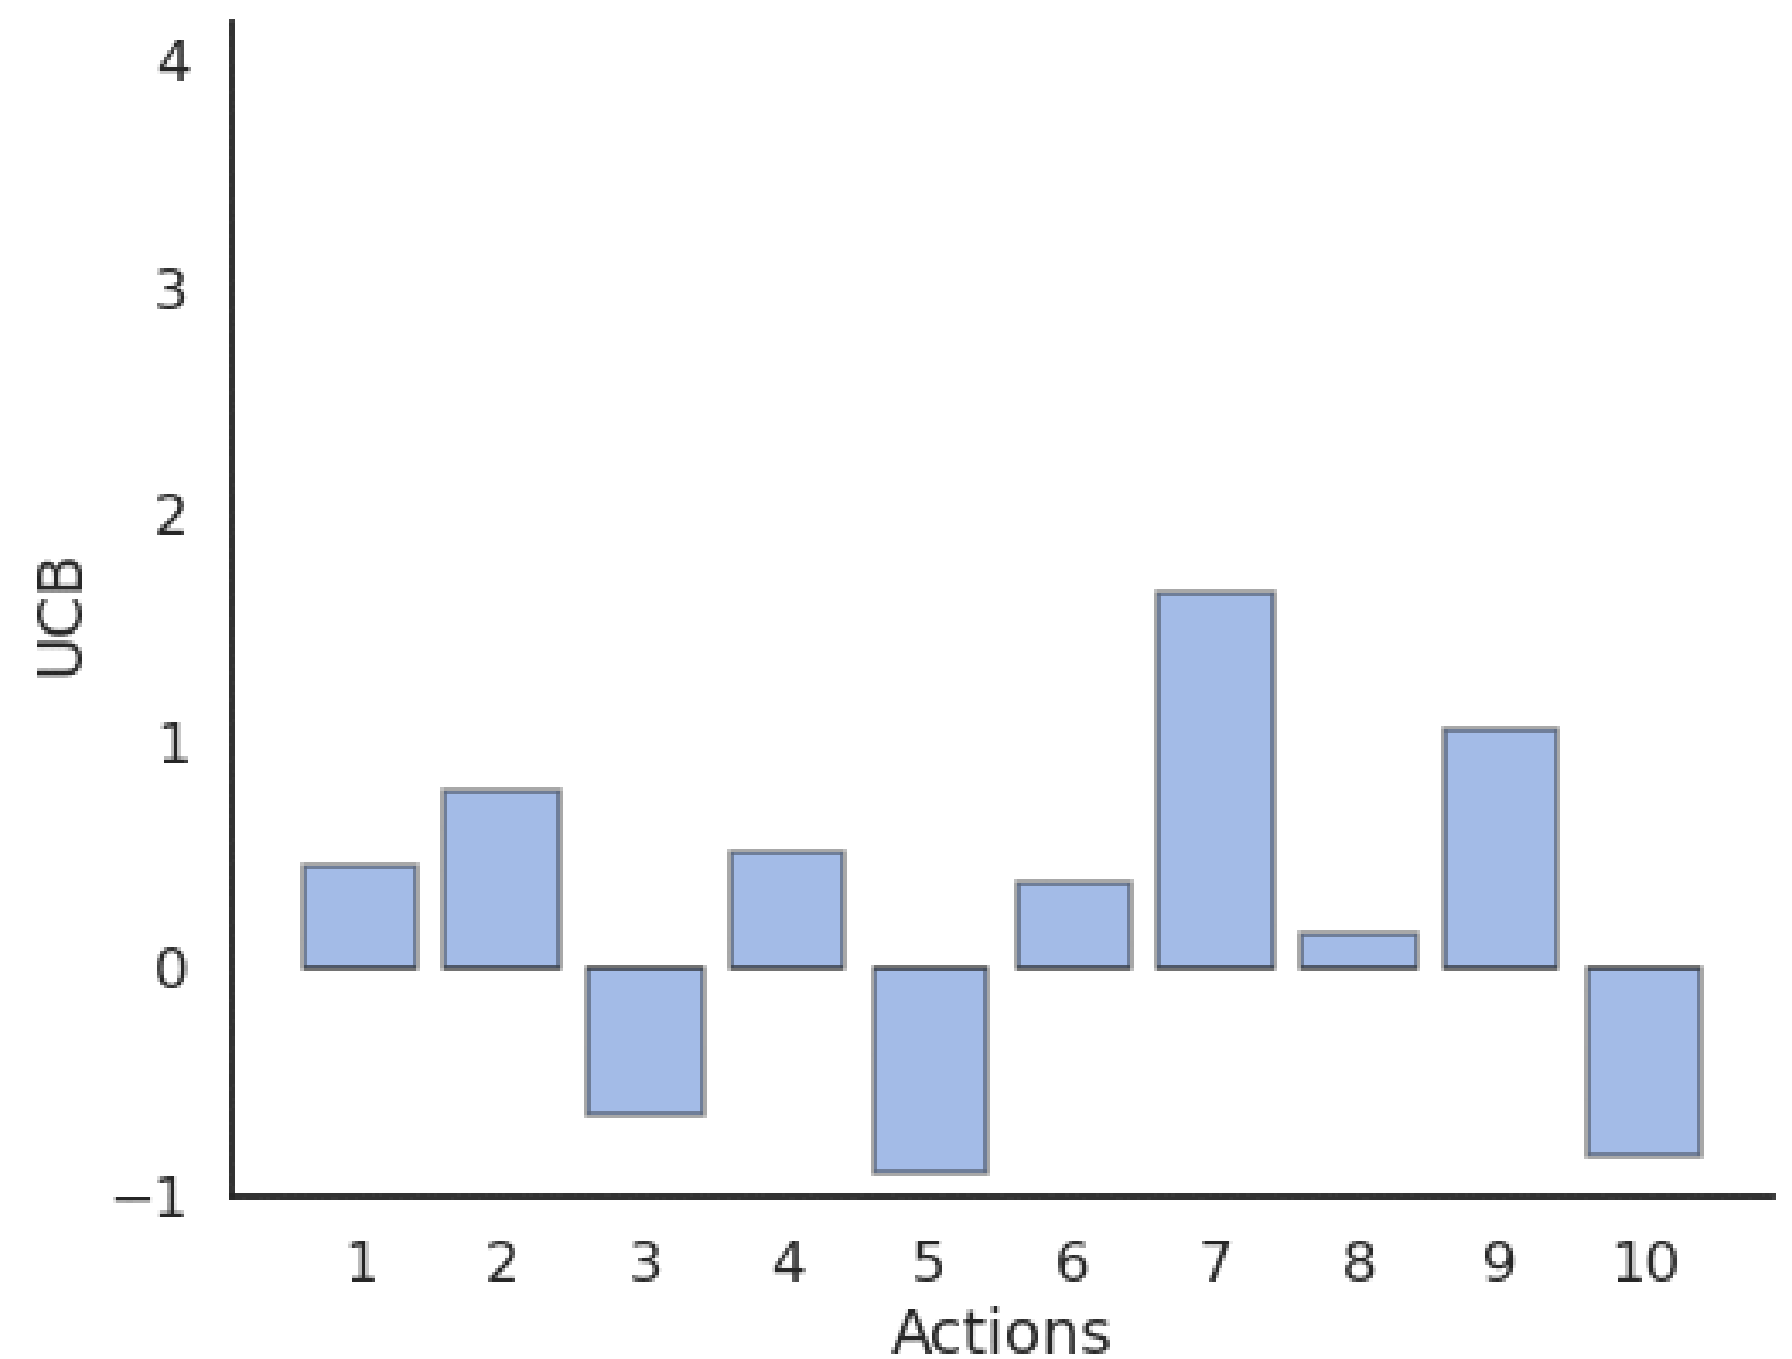
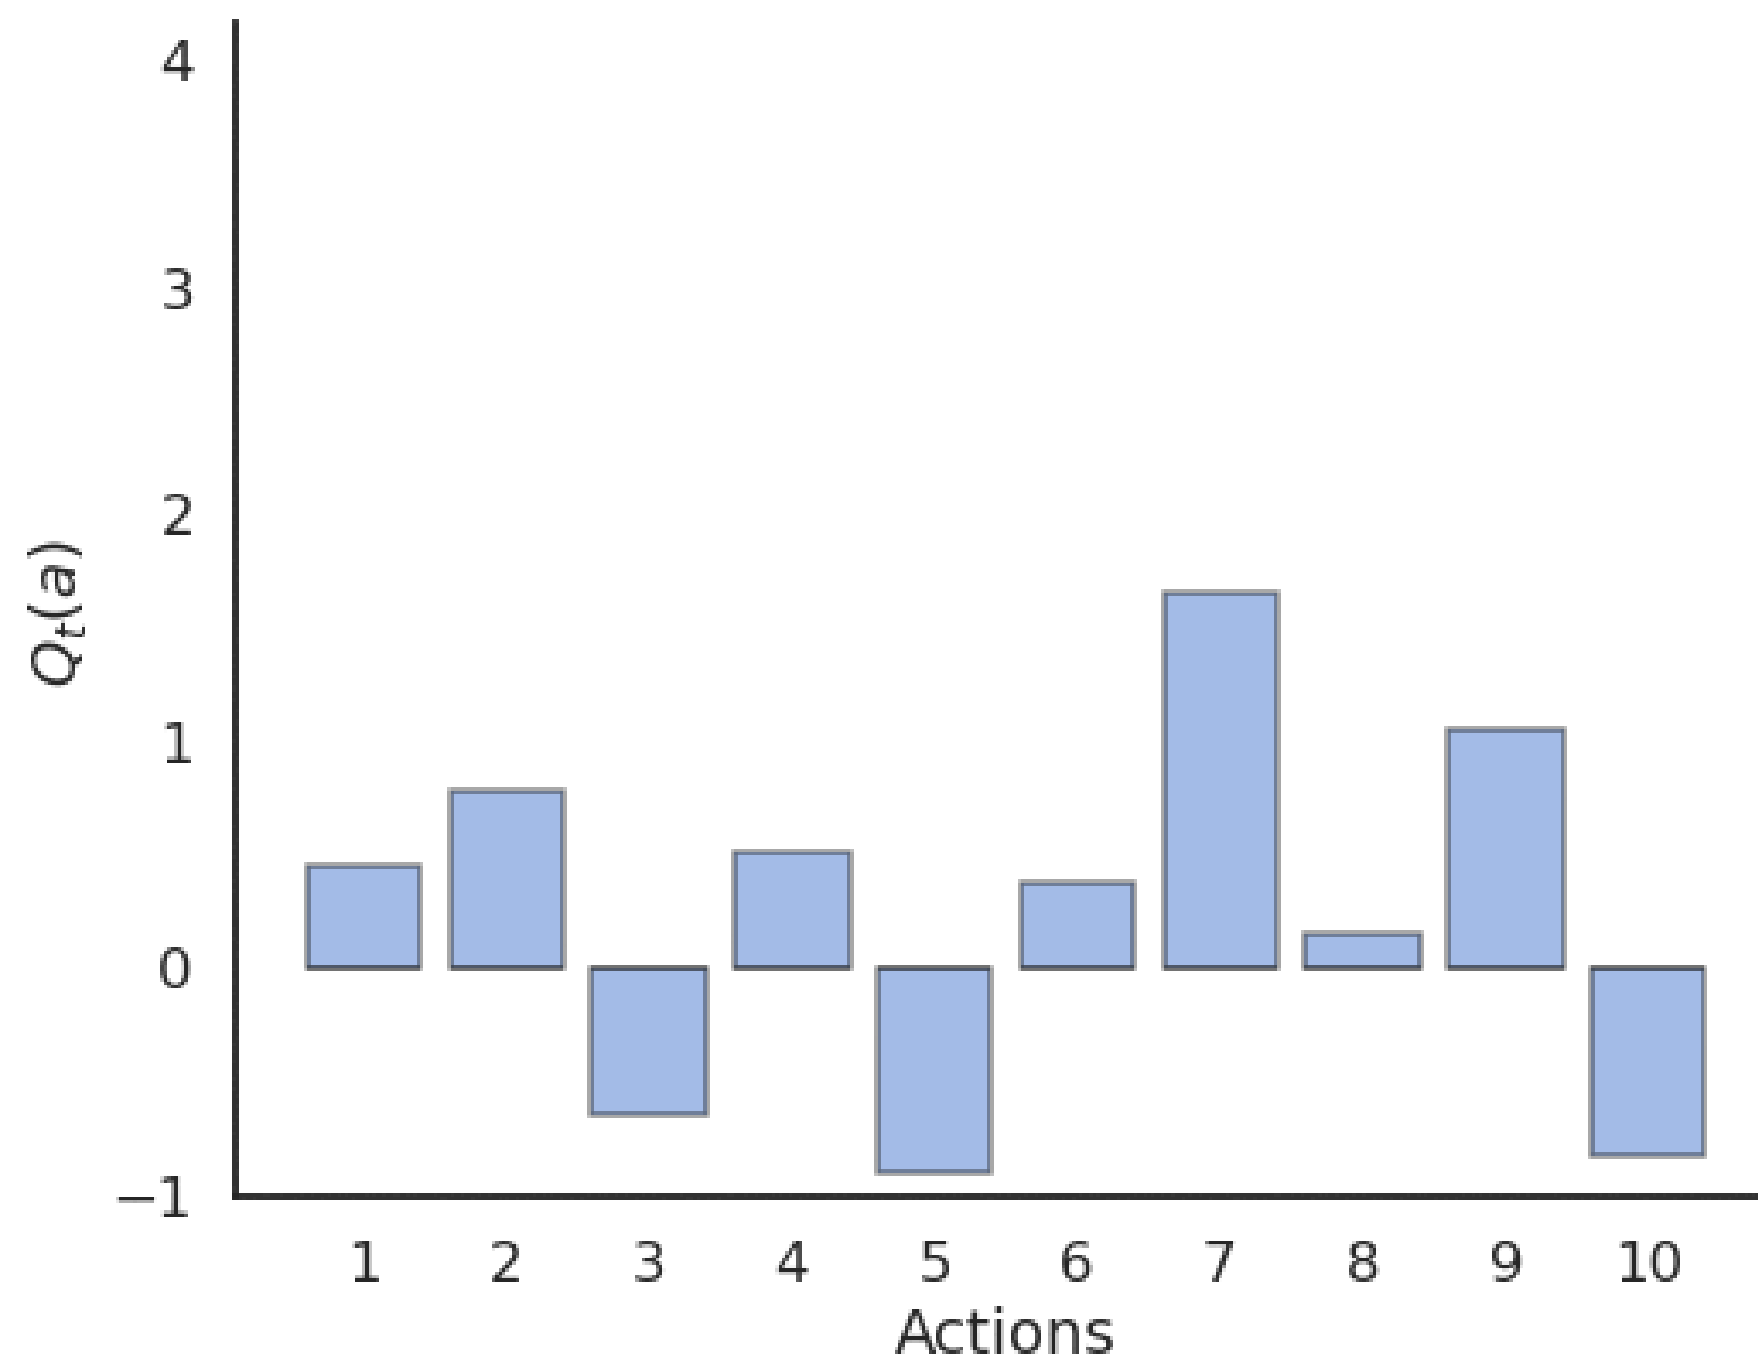
$$a_t^* = \operatorname{argmax}_a \left[ Q_t(a) + c \sqrt{\frac{\ln t}{N_t(a)}} \right]$$

- $Q_t(a)$  is the current estimated value of  $a$  and  $N_t(a)$  is the number of times the action  $a$  has already been selected.
- It realizes a balance between trusting the estimates  $Q_t(a)$  and exploring uncertain actions which have not been explored much yet.
- The term  $\sqrt{\frac{\ln t}{N_t(a)}}$  is an estimate of the variance of  $Q_t(a)$ . The sum of both terms is an **upper-bound** of the true value  $\mu + \sigma$ .
- When an action has not been explored much yet, the uncertainty term will dominate and the action be explored, although its estimated value might be low.
- When an action has been sufficiently explored, the uncertainty term goes to 0 and we greedily follow  $Q_t(a)$ .

# Upper-Confidence-Bound action selection

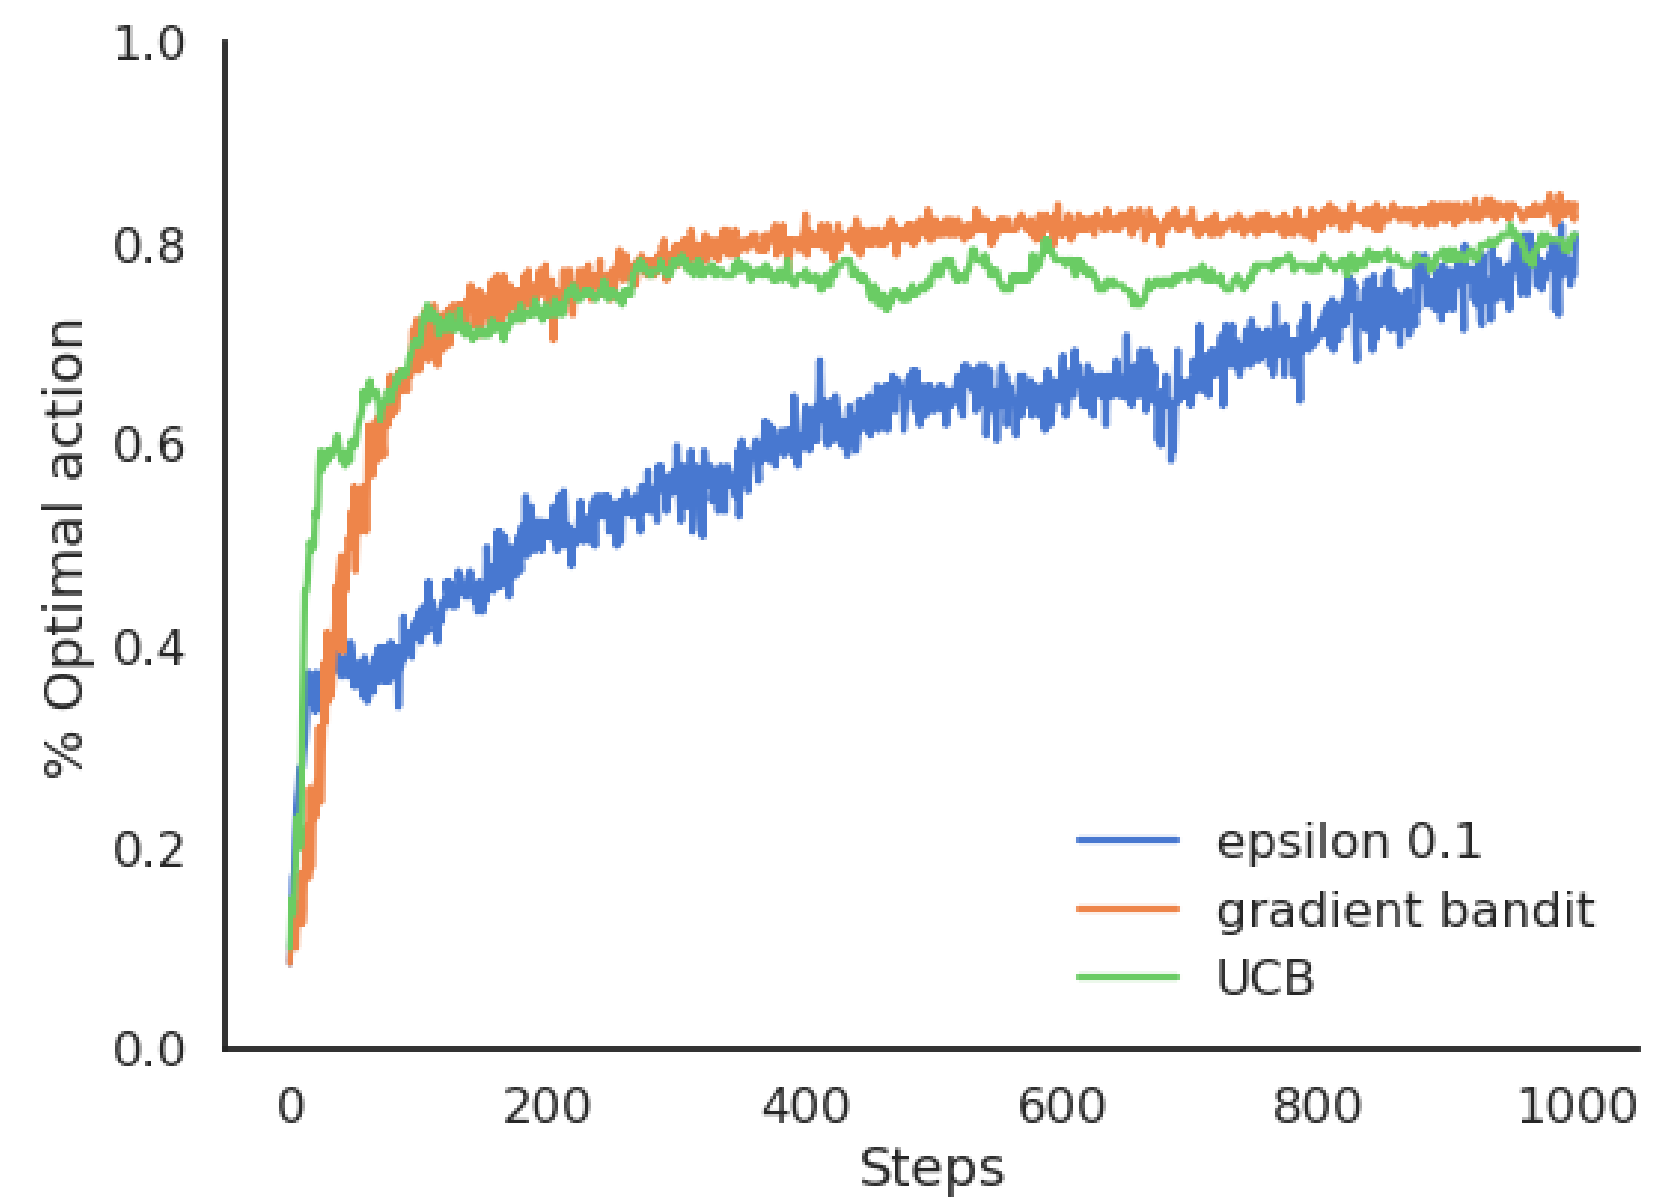
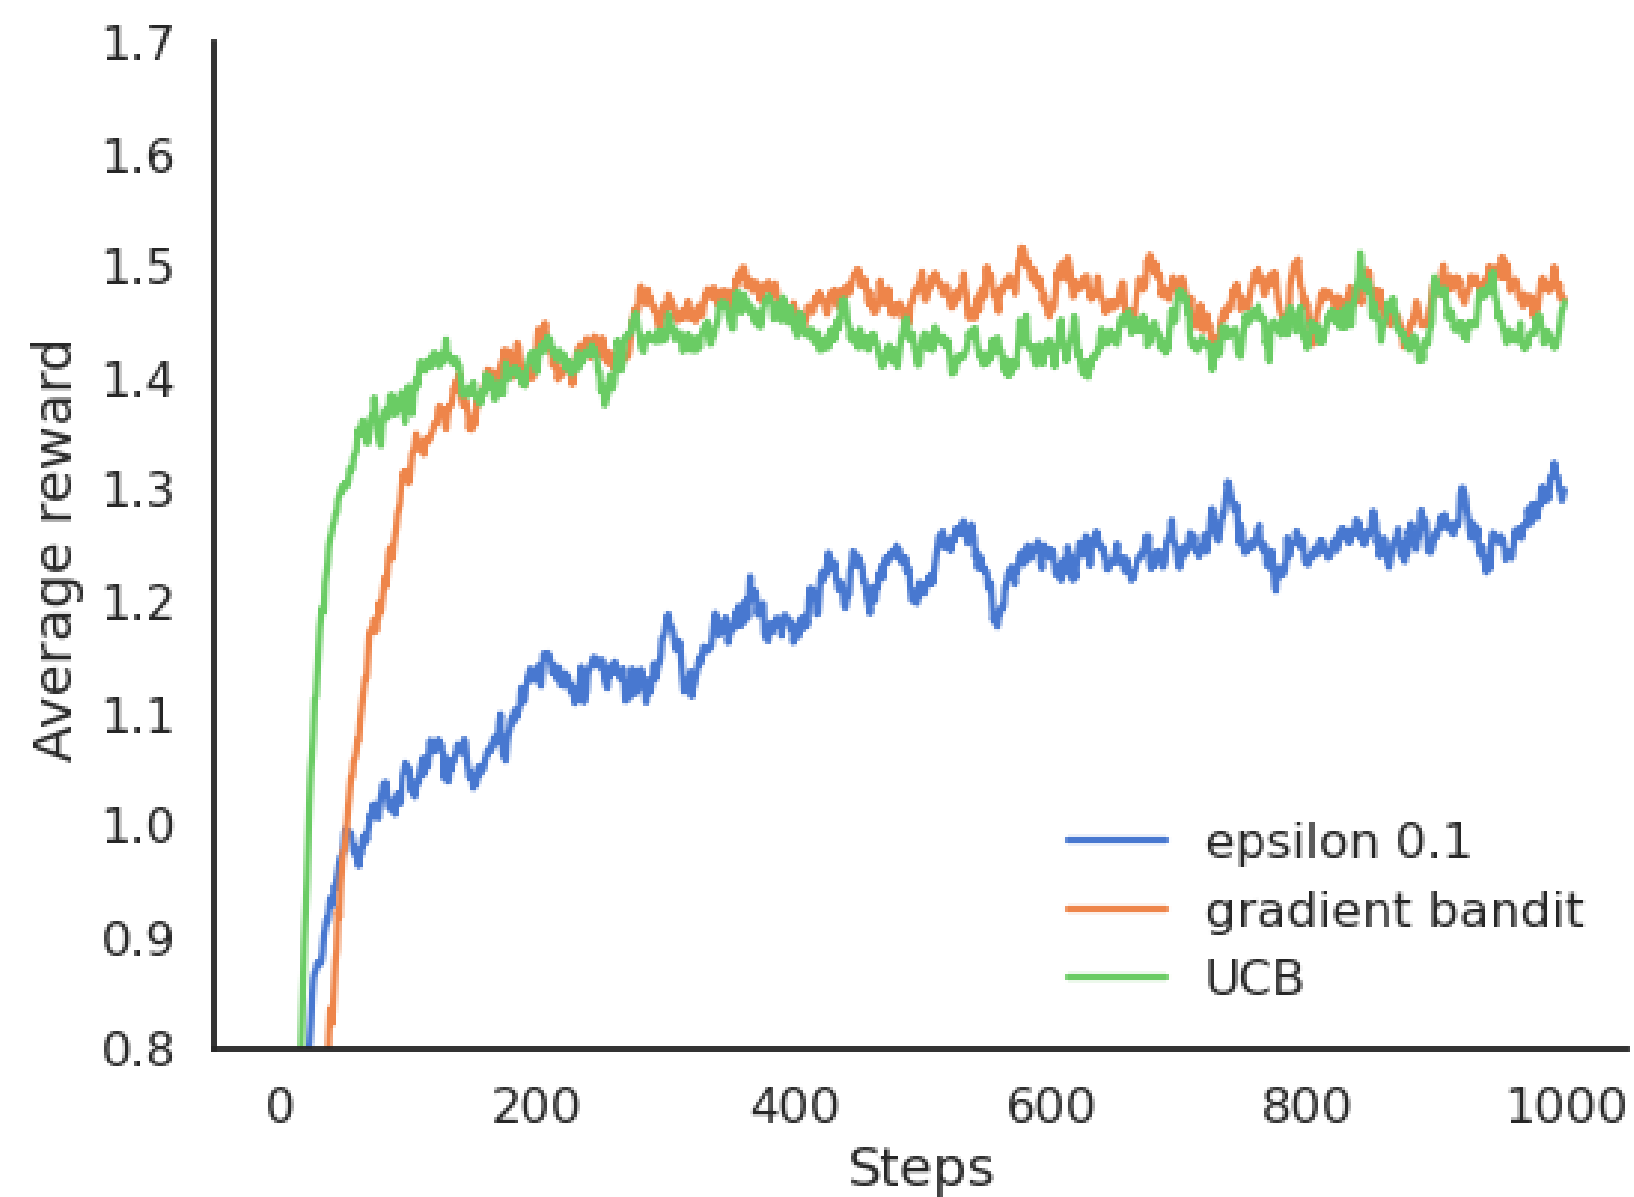
- The **exploration-exploitation** trade-off is automatically adjusted by counting visits to an action.

$$a_t^* = \operatorname{argmax}_a \left[ Q_t(a) + c \sqrt{\frac{\ln t}{N_t(a)}} \right]$$



# Upper-Confidence-Bound action selection

- The “smart” exploration of UCB allows to find the optimal action faster.



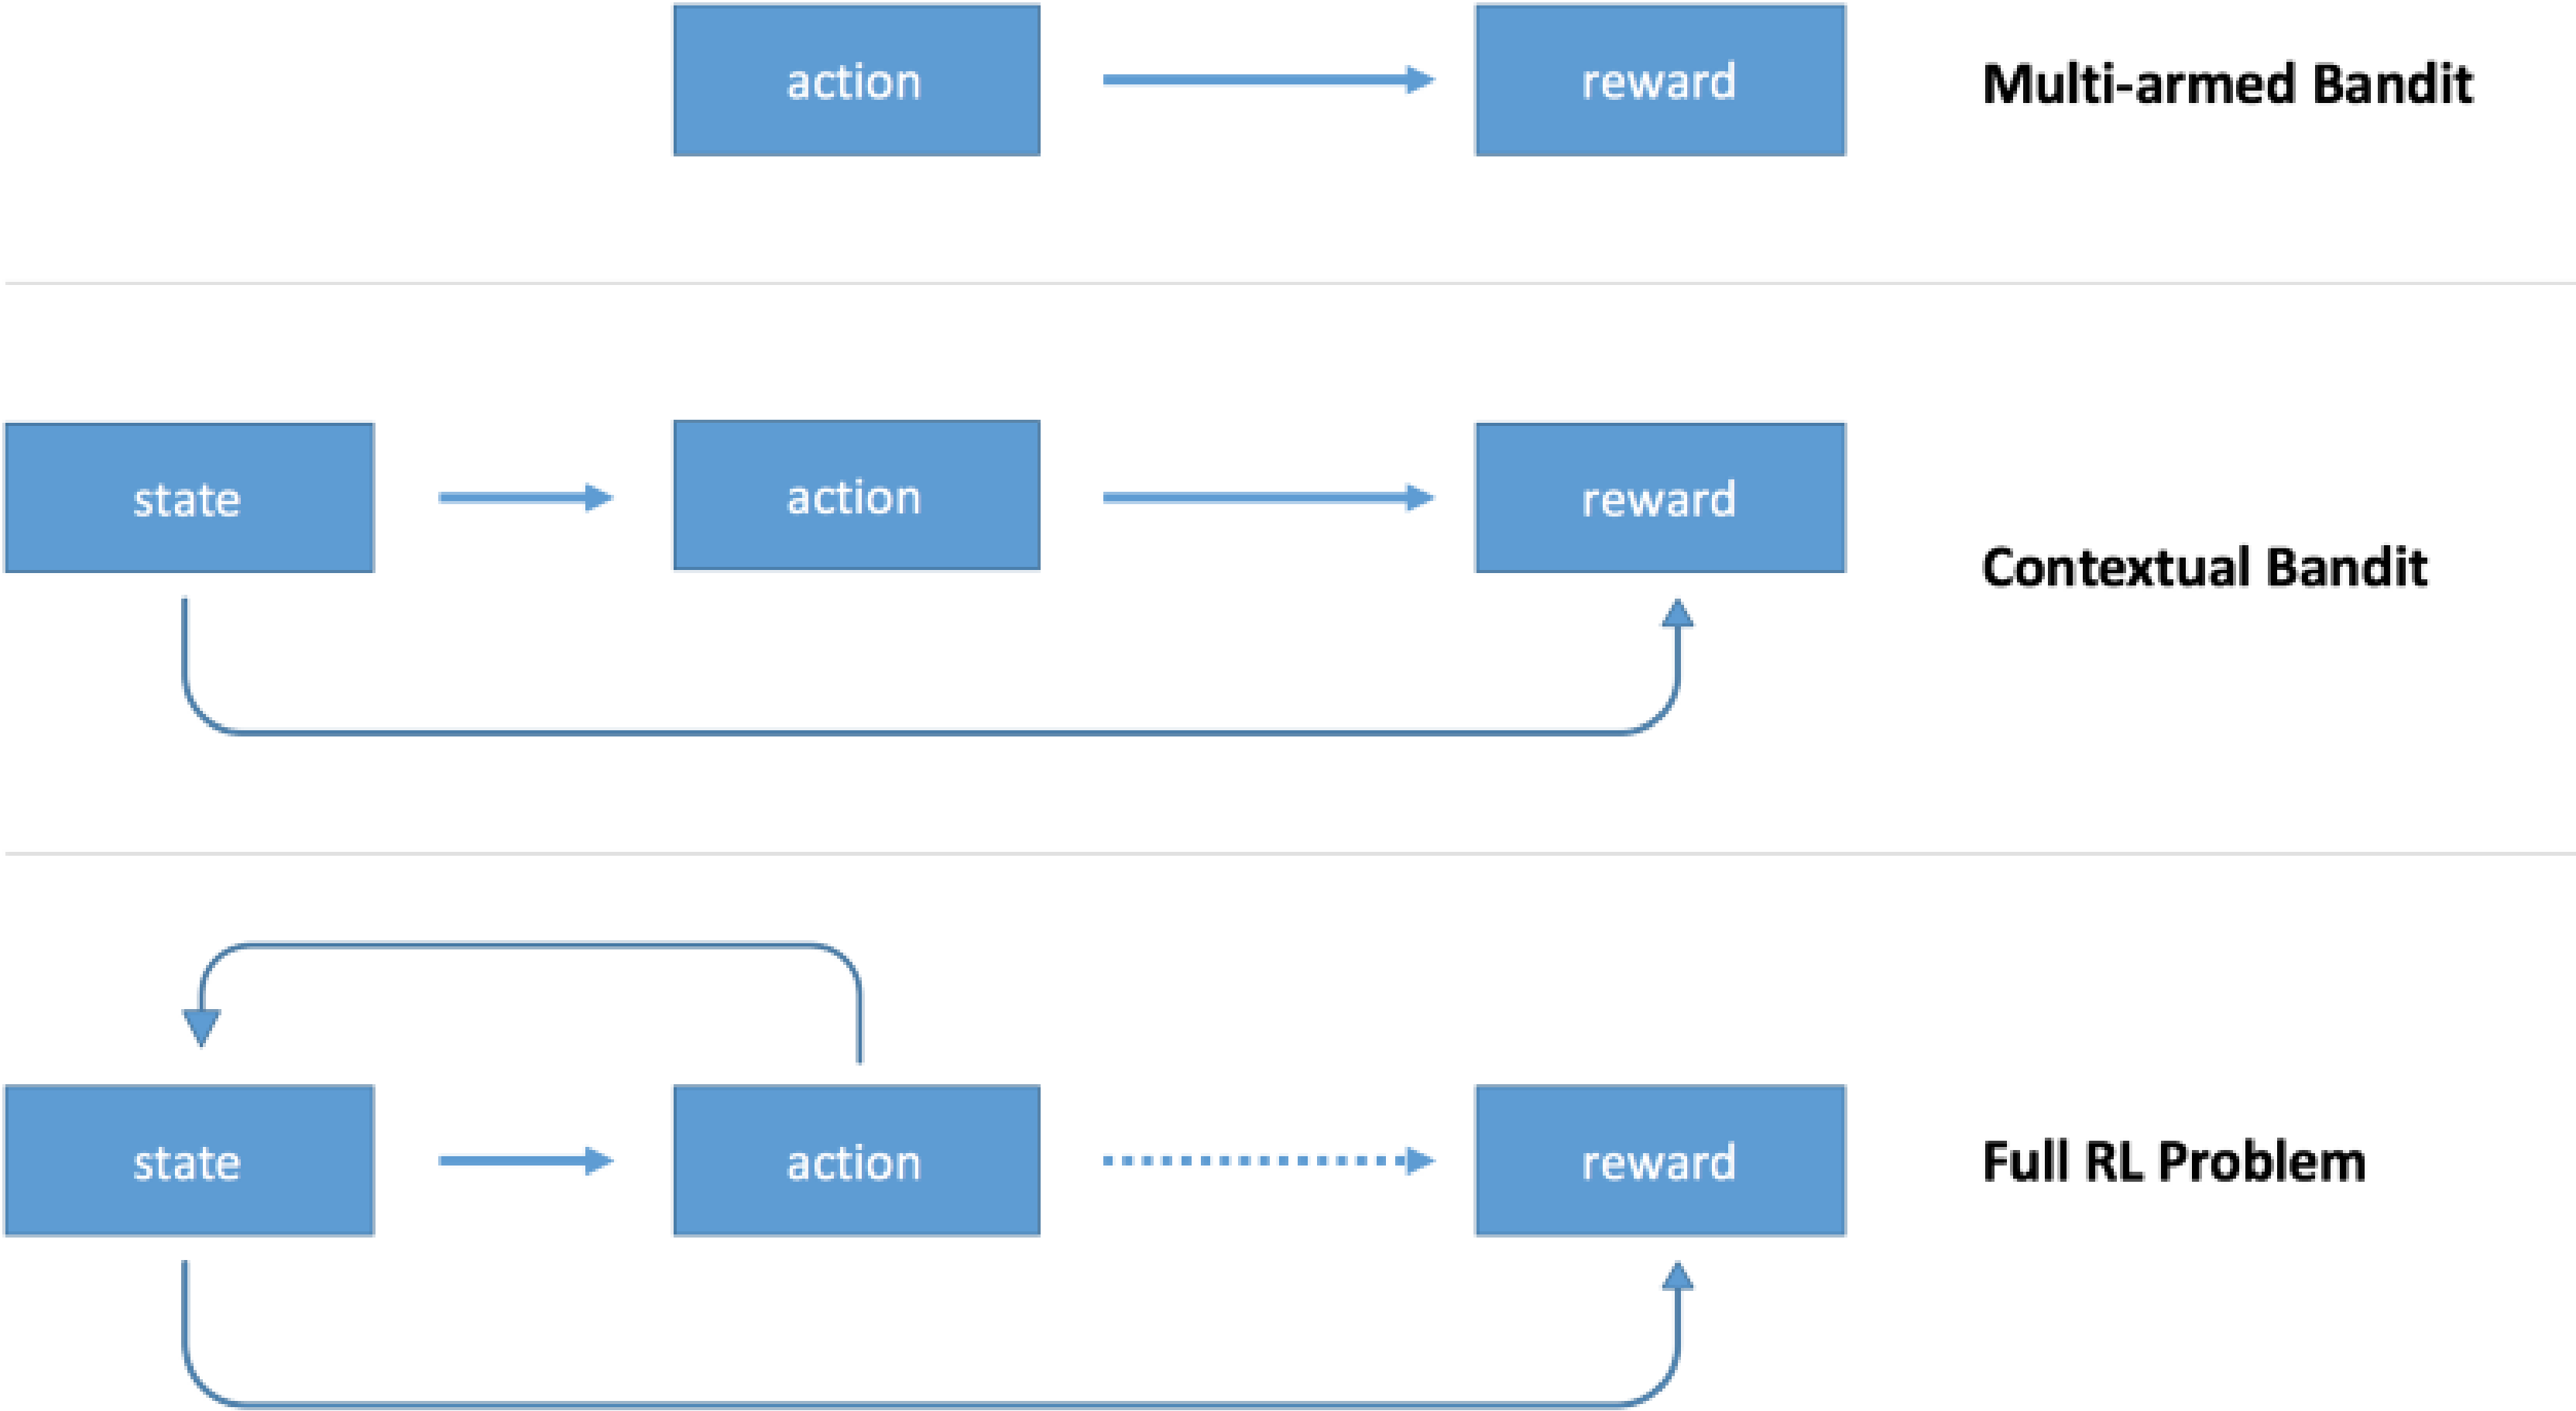
## Summary of evaluative feedback methods

- Greedy,  $\epsilon$ -greedy, softmax, reinforcement comparison, gradient bandit and UCB all have their own advantages and disadvantages depending on the type of problem: stationary or not, high or low reward variance, etc...
- These simple techniques are the most useful ones for bandit-like problems: more sophisticated ones exist, but they either make too restrictive assumptions, or are computationally intractable.
- Take home messages:
  1. RL tries to **estimate values** based on sampled rewards.
  2. One has to balance **exploitation and exploration** throughout learning with the right **action selection scheme**.
  3. Methods exploring more find **better policies**, but are initially slower.



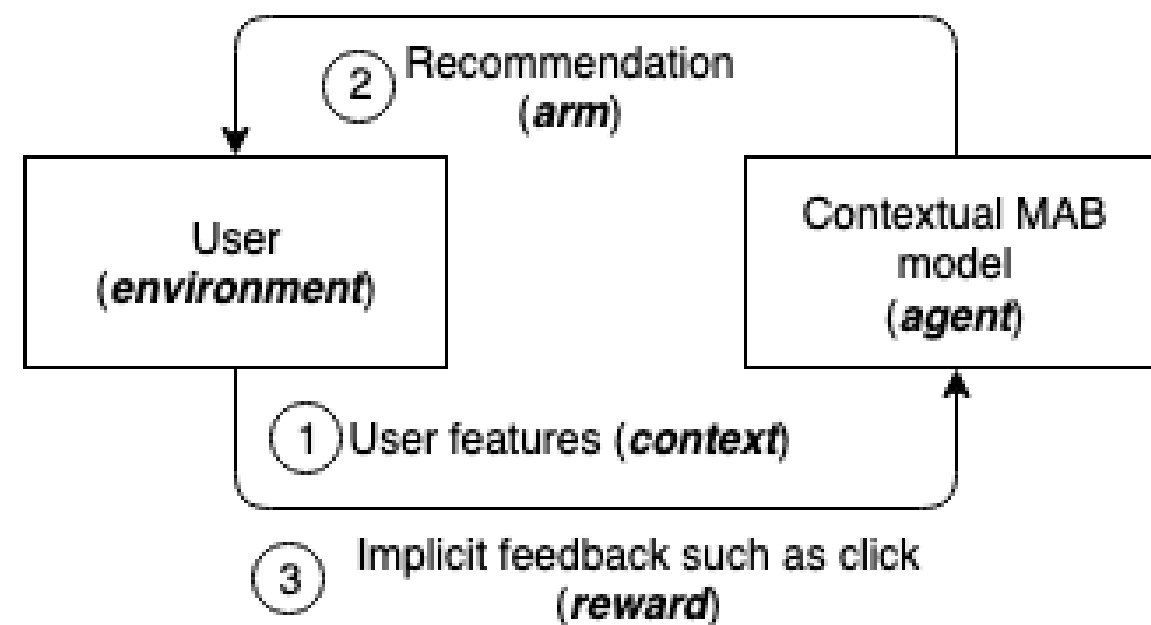
## 4 - Contextual bandits

# Contextual bandits



Source: <https://medium.com/emergent-future/simple-reinforcement-learning-with-tensorflow-part-1-5-contextual-bandits-bff01d1aad9c>

# Contextual bandits



Source: <https://aws.amazon.com/blogs/machine-learning/power-contextual-bandits-using-continual-learning-with-amazon-sagemaker-rl/>

## Recommender systems:

- Actions: advertisements.
- Context: user features / identity.
- Reward: user clicked on the ad.

- Some efficient algorithms have been developed recently, for example:

- In contextual bandits, the obtained rewards do not only depend on the action  $a$ , but also on the **state** or **context**  $s$ :

$$r_{t+1} \sim r(s, a)$$

- For example, the n-armed bandit could deliver rewards with different probabilities depending on:
  - who plays.
  - the time of the year.
  - the availability of funds in the casino.
- The problem is simply to estimate  $Q(s, a)$  instead of  $Q(a)$ ...

Agarwal, A., Hsu, D., Kale, S., Langford, J., Li, L., and Schapire, R. E. (2014). Taming the Monster: A Fast and Simple Algorithm for Contextual Bandits. in Proceedings of the 31 st International Conference on Machine Learning (Beijing, China), 9. <http://proceedings.mlr.press/v32/agarwalb14.pdf>