# Deep Reinforcement Learning

Deep learning

Julien Vitay

Professur für Künstliche Intelligenz - Fakultät für Informatik

1 - Artificial neural networks

Artificial neural networks

- An **artificial neural network** (ANN) is a cascade of **fully-connected** (FC) layers of artificial neurons.

- Each layer k transforms an input vector $\mathbf{h}_{k-1}$ into an output vector $\mathbf{h}_k$ using a weight matrix W_k, a bias vector $\mathbf{b}_k$ and an activation function $f()$.

$$\mathbf{h}_k = f(W_k \times \mathbf{h}_{k-1} + \mathbf{b}_k)$$

- Overall, ANNs are **non-linear parameterized function estimators** from the inputs $\mathbf{x}$ to the outputs $\mathbf{y}$ with parameters θ (all weight matrices and biases).

$$\mathbf{y} = F_\theta(\mathbf{x})$$

Loss functions

- ANNs can be used for both **regression** (continuous outputs) and **classification** (discrete outputs) tasks.

- In supervised learning, we have a fixed **training set** $\mathcal{D}$ of N samples $(\mathbf{x}_t, \mathbf{t}_i)$, where t_i is the **desired output** or **target**.

- **Regression:**

  - The output layer uses a **linear** activation function: $f(x) = x$

  - The network minimizes the **mean square error** (mse) of the model on the training set:

  $$\mathcal{L}(\theta) = \mathbb{E}_{\mathbf{x},\mathbf{t} \in \mathcal{D}}[||\mathbf{t} - \mathbf{y}||^2]$$

- **Classification:**

  - The output layer uses the **softmax** operator to produce a probability distribution: $y_j = \frac{e^{z_j}}{\sum_k e^{z_k}}$

  - The network minimizes the **cross-entropy** or **negative log-likelihood** of the model on the training set:

  $$\mathcal{L}(\theta) = \mathbb{E}_{\mathbf{x},\mathbf{t} \in \mathcal{D}}[-\mathbf{t} \, \log \mathbf{y}]$$

Cross-entropy

- The cross-entropy between two probability distributions X and Y measures their similarity:

$$H(X, Y) = \mathbb{E}_{x \sim X}\left[-\log P(Y = x)\right]$$

- Are samples from X likely under Y?

- Minimizing the cross-entropy makes the two distributions equal almost anywhere.

samples from X

Cross-entropy

- In supervised learning, the targets **t** are fixed **one-hot encoded vectors**.

$$\mathcal{L}(\theta) = \mathbb{E}_{\mathbf{x},\mathbf{t} \in \mathcal{D}}[-\sum_j t_j \log y_j]$$

- But it could be any target distribution.

Backpropagation

- In both cases, we want to minimize the loss function by applying **Stochastic Gradient Descent** (SGD) or a variant (Adam, RMSprop).
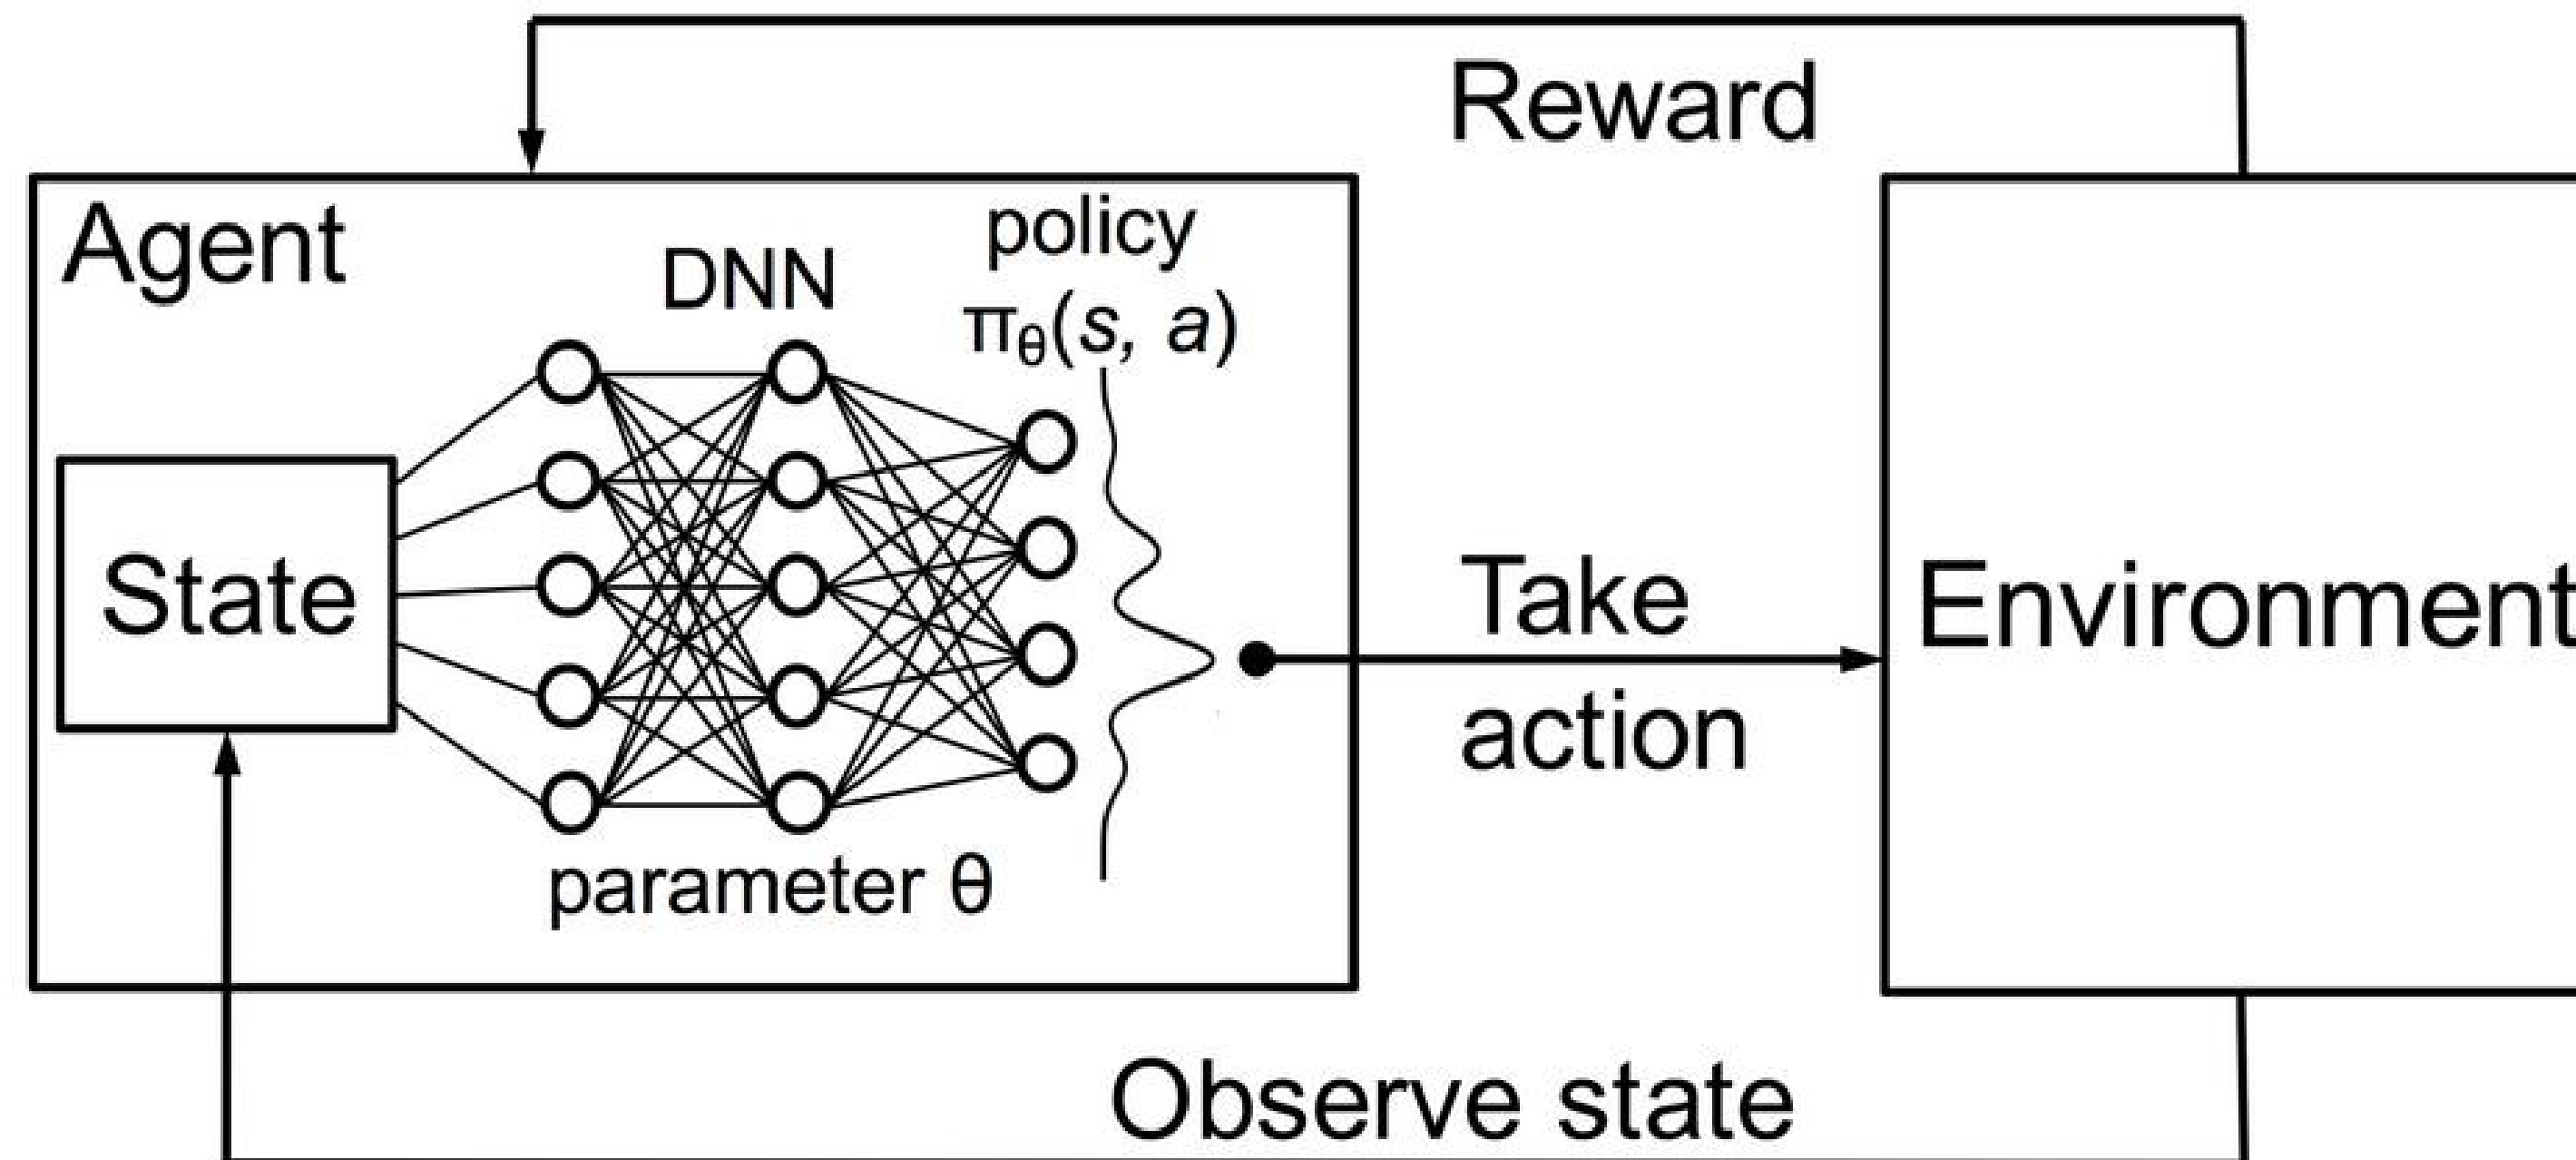
$$\Delta\theta = -\eta \, \nabla_\theta \mathcal{L}(\theta)$$

- The question is how to compute the **gradient of the loss function** w.r.t the parameters θ.

- For both the mse and cross-entropy loss functions, we have:

$$\nabla_\theta \mathcal{L}(\theta) = \mathbb{E}_{\mathcal{D}}\big[-(\mathbf{t} - \mathbf{y}) \, \nabla_\theta \, \mathbf{y}\big]$$

- There is an algorithm to compute efficiently the gradient of the output w.r.t the parameters: **backpropagation** (see Neurocomputing).

- In deep RL, we do not care about backprop: tensorflow or pytorch do it for us.

Rumelhart, Hinton, Williams (1986). "Learning representations by back-propagating errors". Nature.

Components of neural networks

- There are three aspects to consider when building a neural network:

1. **Architecture:** how many layers, what type of layers, how many neurons, etc.

    - Task-dependent: each RL task will require a different architecture. Not our focus.

2. **Loss function:** what should the network do?

    - Central to deep RL!

3. **Update rule** how should we update the parameters θ to minimize the loss function? SGD, backprop.

    - Not really our problem, but see *natural gradients* later.

2 - Convolutional neural networks

Convolutional layers

- When using images as inputs, **fully-connected networks** (FCN) would have too many weights:

  - Slow.

  - Overfitting.

- **Convolutional layers** reduce the number of weights by **reusing** weights at different locations.

  - Principle of a convolution.

  - Fast and efficient.

Convolutional layers

- A **convolutional layer** extracts **features** of its inputs.

- d filters are defined with very small sizes (3x3, 5x5...).

- Each filter is convoluted over the input image (or the previous layer) to create a **feature map**.

- The set of d feature maps becomes a new 3D structure: a **tensor**.

- If the input image is 32x32x3, the resulting tensor will be 32x32xd.

- The convolutional layer has only very few parameters: each feature map has 3x3 values in the filter and a bias, i.e. 10 parameters.

- The convolution operation is **differentiable**: backprop will work.

Source: https://github.com/vdumoulin/conv_arithmetic

Max-pooling

Source: http://cs231n.github.io/convolutional-networks/

- The number of elements in a convolutional layer is still too high. We need to reduce the spatial dimension of a convolutional layer by **downsampling** it.

- For each feature, a **max-pooling** layer takes the maximum value of a feature for each subregion of the image (mostly 2x2).

- Pooling allows translation invariance: the same input pattern will be detected whatever its position in the input image.

- Max-pooling is also differentiable.

Convolutional neural networks

INPUT 32x32 · C1: feature maps 6@28x28 · S2: f. maps 6@14x14 · C3: f. maps 16@10x10 · S4: f. maps 16@5x5 · C5: layer 120 · F6: layer 84 · OUTPUT 10 · Convolutions · Subsampling · Convolutions · Subsampling · Full connection · Full connection · Gaussian connections

- A **convolutional neural network** (CNN) is a cascade of convolution and pooling operations, extracting layer by layer increasingly complex features.

- The spatial dimensions decrease after each pooling operation, but the number of extracted features increases after each convolution.

- One usually stops when the spatial dimensions are around 7x7.

- The last layers are fully connected. Can be used for regression and classification depending on the output layer and the loss function.

- Training a CNN uses backpropagation all along: the convolution and pooling operations are differentiable.

Convolutional neural networks

INPUT   CONVOLUTION + RELU   POOLING   CONVOLUTION + RELU   POOLING   FLATTEN   FULLY CONNECTED   SOFTMAX

FEATURE LEARNING   **CLASSIFICATION**

Source: https://towardsdatascience.com/a-comprehensive-guide-to-convolutional-neural-networks-the-eli5-way-3bd2b1164a53
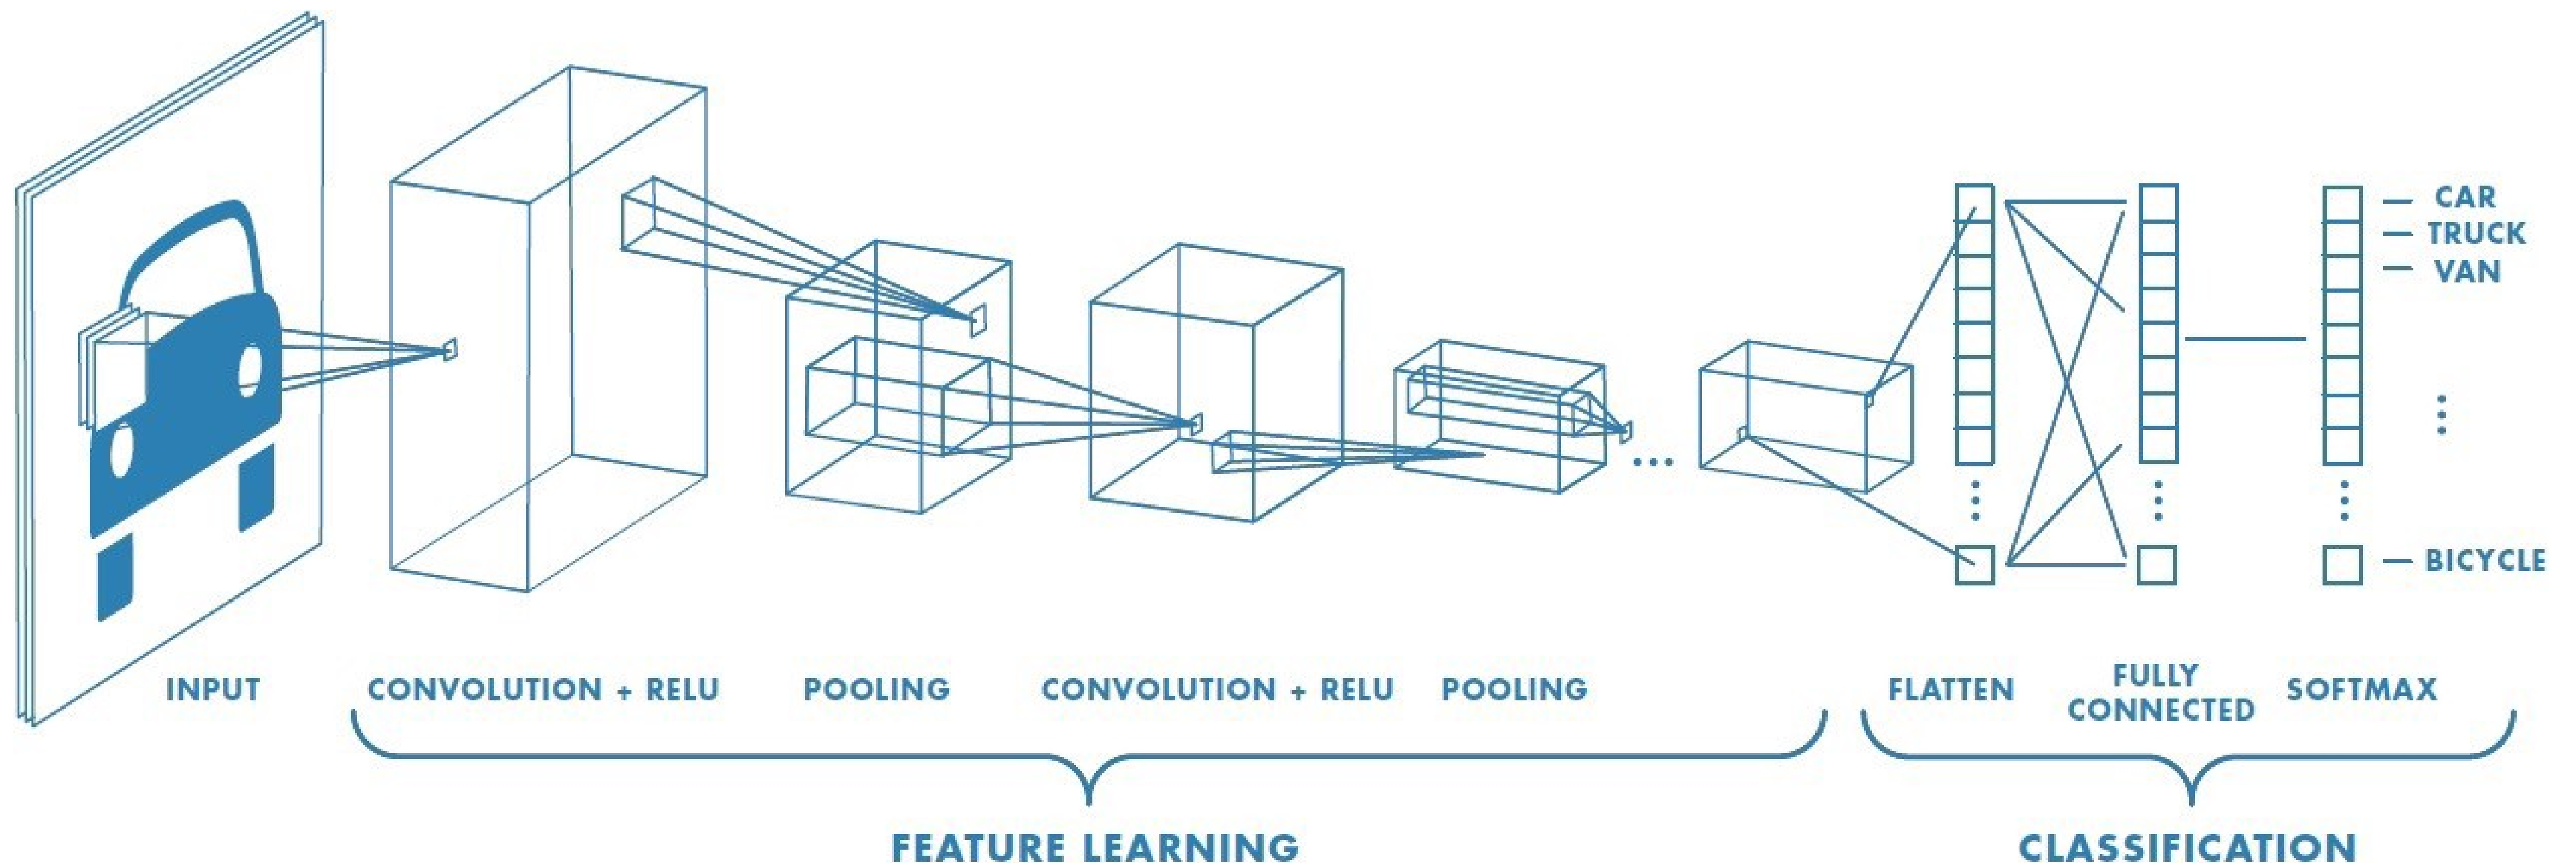
- The only thing we need to know is that CNNs are non-linear function approximators that work well with images.

$$\mathbf{y} = F_\theta(\mathbf{x})$$

- The conv layers **extract complex features** from the images through learning.

- The last FC layers allow to approximate values (regression) or probability distributions (classification).

3 - Autoencoders

Autoencoders

encode →   decode →

input   hidden   output

$q_\phi(z|x)$ → $p_\theta(x|z)$ →

x → $\tilde{x}$

- The problem with FCN and CNN is that they **extract features** in supervised learning tasks.

  - Need for a lot of annotated data (image, label).

- **Autoencoders** allows **unsupervised learning**:

  - They only need inputs (images).

- Their task is to **reconstruct** the input:

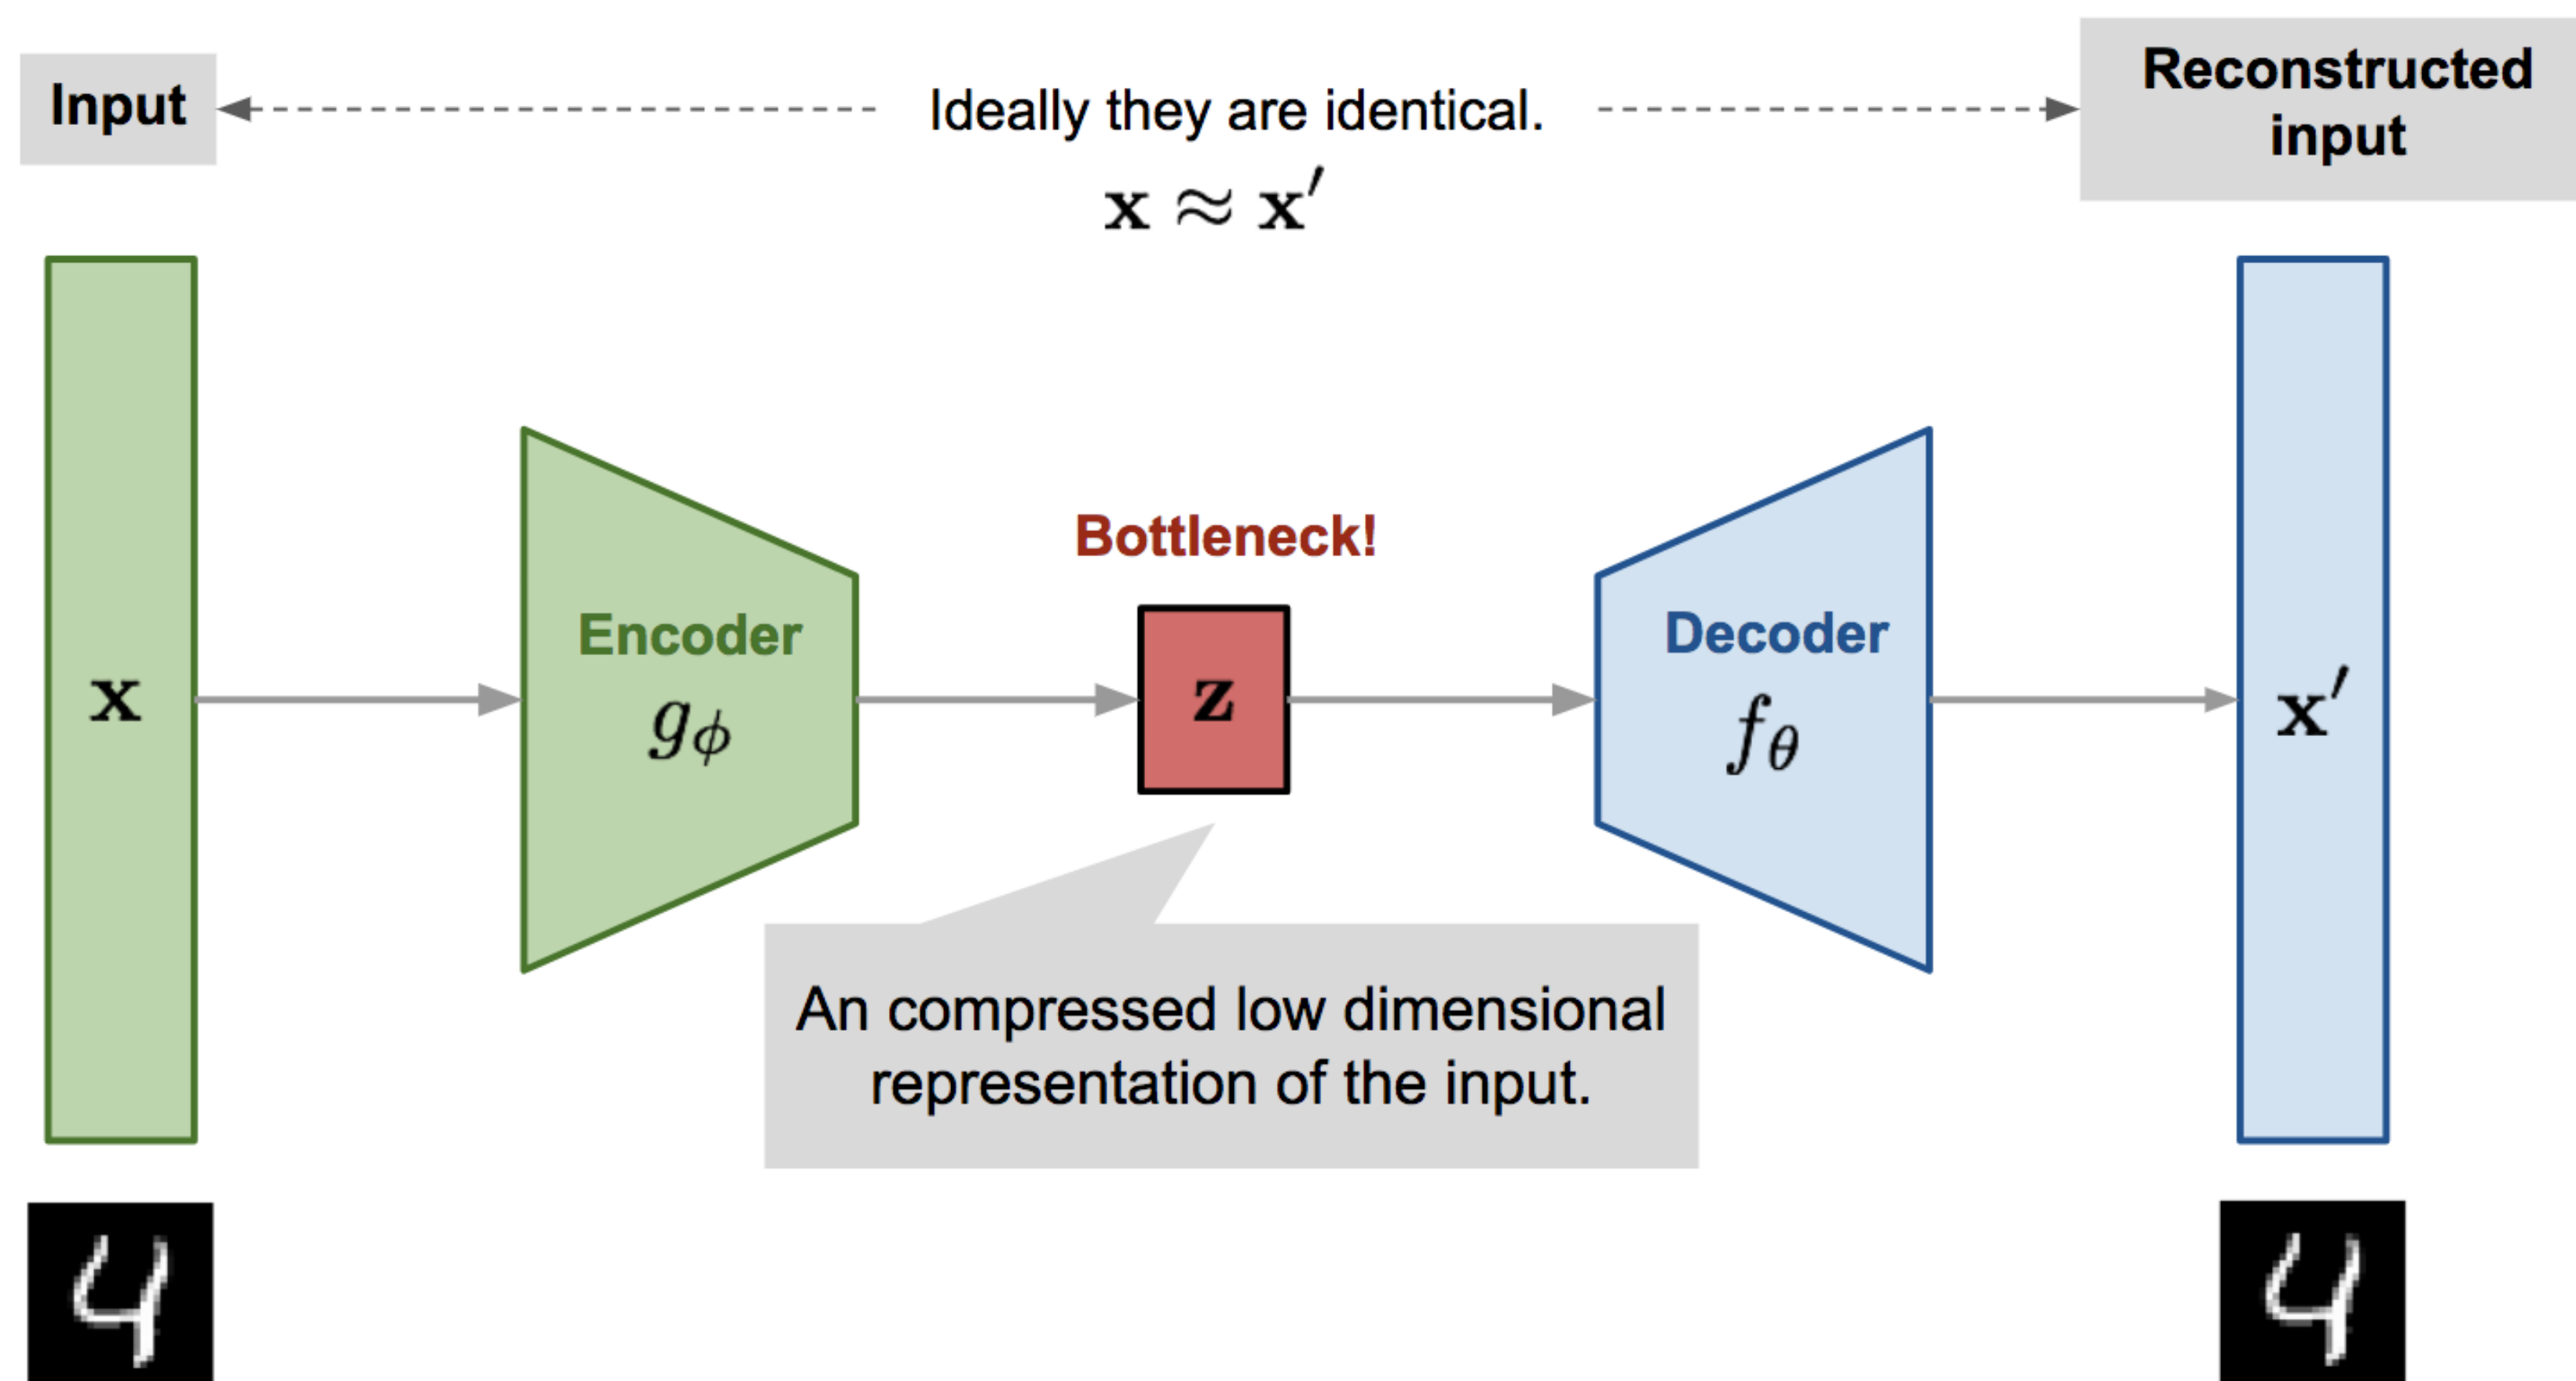$$\mathbf{y} = \tilde{\mathbf{x}} \approx \mathbf{x}$$

- The **reconstruction loss** is simply the **mse** between the input and its reconstruction.

$$\mathcal{L}_{\text{autoencoder}}(\theta) = \mathbb{E}_{\mathbf{x} \in \mathcal{D}}[||\tilde{\mathbf{x}} - \mathbf{x}||^2]$$

- Apart from the loss function, they are trained as regular NNs.
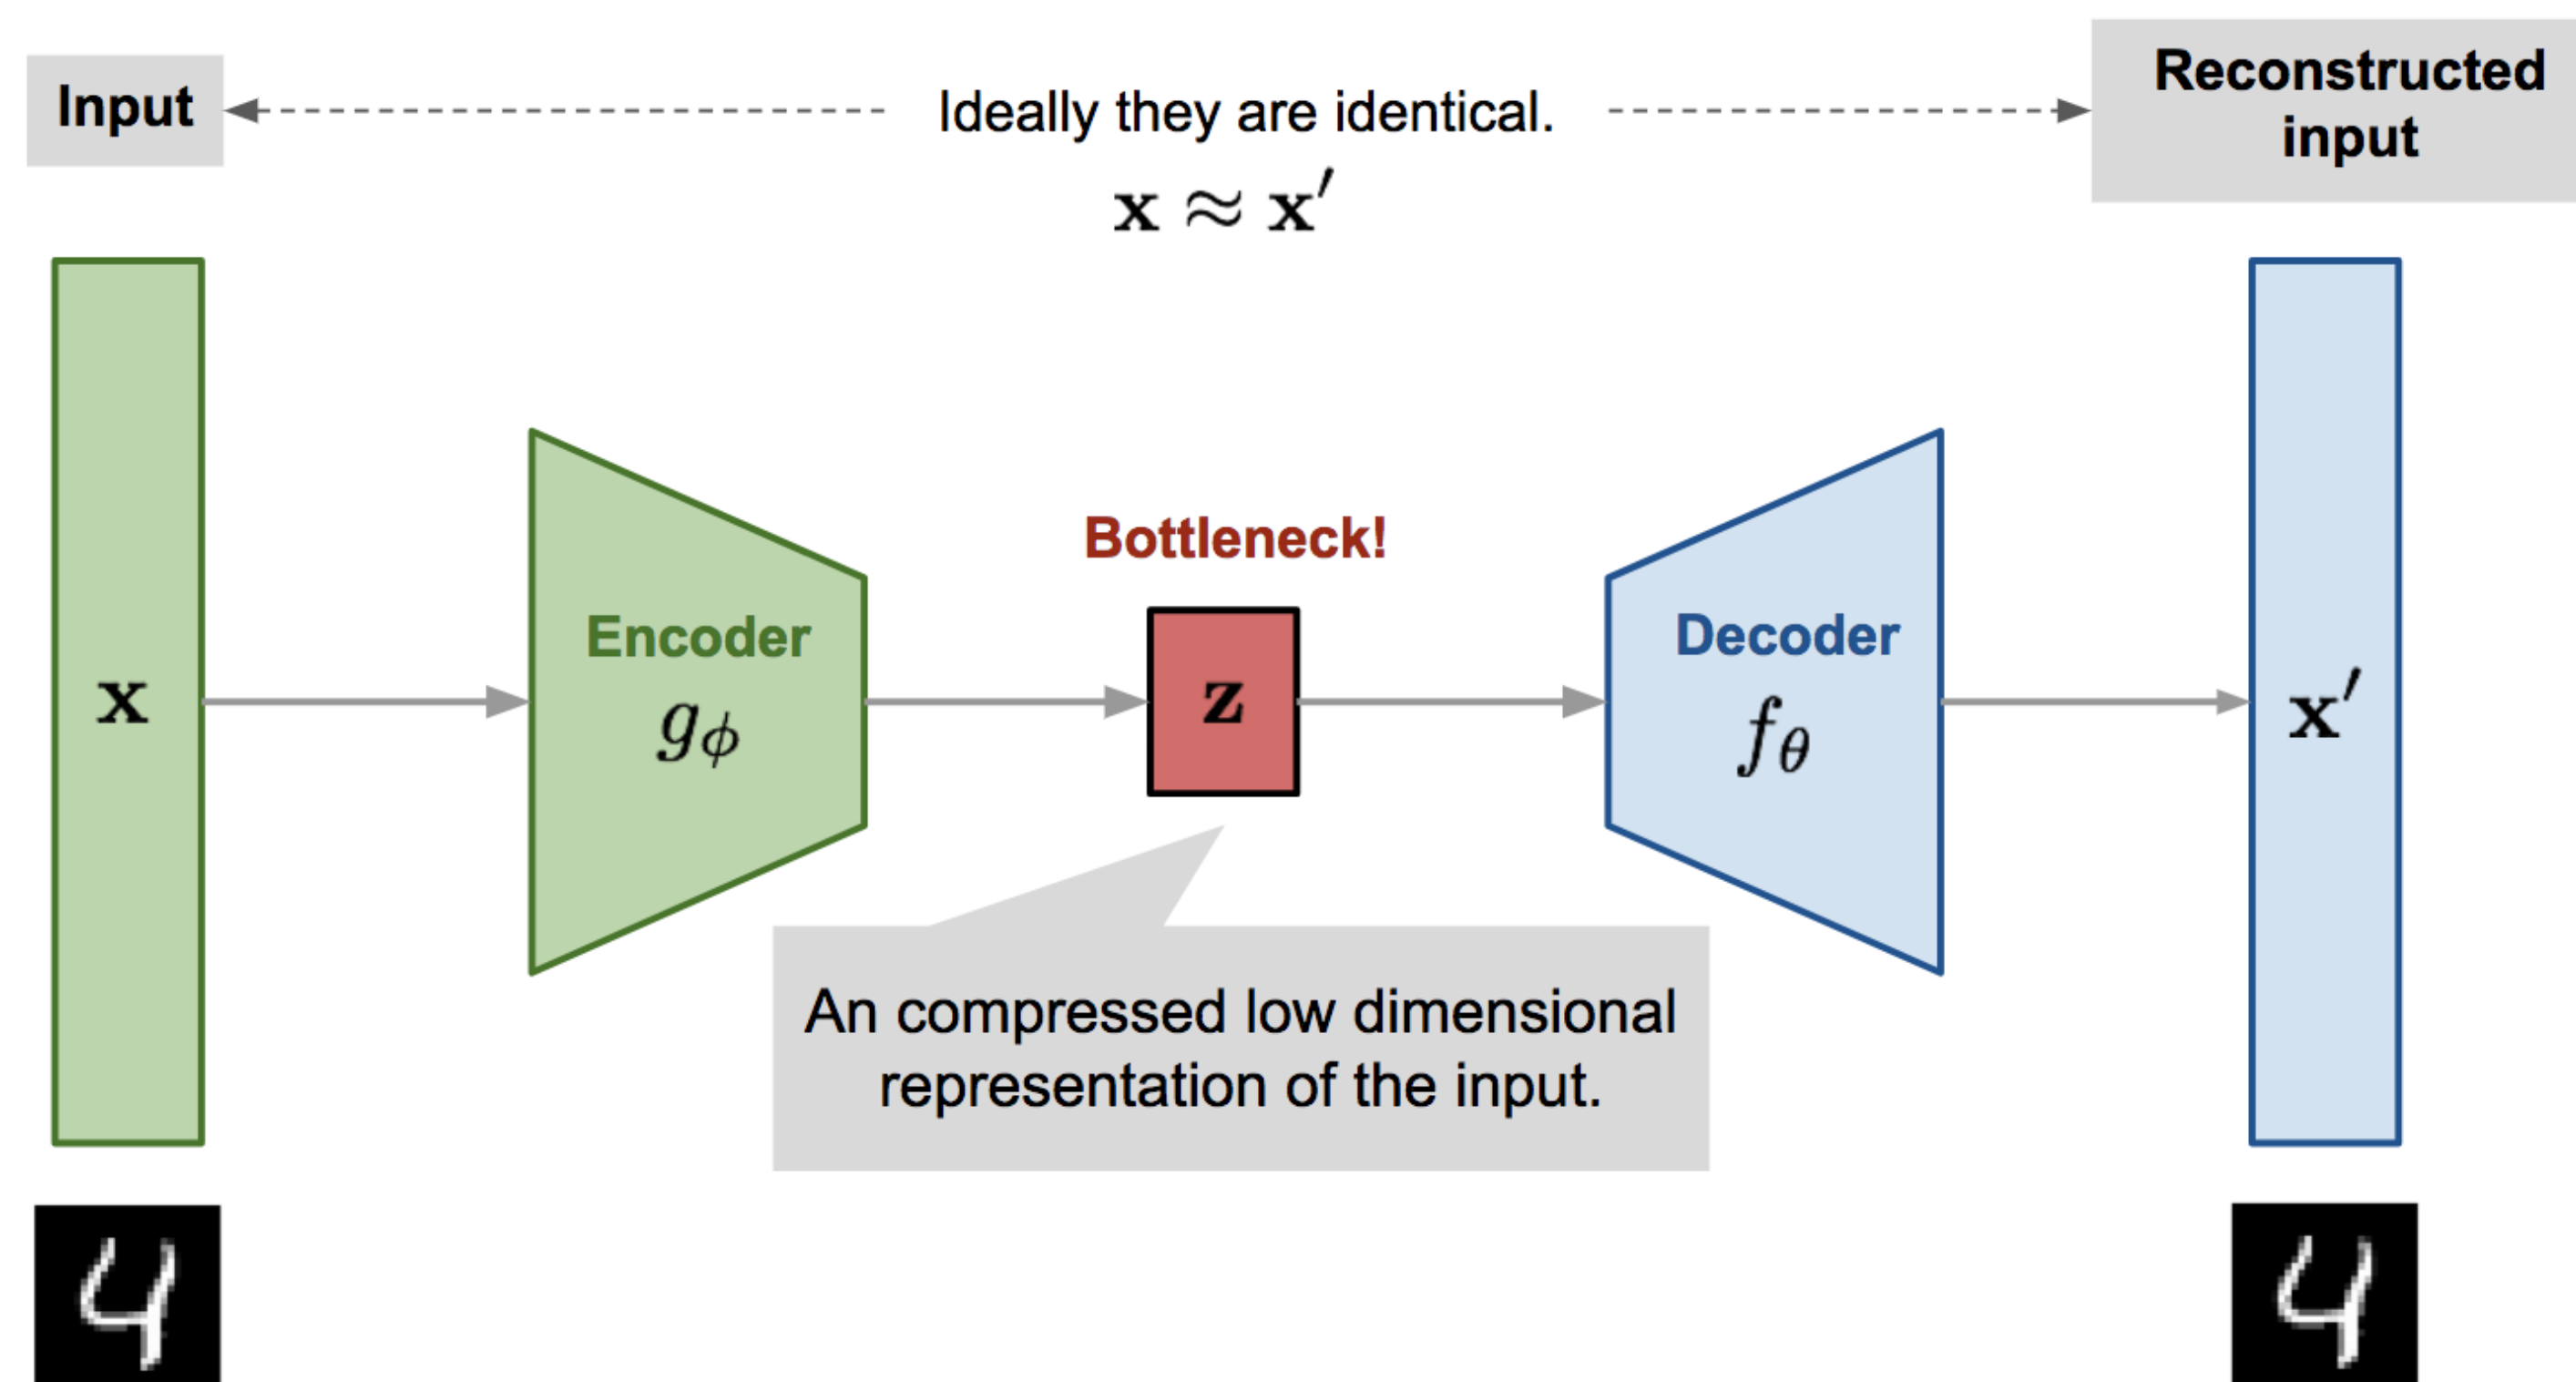
Autoencoders

- Autoencoders consists of:

  - the **encoder**: from the input $\mathbf{x}$ to the **latent space** $\mathbf{z}$.

  - the **decoder**: from the latent space $\mathbf{z}$ to the reconstructed input $\tilde{\mathbf{x}}$.

Source: https://lilianweng.github.io/lil-log/2018/08/12/from-autoencoder-to-beta-vae.html
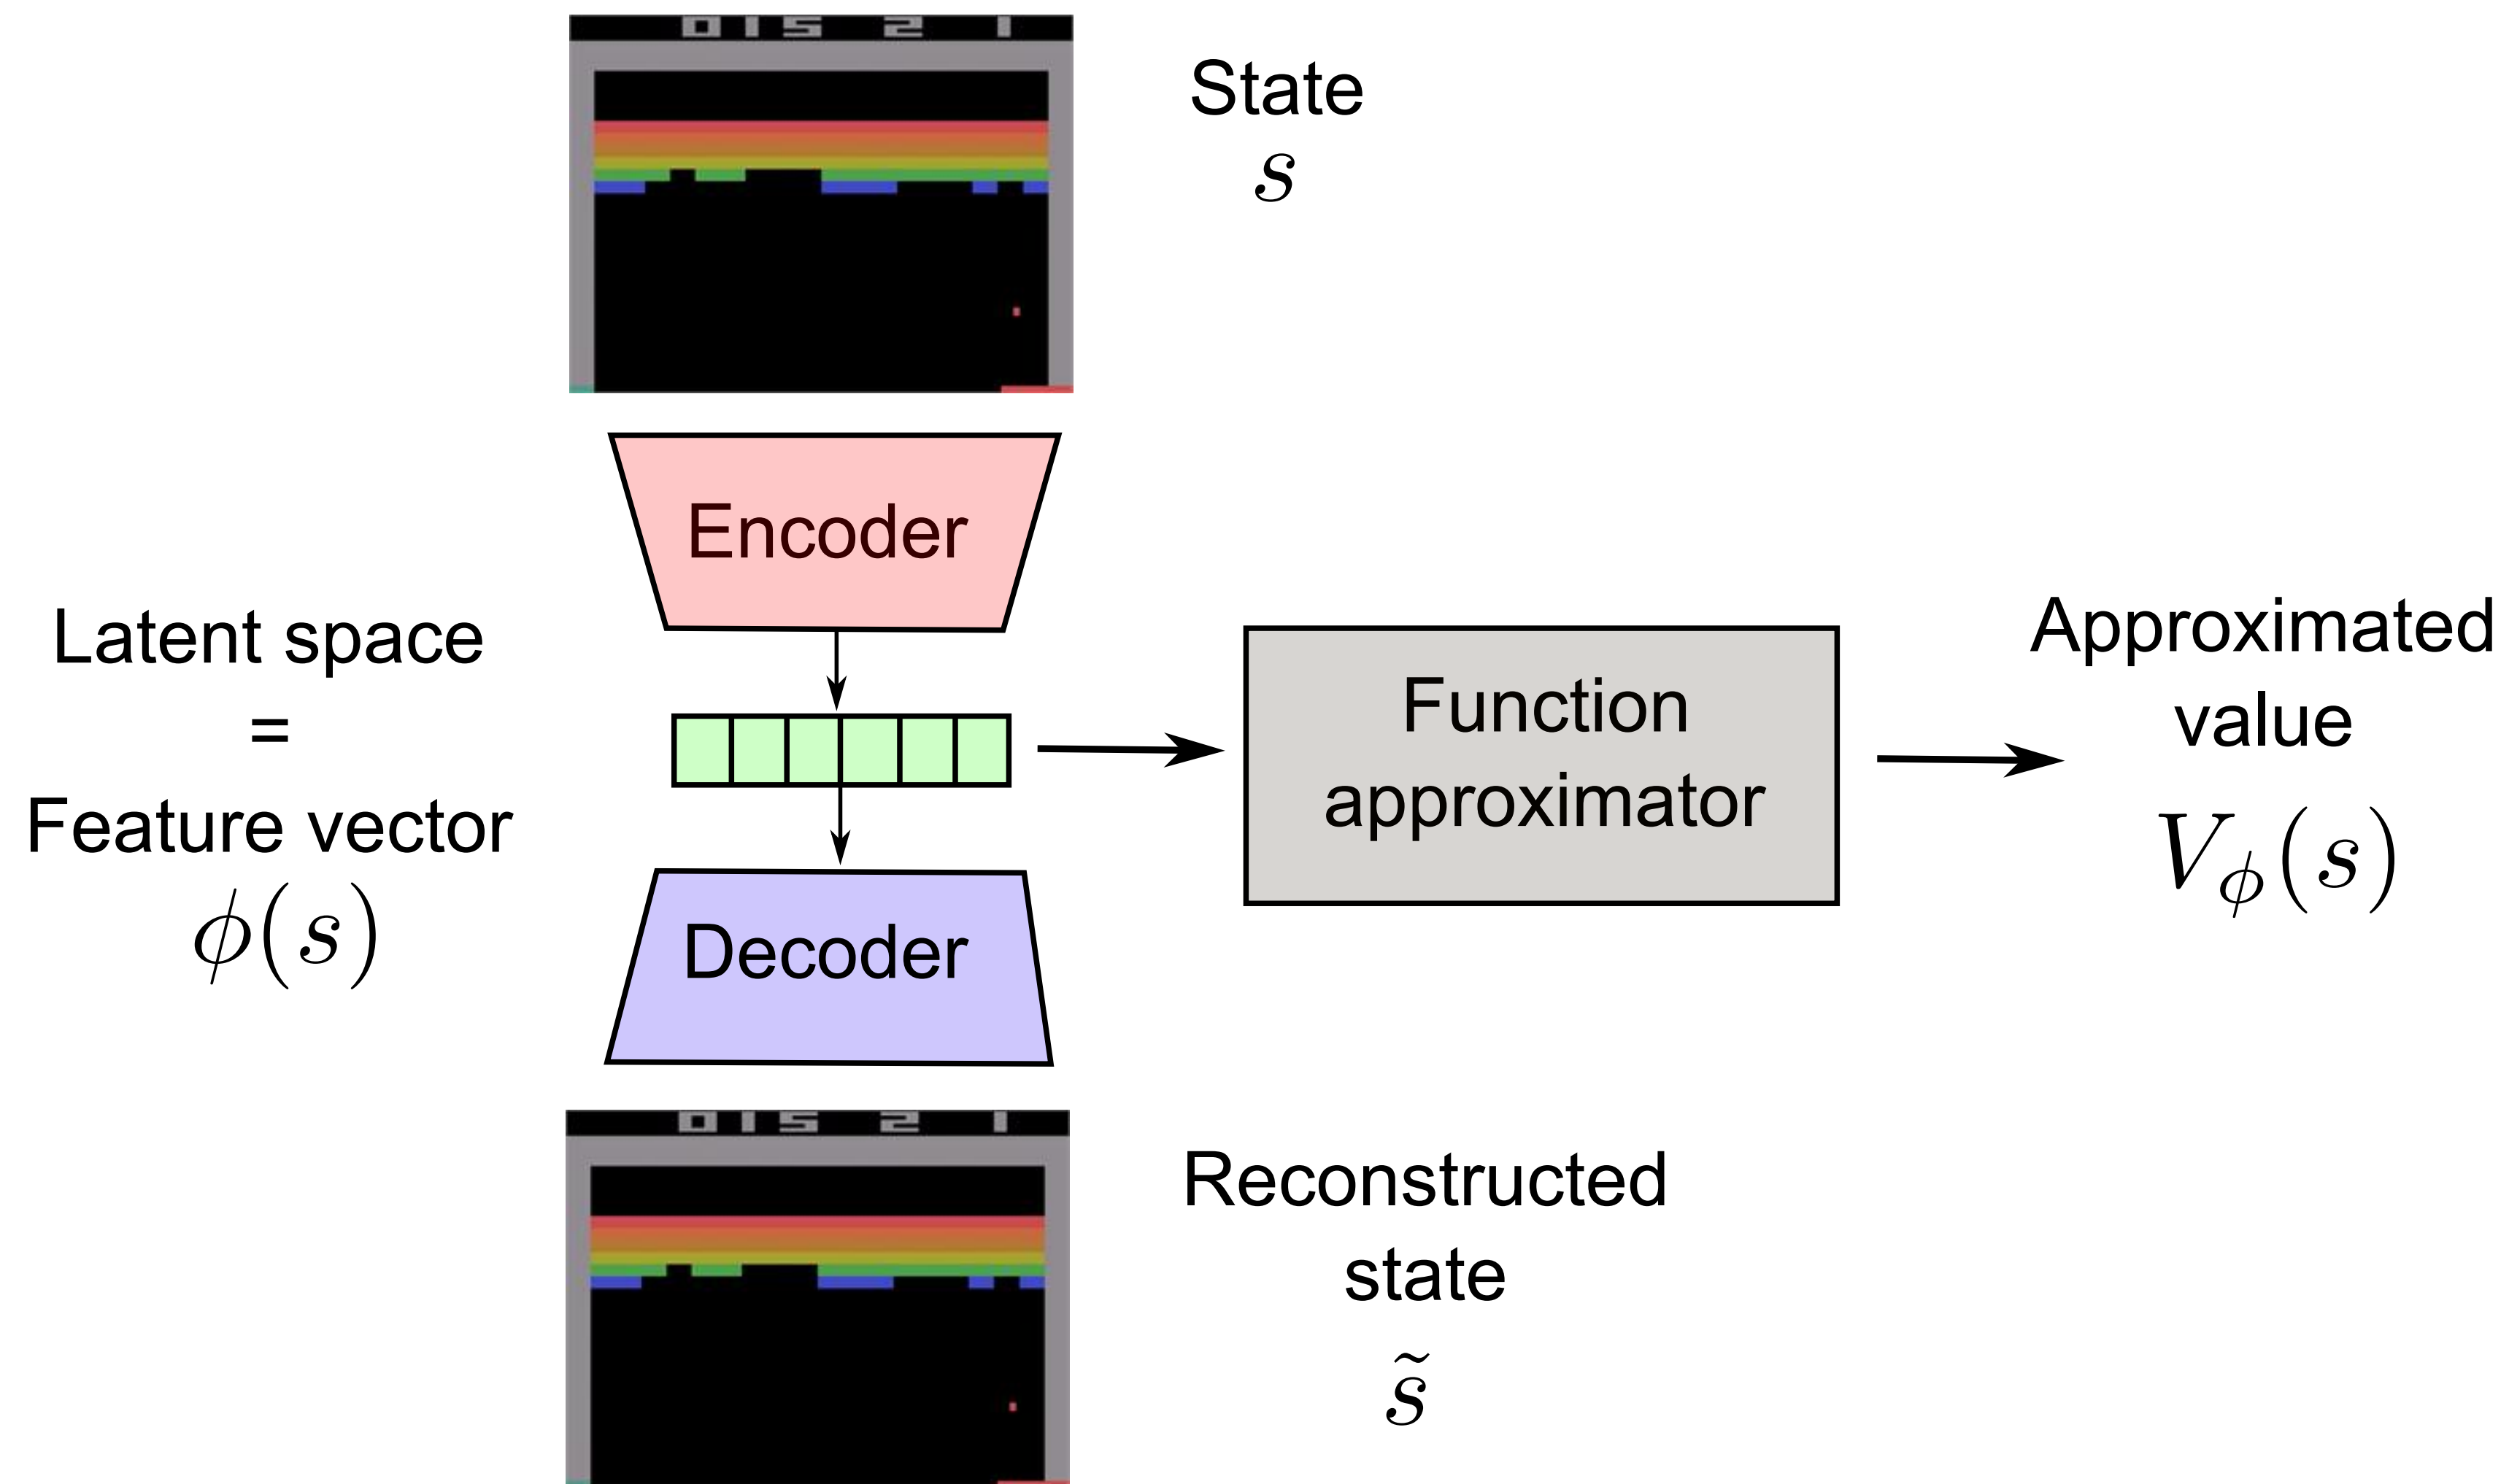
Autoencoders

- The **latent space z** is a **compressed representation** (bottleneck) of the inputs **x**.

- It has to learn to compress efficiently the inputs without losing too much information, in order to reconstruct the inputs.

  - Dimensionality reduction.

  - Unsupervised feature extraction.

Source: https://lilianweng.github.io/lil-log/2018/08/12/from-autoencoder-to-beta-vae.html

Autoencoders in deep RL

- In deep RL, we can construct the feature vector with an autoencoder.

- The autoencoder can be trained offline with a random agent or online with the current policy (auxiliary loss).

State
s

Encoder

Latent space
=
Feature vector
$\phi(s)$

Function approximator

Approximated value
$V_\phi(s)$

Decoder

Reconstructed state
$\tilde{s}$

4 - Recurrent neural networks

Recurrent neural networks

- FCN, CNN and AE are **feedforward neural networks**: they transform an input $\mathbf{x}$ into an output $\mathbf{y}$:

$$\mathbf{y} = F_\theta(\mathbf{x})$$

- If you present a sequence of inputs $\mathbf{x}_0, \mathbf{x}_1, \ldots, \mathbf{x}_t$ to a feedforward network, the outputs will be independent from each other:

$$\mathbf{y}_0 = F_\theta(\mathbf{x}_0)$$

$$\mathbf{y}_1 = F_\theta(\mathbf{x}_1)$$

$$\ldots$$

$$\mathbf{y}_t = F_\theta(\mathbf{x}_t)$$

- The output $\mathbf{y}_t$ does **not** depend on the history of inputs $\mathbf{x}_0, \mathbf{x}_1, \ldots, \mathbf{x}_{t-1}$.

Recurrent neural networks

- This not always what you want.

- If your inputs are frames of a video, the correct response at time t might also depend on previous frames.

Source: https://srirangatarun.wordpress.com/2018/07/09/video-frame-prediction-with-keras/

- The task of the NN could be to explain what happens at each frame.

- As we saw, a single frame is often not enough to predict the future (**Markov property**).

# Recurrent neural networks

Source: C. Olah

- A **recurrent neural network** (RNN) uses it previous output as an additional input (*context*).

- All vectors have a time index t denoting the time at which this vector was computed.

- The input vector at time t is $\mathbf{x}_t$, the output vector is $\mathbf{h}_t$:

$$\mathbf{h}_t = f(W_x \times \mathbf{x}_t + W_h \times \mathbf{h}_{t-1} + \mathbf{b})$$

- The input $\mathbf{x}_t$ and previous output $\mathbf{h}_{t-1}$ are multiplied by **learnable weights**:

  - W_x is the input weight matrix.

  - W_h is the recurrent weight matrix.

Recurrent neural networks
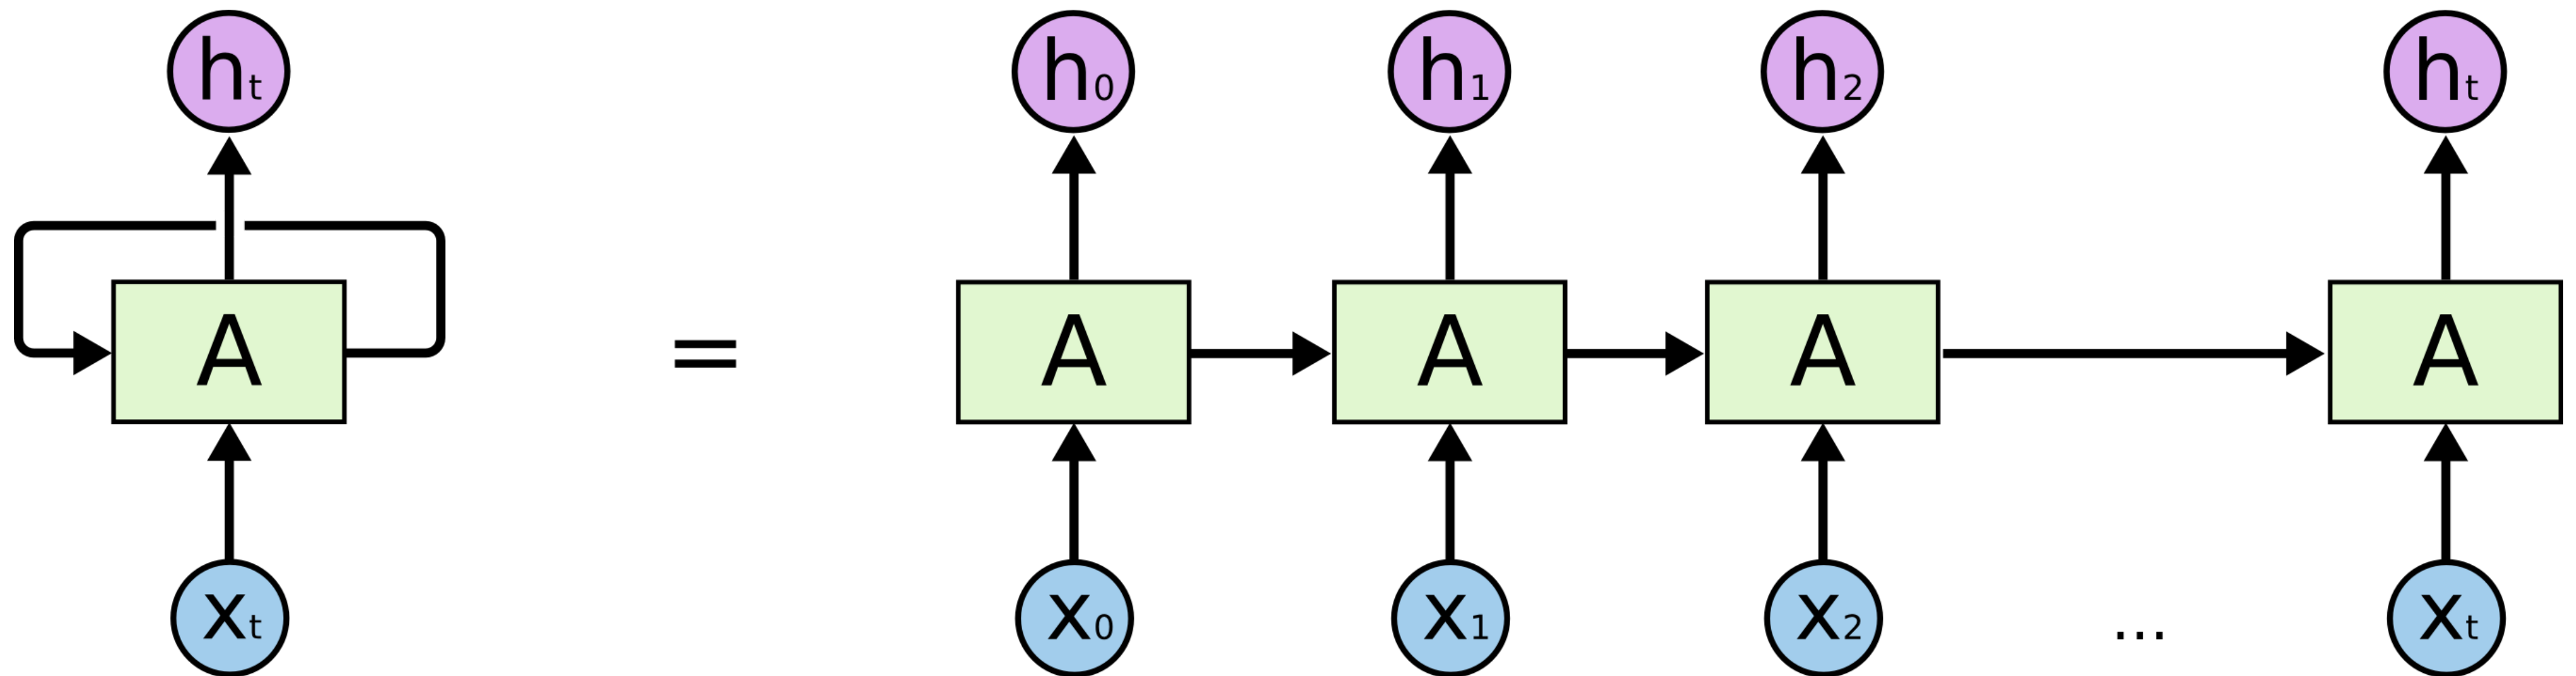
Source: C. Olah

- This is equivalent to a deep neural network taking the whole history $\mathbf{x}_0, \mathbf{x}_1, \ldots, \mathbf{x}_t$ as inputs, but reusing weights between two time steps.

- The weights are trainable using **backpropagation through time** (BPTT).

- A RNN can learn the **temporal dependencies** between inputs.

LSTM cell

- A popular variant of RNN is **LSTM** (long short-term memory).

- In addition to the input $\mathbf{x}_t$ and output $\mathbf{h}_t$, it also has a **state** (or **memory** or **context**) $\mathbf{C}_t$ which is maintained over time.

- It also contains three multiplicative **gates**:

  - The **input gate** controls which inputs should enter the memory.

  - The **forget gate** controls which memory should be forgotten.

  - The **output gate** controls which part of the memory should be used to produce the output.

Source: C. Olah

RNN in RL

Source: https://deepmind.com/blog/article/capture-the-flag-science

- An obvious use case of RNNs in deep RL is for POMDP (partially observable MDP).
- If the individual states s_t do not have the Markov property, the output of a LSTM does:
  - The output of the RNN is a representation of the complete history $s_0, s_1, \ldots, s_t$.
- We can apply RL on the output of a RNN and solve POMDPs for free!