

UNIVERSITY OF TECHNOLOGY
IN THE EUROPEAN CAPITAL OF CULTURE
CHEMNITZ

Deep Reinforcement Learning

Learned world models

Julien Vitay

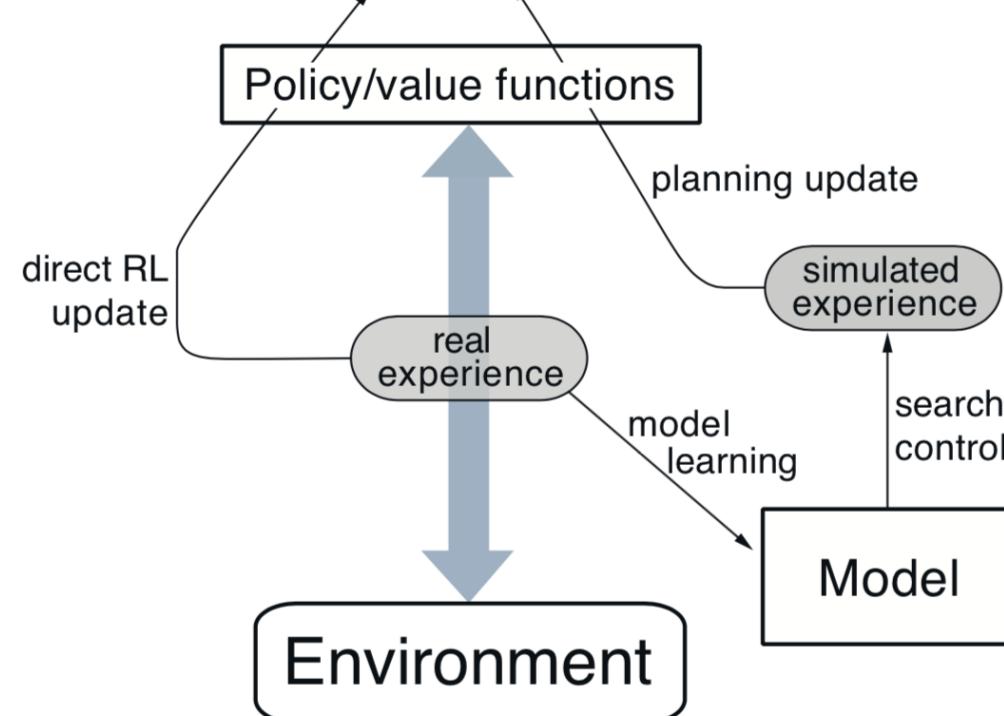
Professur für Künstliche Intelligenz - Fakultät für Informatik

<https://tu-chemnitz.de/informatik/KI/edu/deeprl>

Model-based RL algorithms with learned models

Model-based augmented model-free (MBMF)

- Dyna-Q: the model **generates** imaginary transitions/rollouts that are used to train a model-free algorithm.




Source: <https://towardsdatascience.com/reinforcement-learning-model-based-planning-methods-5e99cae0abb8>

- NAF: Normalized advantage functions (Gu et al., 2016)
- I2A: Imagination-augmented agents (Weber et al., 2017)
- MBVE: model-based value estimation (Feinberg et al., 2018)

Model-based planning

- MPC: the learned model is used to **plan** actions that maximize the RL objective.



Source: <https://arxiv.org/abs/1901.03737>

- TDM: Temporal difference models (Pong et al., 2018)
- World models (Ha and Schmidhuber, 2018)
- PlaNet (Hafner et al., 2019)
- Dreamer (Hafner et al., 2020)

1 - I2A - Imagination-augmented agents

Imagination-Augmented Agents for Deep Reinforcement Learning

Théophane Weber* Sébastien Racanière* David P. Reichert* Lars Buesing
Arthur Guez Danilo Rezende Adria Puigdomènech Badia Oriol Vinyals
Nicolas Heess Yujia Li Razvan Pascanu Peter Battaglia
Demis Hassabis David Silver Daan Wierstra
DeepMind

<https://deepmind.com/blog/article/agents-imagine-and-plan>

I2A - Imagination-augmented agents

- I2A is a **model-based augmented model-free method**: it trains a MF algorithm (A3C) with the help of **rollouts** generated by a MB model.





Figure 3: *Random examples of procedurally generated Sokoban levels.* The player (green sprite) needs to push all 4 boxes onto the red target squares to solve a level, while avoiding irreversible mistakes. Our agents receive sprite graphics (shown above) as observations.

- They showcase their algorithm on the puzzle environment **Sokoban**, where you need to move boxes to specified locations.
- Sokoban is a quite hard game, as actions are irreversible (you can get stuck) and the solution requires many actions (sparse rewards).
- MF methods are bad at this game as they learn through trials-and-(many)-errors.

Sokoban



I2A - Imagination-augmented agents

- The **model** learns to predict the next frame and the next reward based on the four last frames and the chosen action.

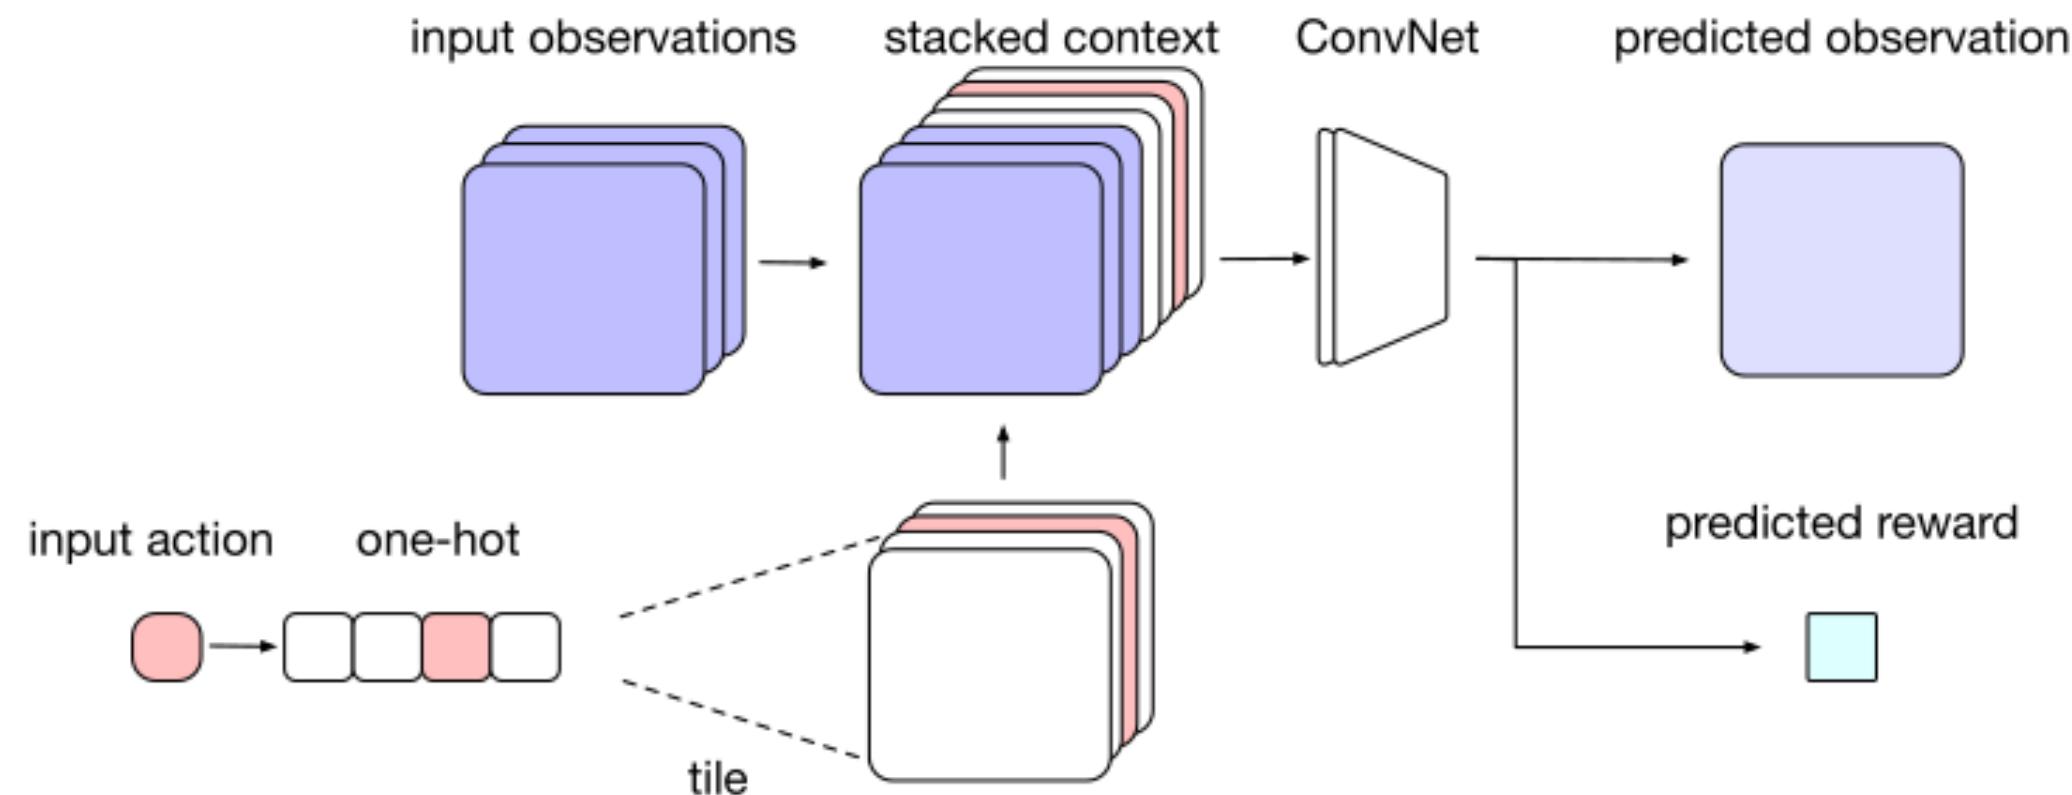
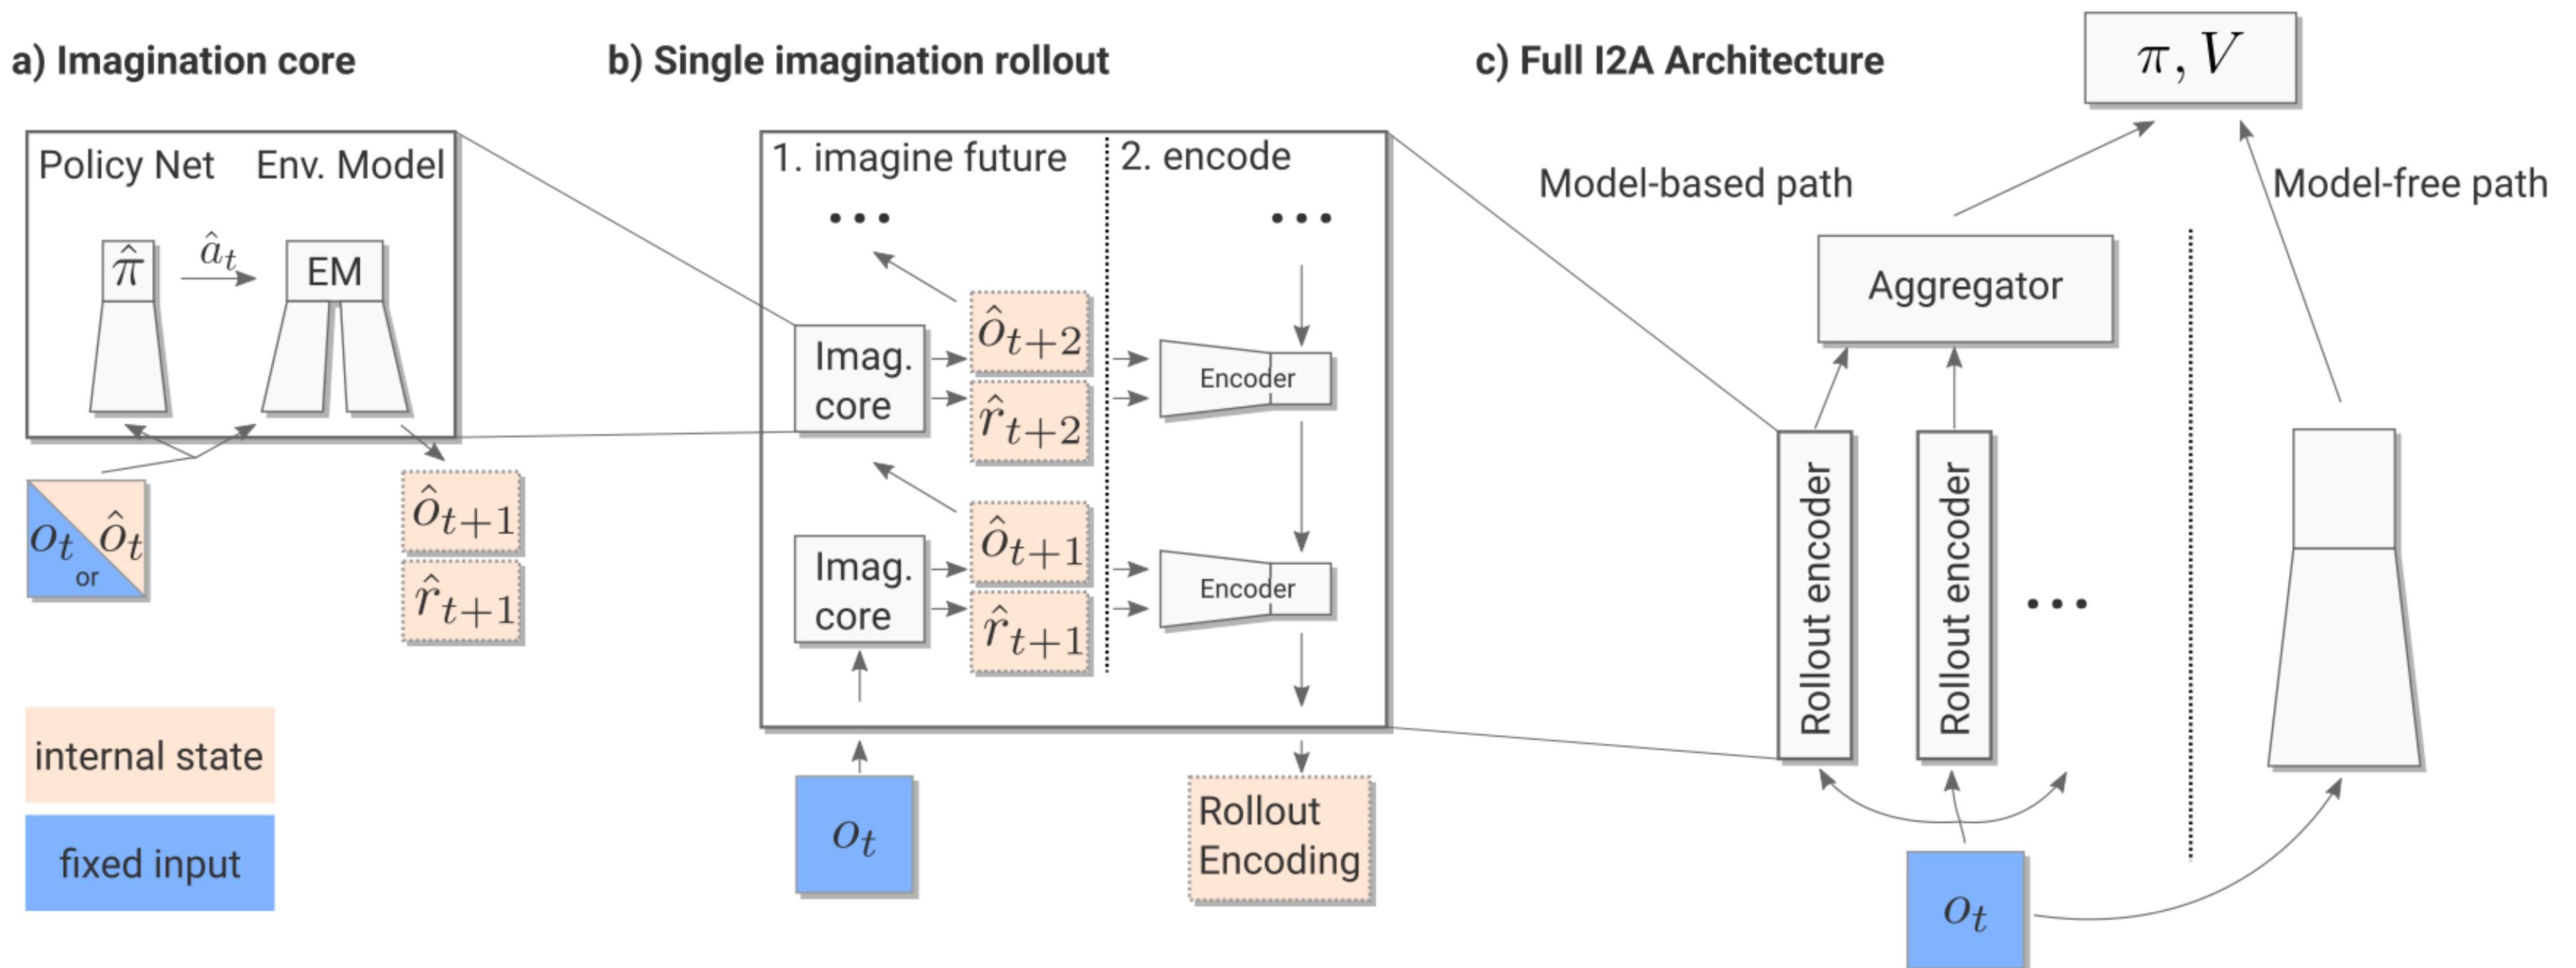


Figure 2: *Environment model*. The input action is broadcast and concatenated to the observation. A convolutional network transforms this into a pixel-wise probability distribution for the output image, and a distribution for the reward.


- It is a **convolutional autoencoder**, taking additionally an action a as input and predicting the next reward.
- It can be pretrained using a random policy, and later fine-tuned during training.

I2A - Imagination-augmented agents



I2A - Imagination-augmented agents


a) Imagination core



- The **imagination core** is composed of the environment model $M(s, a)$ and a **rollout policy** $\hat{\pi}$.
- As Sokoban is a POMDP (partially observable), the notation uses **observation** o_t instead of states s_t , but it does not really matter here.
- The **rollout policy** $\hat{\pi}$ is a simple and fast policy. It does not have to be the trained policy π .
- It could even be a random policy, or a pretrained policy using for example A3C directly.
- In I2A, it is a **distilled policy** from the trained policy π (see later).
- Take home message: given the current observation o_t and a policy $\hat{\pi}$, we can predict the next observation \hat{o}_{t+1} and the next reward \hat{r}_{t+1} .


I2A - Imagination-augmented agents

b) Single imagination rollout



- The **imagination rollout module** uses the imagination core to predict iteratively the next τ frames and rewards using the current frame o_t and the rollout policy:
$$o_t \rightarrow \hat{o}_{t+1} \rightarrow \hat{o}_{t+2} \rightarrow \dots \rightarrow \hat{o}_{t+\tau}$$
- The τ frames and rewards are passed **backwards** to a convolutional LSTM (from $t + \tau$ to t) which produces an embedding / encoding of the rollout.
- The output of the imagination rollout module is a vector e_i (the final state of the LSTM) representing the whole rollout, including the (virtually) obtained rewards.
- Note that because of the stochasticity of the rollout policy $\hat{\pi}$, different rollouts can lead to different encoding vectors.

I2A - Imagination-augmented agents



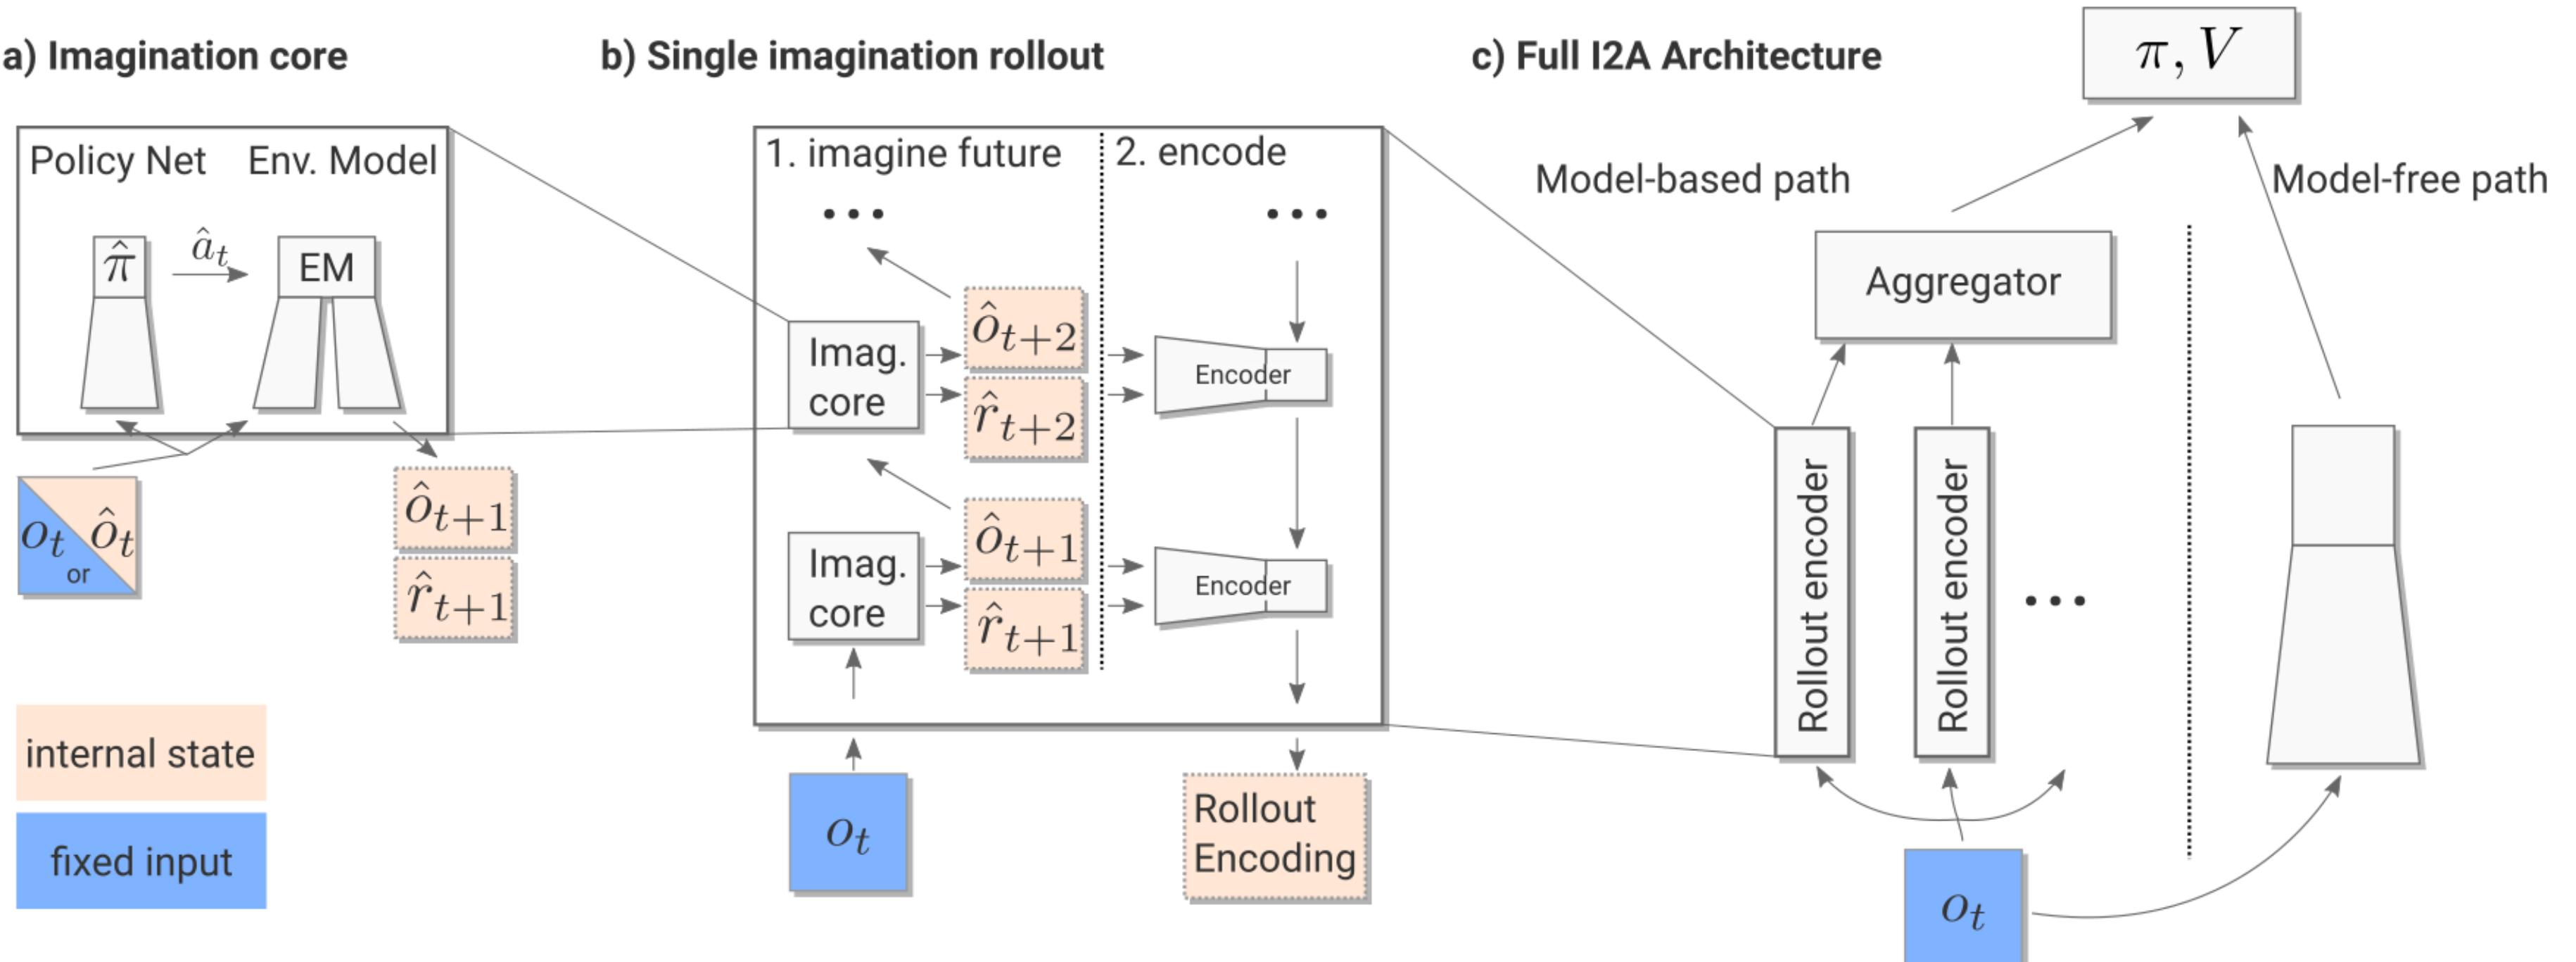
- For the current observation o_t , we then generate one **rollout** per possible action (5 in Sokoban):
 - What would happen if I do action 1?
 - What would happen if I do action 2?
 - etc.
- The resulting vectors are concatenated to the output of **model-free** path (a convolutional neural network taking the current observation as input).
- Altogether, we have a huge NN with weights θ (model, encoder, MF path) producing an input s_t to the **A3C** module.
- We can then learn the policy π and value function V based on this input to maximize the returns:

$$\nabla_{\theta} \mathcal{J}(\theta) = \mathbb{E}_{s_t \sim \rho_{\theta}, a_t \sim \pi_{\theta}} [\nabla_{\theta} \log \pi_{\theta}(s_t, a_t) (\sum_{k=0}^{n-1} \gamma^k r_{t+k+1} + \gamma^n V_{\varphi}(s_{t+n}) - V_{\varphi}(s_t))]$$

$$\mathcal{L}(\varphi) = \mathbb{E}_{s_t \sim \rho_{\theta}, a_t \sim \pi_{\theta}} [(\sum_{k=0}^{n-1} \gamma^k r_{t+k+1} + \gamma^n V_{\varphi}(s_{t+n}) - V_{\varphi}(s_t))^2]$$

I2A - Imagination-augmented agents

- The complete architecture may seem complex, but everything is differentiable so we can apply backpropagation and train the network **end-to-end** using multiple workers.
- It is the A3C algorithm (MF), but **augmented** by MB rollouts, i.e. with explicit information about the future.

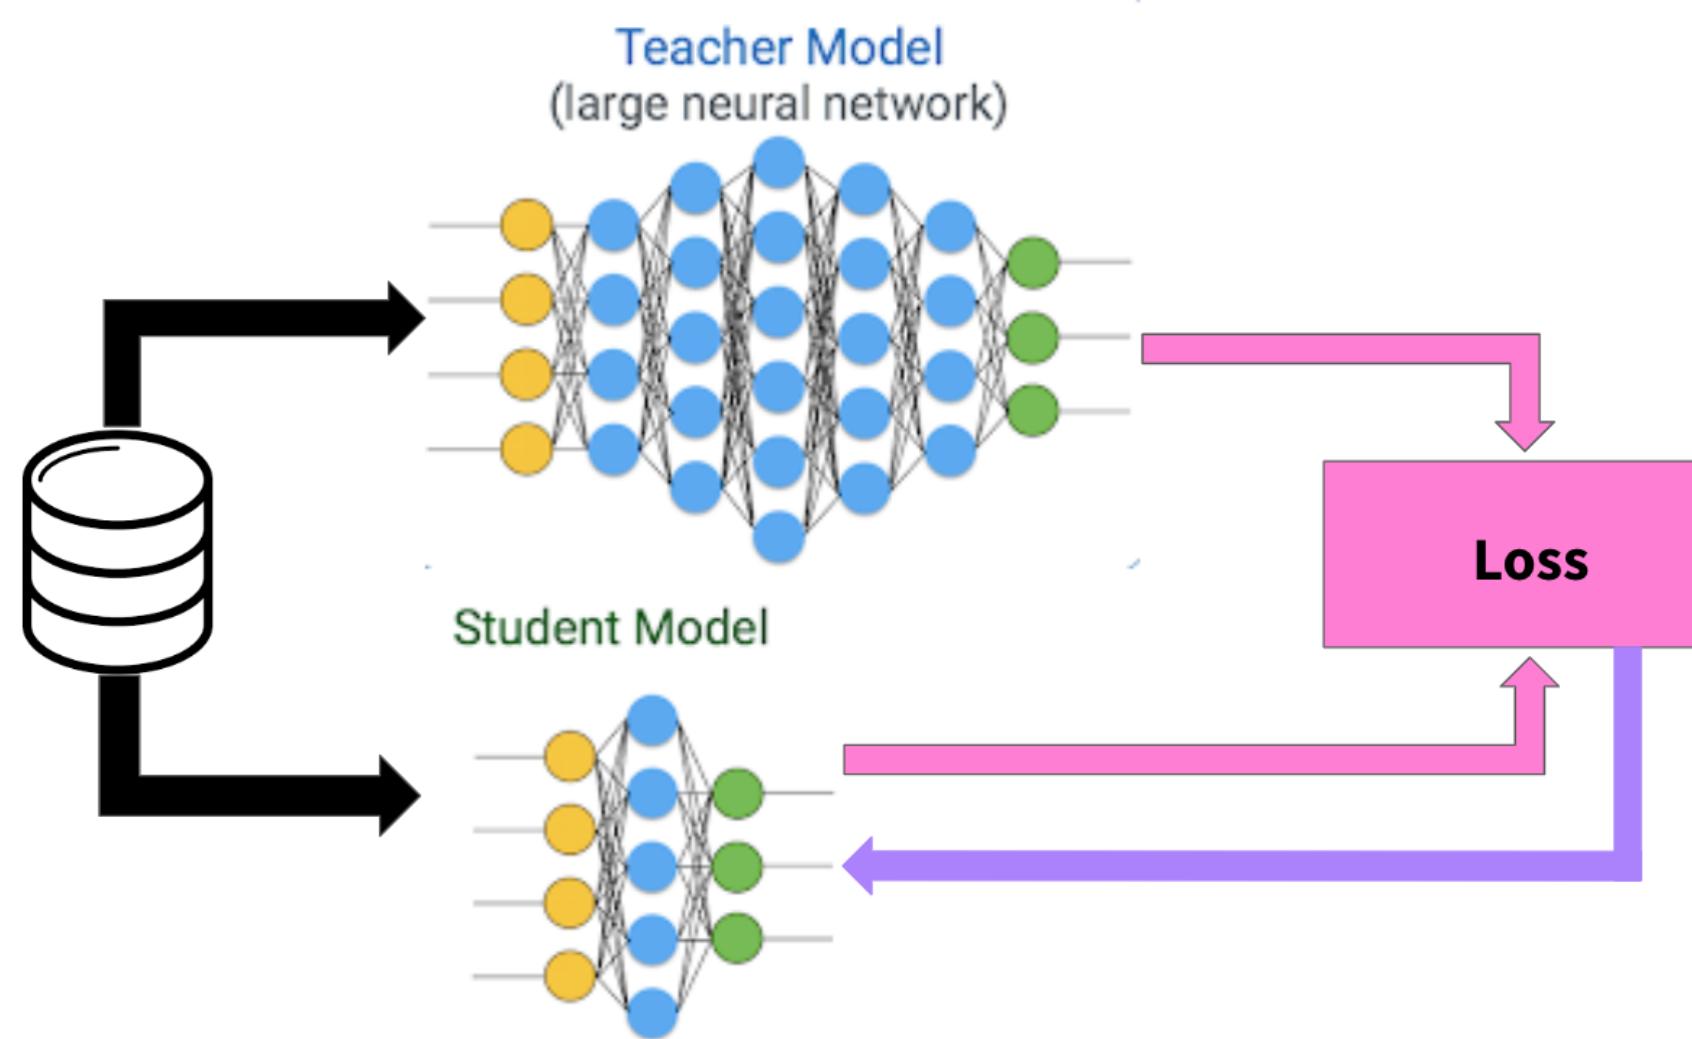


Policy distillation

- The **rollout policy** $\hat{\pi}$ is trained using **policy distillation** of the trained policy π .
- The small rollout policy network with weights $\hat{\theta}$ tries to copy the outputs $\pi(s, a)$ of the bigger policy network (A3C).
- This is a supervised learning task: just minimize the KL divergence between the two policies:


$$\mathcal{L}(\hat{\theta}) = \mathbb{E}_{s,a}[D_{\text{KL}}(\hat{\pi}(s, a) || \pi(s, a))]$$

- As the network is smaller, it won't be as good as π , but its learning objective is easier.



Distral : distill and transfer learning

- FYI: distillation can be used to ensure generalization over different environments.
- Each learning algorithms learns its own task, but tries not to diverge too much from a **shared policy**, which turns out to be good at all tasks.



I2A - Imagination-augmented agents

- Unsurprisingly, I2A performs better than A3C on Sokoban.
- The deeper the rollout, the better.




Figure 4: *Sokoban learning curves*. *Left*: training curves of I2A and baselines. Note that I2A use additional environment observations to pretrain the environment model, see main text for discussion. *Right*: I2A training curves for various values of imagination depth.

I2A - Imagination-augmented agents

- The model does not even have to be perfect: the MF path can compensate for imperfections.

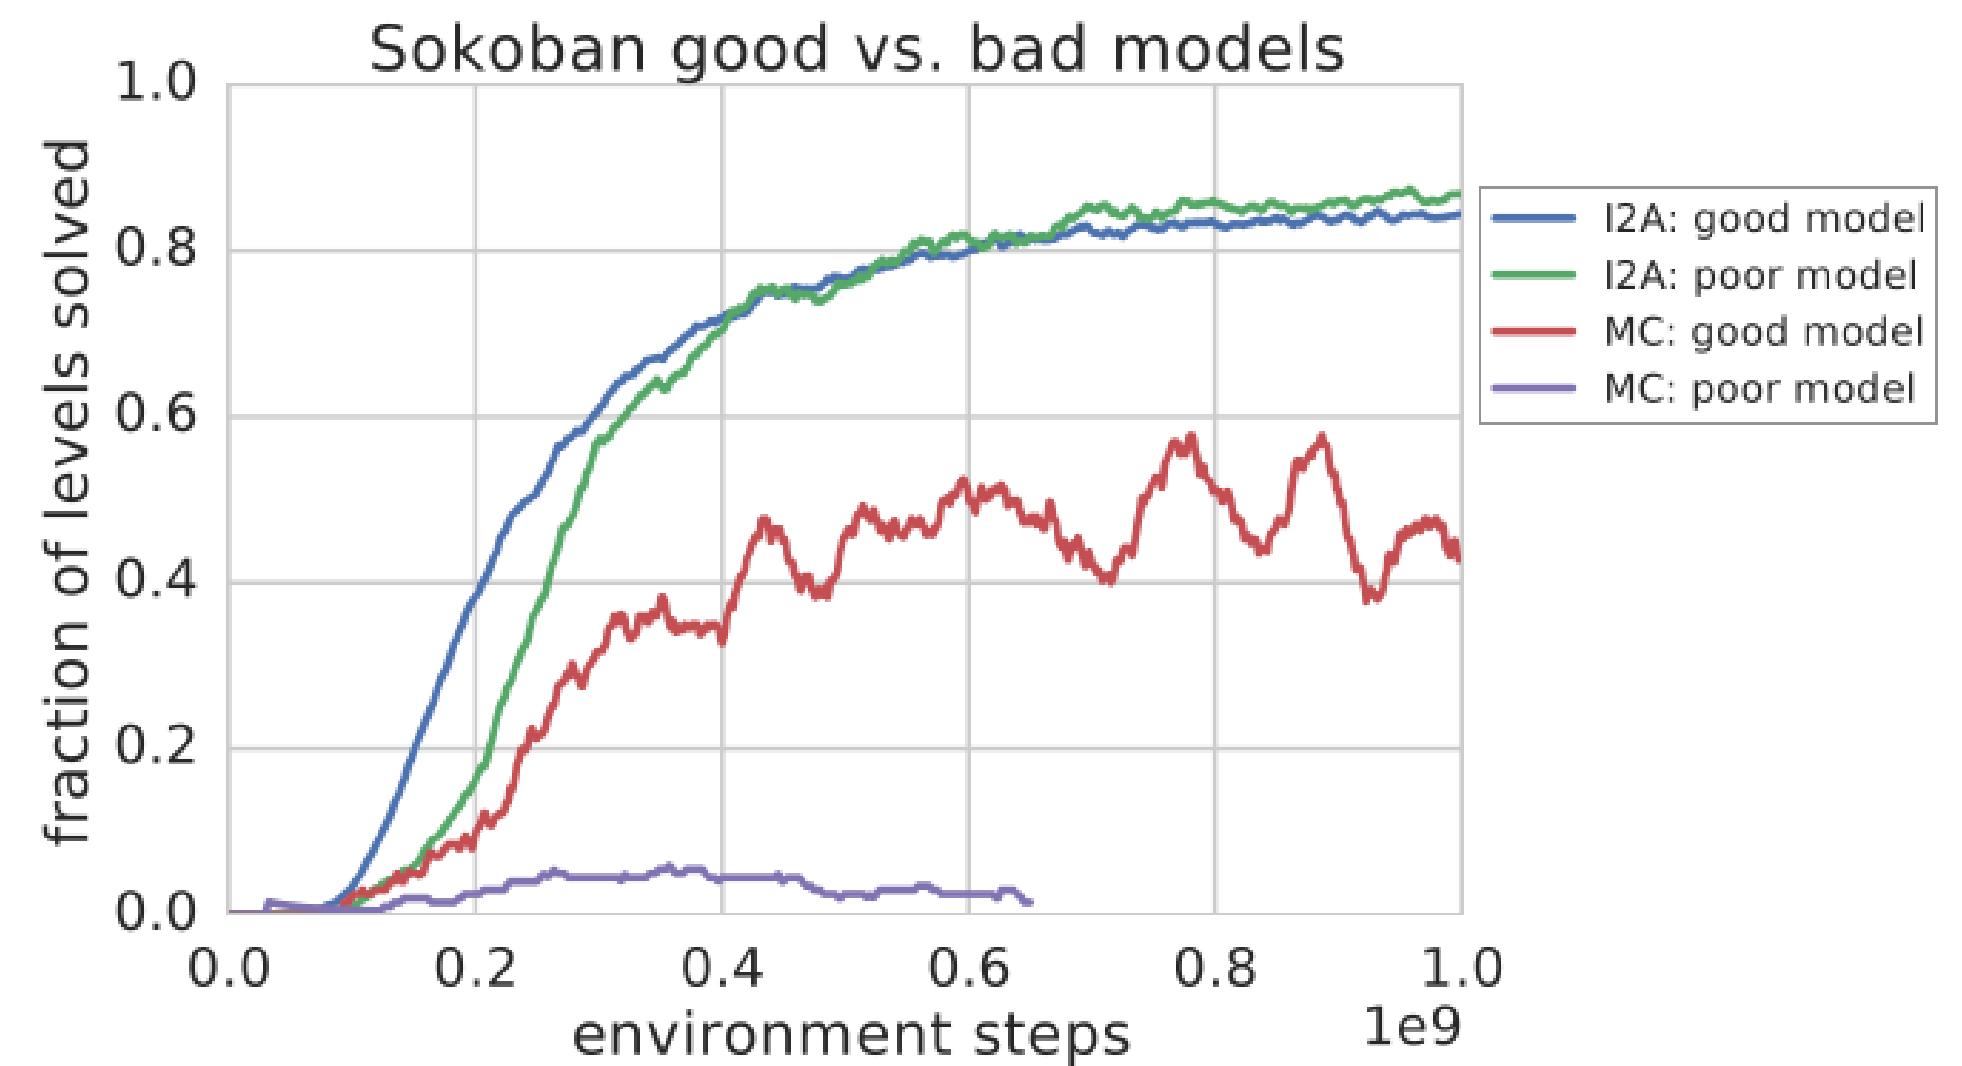
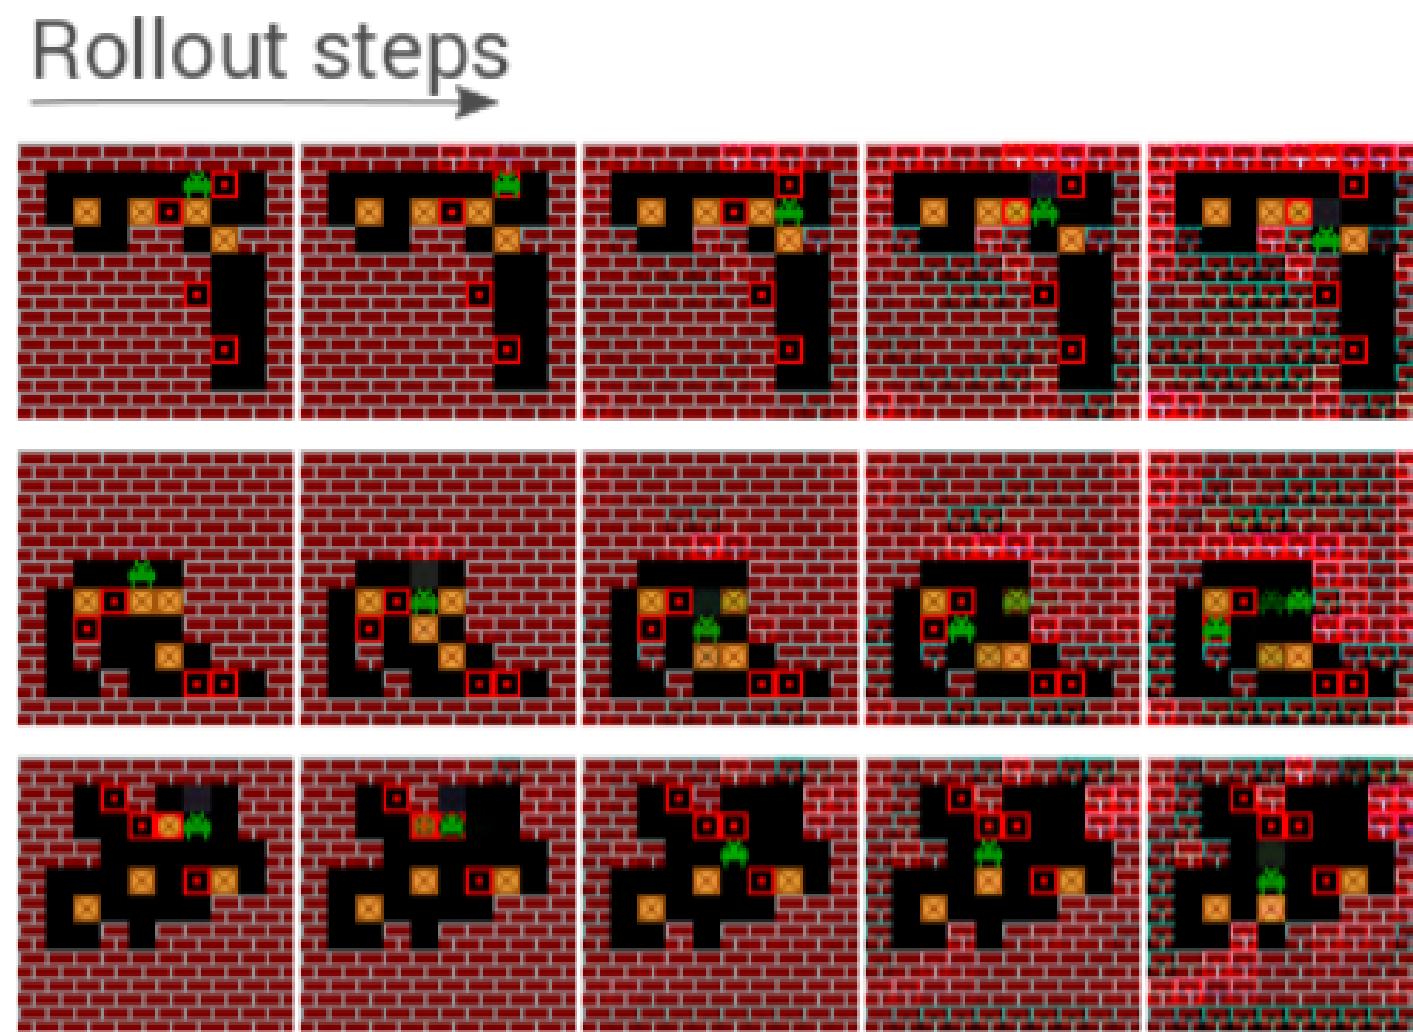
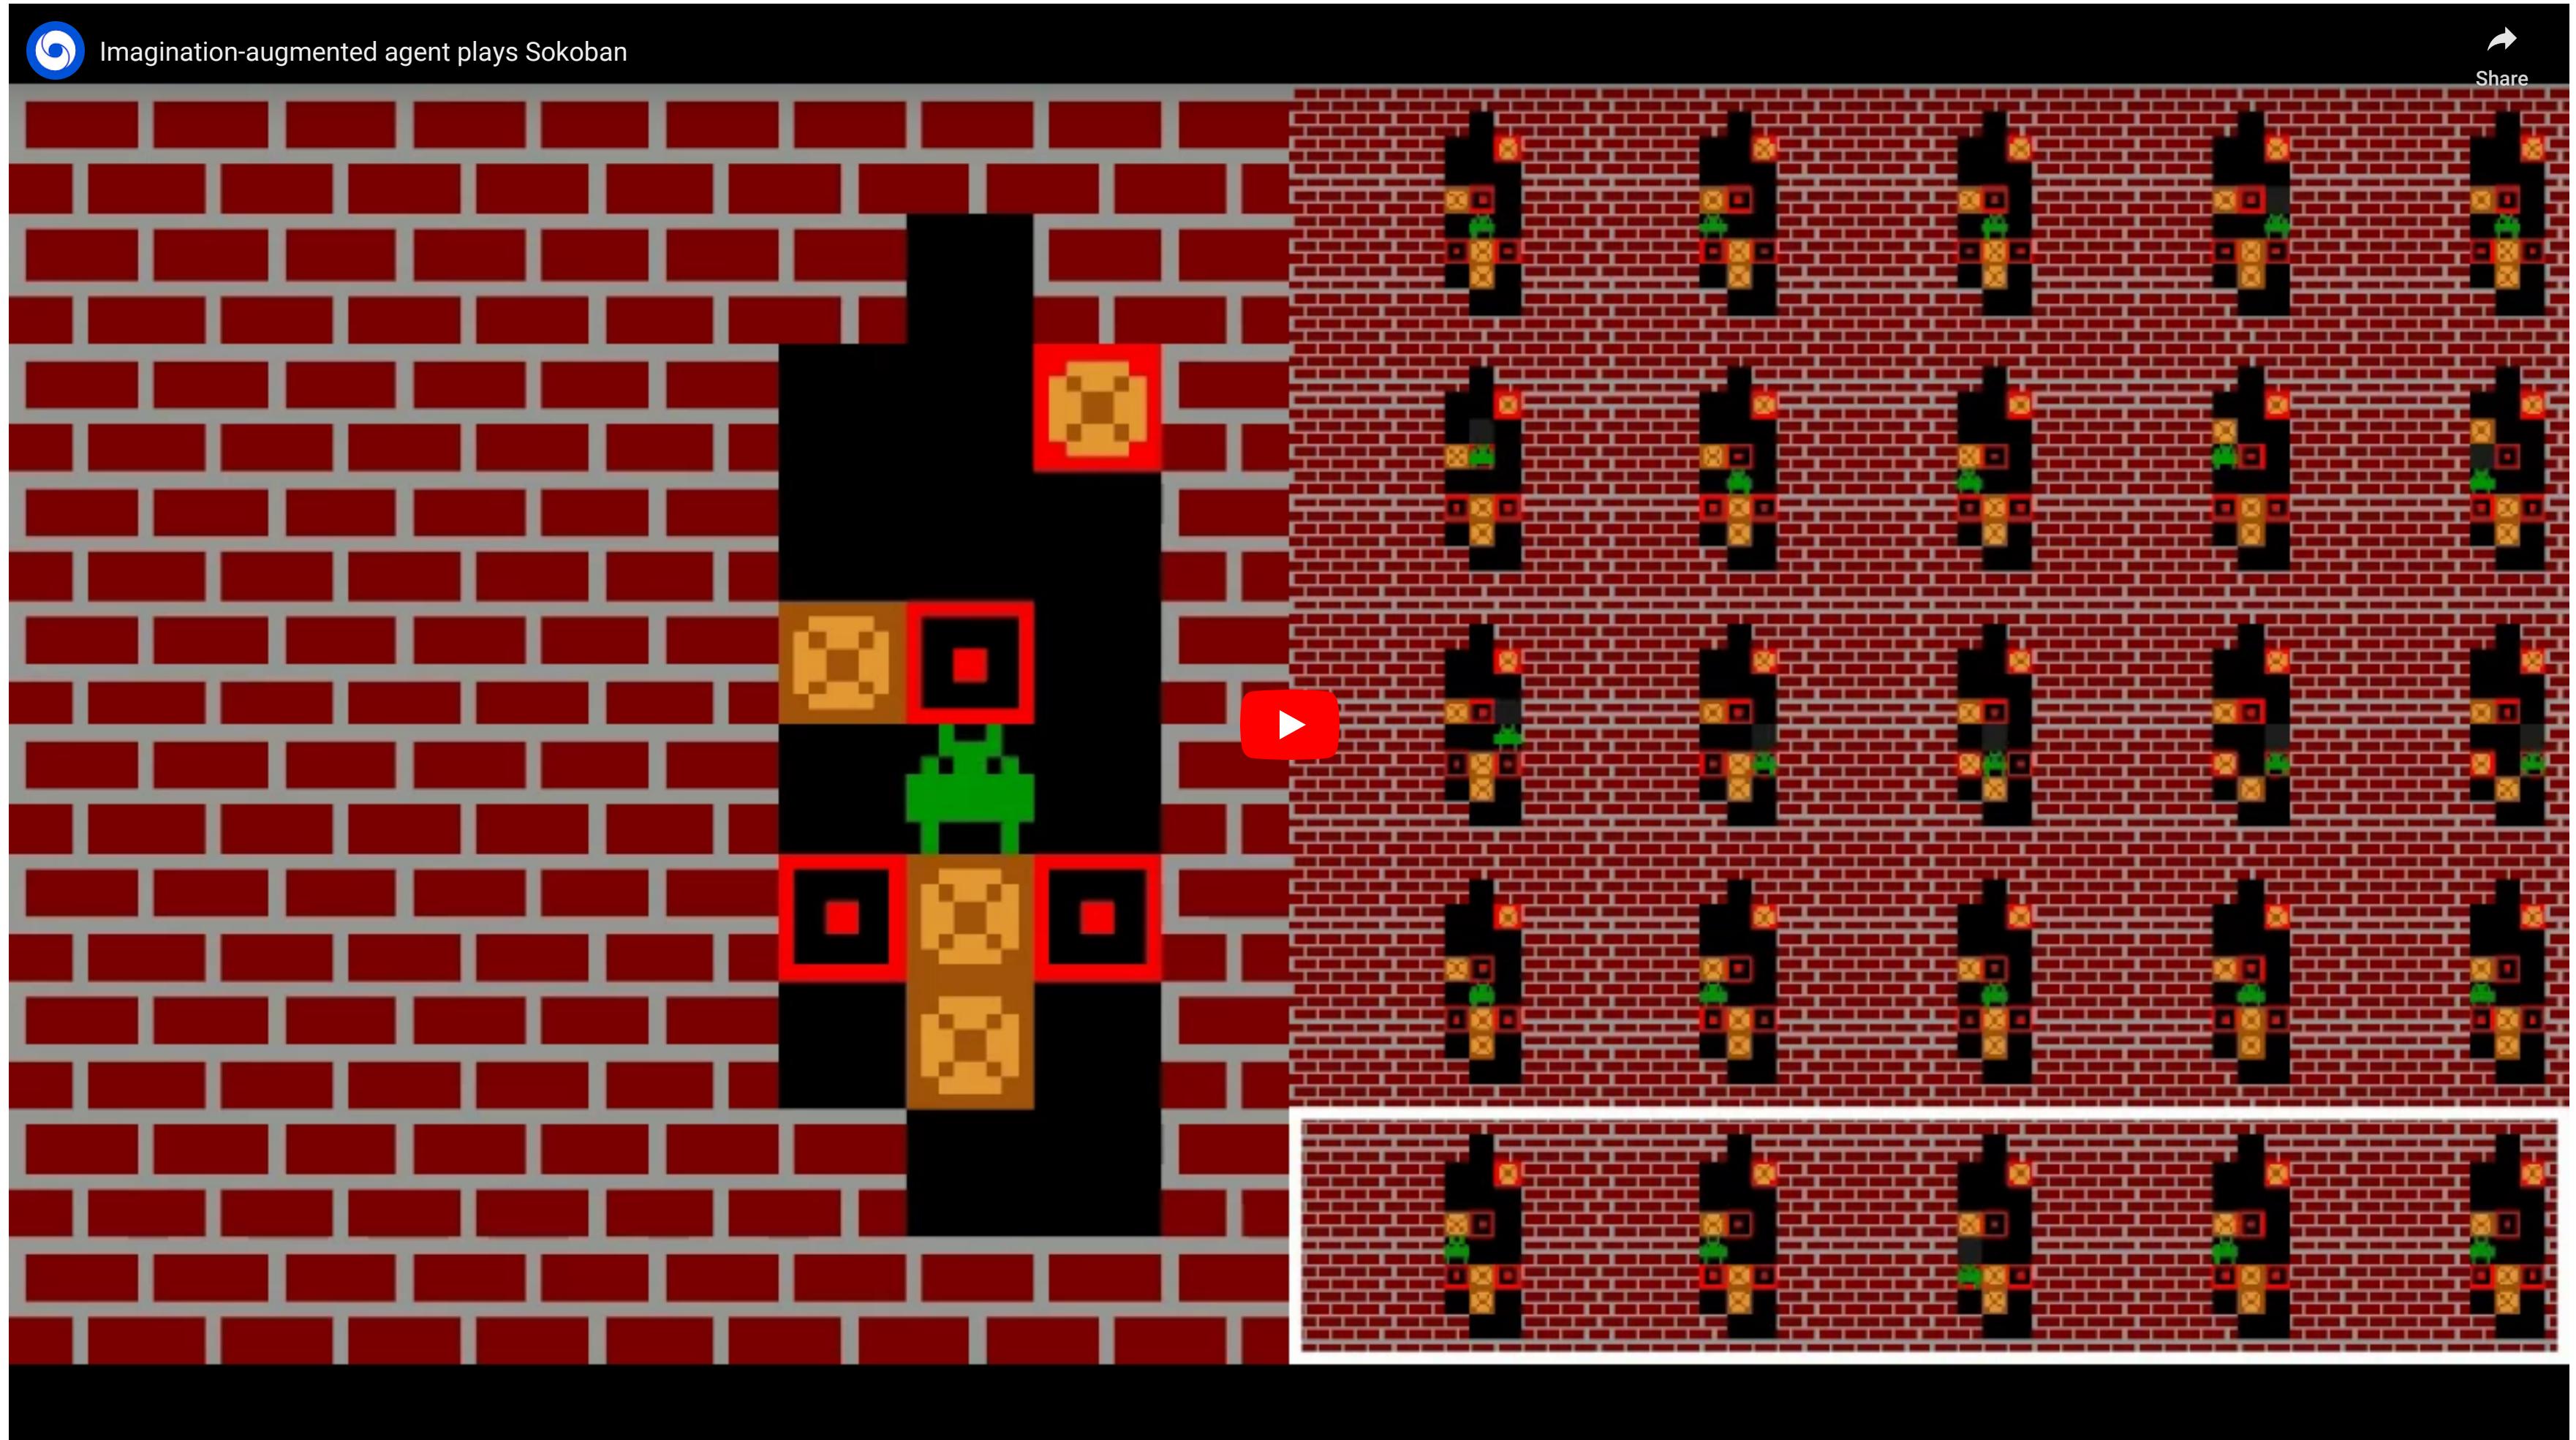


Figure 5: *Experiments with a noisy environment model.* *Left:* each row shows an example 5-step rollout after conditioning on an environment observation. Errors accumulate and lead to various artefacts, including missing or duplicate sprites. *Right:* comparison of Monte-Carlo (MC) search and I2A when using either the accurate or the noisy model for rollouts.

I2A - Sokoban



2 - Temporal difference models - TDM

TEMPORAL DIFFERENCE MODELS: MODEL-FREE DEEP RL FOR MODEL-BASED CONTROL

Vitchyr Pong*

University of California, Berkeley
vitchyr@berkeley.edu

Shixiang Gu*

University of Cambridge
Max Planck Institute
Google Brain
sg717@cam.ac.uk

Murtaza Dalal


University of California, Berkeley
mdalal@berkeley.edu

Sergey Levine

University of California, Berkeley
svlevine@eecs.berkeley.edu

TDM

- One problem with model-based planning is the **discretization time step** (difference between t and $t + 1$).
- It is determined by the action rate: how often a different action a_t has to be taken.
- In robotics, it could be below the millisecond, leading to very long trajectories in terms of steps.



- If you want to go from Berkeley to the Golden State bridge with your bike, planning over leg movements will be very expensive (long horizon).
- A solution is **multiple steps ahead planning**. Instead of learning a one-step model:

$$s_{t+1} = f_\theta(s_t, a_t)$$

one learns to predict the state achieved in T steps using the current policy:

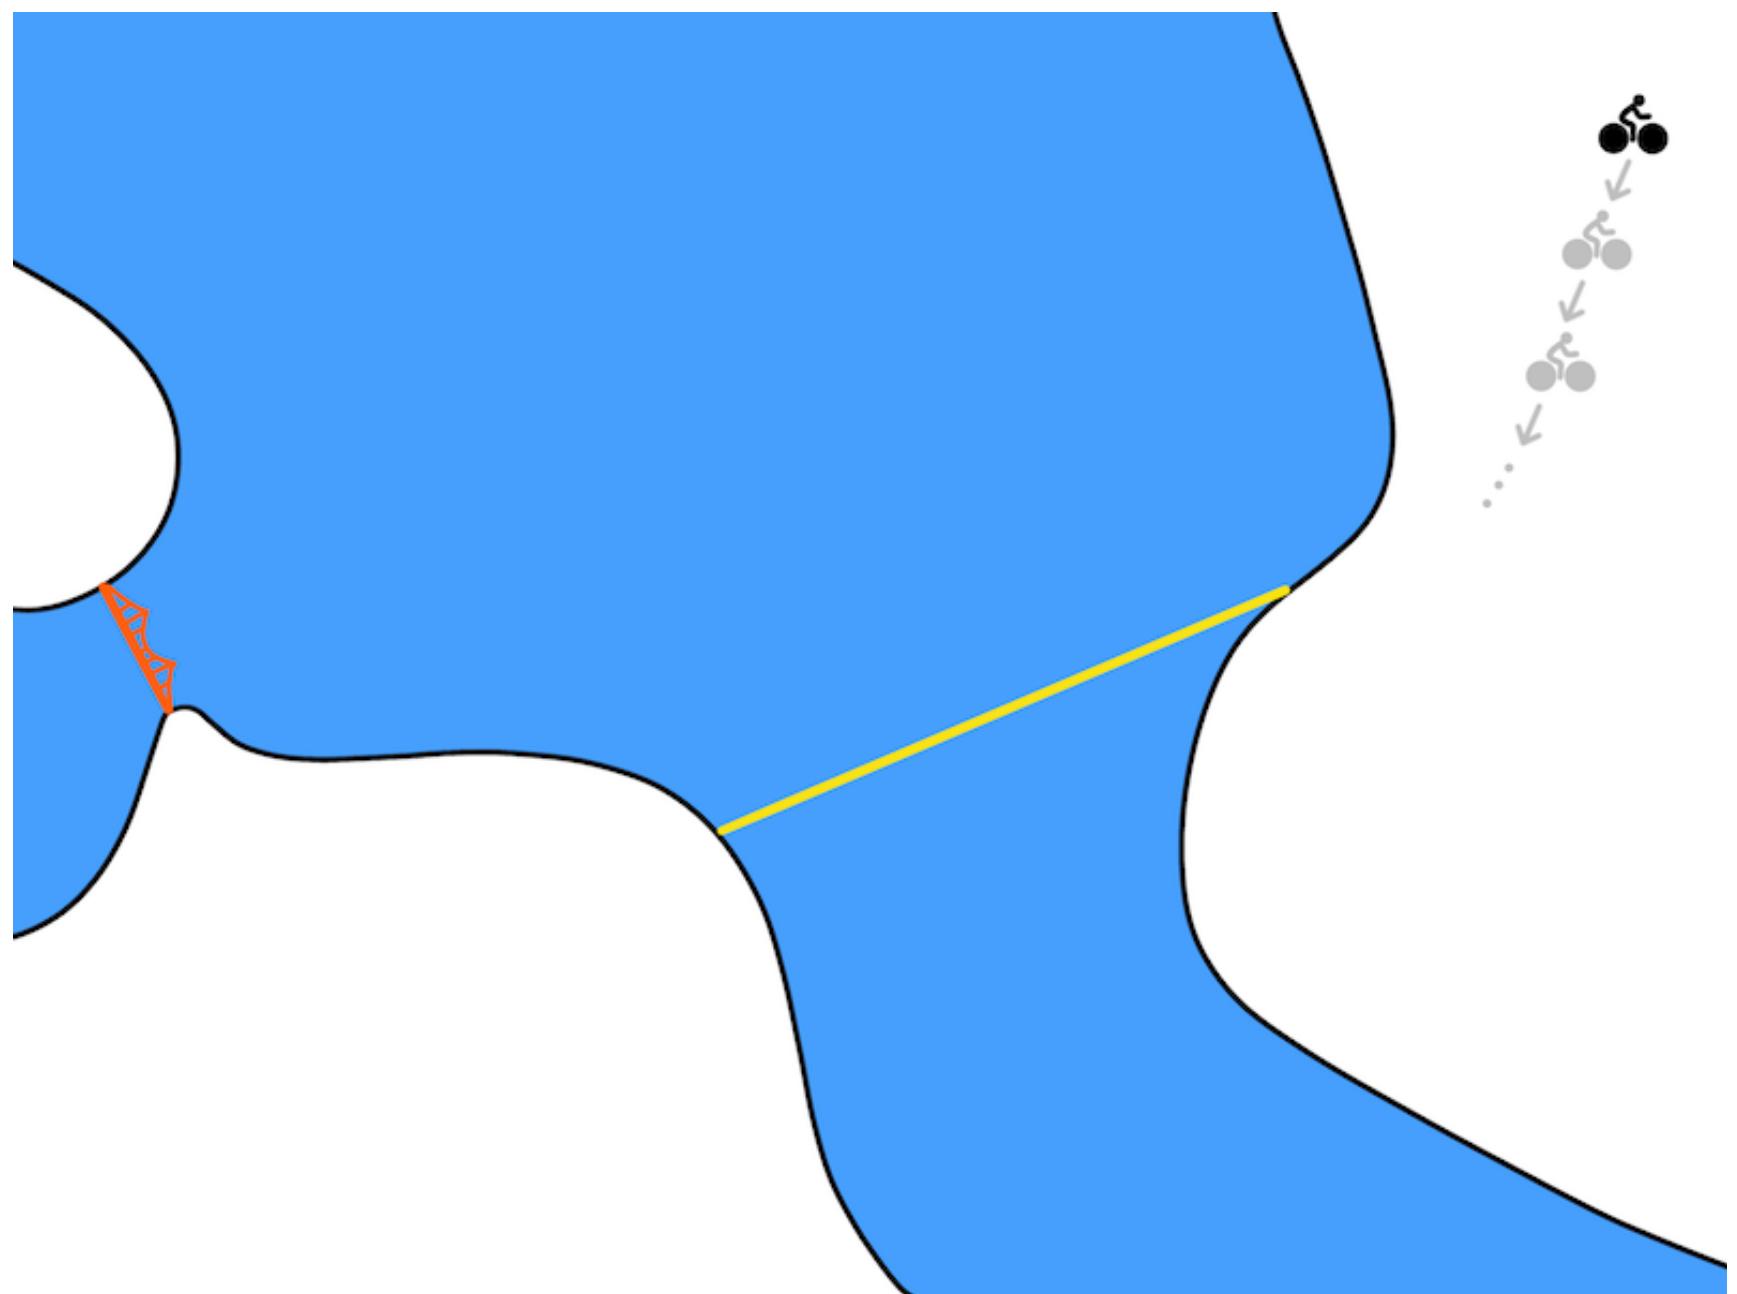
$$s_{t+T} = f_\theta(s_t, a_t, \pi)$$

Source: <https://bairblog.github.io/2018/04/26/tdm/>

- Planning and acting occur at different time scales.

TDM

- A problem with RL in general is how to define the **reward function**.



- If your goal is to travel from Berkeley to the Golden State bridge, which reward function should you use?
 - +1 at the bridge, 0 otherwise (sparse).
 - +100 at the bridge, -1 otherwise (sparse).
 - minus the distance to the bridge (dense).
- **Goal-conditioned RL** defines the reward function using the distance between the achieved state s_{t+1} and a **goal state** s_g :


$$r(s_t, a_t, s_{t+1}) = -\|s_{t+1} - s_g\|$$

Source: <https://bairblog.github.io/2018/04/26/tdm/>

- An action is good if it brings the agent closer to its goal.
- The Euclidean distance works well for the biking example (e.g. using a GPS), but the metric can be adapted to the task.

Goal-conditioned RL

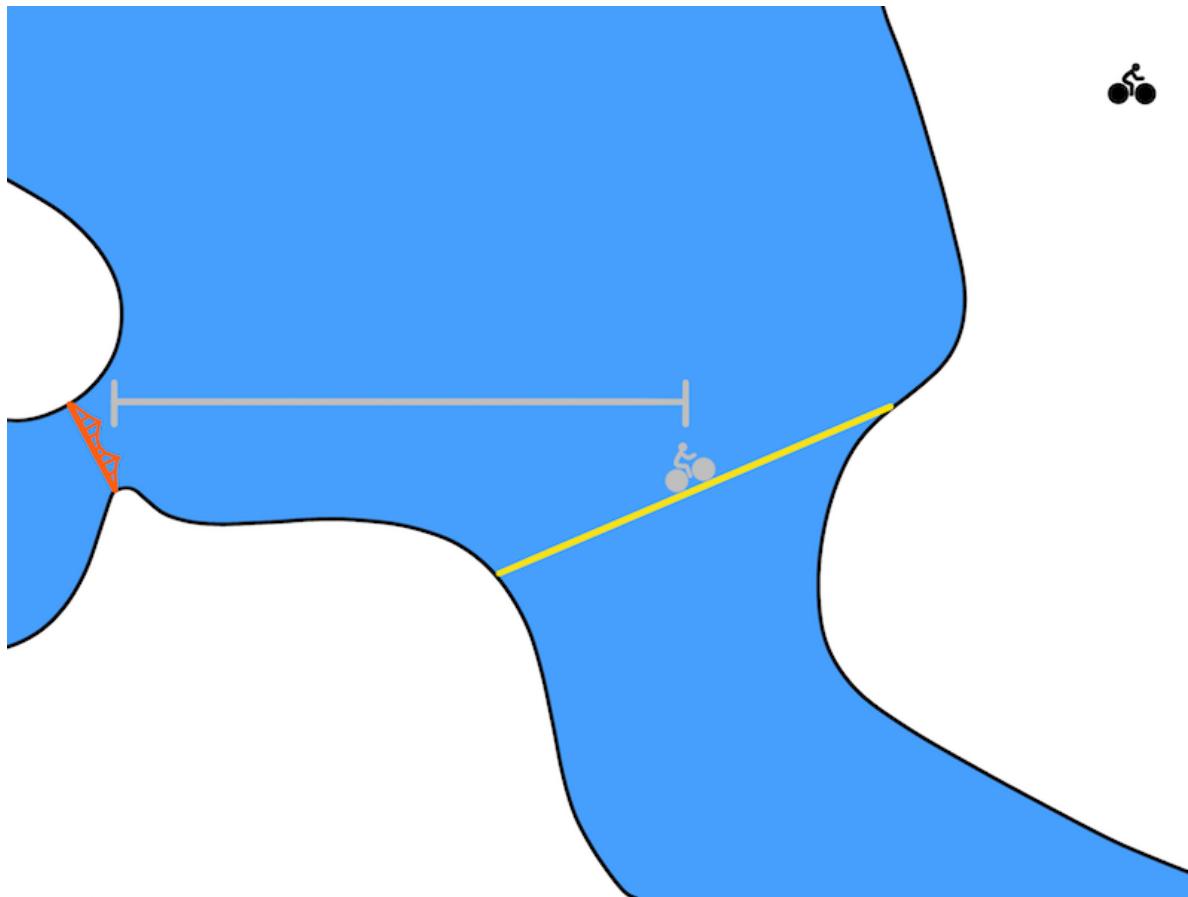
- One advantage is that you can learn multiple “tasks” at the same time with a single policy, not the only one hard-coded in the reward function.
- Another advantage is that it makes a better use of exploration by learning from mistakes: **hindsight experience replay** (HER, Andrychowicz et al., 2017).
- If your goal is to reach s_g but the agent generates a trajectory landing in $s_{g'}$, you can learn that this trajectory is good way to reach $s_{g'}$!
- In football, if you try to score a goal but end up doing a pass to a teammate, you can learn that this was a bad shot **and** a good pass.
- HER is a model-based method: you implicitly learn a model of the environment by knowing how to reach any position.
- Exploration never fails: you always learn to do something, even if this was not your original goal.
- The principle of HER can be used in all model-free methods: DQN, DDPG, etc.



Source: <https://openai.com/blog/ingredients-for-robotics-research/>

TDM

- Using the goal-conditioned reward function $r(s_t, a_t, s_{t+1}) = -||s_{t+1} - s_g||$, how can we learn?



- TDM introduces goal-conditioned Q-value with a horizon T : $Q(s, a, s_g, T)$.
- The Q-value of an action should denote **how close** we will be from the goal s_g in T steps.
- If we can estimate these Q-values, we can use a planning algorithm such as MPC to find the action that will bring us closer to the goal easily:

$$a^* = \arg \max_{a_t} r(s_{t+T}, a_{t+T}, s_{t+T+1})$$

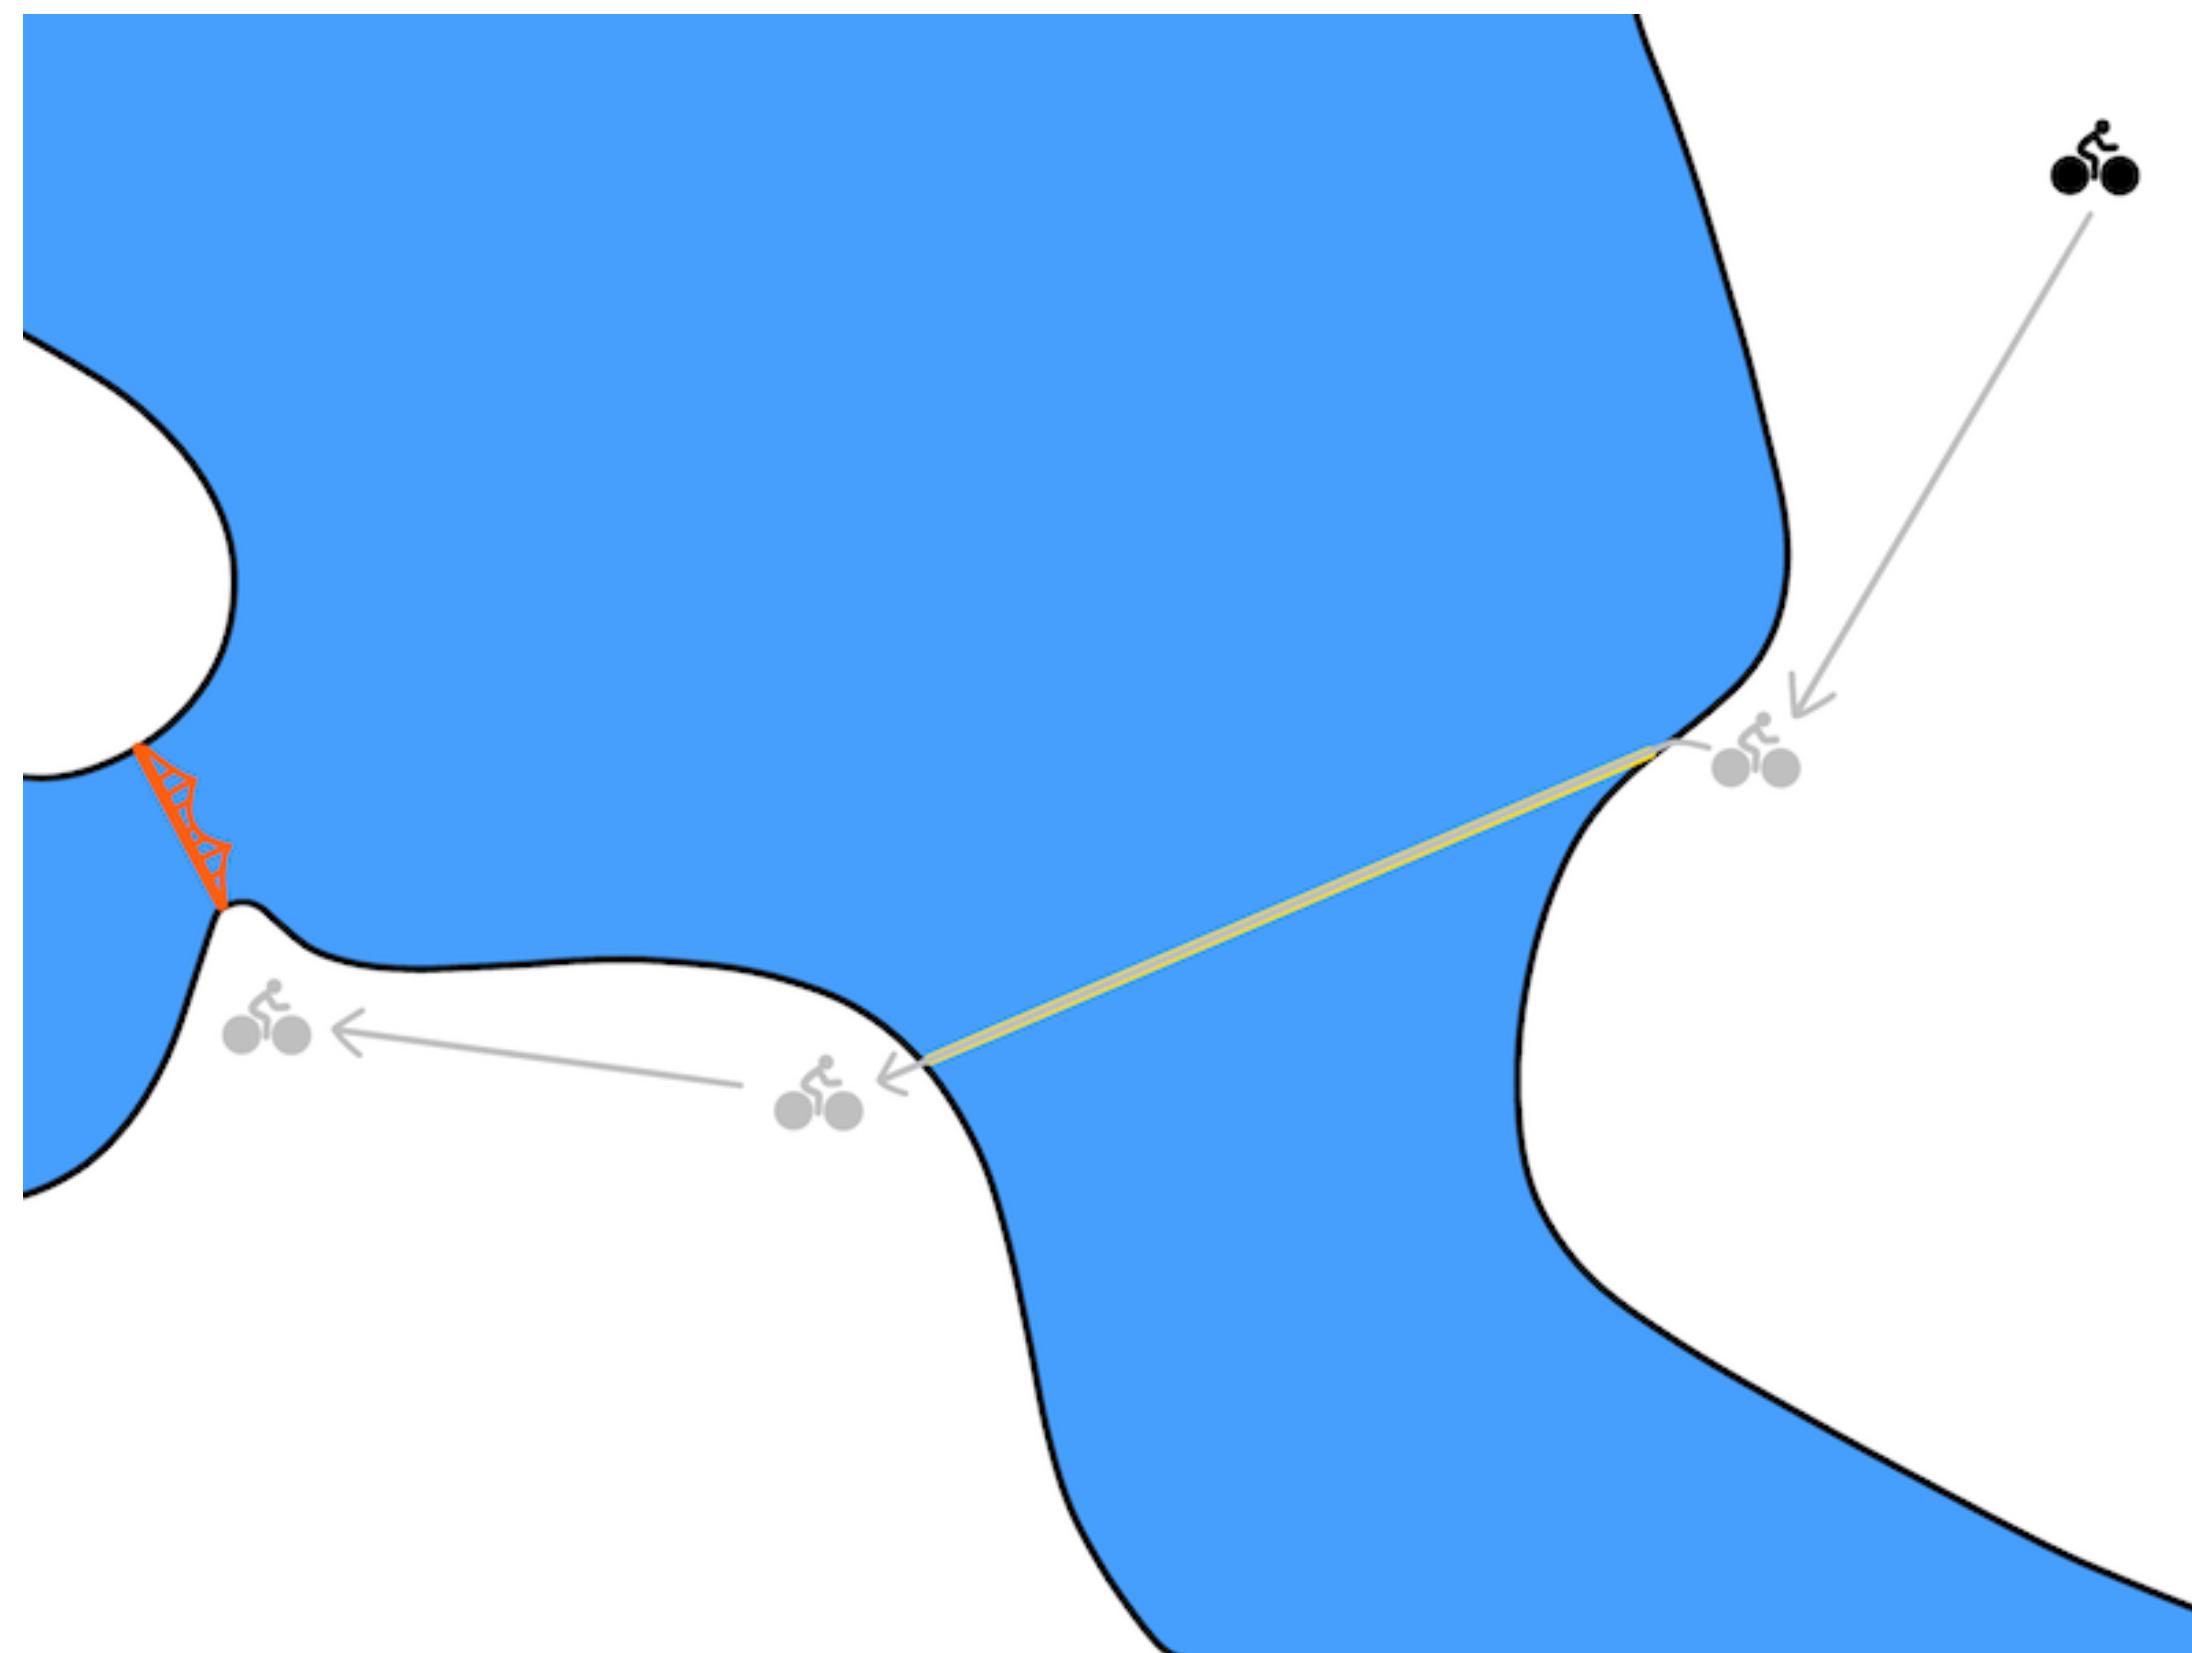
- This corresponds to planning T steps ahead; which action should I do now in order to be close to the goal in T steps?



Source: <https://bairblog.github.io/2018/04/26/tdm/>

TDM

- If the horizon T is well chosen, we only need to plan over a small number of intermediary positions, not over each possible action.
- TDM is model-free on each subgoal, but model-based on the whole trajectory.



Source: <https://bairblog.github.io/2018/04/26/tdm/>

TDM

- How can we learn the goal-conditioned Q-values $Q(s, a, s_g, T)$ with a **model**?
- TDM introduces a recursive relationship for the Q-values:

$$Q(s, a, s_g, T) = \begin{cases} \mathbb{E}_{s'}[r(s, a, s')] \text{ if } T = 0 \\ \mathbb{E}_{s'}[\max_a Q(s', a, s_g, T - 1)] \text{ otherwise.} \end{cases}$$
$$= \mathbb{E}_{s'}[r(s, a, s') \mathbf{1}(T = 0) + \max_a Q(s', a, s_g, T - 1) \mathbf{1}(T \neq 0)]$$

- If we plan over $T = 0$ steps, i.e. immediately after the action (s, a) , the Q-value is the remaining distance to the goal from the next state s' .
- Otherwise, it is the Q-value of the greedy action in the next state s' with an horizon $T - 1$ (one step shorter).
- This allows to learn the Q-values from **single transitions** (s_t, a_t, s_{t+1}) :
 - with $T = 0$, the target is the remaining distance to the goal.
 - with $T > 0$, the target is the Q-value of the next action at a shorter horizon.

TDM

- The critic learns to minimize the prediction error **off-policy**:

$$\mathcal{L}(\theta) = \mathbb{E}_{s_t, a_t, s_{t+1} \in \mathcal{D}} [(r(s_t, a_t, s_{t+1}) \mathbf{1}(T = 0) + \max_a Q(s_{t+1}, a, s_g, T - 1) \mathbf{1}(T \neq 0) - Q(s_t, a_t, s_g, T))]$$

- This is a model-free Q-learning-like update rule, that can be learned by any off-policy value-based algorithm (DQN, DDPG) and an experience replay memory.
- The cool trick is that, with a single transition (s_t, a_t, s_{t+1}) , you can train the critic with:
 - different horizons T , e.g. between 0 and T_{\max} .
 - different goals s_g . You can sample any achievable state as a goal, including the “true” s_{t+T} (hindsight).
- You do not only learn to reach s_g , but any state! TDM learns a lot of information from a single transition, so it has a very good sample complexity.

Summary of TDM

- TDM learns to break long trajectories into finite horizons (model-based planning) by learning model-free (Q-learning updates).
- The critic learns how good an action (s, a) is order to reach a state s_g in T steps.

$$Q(s, a, s_g, T) = \mathbb{E}_{s'} [r(s, a, s') \mathbf{1}(T = 0) + \max_a Q(s', a, s_g, T - 1) \mathbf{1}(T \neq 0)]$$


- The actor uses MPC planning to iteratively select actions that bring us closer to the goal in T steps:

$$a_t = \arg \max_a Q(s_t, a, s_g, T)$$


- The argmax can be estimated via sampling.
- TDM is a model-based method in disguise: it does predict the next state directly, but how much closer it will be to the goal via Q-learning.

TDM results

- For problems where the model is easy to learn, the performance of TDM is on par with model-based methods (MPC).



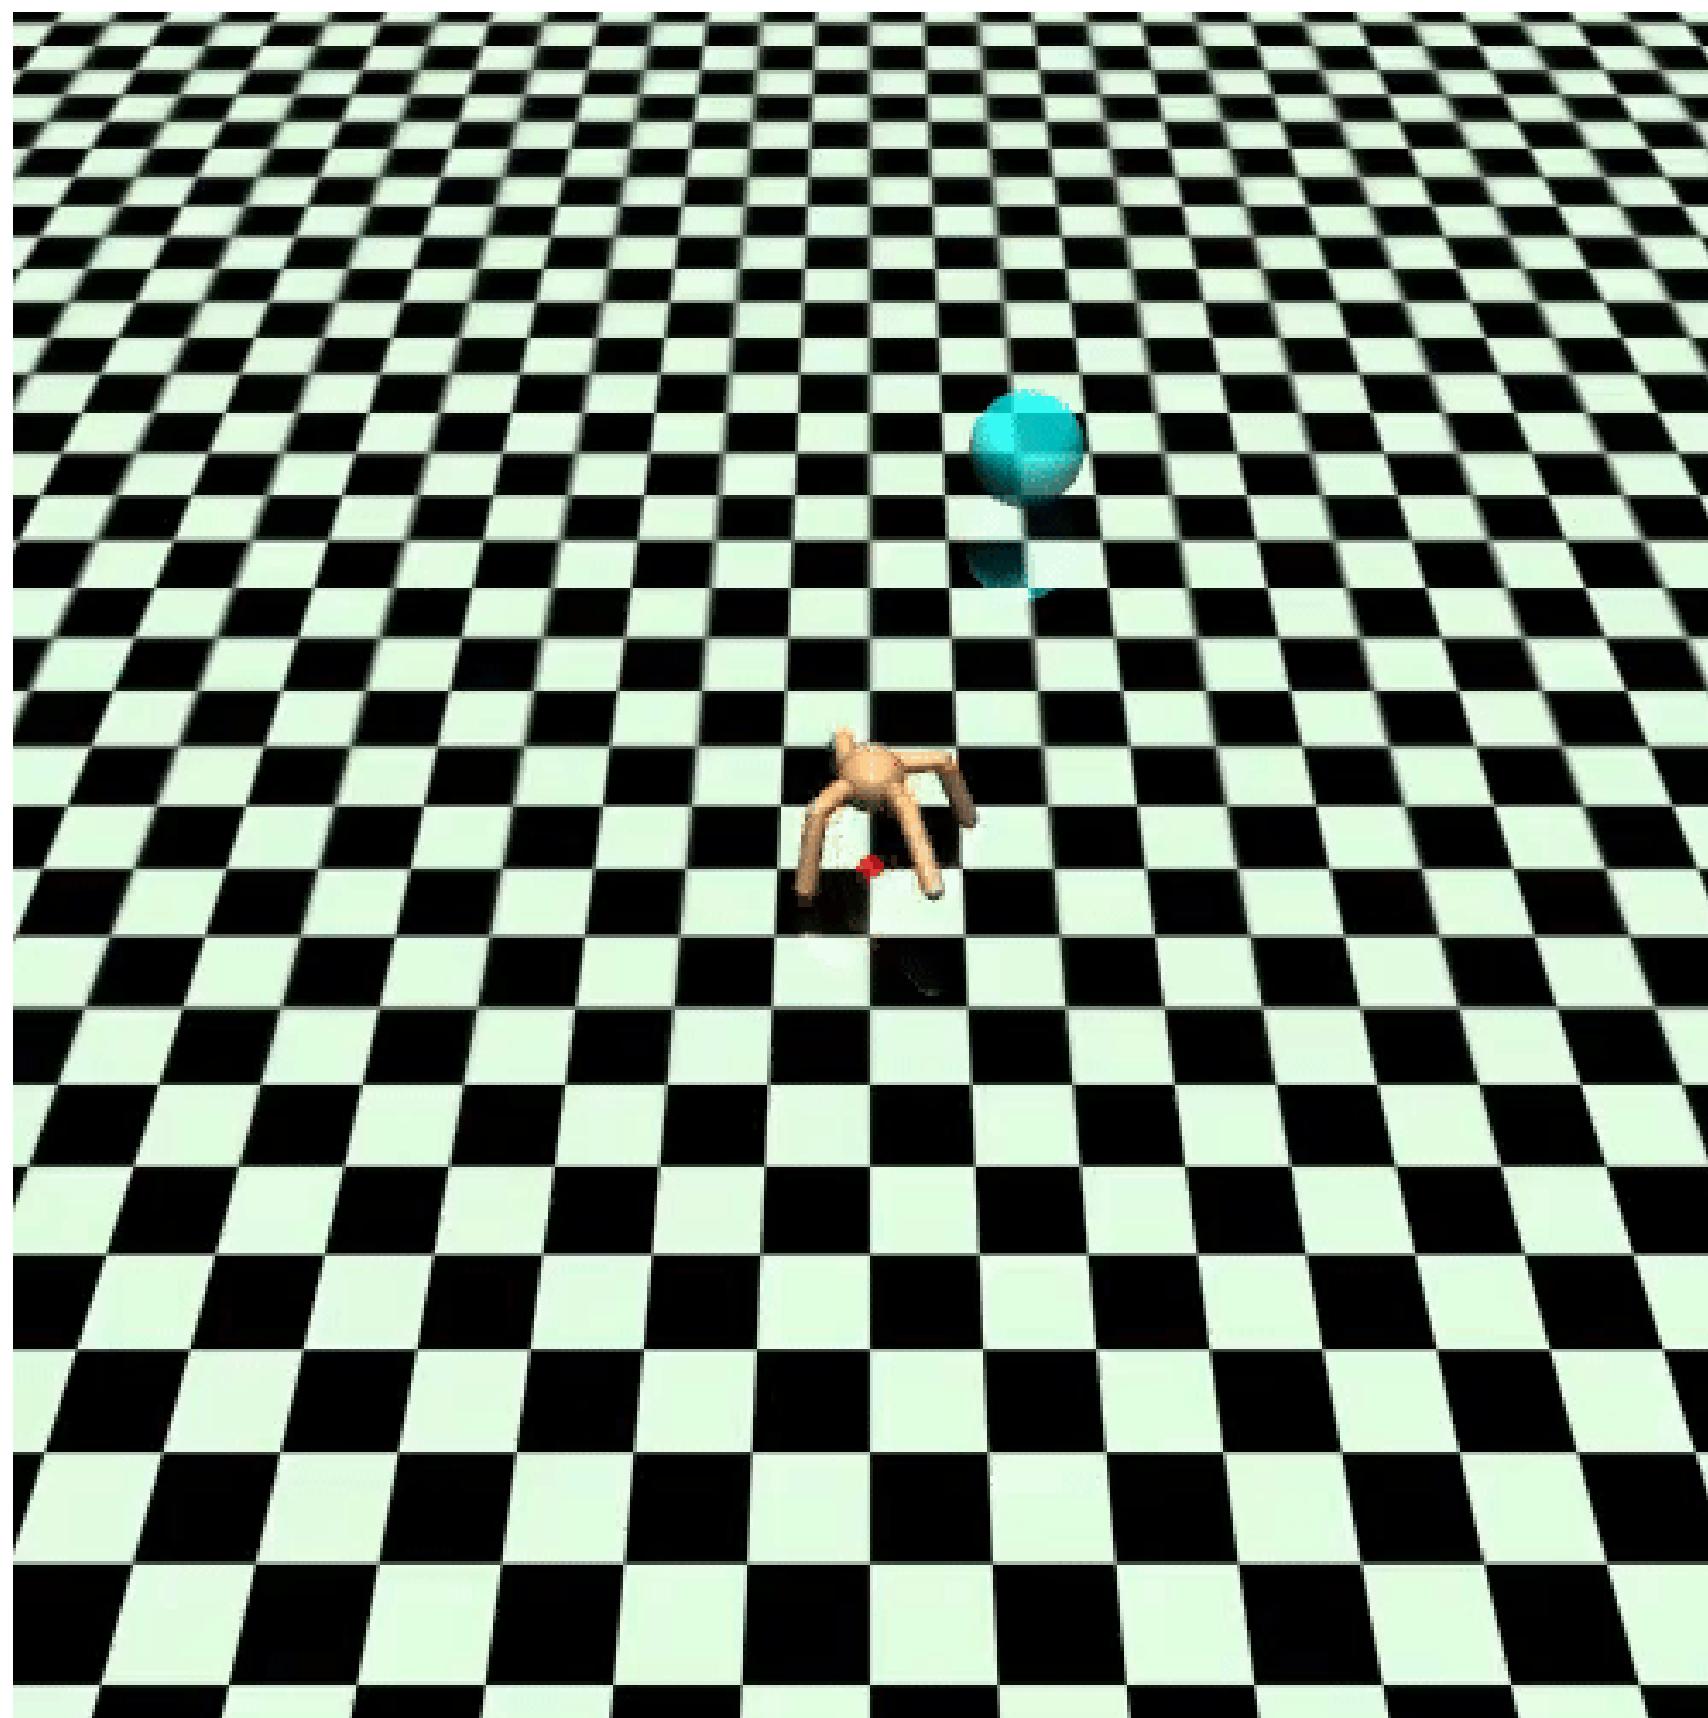
- Model-free methods have a much higher sample complexity.
- TDM learns much more from single transitions.



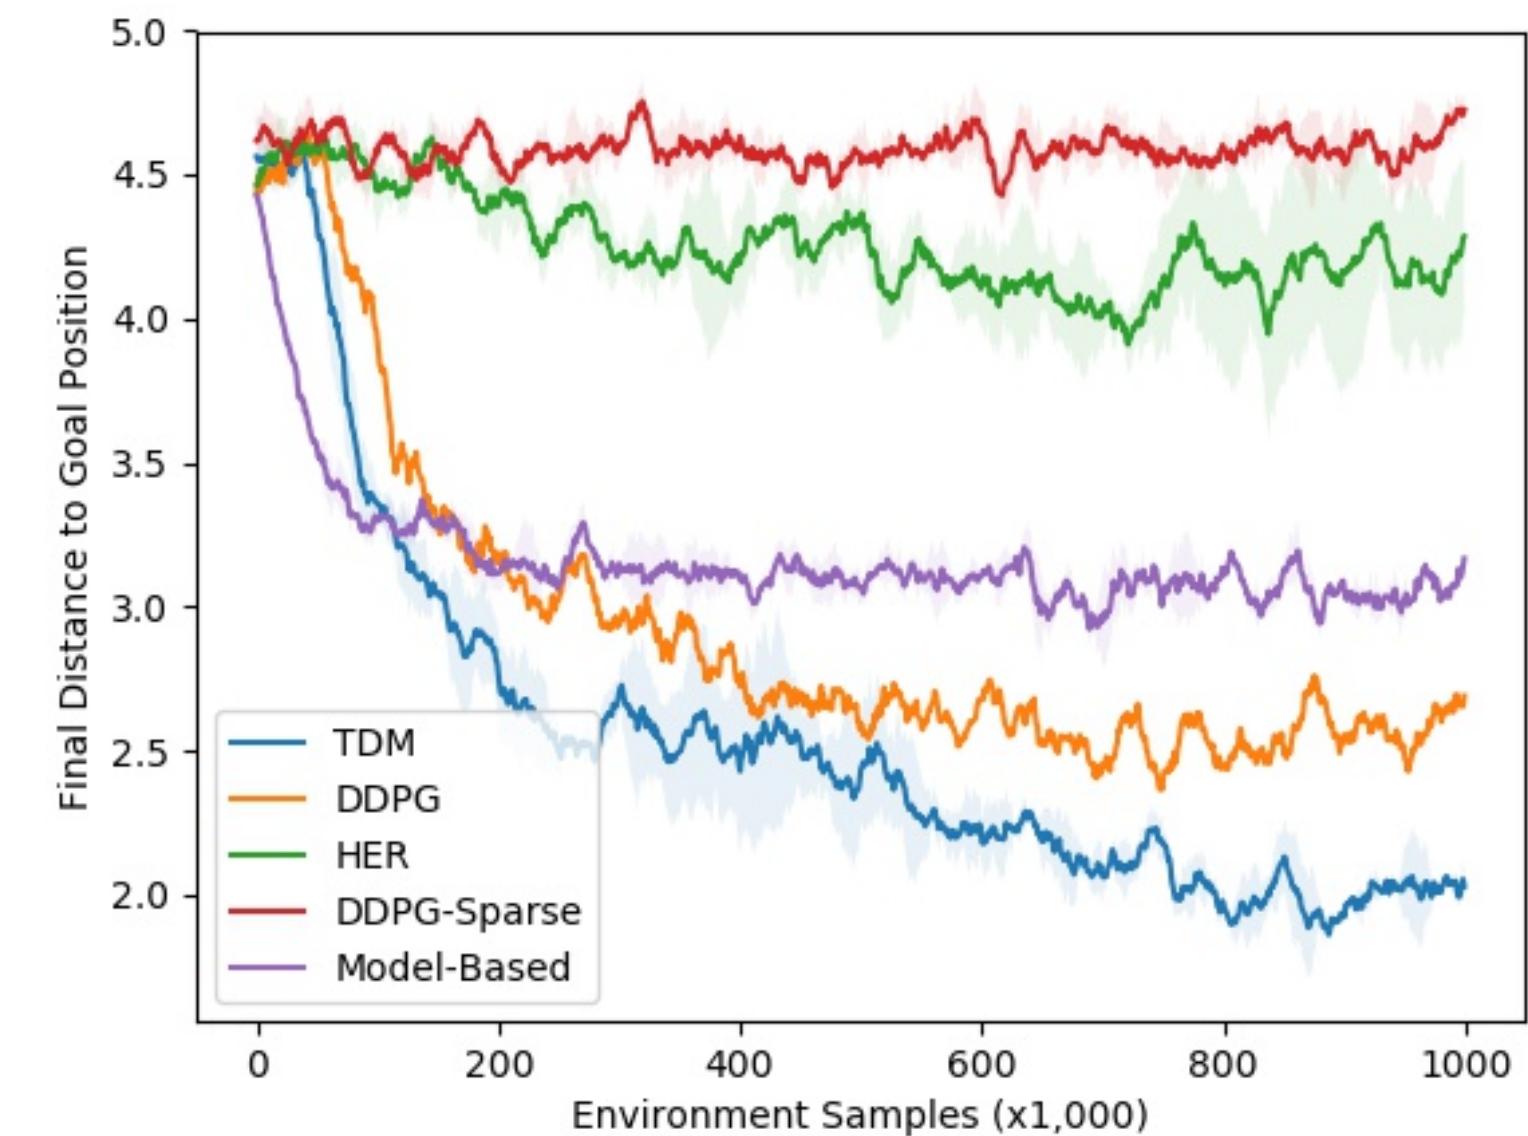
Source: <https://bairblog.github.io/2018/04/26/tdm/>

TDM results

- For problems where the model is complex to learn, the performance of TDM is on par with model-free methods (DDPG).



- Model-based methods suffer from model imprecision on long horizons.
- TDM plans over shorter horizons T .



Source: <https://bairblog.github.io/2018/04/26/tdm/>

3 - World models


World Models

David Ha¹ Jürgen Schmidhuber^{2,3}

<https://worldmodels.github.io/>

World models

- The core idea of **world models** is to explicitly separate the **world model** (what will happen next) from the **controller** (how to act).
- Deep RL NN are usually small, as rewards do not contain enough information to train huge networks.



<https://worldmodels.github.io/>

World models

- A huge **world model** can be efficiently trained by supervised or unsupervised methods.
- A small **controller** should not need too many trials if its input representations are good.

At each time step, our agent receives an **observation** from the environment.

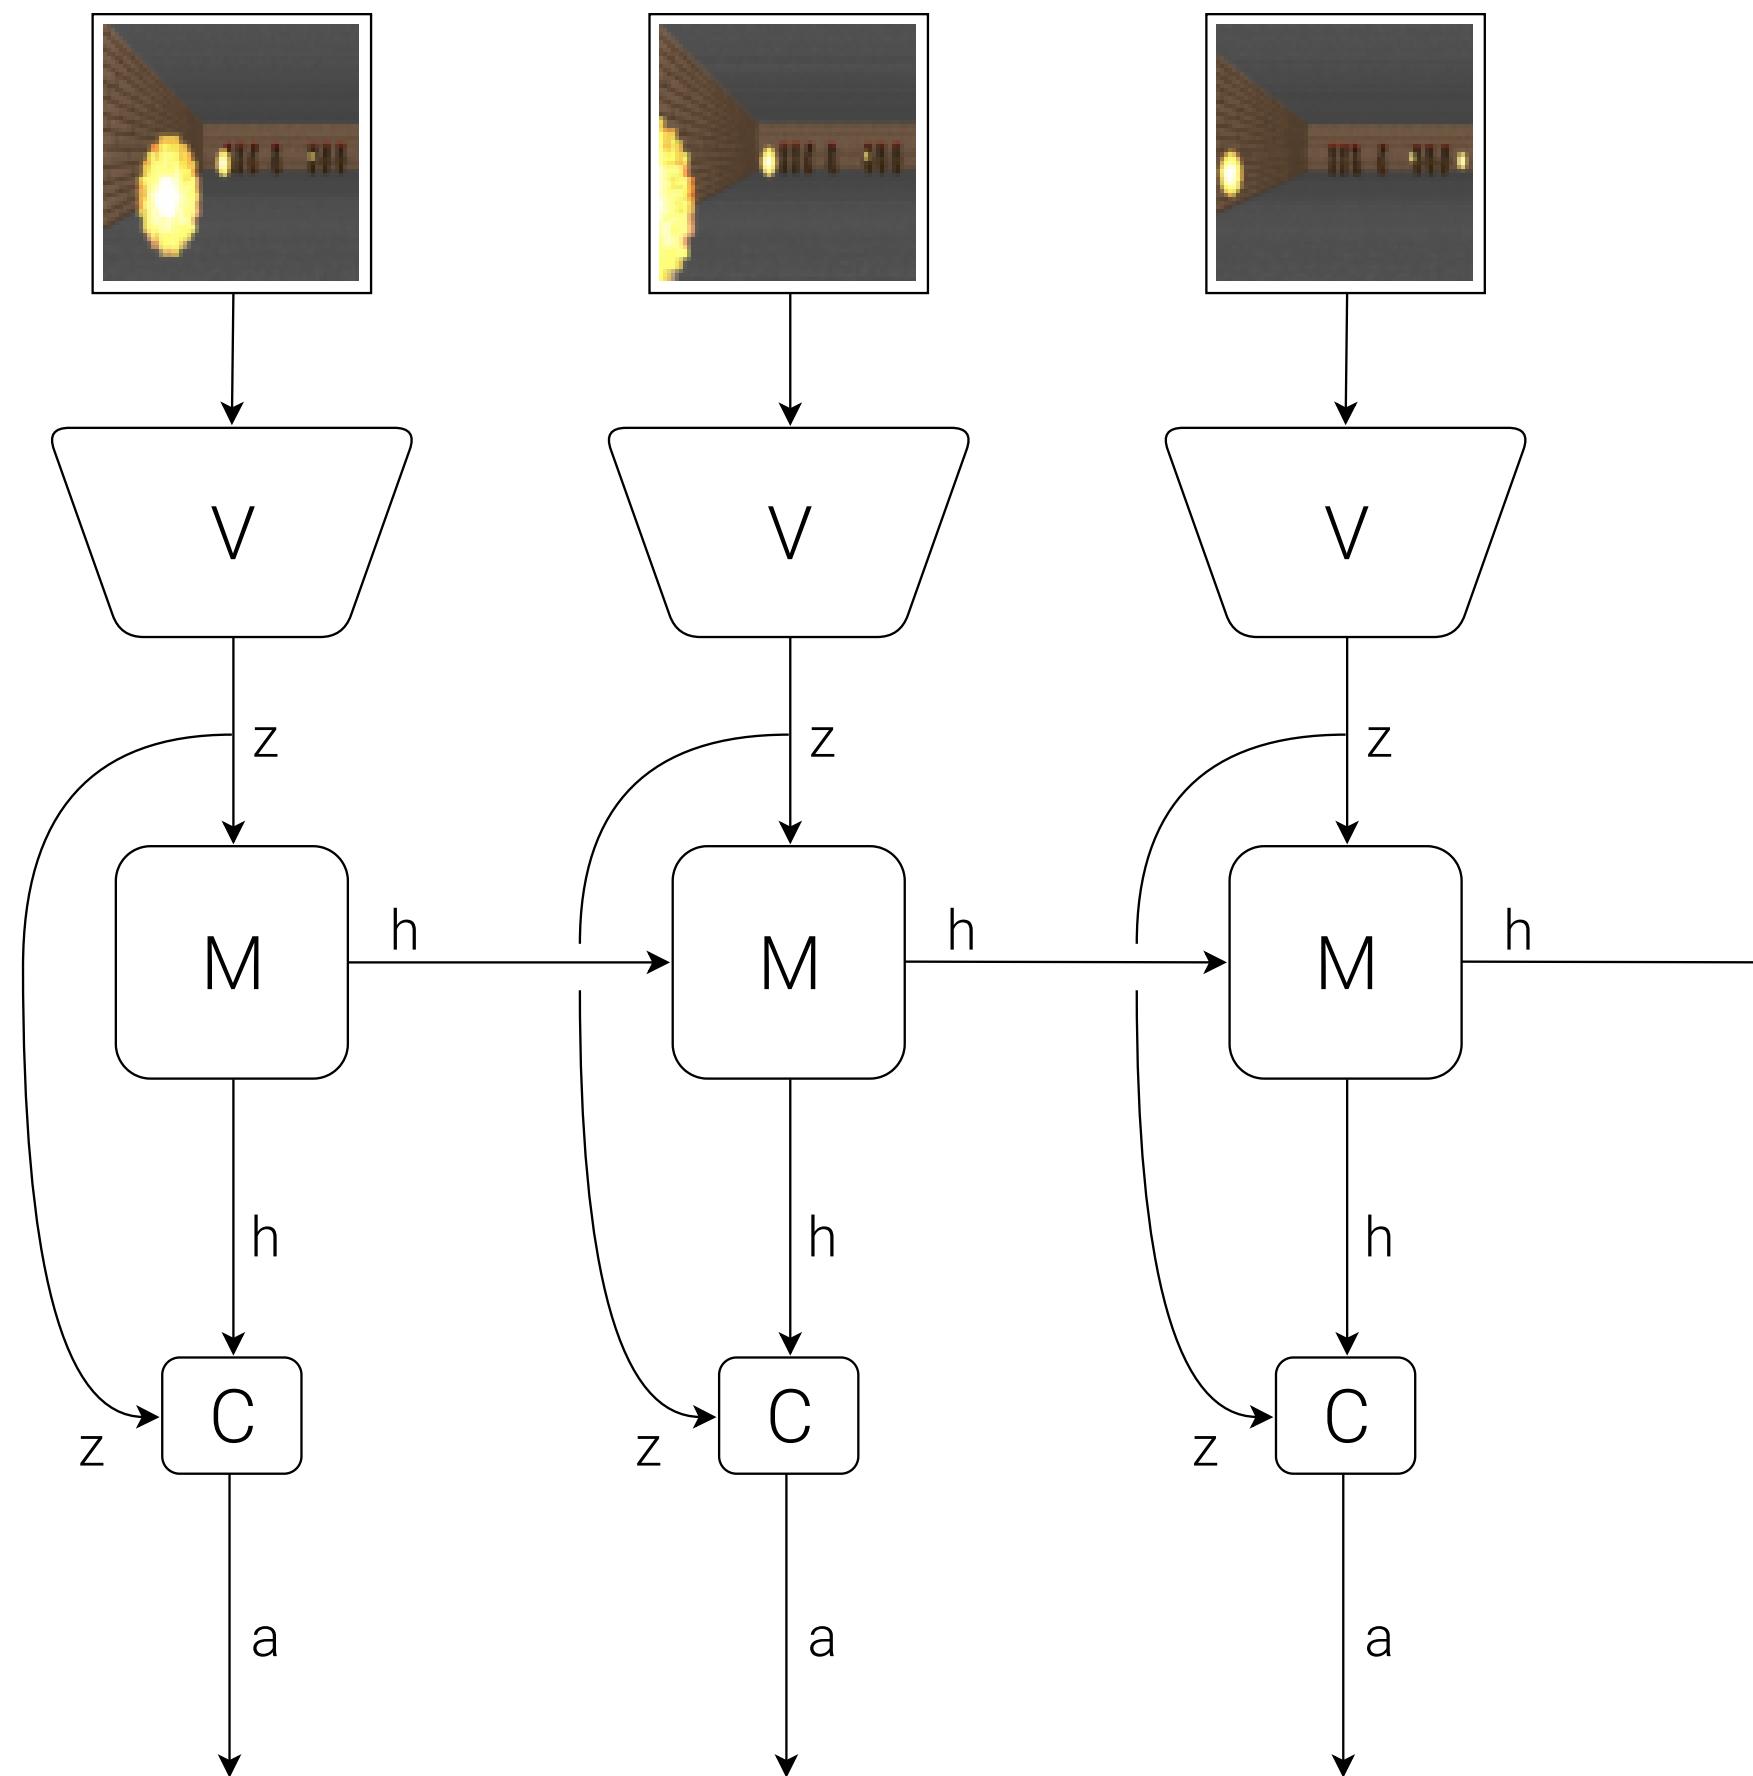
World Model

The **Vision Model (V)** encodes the high-dimensional observation into a low-dimensional latent vector.

The **Memory RNN (M)** integrates the historical codes to create a representation that can predict future states.

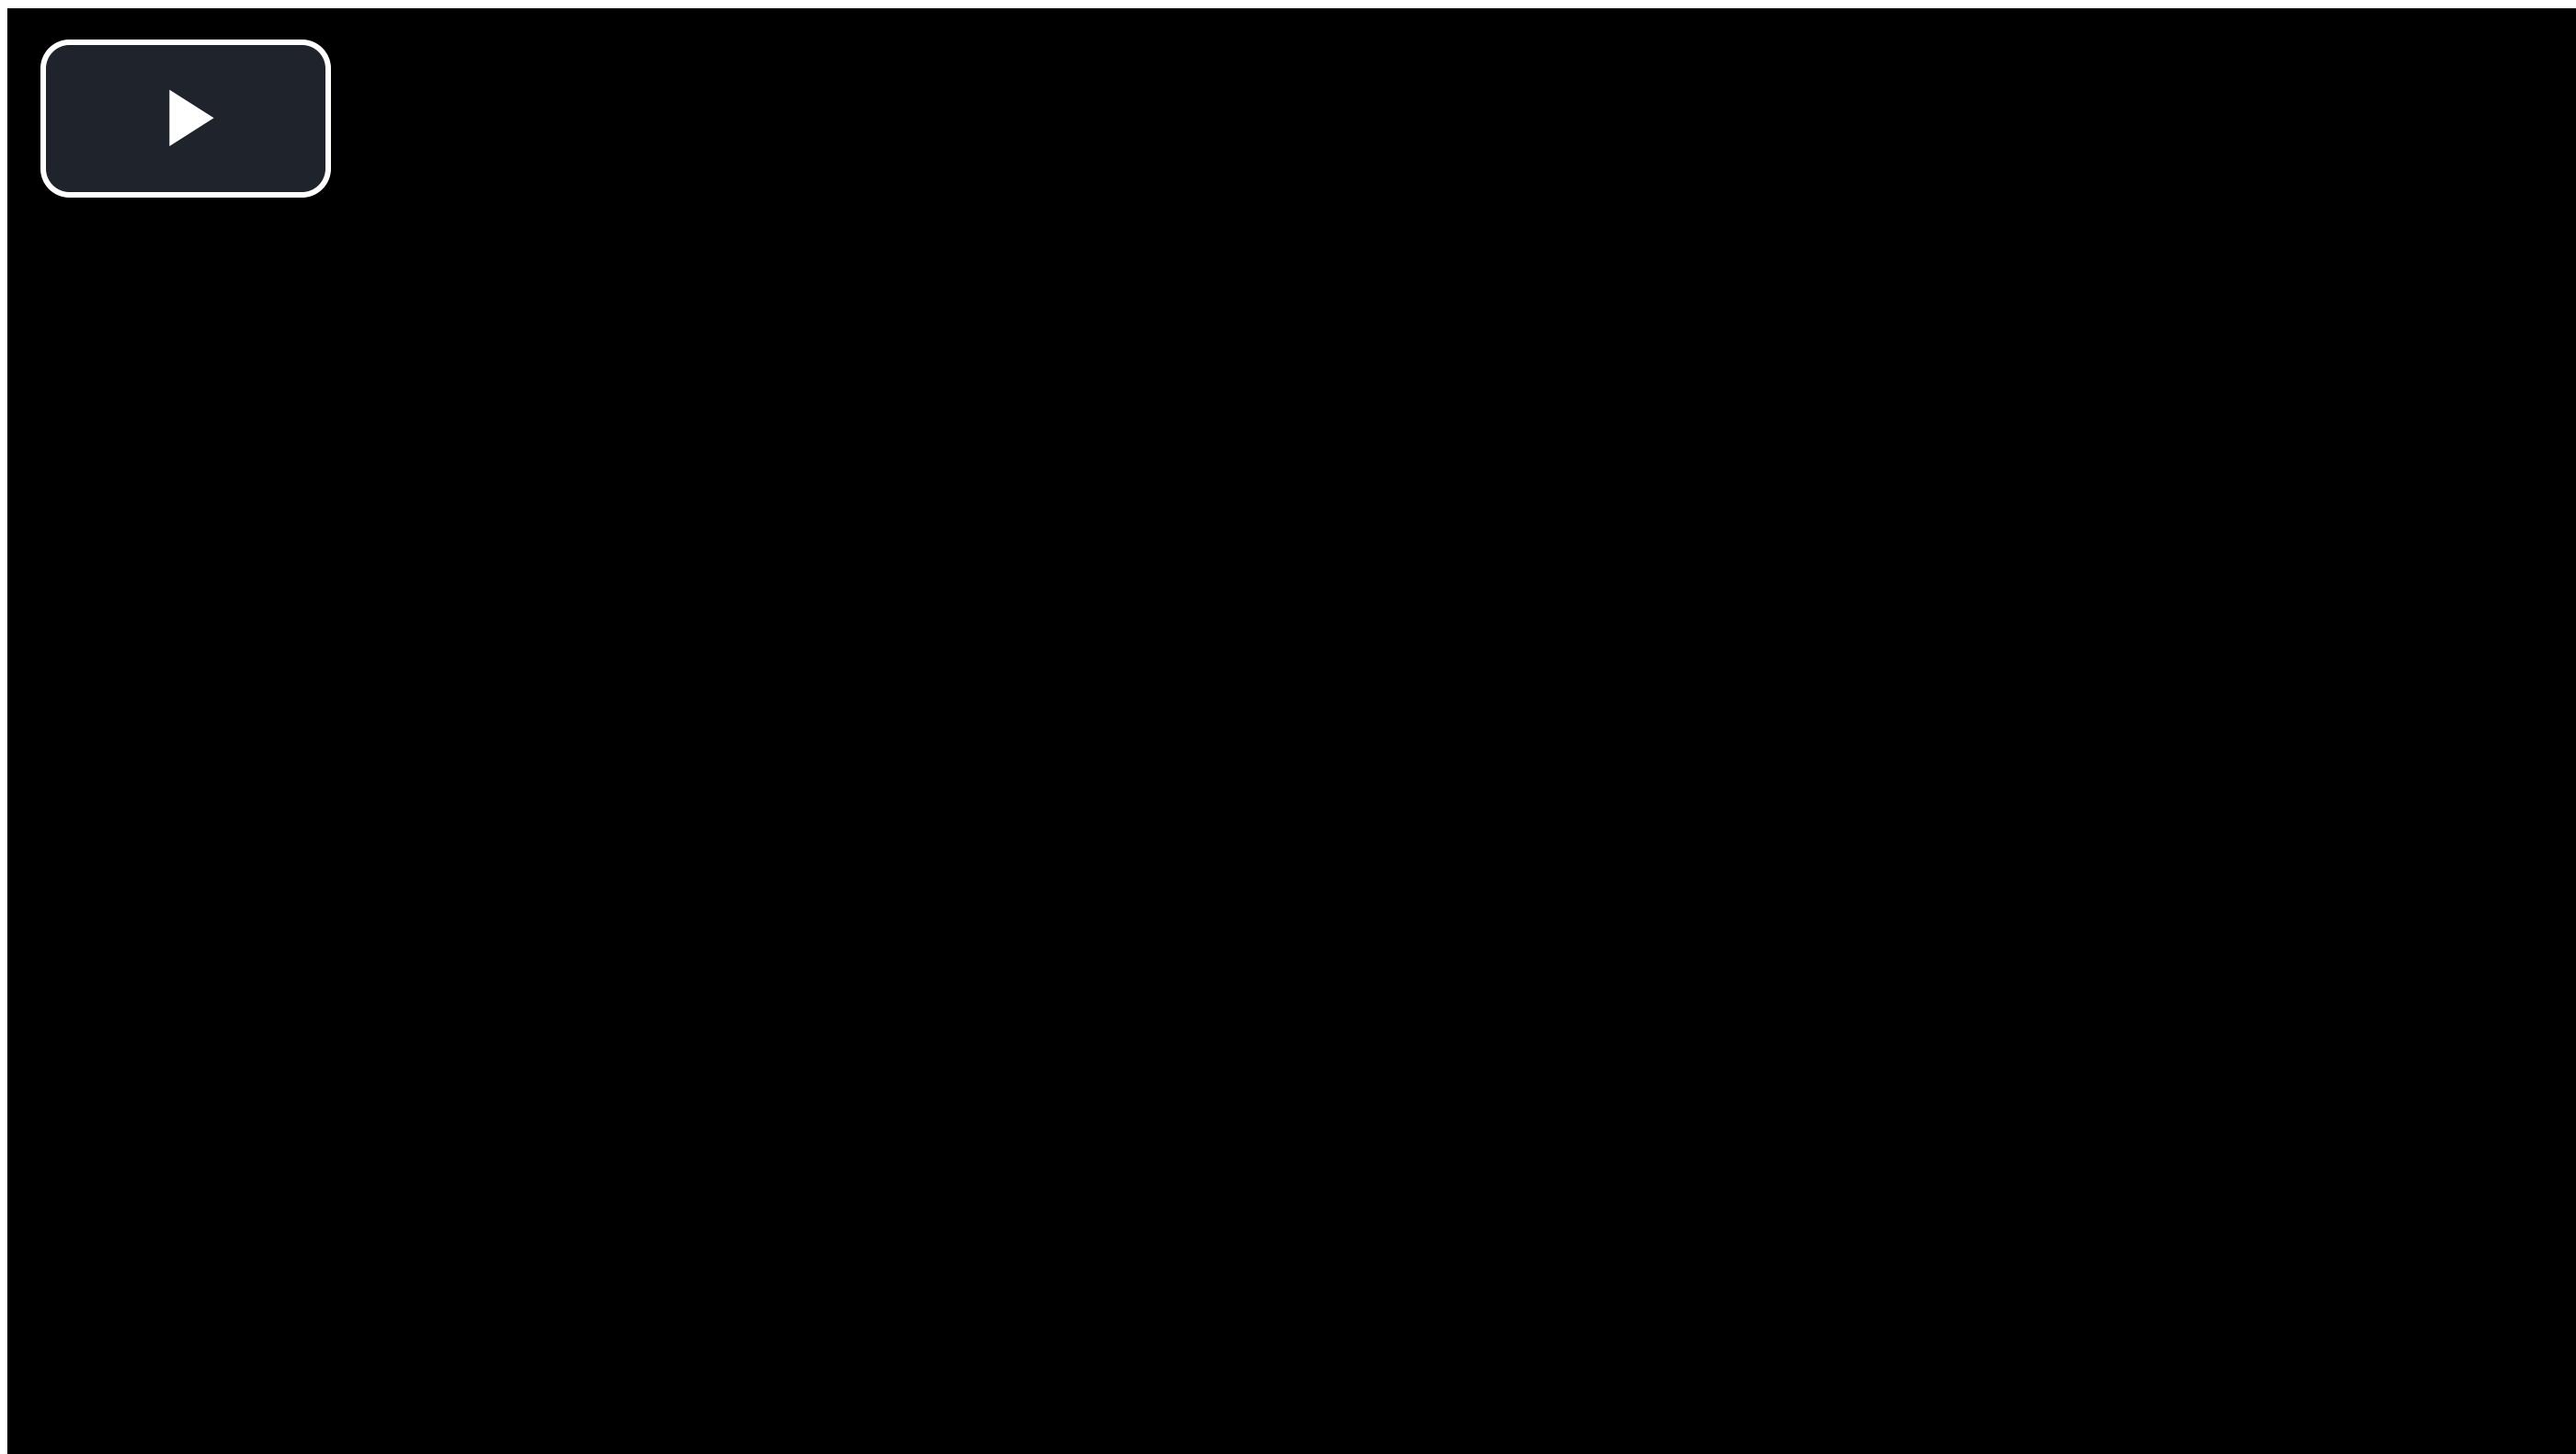
A small **Controller (C)** uses the representations from both **V** and **M** to select good actions.

The agent performs **actions** that go back and affect the environment.




<https://worldmodels.github.io/>

The Vizdoom Take Cover environment



World models

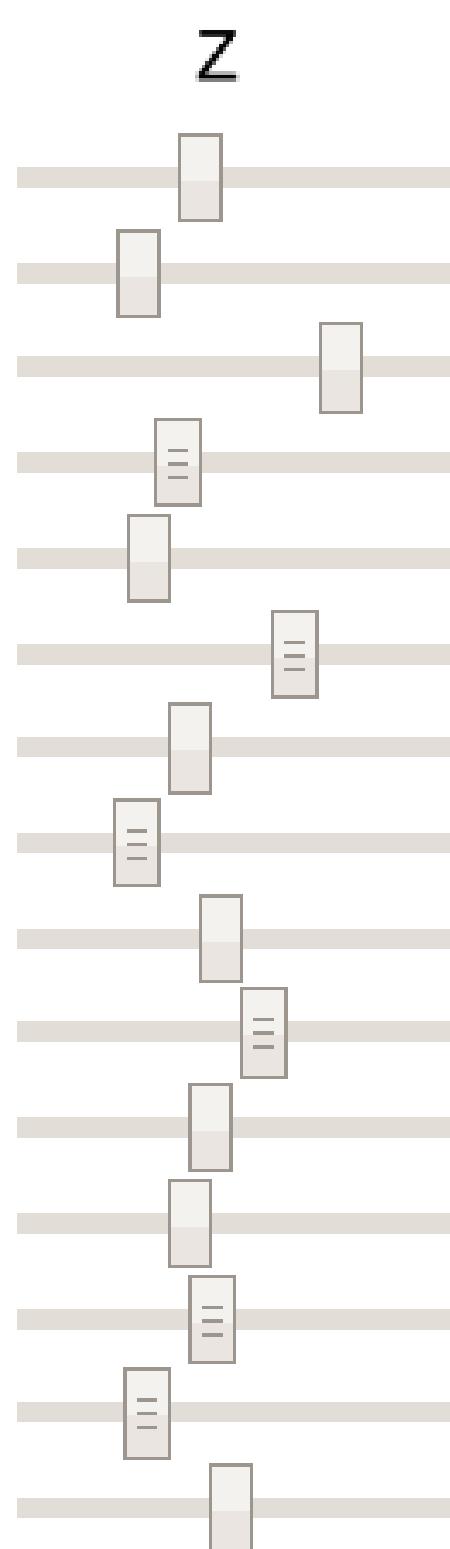
- The vision module V is trained as a **variational autoencoder** (VAE) on single frames of the game.
- The latent vector \mathbf{z}_t contains a compressed representation of the frame \mathbf{o}_t .



<https://worldmodels.github.io/>

World models

Screenshot Image



Reconstruction




[Load Random Screenshot](#)

[Randomize Z](#)

- Go to <https://worldmodels.github.io/> for an interactive demo.

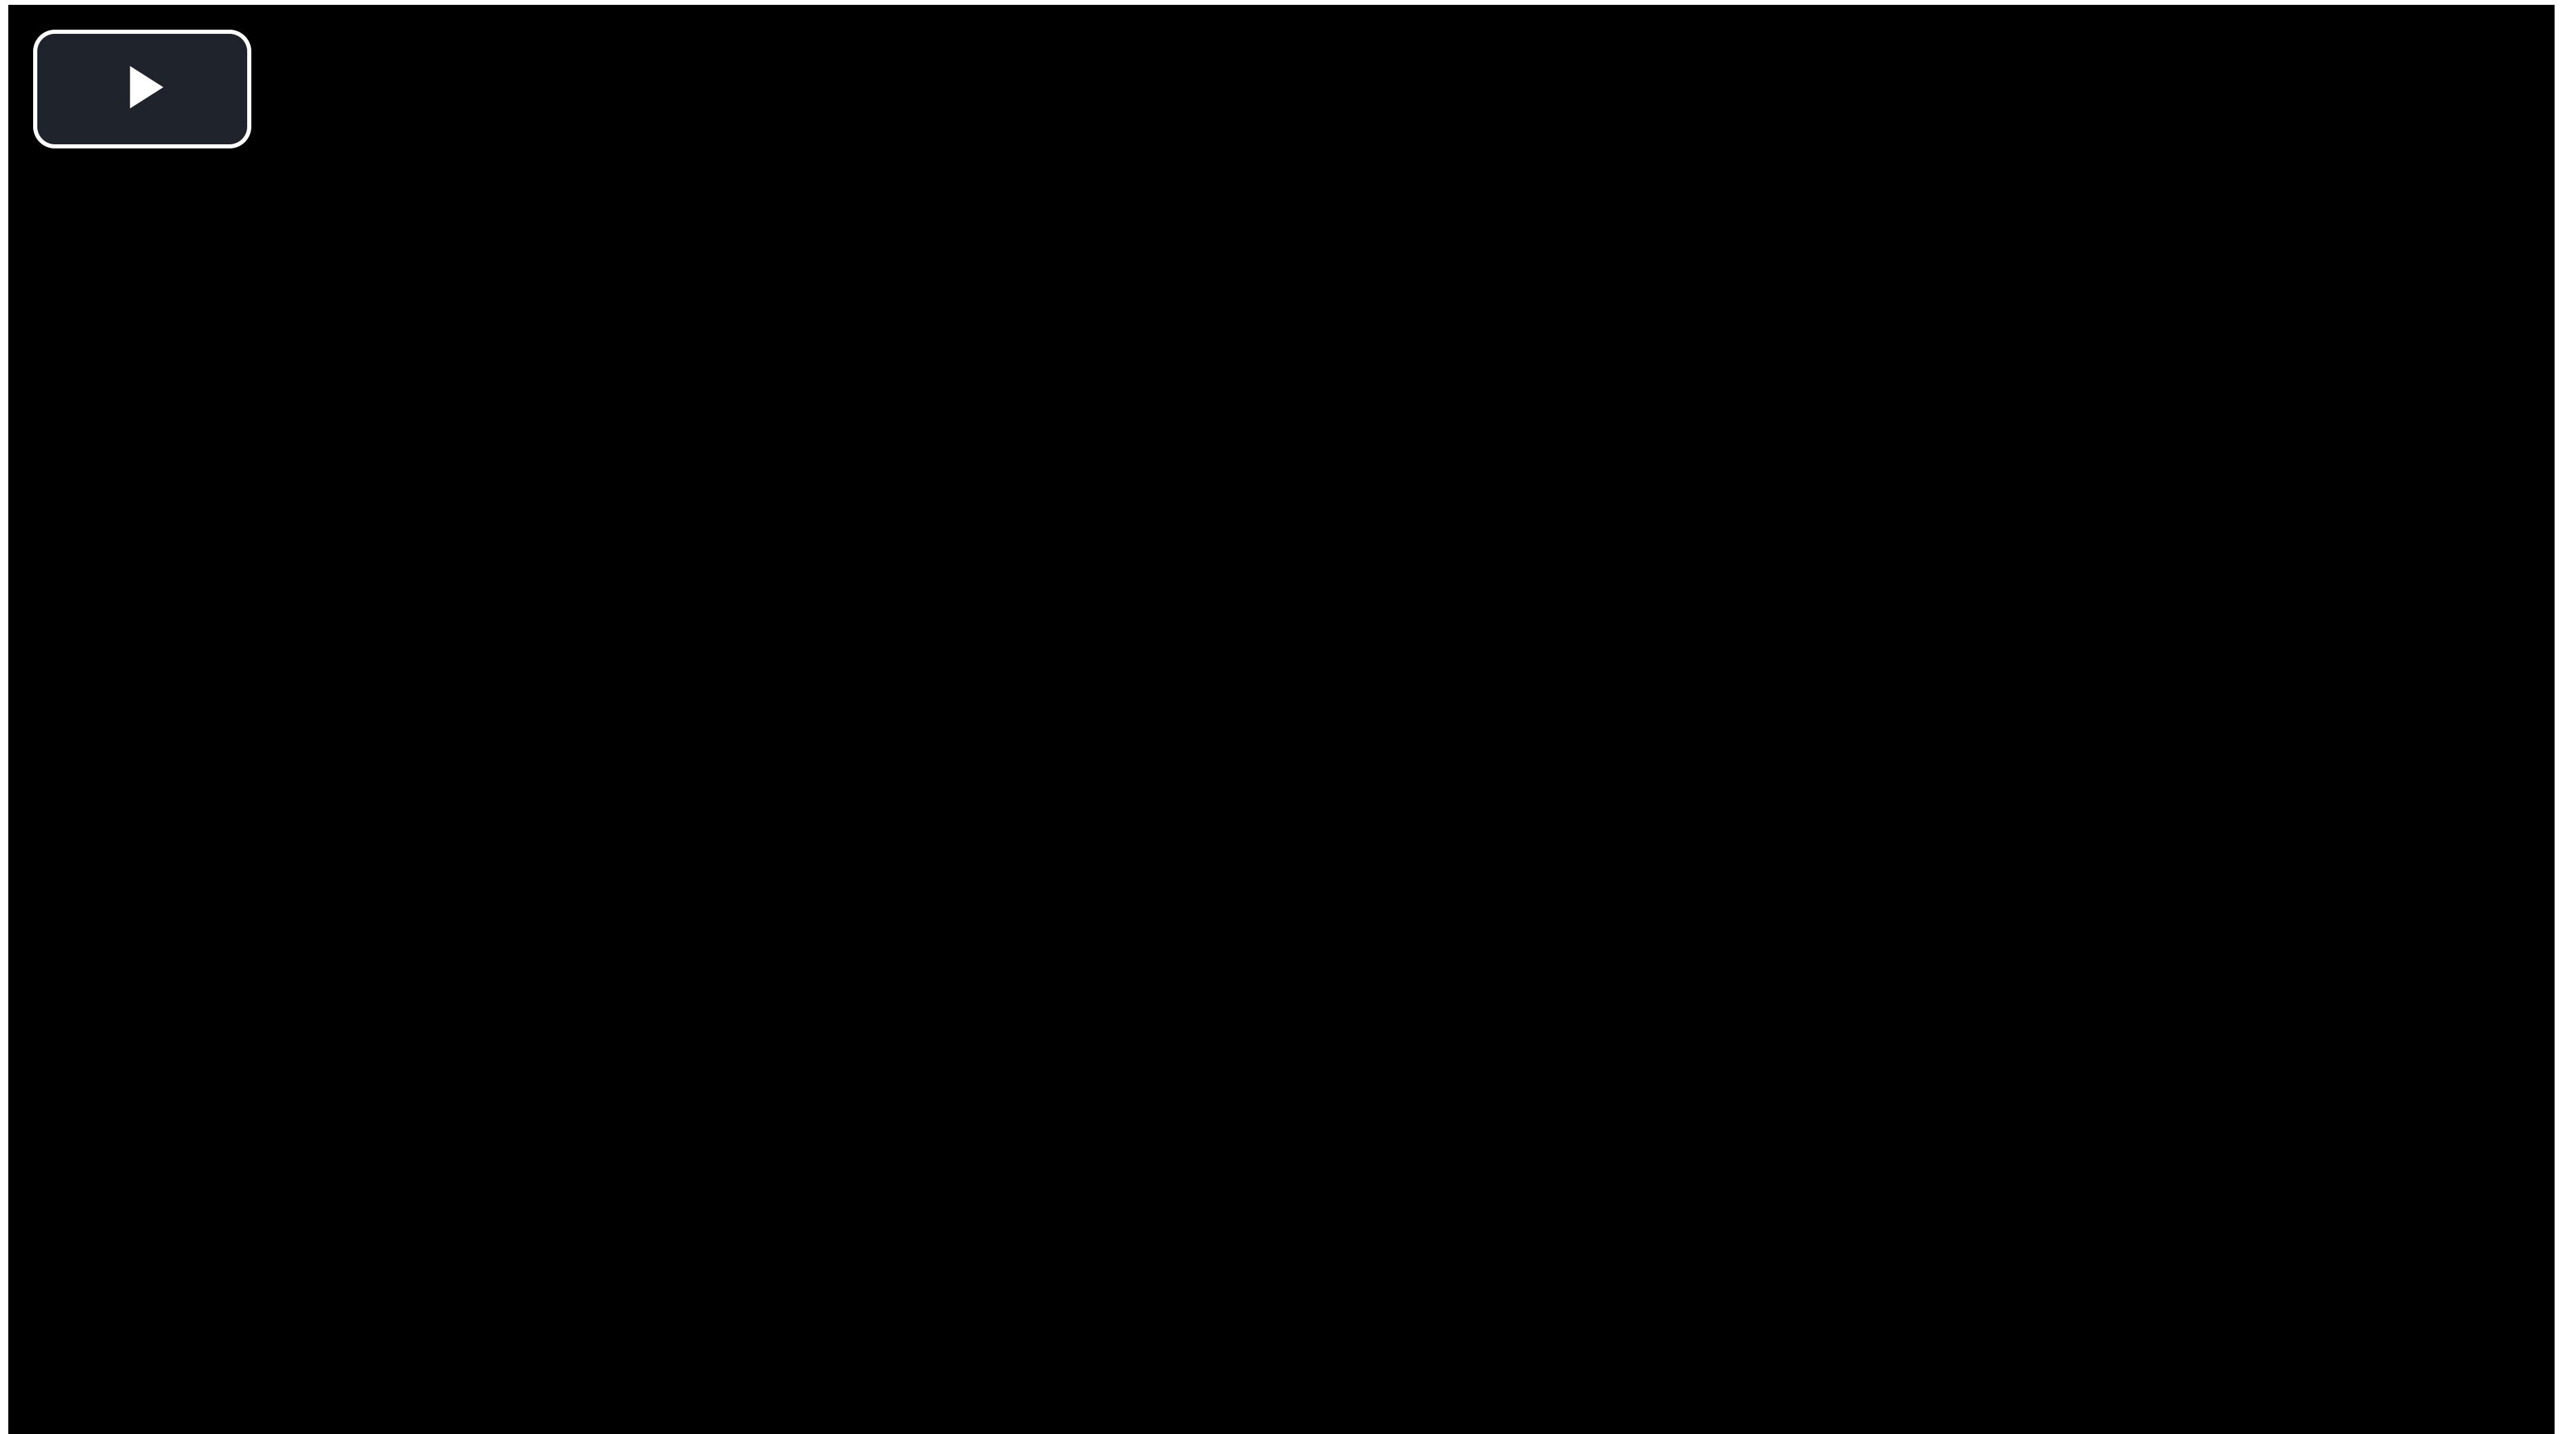
World models

- The sequence of latent representations $\mathbf{z}_0, \dots, \mathbf{z}_t$ in a game is fed to a LSTM layer together with the actions a_t to compress what happens over time.
- A **Mixture Density Network (MDN)** is used to predict the **distribution** of the next latent representations $P(\mathbf{z}_{t+1} | a_t, \mathbf{h}_t, \dots, \mathbf{z}_t)$.
- The RNN-MDN architecture has been used successfully in the past for sequence generation problems such as generating handwriting and sketches (Sketch-RNN).



<https://worldmodels.github.io/>

Sketch-RNN



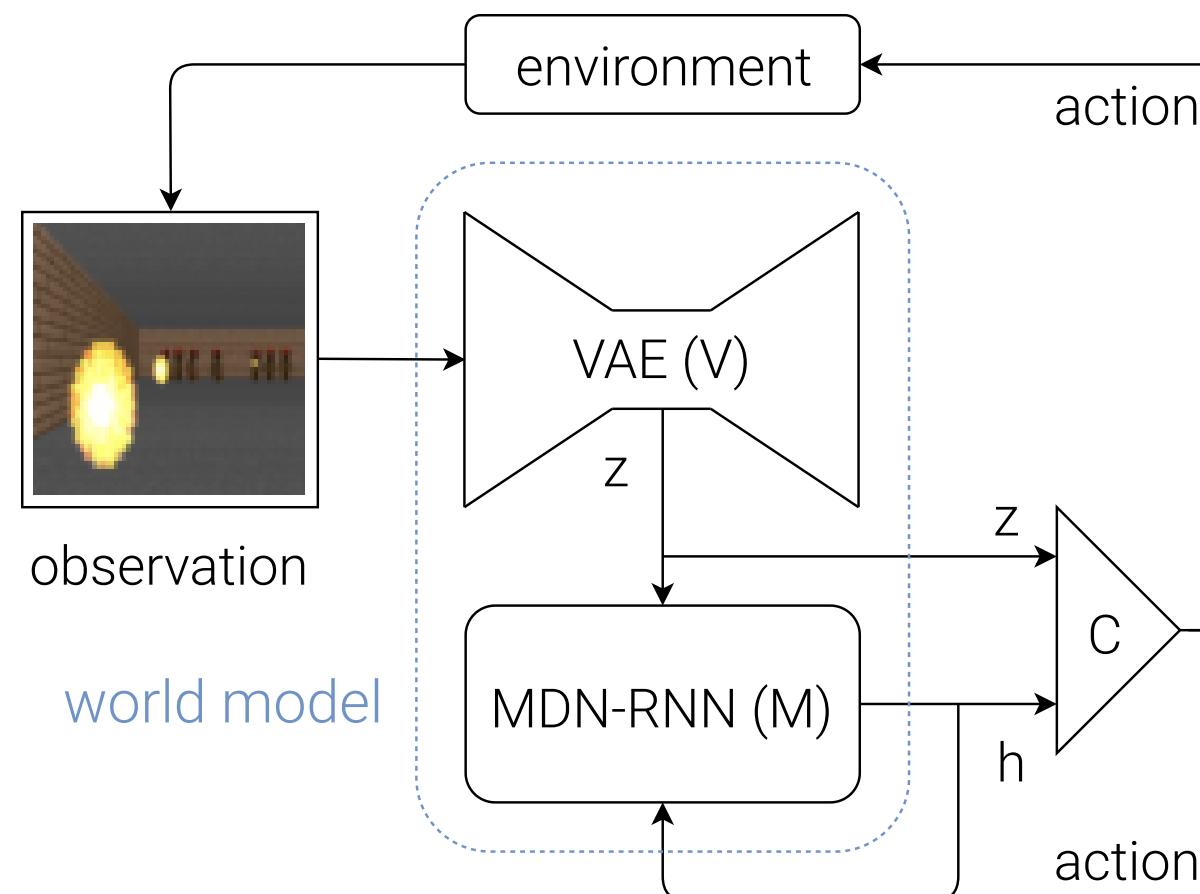
<https://magenta.tensorflow.org/sketch-rnn-demo>

World models

- The last step is the **controller**. It takes a latent representation \mathbf{z}_t and the current hidden state of the LSTM \mathbf{h}_t as inputs and selects an action **linearly**:

$$a_t = \tanh(W[\mathbf{z}_t, \mathbf{h}_t] + b)$$

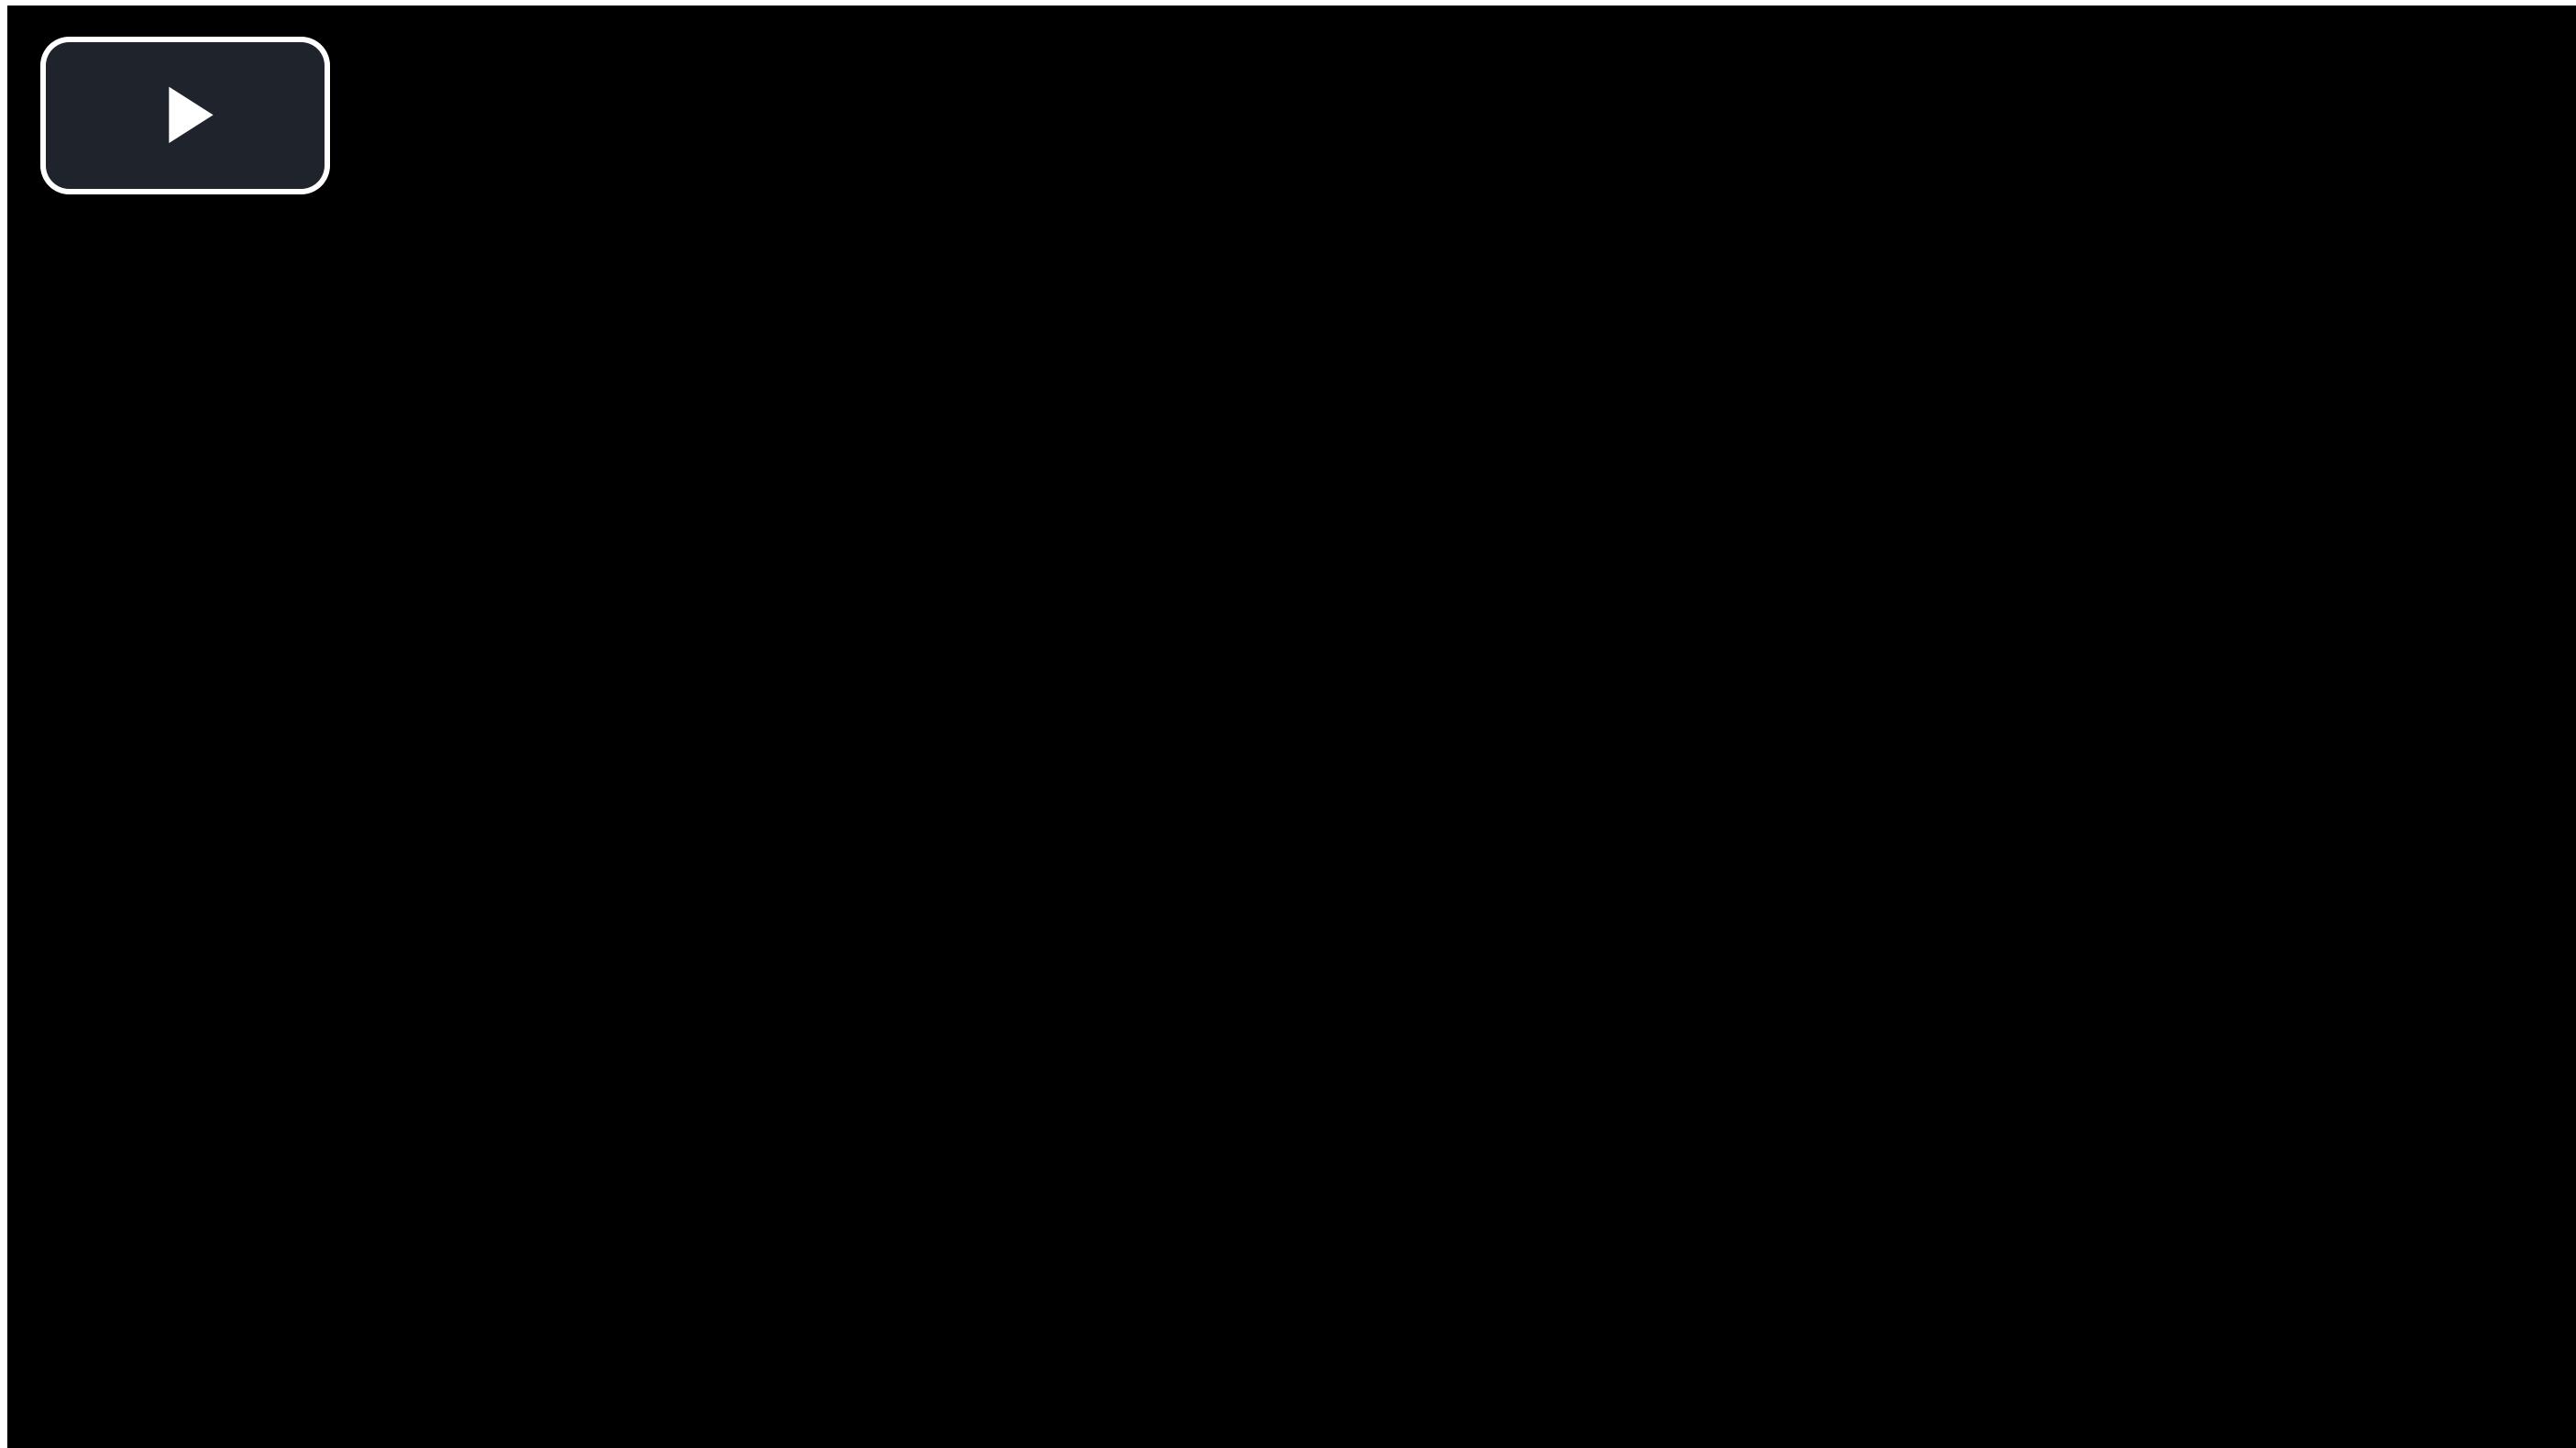
- A RL actor cannot get simpler as that...



<https://worldmodels.github.io/>


- The controller is not even trained with RL: it uses a genetic algorithm, the Covariance-Matrix Adaptation Evolution Strategy (CMA-ES), to find the output weights that maximize the returns.
- The world model is trained by classical supervised learning using a random agent before learning.

World models : car racing




World models : car racing

- Below is the input of the VAE and the reconstruction.
- The reconstruction does not have to be perfect as long as the latent space is informative.



World models : car racing

- Controller seeing only \mathbf{z}_t .
- Controller seeing both \mathbf{z}_t and \mathbf{h}_t .



- Having access to a full rollout of the future leads to more stable driving.

World models

Algorithm:

1. Collect 10,000 rollouts from a random policy.
2. Train VAE (V) to encode each frame into a latent vector $\mathbf{z} \in \mathcal{R}^{32}$.
3. Train MDN-RNN (M) to model $P(\mathbf{z}_{t+1} | a_t, \mathbf{h}_t, \dots, \mathbf{z}_t)$.
4. Evolve Controller (C) to maximize the expected cumulative reward of a rollout.

Parameters for car racing:

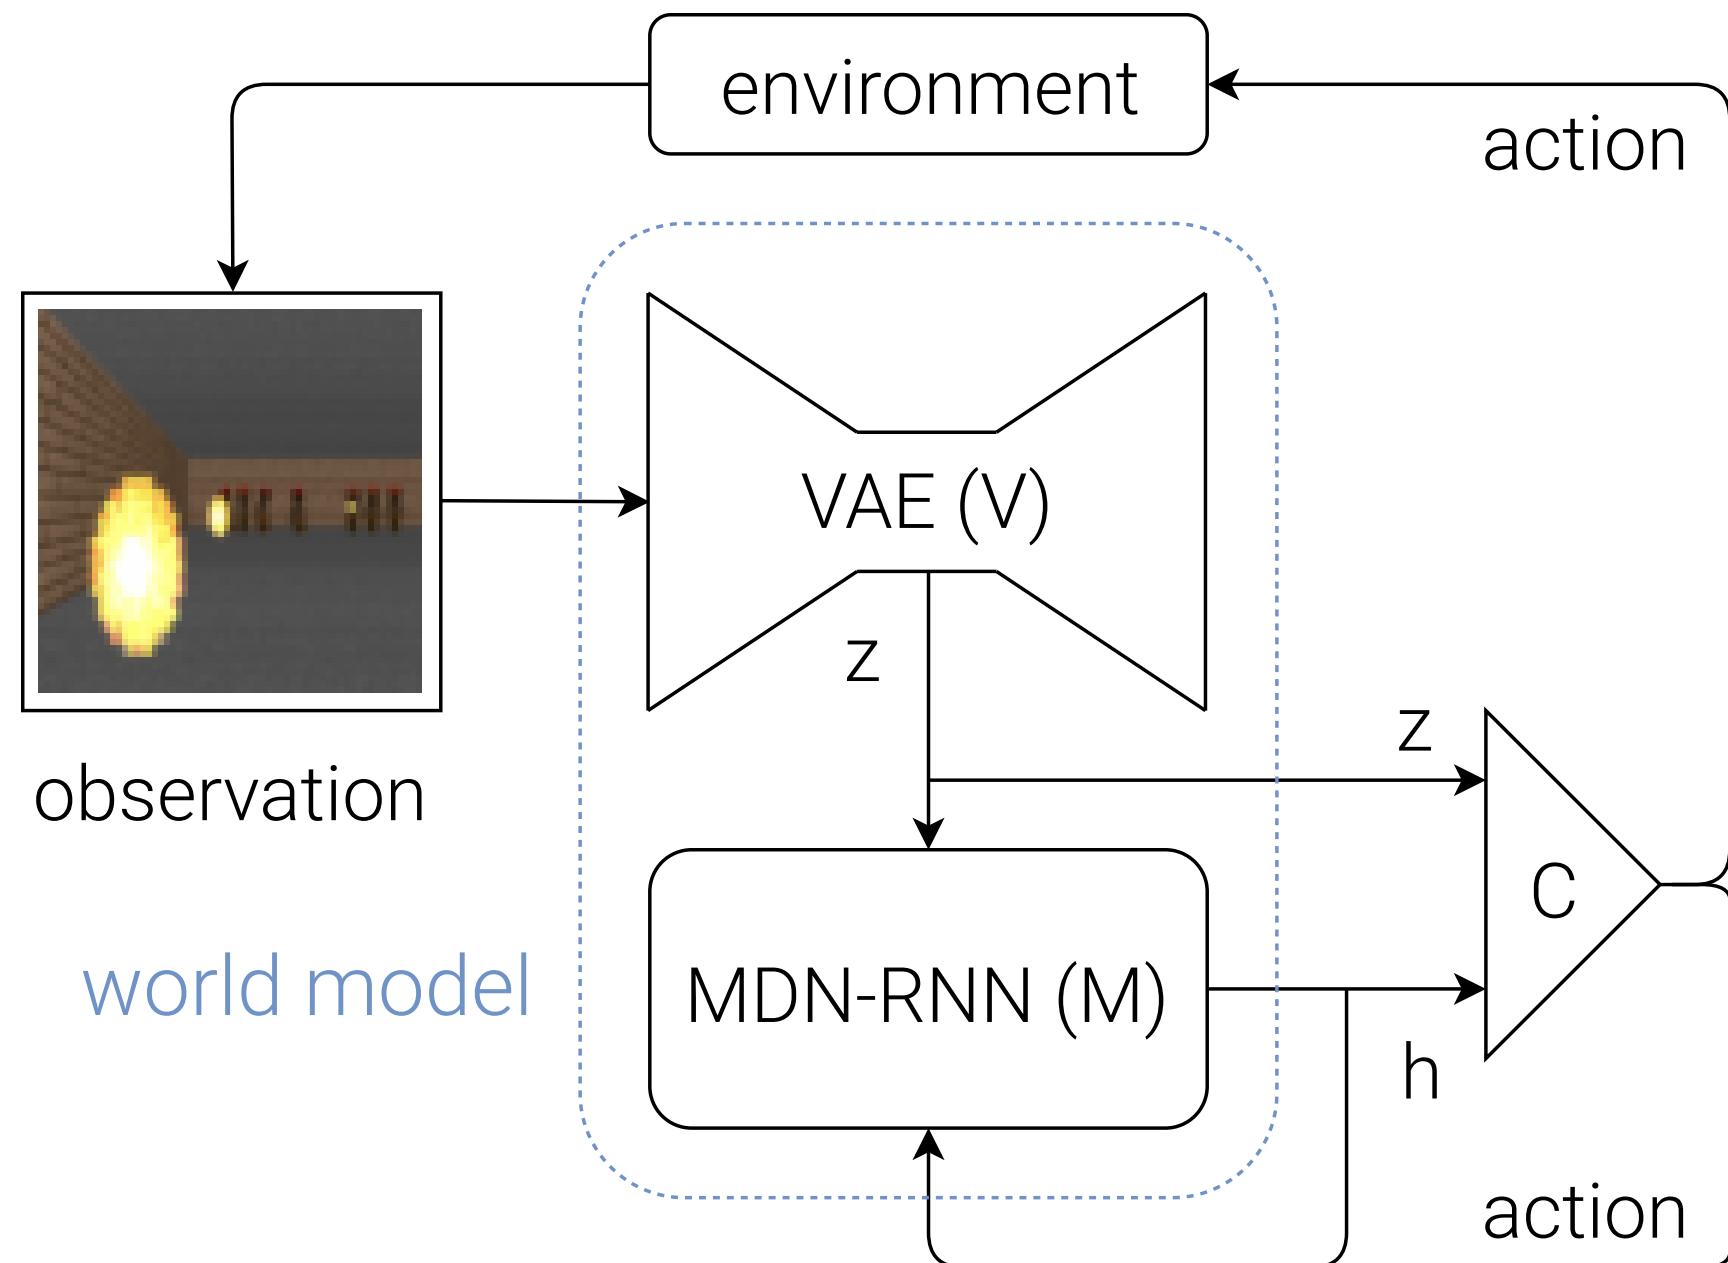
Model	Parameter Count
VAE	4,348,547
MDN-RNN	422,368
Controller	867

World models : car racing

Method	Average Score over 100 Random Tracks
DQN [53]	343 ± 18
A3C (continuous) [52]	591 ± 45
A3C (discrete) [51]	652 ± 10
ceobillionaire's algorithm (unpublished) [47]	838 ± 11
V model only, z input	632 ± 251
V model only, z input with a hidden layer	788 ± 141
Full World Model, z and h	906 ± 21

<https://worldmodels.github.io/>

World models



- The **world model V+M** is learned **offline** with a random agent, using unsupervised learning.
- The **controller C** has few weights (1000) and can be trained by evolutionary algorithms, not even RL.
- The network can even learn by playing entirely in its **own imagination** as the world model can be applied on itself and predict all future frames.
- It just need to additionally predict the reward.
- The learned policy can then be transferred to the real environment.

<https://worldmodels.github.io/>

4 - Deep Planning Network - PlaNet

Learning Latent Dynamics for Planning from Pixels

Danijar Hafner^{1,2} **Timothy Lillicrap³** **Ian Fischer⁴** **Ruben Villegas^{1,5}**
David Ha¹ **Honglak Lee¹** **James Davidson¹**


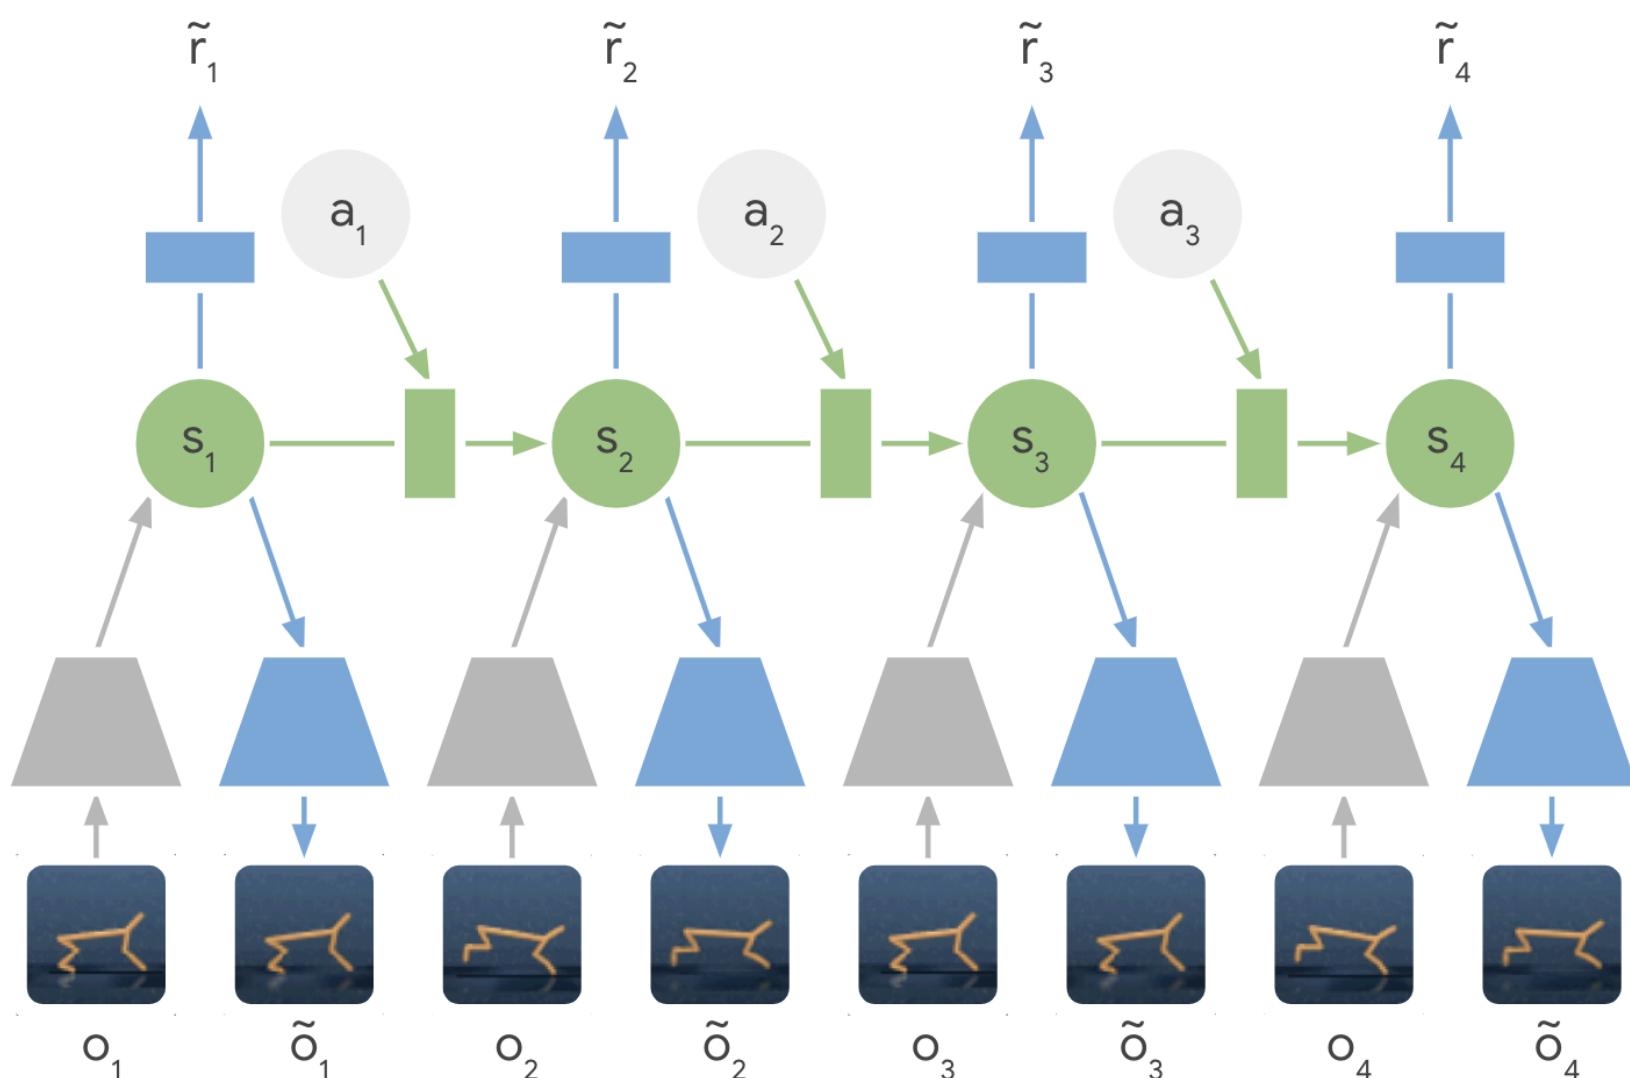
PlaNet

- PlaNet extends the idea of World models by learning the model together with the policy (**end-to-end**).
- It learns a **latent dynamics model** that takes the past observations o_t into account (needed for POMDPs):

$$s_t, r_{t+1}, \hat{o}_t = f(o_t, a_t, s_{t-1})$$

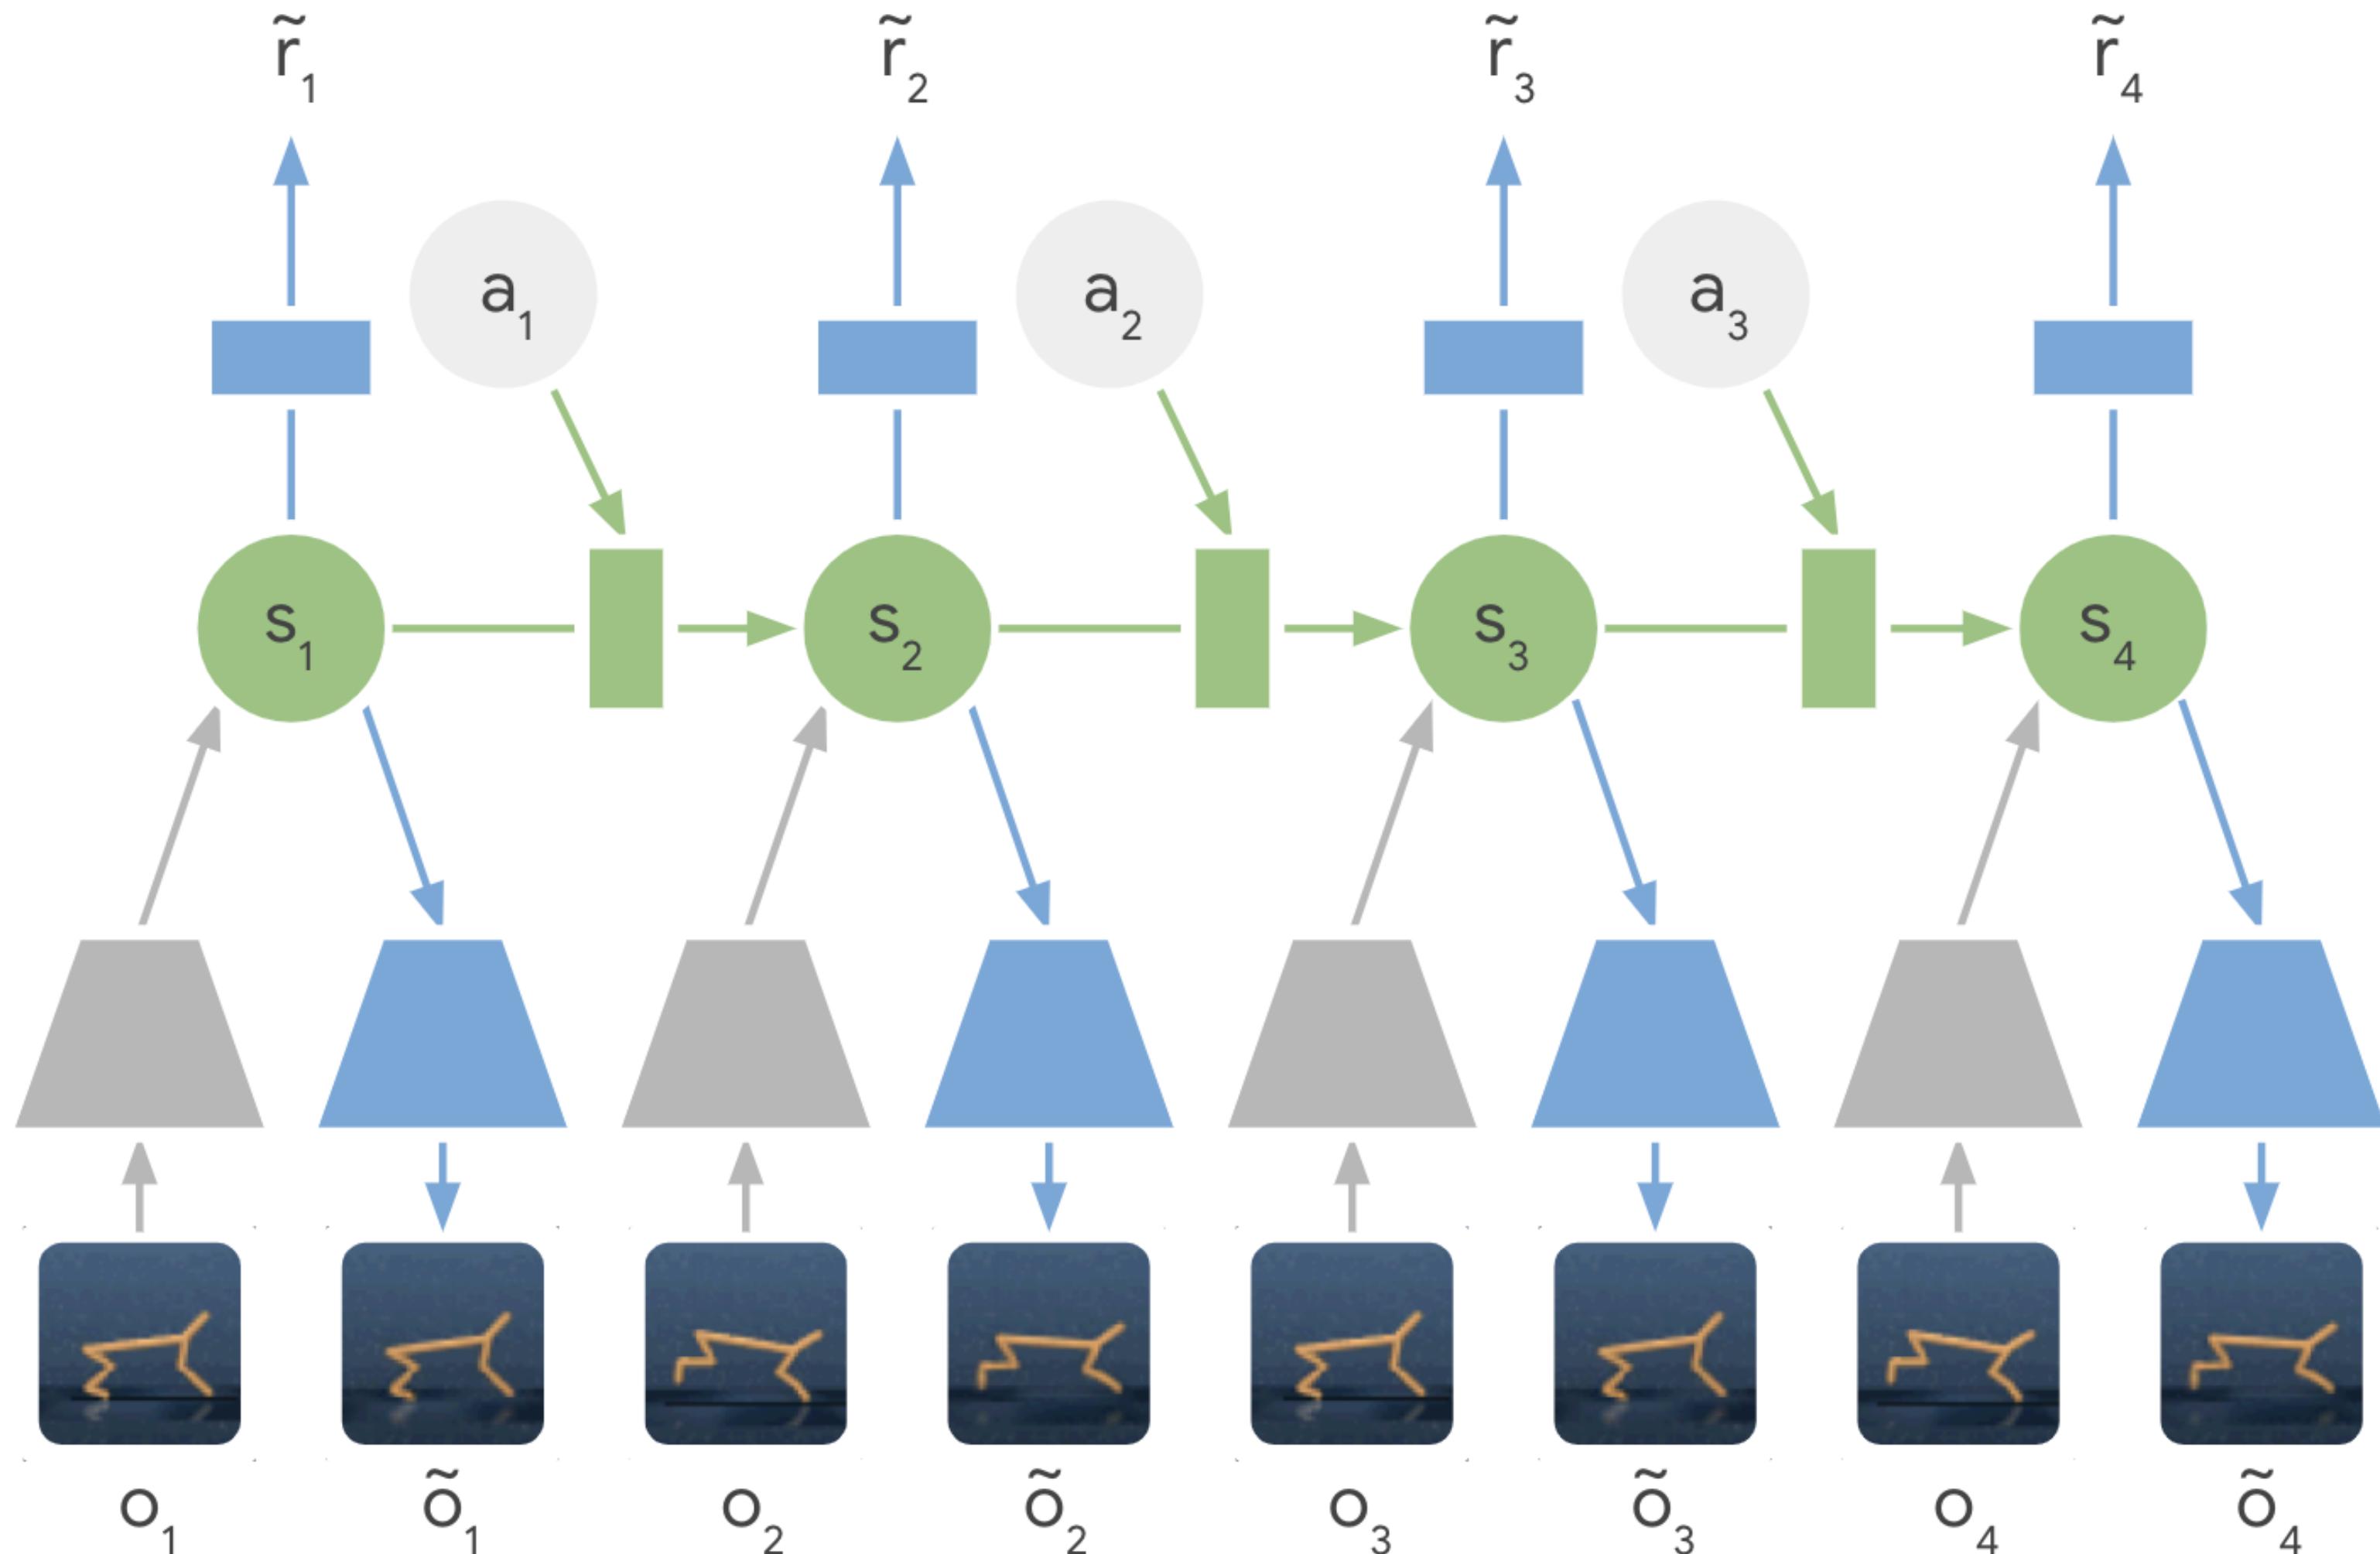
and plans in the latent space using multiple rollouts:

$$a_t = \arg \max_a \mathbb{E}[R(s_t, a, s_{t+1}, \dots)]$$



Source: <https://planetrl.github.io/>

PlaNet: latent dynamics model



Source: <https://ai.googleblog.com/2019/02/introducing-planet-deep-planning.html>

PlaNet: latent dynamics model

- The latent dynamics model is a sequential variational autoencoder learning concurrently:

- An **encoder** from the observation o_t to the latent space s_t .

$$q(s_t | o_t)$$

- A **decoder** from the latent space to the reconstructed observation \hat{o}_t .

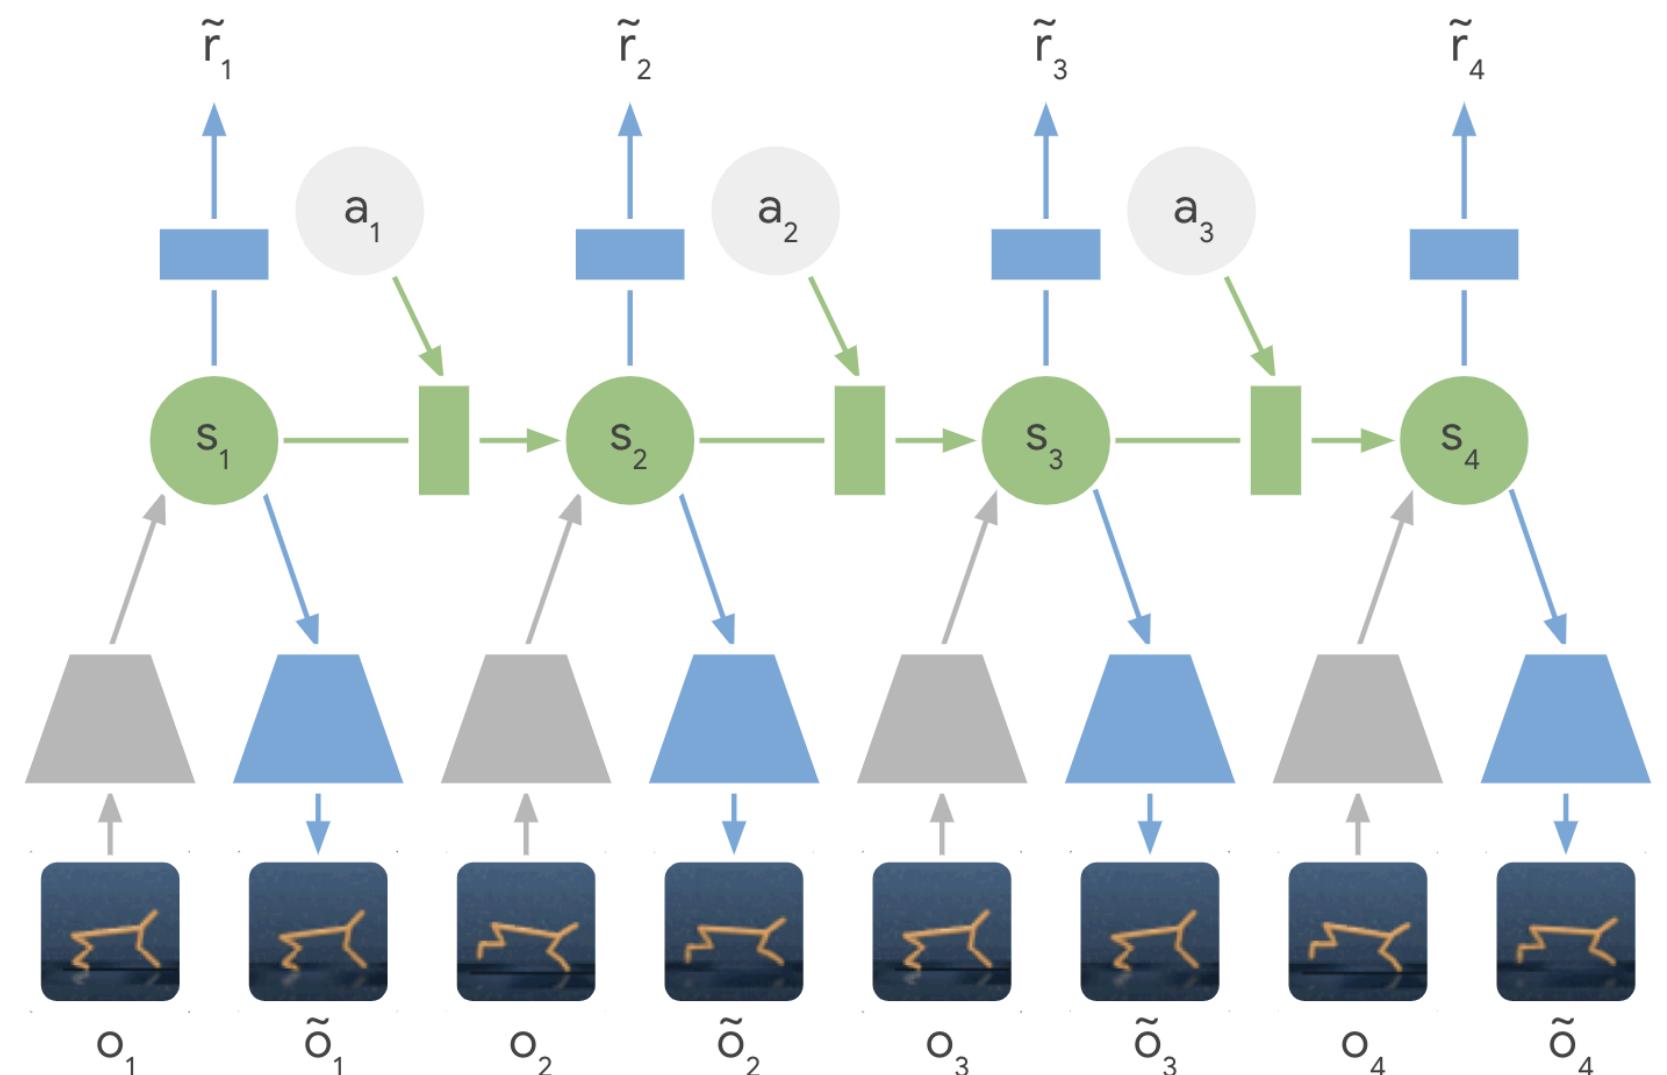
$$p(\hat{o}_t | s_t)$$

- A **transition model** to predict the next latent representation given an action.

$$p(s_{t+1} | s_t, a_t)$$

- A **reward model** predicting the immediate reward.


$$p(r_t | s_t)$$



Source: <https://ai.googleblog.com/2019/02/introducing-planet-deep-planning.html>

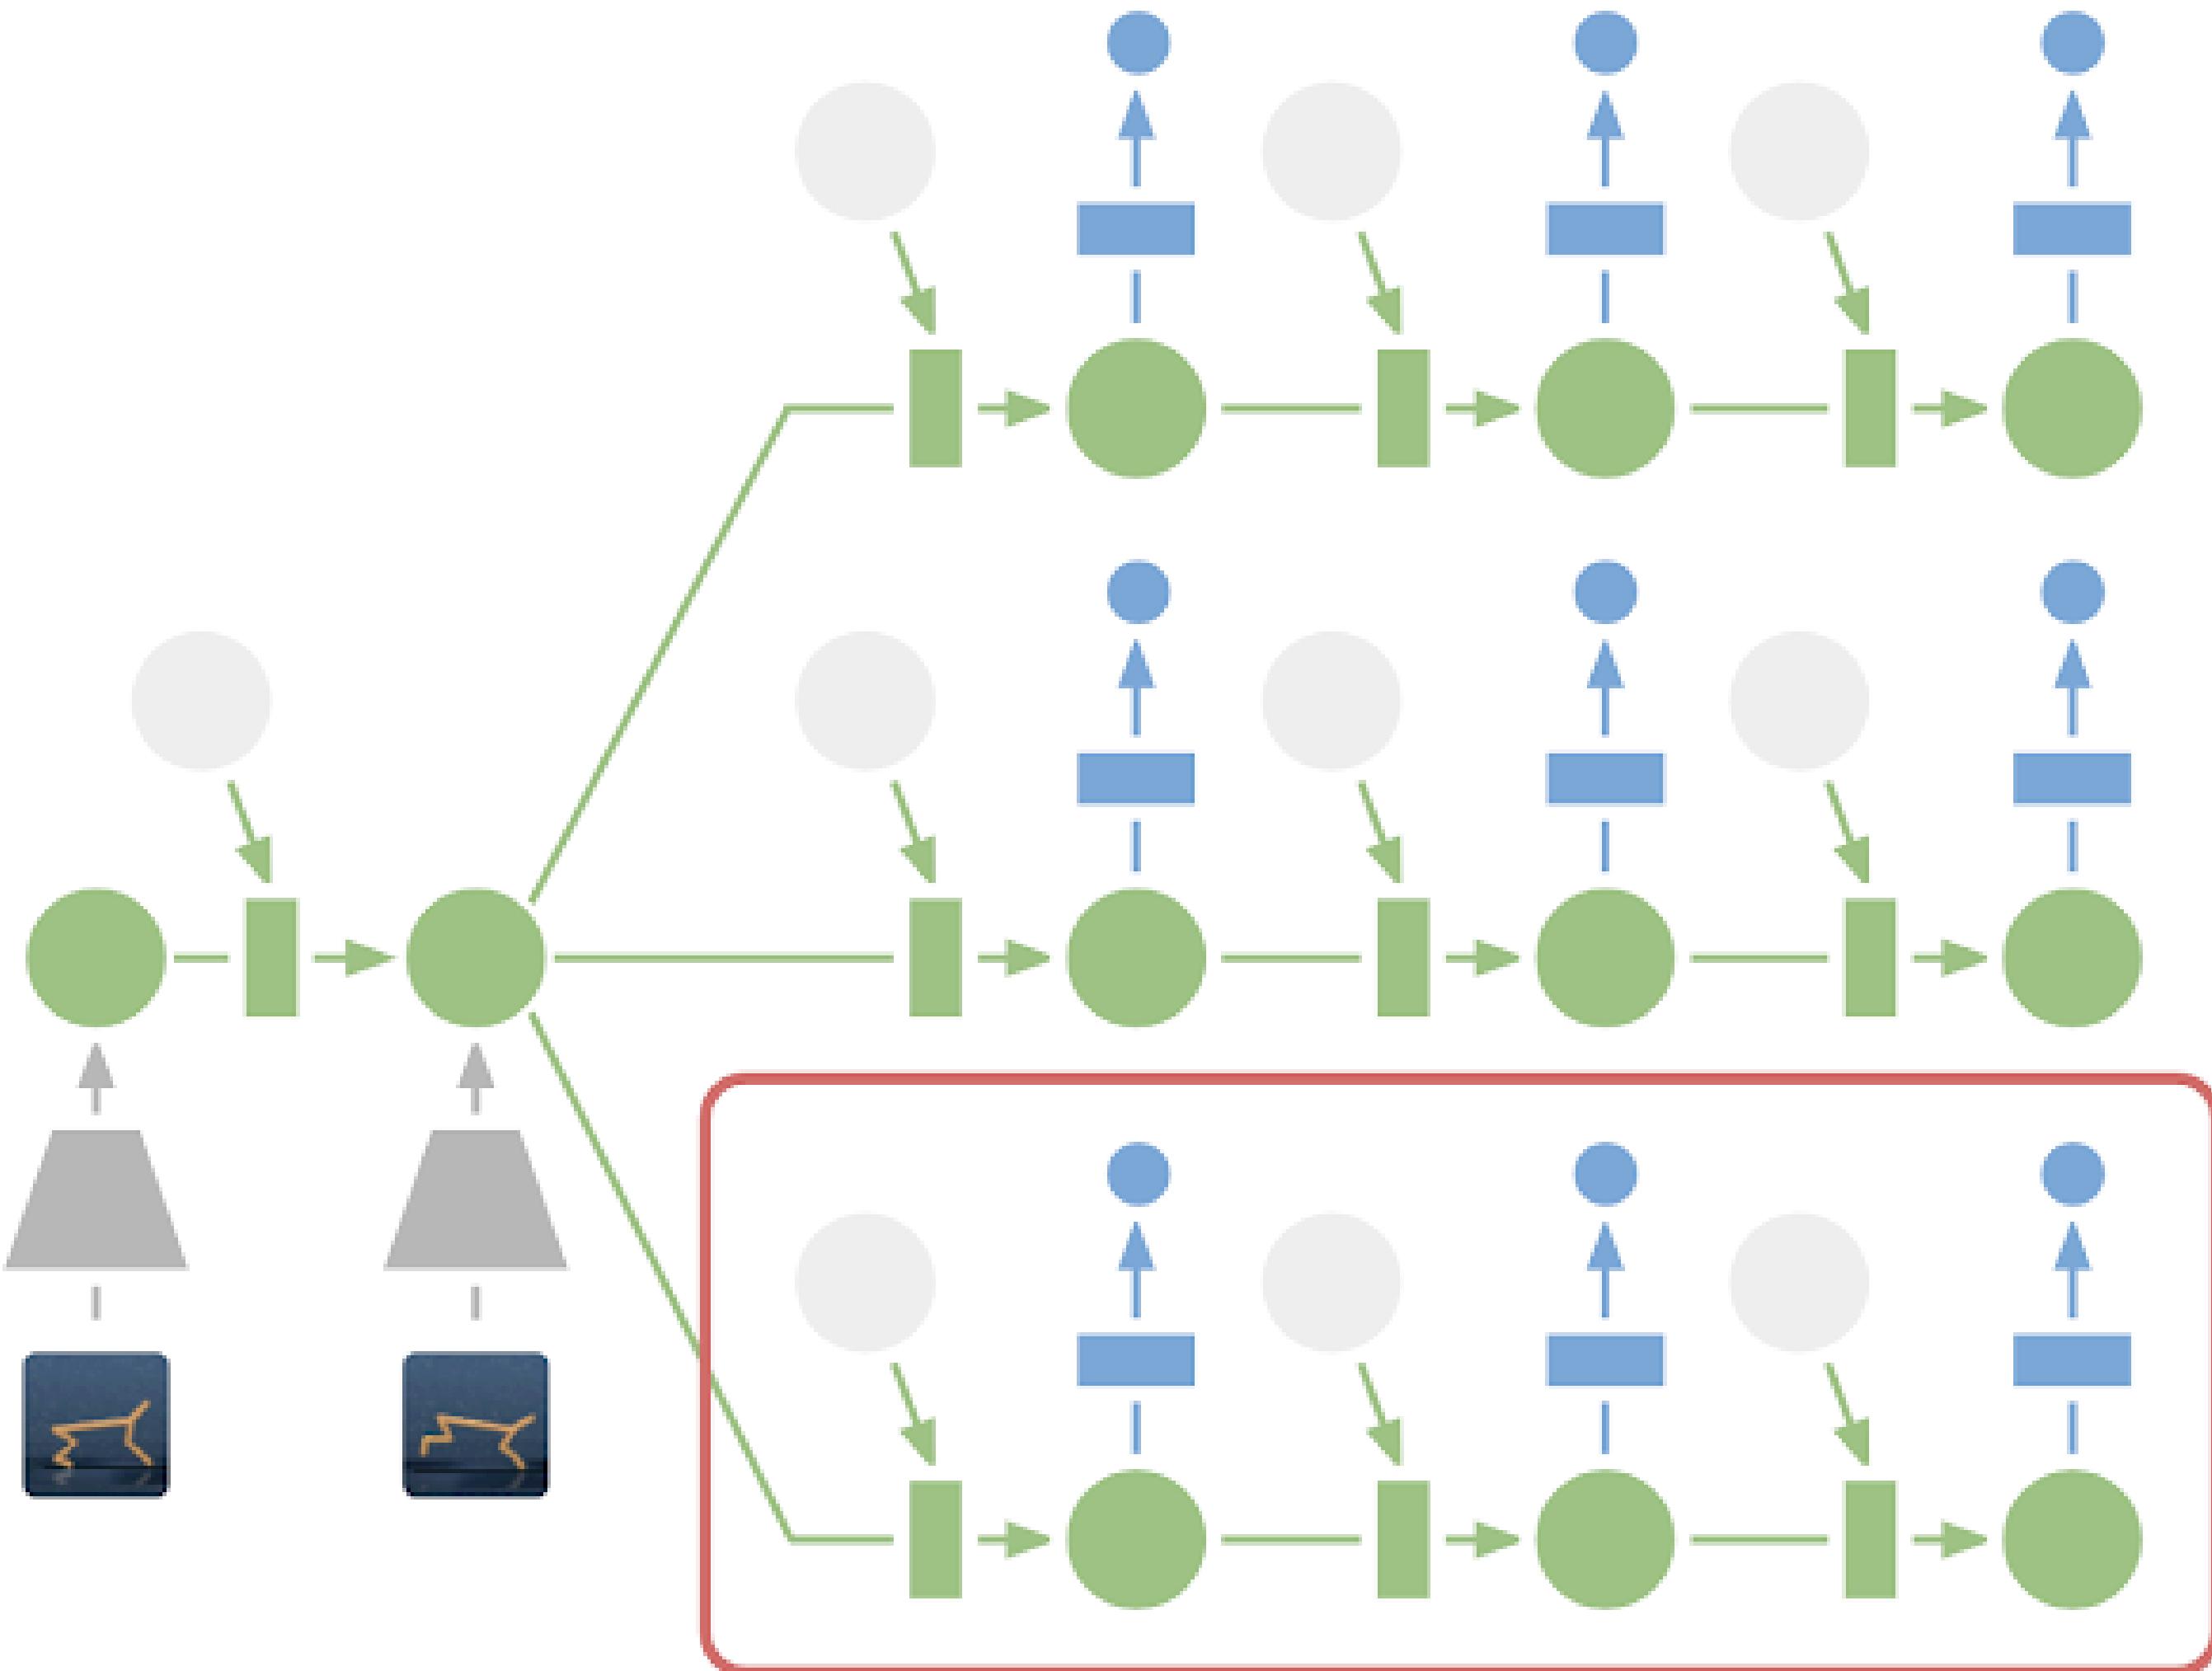
PlaNet: latent dynamics model

- Training sequences $(o_1, a_1, o_2, \dots, o_T)$ can be generated **off-policy** (e.g. from demonstrations) or on-policy.



Source: <https://ai.googleblog.com/2020/03/introducing-dreamer-scalable.html>

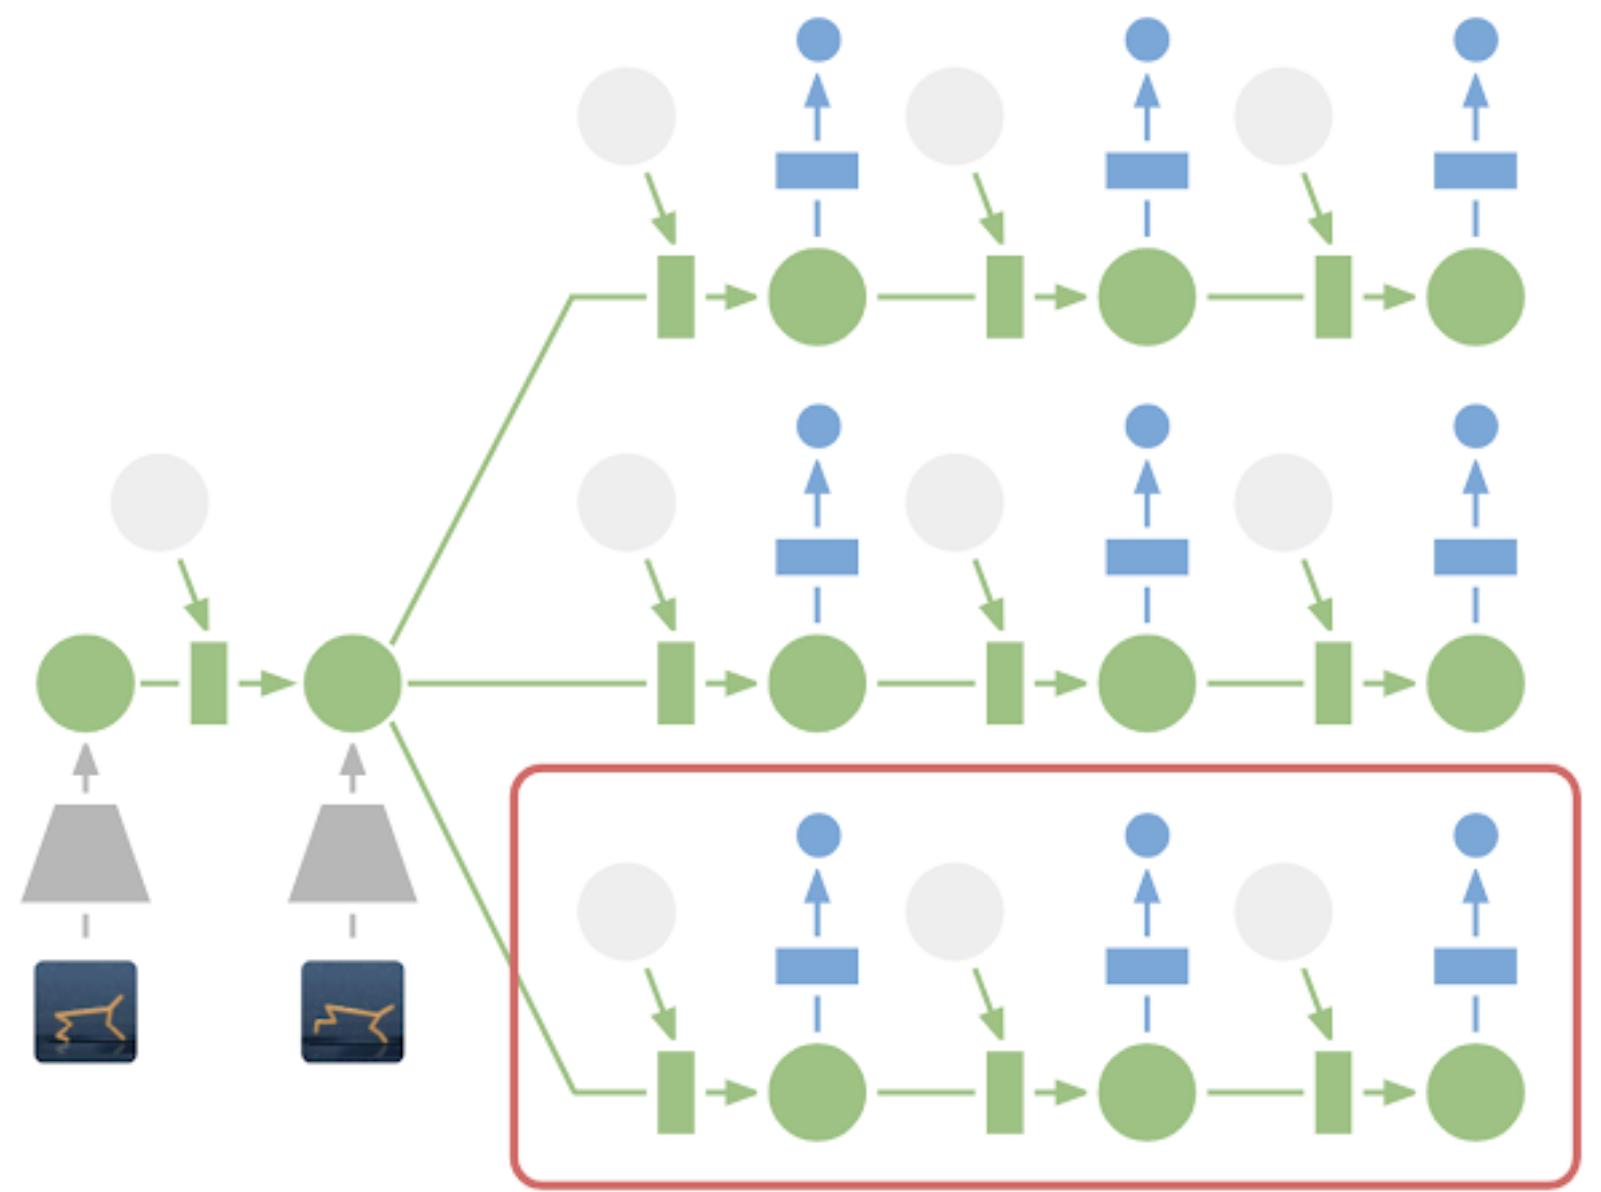
PlaNet: latent space planning



Source: <https://ai.googleblog.com/2019/02/introducing-planet-deep-planning.html>

PlaNet: latent space planning

- From a single observation o_t encoded into s_t , 10000 rollouts are generated using **random sampling**.
- A belief over action sequences is updated using the **cross-entropy method** (CEM) in order to restrict the search.
- The first action of the sequence with the highest estimated return (reward model) is executed.
- At the next time step, planning starts from scratch: Model Predictive Control.
- There is no actor in PlaNet, only a transition model used for planning.



Source: <https://ai.googleblog.com/2019/02/introducing-planet-deep-planning.html>


PlaNet results

- Planet learns continuous image-based control problems in 2000 episodes, where D4PG needs 50 times more.



PlaNet results

- The latent dynamics model can learn 6 control tasks **at the same time**.
- As there is no actor, but only a planner, the same network can control all agents!



Source: <https://ai.googleblog.com/2019/02/introducing-planet-deep-planning.html>

5 - Dreamer

Published as a conference paper at ICLR 2020

DREAM TO CONTROL: LEARNING BEHAVIORS BY LATENT IMAGINATION

Danijar Hafner *

University of Toronto

Google Brain

Timothy Lillicrap

DeepMind

Jimmy Ba

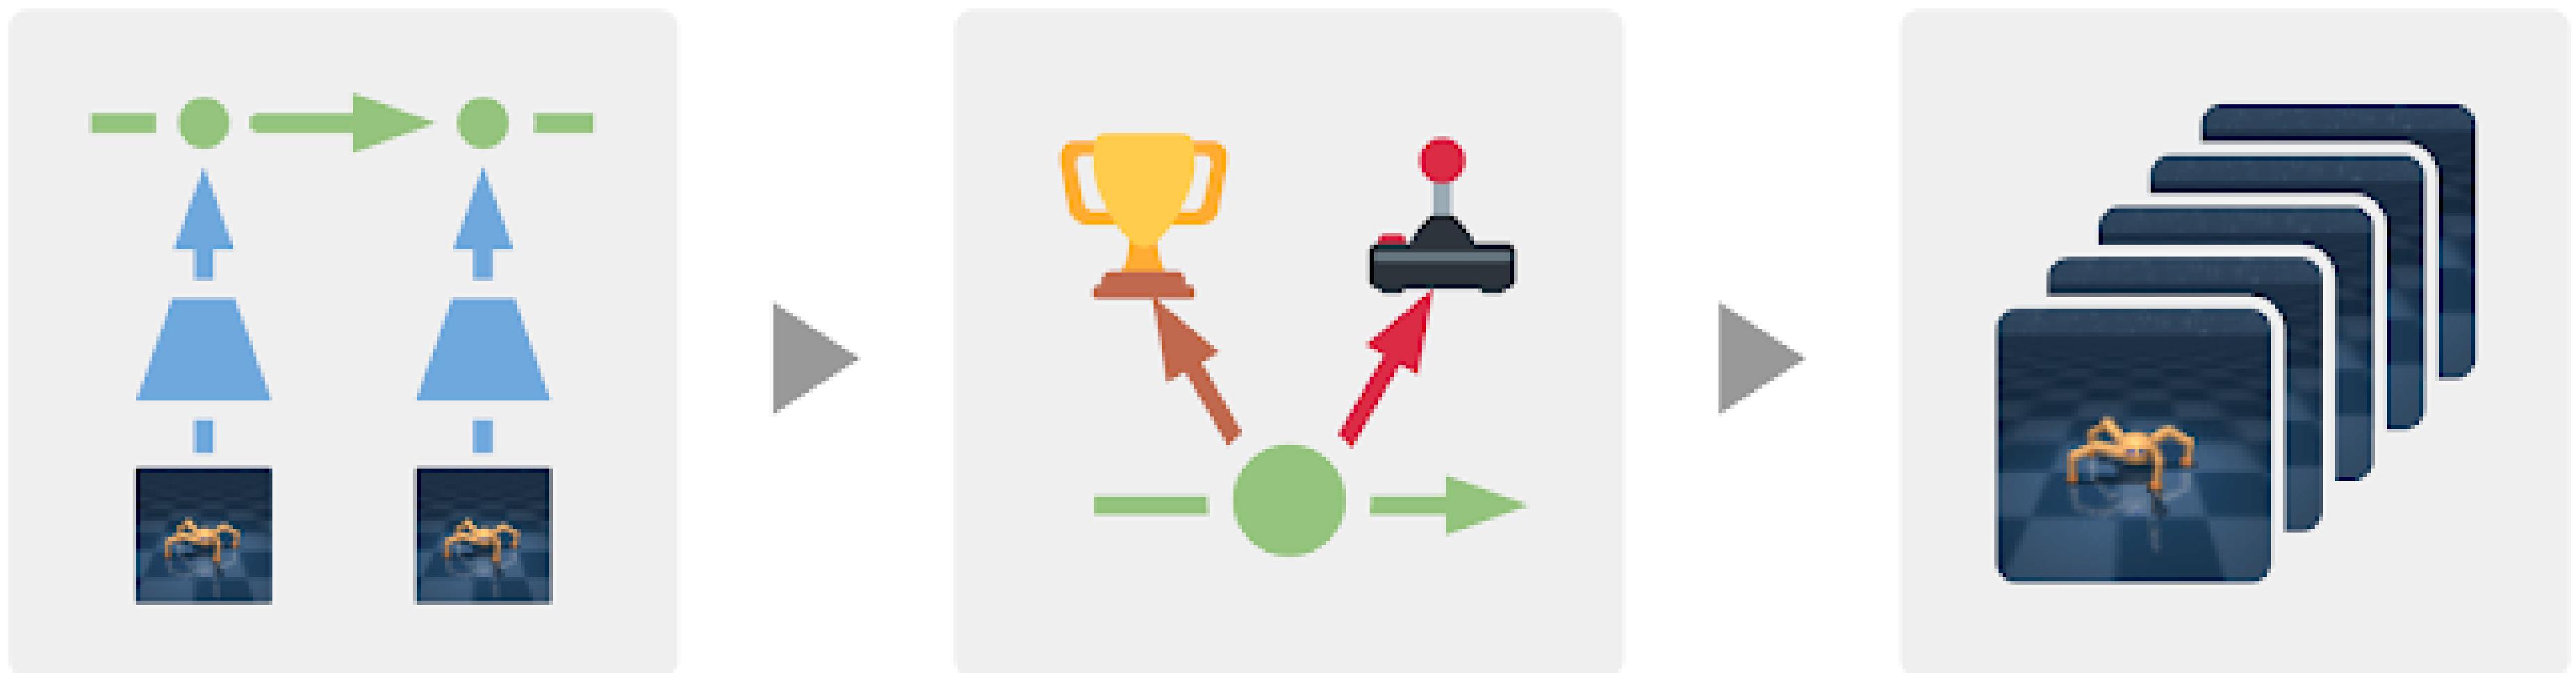
University of Toronto

Mohammad Norouzi

Google Brain

Dreamer

- Dreamer extends the idea of PlaNet by additionally **training an actor** instead of using a MPC planner.
- The latent dynamics model is the same RSSM architecture.
- Training a “model-free” actor on imaginary rollouts instead of MPC planning should reduce the computational time.



World Model
Learning


Learning Value and
Actor Networks

Environment
Interaction

Source: <https://ai.googleblog.com/2020/03/introducing-dreamer-scalable.html>

Dreamer: latent dynamics model

- The latent dynamics model is the same as in PlaNet, learning from past experiences.



Source: <https://ai.googleblog.com/2020/03/introducing-dreamer-scalable.html>

Dreamer: behavior module

- The behavior module learns to predict the value of a state $V_\varphi(s)$ and the policy $\pi_\theta(s)$ (actor-critic).
- It is trained **in imagination** in the latent space using the reward model for the immediate rewards (to compute returns) and the transition model for the next states.



o_1

Source: <https://ai.googleblog.com/2020/03/introducing-dreamer-scalable.html>

- The current observation o_t is encoded into a state s_t , the actor selects an action a_t , the transition model predicts s_{t+1} , the reward model predicts r_{t+1} , the critic predicts $V_\varphi(s_t)$.
- At the end of the sequence, we apply **backpropagation-through-time** to train the actor and the critic.

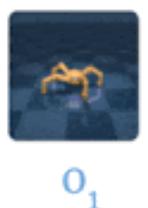
Dreamer: behavior module

- The **critic** $V_\varphi(s_t)$ is trained on the imaginary sequence $(s_t, a_t, r_{t+1}, s_{t+1}, \dots, s_T)$ to minimize the prediction error with the λ -return:

$$R_t^\lambda = (1 - \lambda) \sum_{n=1}^{T-t-1} \lambda^{n-1} R_t^n + \lambda^{T-t-1} R_t$$

- The **actor** $\pi_\theta(s_t, a_t)$ is trained on the sequence to maximize the sum of the value of the future states:


$$\mathcal{J}(\theta) = \mathbb{E}_{s_t, a_t \sim \pi_\theta} \left[\sum_{t'=t}^T V_\varphi(s_{t'}) \right]$$




Source: <https://ai.googleblog.com/2020/03/introducing-dreamer-scalable.html>

Dreamer


- The main advantage of training an actor is that we need only one rollout when training it: backpropagation maximizes the expected returns.
- When acting, we just need to encode the history of the episode in the latent space, and the actor becomes model-free!



(a) Learn dynamics from experience




(b) Learn behavior in imagination



(c) Act in the environment


Dreamer results

- Dreamer beats model-free and model-based methods on 20 continuous control tasks.



Sparse Cartpole Acrobot Swingup Hopper Hop Walker Run Quadruped Run

■ Dreamer (5e6 steps) ■ PlaNet (5e6 steps) ■ D4PG (1e8 steps) ■ A3C (1e8 steps, proprio)



Source: <https://ai.googleblog.com/2020/03/introducing-dreamer-scalable.html>

Dreamer results

- It also learns Atari and Deepmind lab video games, sometimes on par with Rainbow or IMPALA!

