



UNIVERSITY OF TECHNOLOGY  
IN THE EUROPEAN CAPITAL OF CULTURE  
CHEMNITZ

# Deep Reinforcement Learning


Dynamic Programming

Julien Vitay

Professur für Künstliche Intelligenz - Fakultät für Informatik

<https://tu-chemnitz.de/informatik/KI/edu/deepri>

# Dynamic Programming (DP)



- Dynamic Programming (DP) iterates over two steps:

## 1. Policy evaluation

- For a given policy  $\pi$ , the value of all states  $V^\pi(s)$  or all state-action pairs  $Q^\pi(s, a)$  is calculated based on the Bellman equations:

$$V^\pi(s) = \sum_{a \in \mathcal{A}(s)} \pi(s, a) \sum_{s' \in \mathcal{S}} p(s'|s, a) [r(s, a, s') + \gamma V^\pi(s')]$$

## 2. Policy improvement

- From the current estimated values  $V^\pi(s)$  or  $Q^\pi(s, a)$ , a new **better** policy  $\pi$  is derived.

- After enough iterations, the policy converges to the **optimal policy** (if the states are Markov).
- Two main algorithms: **policy iteration** and **value iteration**.

# 1 - Policy iteration

# Policy evaluation

- Bellman equation for the state  $s$  and a fixed policy  $\pi$ :

$$V^\pi(s) = \sum_{a \in \mathcal{A}(s)} \pi(s, a) \sum_{s' \in \mathcal{S}} p(s'|s, a) [r(s, a, s') + \gamma V^\pi(s')]$$

- Let's note  $\mathcal{P}_{ss'}^\pi$  the transition probability between  $s$  and  $s'$  (dependent on the policy  $\pi$ ) and  $\mathcal{R}_s^\pi$  the expected reward in  $s$  (also dependent):

$$\mathcal{P}_{ss'}^\pi = \sum_{a \in \mathcal{A}(s)} \pi(s, a) \sum_{s' \in \mathcal{S}} p(s'|s, a)$$

$$\mathcal{R}_s^\pi = \sum_{a \in \mathcal{A}(s)} \pi(s, a) \sum_{s' \in \mathcal{S}} p(s'|s, a) r(s, a, s')$$

- The Bellman equation becomes:

$$V^\pi(s) = \mathcal{R}_s^\pi + \gamma \sum_{s' \in \mathcal{S}} \mathcal{P}_{ss'}^\pi V^\pi(s')$$

- As we have a fixed policy during the evaluation (Markov Reward Process), the Bellman equation is simplified.

# Policy evaluation

- Let's now put the Bellman equations in a matrix-vector form.

$$V^\pi(s) = \mathcal{R}_s^\pi + \gamma \sum_{s' \in \mathcal{S}} \mathcal{P}_{ss'}^\pi V^\pi(s')$$

- We first define the **vector of state values**  $\mathbf{V}^\pi$ :
- and the **vector of expected reward**  $\mathcal{R}^\pi$ :

$$\mathbf{V}^\pi = \begin{bmatrix} V^\pi(s_1) \\ V^\pi(s_2) \\ \vdots \\ V^\pi(s_n) \end{bmatrix}$$

$$\mathcal{R}^\pi = \begin{bmatrix} \mathcal{R}^\pi(s_1) \\ \mathcal{R}^\pi(s_2) \\ \vdots \\ \mathcal{R}^\pi(s_n) \end{bmatrix}$$

- The **state transition matrix**  $\mathcal{P}^\pi$  is defined as:

$$\mathcal{P}^\pi = \begin{bmatrix} \mathcal{P}_{s_1 s_1}^\pi & \mathcal{P}_{s_1 s_2}^\pi & \dots & \mathcal{P}_{s_1 s_n}^\pi \\ \mathcal{P}_{s_2 s_1}^\pi & \mathcal{P}_{s_2 s_2}^\pi & \dots & \mathcal{P}_{s_2 s_n}^\pi \\ \vdots & \vdots & \vdots & \vdots \\ \mathcal{P}_{s_n s_1}^\pi & \mathcal{P}_{s_n s_2}^\pi & \dots & \mathcal{P}_{s_n s_n}^\pi \end{bmatrix}$$

# Policy evaluation

- You can simply check that:

$$\begin{bmatrix} V^\pi(s_1) \\ V^\pi(s_2) \\ \vdots \\ V^\pi(s_n) \end{bmatrix} = \begin{bmatrix} \mathcal{R}^\pi(s_1) \\ \mathcal{R}^\pi(s_2) \\ \vdots \\ \mathcal{R}^\pi(s_n) \end{bmatrix} + \gamma \begin{bmatrix} \mathcal{P}_{s_1 s_1}^\pi & \mathcal{P}_{s_1 s_2}^\pi & \dots & \mathcal{P}_{s_1 s_n}^\pi \\ \mathcal{P}_{s_2 s_1}^\pi & \mathcal{P}_{s_2 s_2}^\pi & \dots & \mathcal{P}_{s_2 s_n}^\pi \\ \vdots & \vdots & \vdots & \vdots \\ \mathcal{P}_{s_n s_1}^\pi & \mathcal{P}_{s_n s_2}^\pi & \dots & \mathcal{P}_{s_n s_n}^\pi \end{bmatrix} \times \begin{bmatrix} V^\pi(s_1) \\ V^\pi(s_2) \\ \vdots \\ V^\pi(s_n) \end{bmatrix}$$

leads to the same equations as:

$$V^\pi(s) = \mathcal{R}_s^\pi + \gamma \sum_{s' \in \mathcal{S}} \mathcal{P}_{ss'}^\pi V^\pi(s')$$

for all states  $s$ .

- The Bellman equations for all states  $s$  can therefore be written with a matrix-vector notation as:

$$\mathbf{V}^\pi = \mathcal{R}^\pi + \gamma \mathcal{P}^\pi \mathbf{V}^\pi$$

# Policy evaluation

- The Bellman equations for all states  $s$  is:

$$\mathbf{V}^\pi = \mathcal{R}^\pi + \gamma \mathcal{P}^\pi \mathbf{V}^\pi$$

- If we know  $\mathcal{P}^\pi$  and  $\mathcal{R}^\pi$  (dynamics of the MDP for the policy  $\pi$ ), we can simply obtain the state values:

$$(\mathbb{I} - \gamma \mathcal{P}^\pi) \times \mathbf{V}^\pi = \mathcal{R}^\pi$$

where  $\mathbb{I}$  is the identity matrix, what gives:


$$\mathbf{V}^\pi = (\mathbb{I} - \gamma \mathcal{P}^\pi)^{-1} \times \mathcal{R}^\pi$$

- Done!
- **But**, if we have  $n$  states, the matrix  $\mathcal{P}^\pi$  has  $n^2$  elements.
- Inverting  $\mathbb{I} - \gamma \mathcal{P}^\pi$  requires at least  $\mathcal{O}(n^{2.37})$  operations.
- Forget it if you have more than a thousand states ( $1000^{2.37} \approx 13$  million operations).
- In **dynamic programming**, we will use **iterative methods** to estimate  $\mathbf{V}^\pi$ .

# Iterative policy evaluation

- The idea of **iterative policy evaluation** (IPE) is to consider a sequence of consecutive state-value functions which should converge from initially wrong estimates  $V_0(s)$  towards the real state-value function  $V^\pi(s)$ .

$$V_0 \rightarrow V_1 \rightarrow V_2 \rightarrow \dots \rightarrow V_k \rightarrow V_{k+1} \rightarrow \dots \rightarrow V^\pi$$



- The value function at step  $k + 1$   $V_{k+1}(s)$  is computed using the previous estimates  $V_k(s)$  and the Bellman equation transformed into an **update rule**.
- In vector notation:

$$\mathbf{V}_{k+1} = \mathcal{R}^\pi + \gamma \mathcal{P}^\pi \mathbf{V}_k$$

Source: David Silver.

<http://www0.cs.ucl.ac.uk/staff/d.silver/web/Teaching.html>



# Iterative policy evaluation

- Let's start with dummy (e.g. random) initial estimates  $V_0(s)$  for the value of every state  $s$ .
- We can obtain new estimates  $V_1(s)$  which are slightly less wrong by applying once the **Bellman operator**:

$$V_1(s) \leftarrow \sum_{a \in \mathcal{A}(s)} \pi(s, a) \sum_{s' \in \mathcal{S}} p(s'|s, a) [r(s, a, s') + \gamma V_0(s')] \quad \forall s \in \mathcal{S}$$

- Based on these estimates  $V_1(s)$ , we can obtain even better estimates  $V_2(s)$  by applying again the Bellman operator:

$$V_2(s) \leftarrow \sum_{a \in \mathcal{A}(s)} \pi(s, a) \sum_{s' \in \mathcal{S}} p(s'|s, a) [r(s, a, s') + \gamma V_1(s')] \quad \forall s \in \mathcal{S}$$

- Generally, state-value function estimates are improved iteratively through:

$$V_{k+1}(s) \leftarrow \sum_{a \in \mathcal{A}(s)} \pi(s, a) \sum_{s' \in \mathcal{S}} p(s'|s, a) [r(s, a, s') + \gamma V_k(s')] \quad \forall s \in \mathcal{S}$$

- $V_\infty = V^\pi$  is a fixed point of this update rule because of the uniqueness of the solution to the Bellman equation.

# Bellman operator

- The **Bellman operator**  $\mathcal{T}^\pi$  is a mapping between two vector spaces:

$$\mathcal{T}^\pi(\mathbf{V}) = \mathcal{R}^\pi + \gamma \mathcal{P}^\pi \mathbf{V}$$

- If you apply repeatedly the Bellman operator on any initial vector  $\mathbf{V}_0$ , it converges towards the solution of the Bellman equations  $\mathbf{V}^\pi$ .
- Mathematically speaking,  $\mathcal{T}^\pi$  is a  $\gamma$ -contraction, i.e. it makes value functions closer by at least  $\gamma$ :

$$\|\mathcal{T}^\pi(\mathbf{V}) - \mathcal{T}^\pi(\mathbf{U})\|_\infty \leq \gamma \|\mathbf{V} - \mathbf{U}\|_\infty$$


- The **contraction mapping theorem** ensures that  $\mathcal{T}^\pi$  converges to a unique fixed point:
  - existence and uniqueness of the solution of the Bellman equations.

# Backup diagram of IPE

- Iterative Policy Evaluation relies on **full backups**: it backs up the value of ALL possible successive states into the new value of a state.

$$V_{k+1}(s) \leftarrow \sum_{a \in \mathcal{A}(s)} \pi(s, a) \sum_{s' \in \mathcal{S}} p(s'|s, a) [r(s, a, s') + \gamma V_k(s')] \quad \forall s \in \mathcal{S}$$

- Backup diagram**: which other values do you need to know in order to update one value?



- The backups are **synchronous**: all states are backed up in parallel.


$$\mathbf{V}_{k+1} = \mathcal{R}^\pi + \gamma \mathcal{P}^\pi \mathbf{V}_k$$

- The termination of iterative policy evaluation has to be controlled by hand, as the convergence of the algorithm is only at the limit.
- It is good practice to look at the variations on the values of the different states, and stop the iteration when this variation falls below a predefined threshold.

# Iterative policy evaluation

- For a fixed policy  $\pi$ , initialize  $V(s) = 0 \ \forall s \in \mathcal{S}$ .
- **while** not converged:
  - **for** all states  $s$ :
    - $V_{\text{target}}(s) = \sum_{a \in \mathcal{A}(s)} \pi(s, a) \sum_{s' \in \mathcal{S}} p(s'|s, a) [r(s, a, s') + \gamma V(s')]$
  - $\delta = 0$
  - **for** all states  $s$ :
    - $\delta = \max(\delta, |V(s) - V_{\text{target}}(s)|)$
    - $V(s) = V_{\text{target}}(s)$
  - **if**  $\delta < \delta_{\text{threshold}}$ :
    - converged = True

# Dynamic Programming (DP)



- Dynamic Programming (DP) iterates over two steps:

## 1. Policy evaluation

- For a given policy  $\pi$ , the value of all states  $V^\pi(s)$  or all state-action pairs  $Q^\pi(s, a)$  is calculated based on the Bellman equations:

$$V^\pi(s) = \sum_{a \in \mathcal{A}(s)} \pi(s, a) \sum_{s' \in \mathcal{S}} p(s' | s, a) [r(s, a, s') + \gamma V^\pi(s')]$$


## 2. Policy improvement

- From the current estimated values  $V^\pi(s)$  or  $Q^\pi(s, a)$ , a new **better** policy  $\pi$  is derived.

# Policy improvement

- For each state  $s$ , we would like to know if we should deterministically choose an action  $a \neq \pi(s)$  or not in order to improve the policy.
- The value of an action  $a$  in the state  $s$  for the policy  $\pi$  is given by:

$$Q^\pi(s, a) = \sum_{s' \in \mathcal{S}} p(s'|s, a) [r(s, a, s') + \gamma V^\pi(s')]$$



- If the Q-value of an action  $a$  is higher than the one currently selected by the **deterministic** policy:

$$Q^\pi(s, a) > Q^\pi(s, \pi(s)) = V^\pi(s)$$

then it is better to select  $a$  once in  $s$  and thereafter follow  $\pi$ .

- If there is no better action, we keep the previous policy for this state.
- This corresponds to a **greedy** action selection over the Q-values, defining a **deterministic** policy  $\pi(s)$ :

$$\pi(s) \leftarrow \operatorname{argmax}_a Q^\pi(s, a) = \sum_{s' \in \mathcal{S}} p(s'|s, a) [r(s, a, s') + \gamma V^\pi(s')]$$

# Policy improvement

- After the policy improvement, the Q-value of each deterministic action  $\pi(s)$  has increased or stayed the same.

$$\operatorname{argmax}_a Q^\pi(s, a) = \sum_{s' \in \mathcal{S}} p(s'|s, a) [r(s, a, s') + \gamma V^\pi(s')] \geq Q^\pi(s, \pi(s))$$

- This defines an **improved** policy  $\pi'$ , where all states and actions have a higher value than previously.
- **Greedy action selection** over the state value function implements policy improvement:


$$\pi' \leftarrow \text{Greedy}(V^\pi)$$



Greedy policy improvement:


- **for** each state  $s \in \mathcal{S}$ :
  - $\pi(s) \leftarrow \operatorname{argmax}_a \sum_{s' \in \mathcal{S}} p(s'|s, a) [r(s, a, s') + \gamma V^\pi(s')]$

# Policy iteration



- Once a policy  $\pi$  has been improved using  $V^\pi$  to yield a better policy  $\pi'$ , we can then compute  $V^{\pi'}$  and improve it again to yield an even better policy  $\pi''$ .
- The algorithm **policy iteration** successively uses **policy evaluation** and **policy improvement** to find the optimal policy.

$$\pi_0 \xrightarrow{E} V^{\pi_0} \xrightarrow{I} \pi_1 \xrightarrow{E} V^{\pi_1} \xrightarrow{I} \dots \xrightarrow{I} \pi^* \xrightarrow{E} V^*$$




- The **optimal policy** being deterministic, policy improvement can be greedy over the state values.
- If the policy does not change after policy improvement, the optimal policy has been found.



# Policy iteration

- Initialize a deterministic policy  $\pi(s)$  and set  $V(s) = 0 \forall s \in \mathcal{S}$ .
- **while**  $\pi$  is not optimal:
  - **while** not converged: # Policy evaluation
    - **for** all states  $s$ :
      - $V_{\text{target}}(s) = \sum_{a \in \mathcal{A}(s)} \pi(s, a) \sum_{s' \in \mathcal{S}} p(s'|s, a) [r(s, a, s') + \gamma V(s')]$
    - **for** all states  $s$ :
      - $V(s) = V_{\text{target}}(s)$
  - **for** each state  $s \in \mathcal{S}$ : # Policy improvement
    - $\pi(s) \leftarrow \operatorname{argmax}_a \sum_{s' \in \mathcal{S}} p(s'|s, a) [r(s, a, s') + \gamma V^\pi(s')]$
  - **if**  $\pi$  has not changed: **break**

## Small Gridworld example



	1	2	3
4	5	6	7
8	9	10	11
12	13	14	

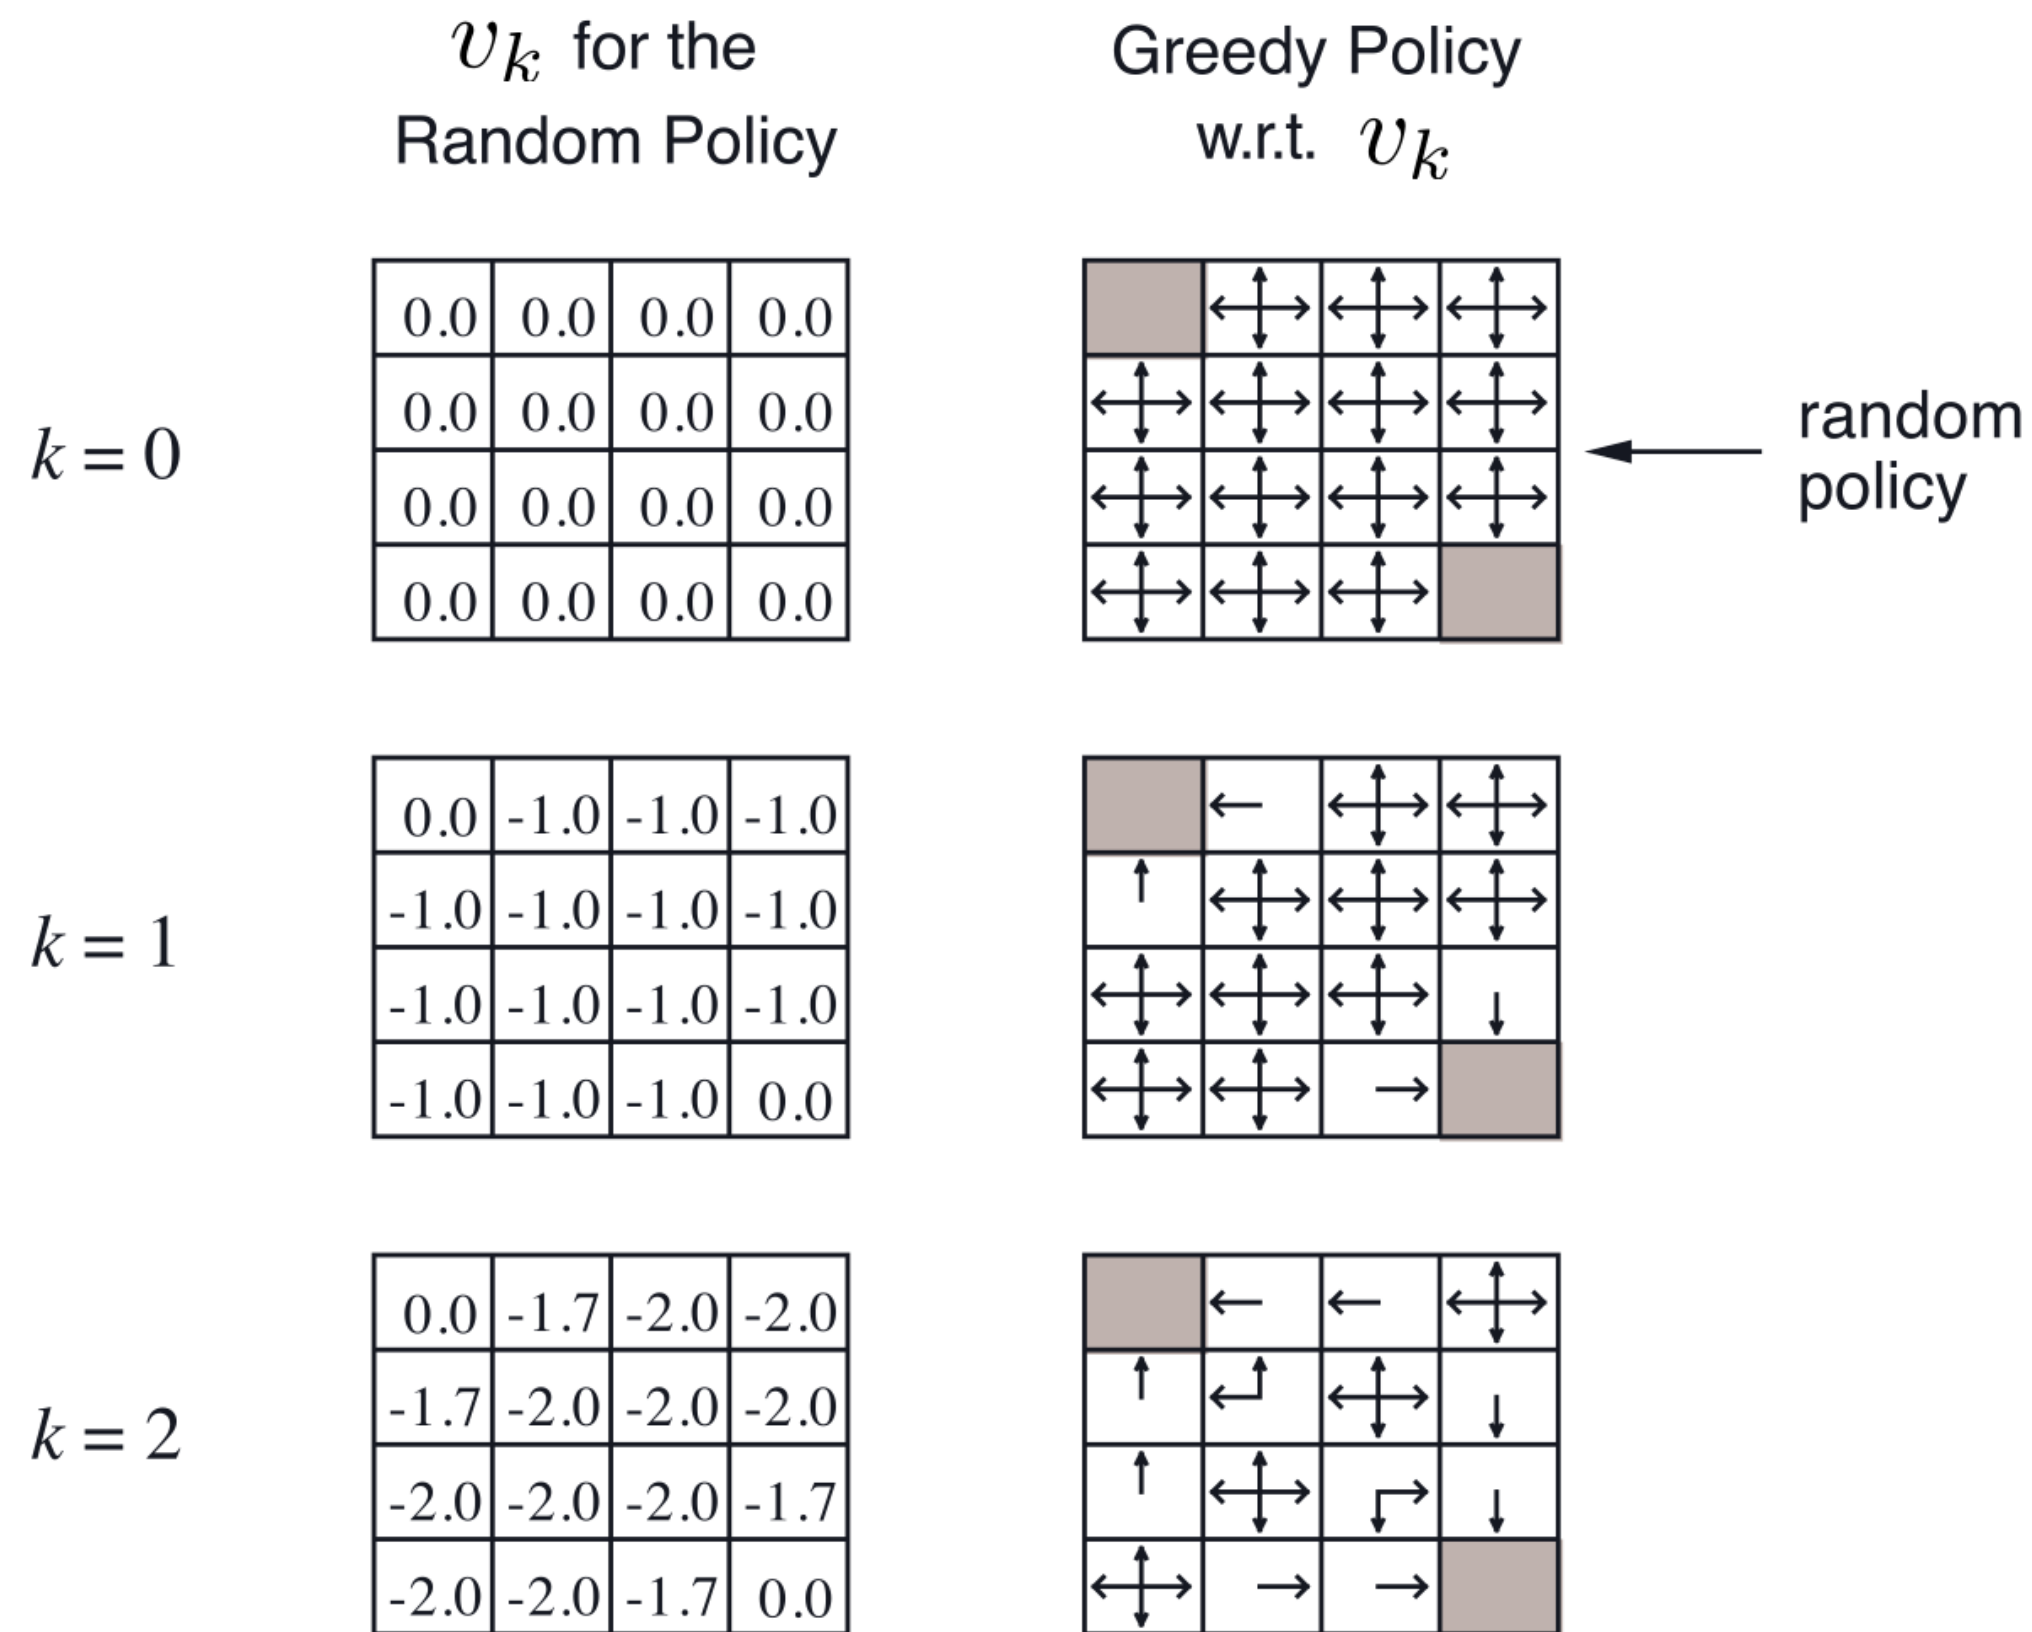
$r = -1$   
on all transitions

Source: David Silver. <http://www0.cs.ucl.ac.uk/staff/d.silver/web/Teaching.html>

- **Gridworld** is an undiscounted MDP (we can take  $\gamma = 1$ ).
- The states are the position in the grid, the actions are up, down, left, right. Transitions to a wall leave in the same state.
- The reward is always -1, except after being in the terminal states in gray ( $r = 0$ ).
- The initial policy is random:

$$\pi(s, \text{up}) = \pi(s, \text{down}) = \pi(s, \text{left}) = \pi(s, \text{right}) = 0.25$$


# Small Gridworld example



- $k = 0$ :
  - The initial values  $V_0$  are set to 0 as the initial policy is random.
- $k = 1$ :
  - The random policy is evaluated: all states get the value of the average immediate reward in that state. -1, except the terminal states (0).
  - The greedy policy is already an improvement over the random policy: adjacent states to the terminal states would decide to go there systematically, as the value is 0 instead of -1.

- $k = 2$ : The previous estimates propagate: states adjacent to the terminal states get a higher value, as there will be less punishments after these states.

# Small Gridworld example



- $k = 3$ :
  - The values continue to propagate.
  - The greedy policy at that step of policy evaluation is already optimal.
- $k > 3$ :
  - The values continue to converge towards the true values.
  - The greedy policy does not change. In this simple example, it is already the optimal policy.

- Two things to notice:
  - There is actually no need to wait until the end of policy evaluation to improve the policy, as the greedy policy might already be optimal.
  - There can be more than one optimal policy: some actions may have the same Q-value: choosing one or other is equally optimal.

## 2 - Value iteration

# Value iteration

- **Policy iteration** can converge in a surprisingly small number of iterations.
- One drawback of *policy iteration* is that it uses a full policy evaluation, which can be computationally exhaustive as the convergence of  $V_k$  is only at the limit and the number of states can be huge.
- The idea of **value iteration** is to interleave policy evaluation and policy improvement, so that the policy is improved after EACH iteration of policy evaluation, not after complete convergence.
- As policy improvement returns a deterministic greedy policy, updating of the value of a state is then simpler:

$$V_{k+1}(s) = \max_a \sum_{s'} p(s'|s, a) [r(s, a, s') + \gamma V_k(s')]$$

- Note that this is equivalent to turning the **Bellman optimality equation** into an update rule.
- Value iteration converges to  $V^*$ , faster than policy iteration, and should be stopped when the values do not change much anymore.


# Value iteration

- Initialize a deterministic policy  $\pi(s)$  and set  $V(s) = 0 \ \forall s \in \mathcal{S}$ .
- **while** not converged:
  - **for** all states  $s$ :
    - $V_{\text{target}}(s) = \max_a \sum_{s' \in \mathcal{S}} p(s'|s, a) [r(s, a, s') + \gamma V(s')]$
  - $\delta = 0$
  - **for** all states  $s$ :
    - $\delta = \max(\delta, |V(s) - V_{\text{target}}(s)|)$
    - $V(s) = V_{\text{target}}(s)$
  - **if**  $\delta < \delta_{\text{threshold}}$ :
    - converged = True


# Comparison of Policy- and Value-iteration

## Full policy-evaluation backup

$$V_{k+1}(s) \leftarrow \sum_{a \in \mathcal{A}(s)} \pi(s, a) \sum_{s' \in \mathcal{S}} p(s' | s, a) [r(s, a, s') + \gamma V_k(s')]$$



$$V_{k+1}(s) \leftarrow \max_{a \in \mathcal{A}(s)} \sum_{s' \in \mathcal{S}} p(s' | s, a) [r(s, a, s') + \gamma V_k(s')]$$





# Asynchronous dynamic programming

- Synchronous DP requires exhaustive sweeps of the entire state set (**synchronous backups**).
  - **while** not converged:
    - **for** all states  $s$ :
      - $V_{\text{target}}(s) = \max_a \sum_{s' \in \mathcal{S}} p(s'|s, a) [r(s, a, s') + \gamma V(s')]$
    - **for** all states  $s$ :
      - $V(s) = V_{\text{target}}(s)$
- Asynchronous DP updates instead each state independently and asynchronously (**in-place**):
  - **while** not converged:
    - Pick a state  $s$  randomly (or following a heuristic).
    - Update the value of this state.

$$V(s) = \max_a \sum_{s' \in \mathcal{S}} p(s'|s, a) [r(s, a, s') + \gamma V(s')]$$

- We must still ensure that all states are visited, but their frequency and order is irrelevant.

# Asynchronous dynamic programming

- Is it possible to select the states to backup intelligently?
- **Prioritized sweeping** selects in priority the states with the largest remaining **Bellman error**:


$$\delta = \left| \max_a \sum_{s' \in \mathcal{S}} p(s'|s, a) [r(s, a, s') + \gamma V(s')] - V(s) \right|$$

- A large Bellman error means that the current estimate  $V(s)$  is very different from the **target**  $y$ :

$$y = \max_a \sum_{s' \in \mathcal{S}} p(s'|s, a) [r(s, a, s') + \gamma V(s')]$$


- States with a high Bellman error should be updated in priority.
- If the Bellman error is small, this means that the current estimate  $V(s)$  is already close to what it should be, there is no hurry in evaluating this state.
- The main advantage is that the DP algorithm can be applied as the agent is actually experiencing its environment (no need for the dynamics of environment to be fully known).

# Efficiency of Dynamic Programming



- Policy-iteration and value-iteration consist of alternations between policy evaluation and policy improvement, although at different frequencies.
- This principle is called **Generalized Policy Iteration** (GPI).
- Finding an optimal policy is polynomial in the number of states and actions:  $\mathcal{O}(n^2 m)$  ( $n$  is the number of states,  $m$  the number of actions).
- However, the number of states is often astronomical, e.g., often growing exponentially with the number of state variables (what Bellman called “**the curse of dimensionality**”).
- In practice, classical DP can only be applied to problems with a few millions of states.

# Curse of dimensionality



Source: <https://medium.com/diogo-menezes-borges/give-me-the-antidote-for-the-curse-of-dimensionality-b14bce4bf4d2>

- If one variable can be represented by 5 discrete values:
  - 2 variables necessitate 25 states,
  - 3 variables need 125 states, and so on...
- The number of states explodes exponentially with the number of dimensions of the problem.