



UNIVERSITY OF TECHNOLOGY
IN THE EUROPEAN CAPITAL OF CULTURE
CHEMNITZ


Deep Reinforcement Learning

Monte Carlo methods

Julien Vitay

Professur für Künstliche Intelligenz - Fakultät für Informatik

Key idea of Reinforcement learning: Generalized Policy Iteration



- RL algorithms iterate over two steps:

1. Policy evaluation

- For a given policy π , the value of all states $V^\pi(s)$ or all state-action pairs $Q^\pi(s, a)$ is calculated, either based on:
 - the Bellman equations (Dynamic Programming)
 - sampled experience (Monte Carlo and Temporal Difference)

2. Policy improvement


- From the current estimated values $V^\pi(s)$ or $Q^\pi(s, a)$, a new **better** policy π is derived.
- After enough iterations, the policy converges to the **optimal policy** (if the states are Markov).

1 - Monte Carlo control

Principle of Monte Carlo (MC) methods

- The value of each state is defined as the mathematical expectation of the return obtained after that state and thereafter following the policy π :

$$V^\pi(s) = \mathbb{E}_{\rho_\pi}(R_t | s_t = s) = \mathbb{E}_{\rho_\pi}\left(\sum_{k=0}^{\infty} \gamma^k r_{t+k+1} | s_t = s\right)$$




- Instead of solving the Bellman equations, **Monte Carlo methods** (MC) approximate this mathematical expectation by **sampling** M trajectories τ_i starting from s and computing the sampling average of the obtained returns:

$$V^\pi(s) = \mathbb{E}_{\rho_\pi}(R_t | s_t = s) \approx \frac{1}{M} \sum_{i=1}^M R(\tau_i)$$

- If you have enough trajectories, the sampling average is an unbiased estimator of the value function.
- The advantage of Monte Carlo methods is that they require only **experience**, not the complete dynamics $p(s' | s, a)$ and $r(s, a, s')$: **model-free**.

Monte Carlo policy evaluation


- The idea of MC policy evaluation is to repeatedly sample **episodes** starting from each possible state s_0 and maintain a **running average** of the obtained returns for each state:



- while** True:
 - Start from an initial state s_0 .
 - Generate a sequence of transitions according to the current policy π until a terminal state s_T is reached.
$$\tau = (s_0, a_0, r_1, s_1, a_1, \dots, s_T)$$
 - Compute the return $R_t = \sum_{k=0}^{\infty} \gamma^k r_{t+k+1}$ for all encountered states s_0, s_1, \dots, s_T .
 - Update the estimated state value $V(s_t)$ of all encountered states using the obtained return:

$$V(s_t) \leftarrow V(s_t) + \alpha (R_t - V(s_t))$$

Monte Carlo policy evaluation of action values



- The same method can be used to estimate Q-values.
- **while** True:
 1. Start from an initial state s_0 .
 2. Generate a sequence of transitions according to the current policy π until a terminal state s_T is reached.

$$\tau = (s_0, a_0, r_1, s_1, a_1, \dots, s_T)$$

3. Compute the return $R_t = \sum_{k=0}^{\infty} \gamma^k r_{t+k+1}$ for all encountered state-action pairs $(s_0, a_0), (s_1, a_1), \dots, (s_{T-1}, a_{T-1})$.
4. Update the estimated action value $Q(s_t, a_t)$ of all encountered state-action pairs using the obtained return:

$$Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha (R_t - Q(s_t, a_t))$$

- There are much more values to estimate (one per state-action pair), but the policy will be easier to derive.


Monte Carlo policy improvement

- After each episode, the state or action values of the visited (s, a) pairs have changed, so the current policy might not be optimal anymore.
- As in DP, the policy can then be improved in a greedy manner:

$$\pi'(s) = \operatorname{argmax}_a Q(s, a)$$

$$= \operatorname{argmax}_a \sum_{s' \in \mathcal{S}} p(s'|s, a) [r(s, a, s') + \gamma V(s')]$$

- Estimating the Q-values allows to act greedily, while estimating the V-values still requires the dynamics $p(s'|s, a)$ and $r(s, a, s')$ (one-step lookahead).
- An approximation of these dynamics can be sufficient.



Monte Carlo control

- **Monte Carlo control** alternates between **MC policy evaluation** and **policy improvement** until the optimal policy is found.
- **while** True:
 1. Select an initial state s_0 .
 2. Generate a sequence of transitions according to the current policy π until a terminal state s_T is reached.

$$\tau = (s_0, a_0, r_1, s_1, a_1, \dots, s_T)$$

3. Compute the return $R_t = \sum_{k=0}^{\infty} \gamma^k r_{t+k+1}$ of all encountered state-action pairs.
4. Update the estimated action value $Q(s_t, a_t)$ of all encountered state-action pairs:

$$Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha (R_t - Q(s_t, a_t))$$


5. For each state s_t in the episode, **improve** the policy:

$$\pi(s_t, a) = \begin{cases} 1 & \text{if } a = \operatorname{argmax} Q(s_t, a) \\ 0 & \text{otherwise.} \end{cases}$$

2 - On-policy Monte Carlo control

How to generate the episodes?

- The problem with MC control is that we need a policy to generate the sample episodes, but it is also that policy that we want to learn.
- We have the same **exploration/exploitation** problem as in bandits:
 - If I trust my estimates too much (**exploitation**), I may miss more interesting solutions by keeping generating the same episodes.
 - If I act randomly (**exploration**), I will find more interesting solutions, but I won't keep doing them.





Source: http://ai.berkeley.edu/lecture_slides.html

Exploration/Exploitation dilemma

- **Exploitation** is using the current estimated values to select the greedy action:
 - The estimated values represent how good we think an action is, so we have to use this value to update the policy.
 - **Exploration** is executing non-greedy actions to try to reduce our uncertainty about the true values:
 - The values are only estimates: they may be wrong so we can not trust them completely.
 - If you only **exploit** your estimates, you may miss interesting solutions.
 - If you only **explore**, you do not use what you know: you act randomly and do not obtain as much reward as you could.
- **You can't exploit all the time; you can't explore all the time.**
- You can never stop exploring; but you can reduce it if your performance is good enough.

Stochastic policies

- **Exploration** can be ensured by forcing the learned policy to be **stochastic**, aka **ϵ -soft**.



- **ϵ -Greedy action selection** randomly selects non-greedy actions with a small probability ϵ :

$$\pi(s, a) = \begin{cases} 1 - \epsilon & \text{if } a = \operatorname{argmax} Q(s, a) \\ \frac{\epsilon}{|\mathcal{A}| - 1} & \text{otherwise.} \end{cases}$$

- **Softmax action selection** uses a Gibbs (or Boltzmann) distribution to represent the probability of choosing the action a in state s :

$$\pi(s, a) = \frac{\exp Q(s, a) / \tau}{\sum_b \exp Q(s, b) / \tau}$$

- ϵ -greedy chooses non-greedy actions randomly, while softmax favors the best alternatives.

On-policy Monte Carlo control

- In **on-policy** control methods, the learned policy has to be ϵ -soft, which means all actions have a probability of at least $\frac{\epsilon}{|\mathcal{A}|}$ to be visited. ϵ -greedy and softmax policies meet this criteria.
- Each sample episode is generated using this policy, which ensures exploration, while the control method still converges towards the optimal ϵ -policy.
- **while** True:

1. Generate an episode $\tau = (s_0, a_0, r_1, \dots, s_T)$ using the current **stochastic** policy π .
2. For each state-action pair (s_t, a_t) in the episode, update the estimated Q-value:

$$Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha (R_t - Q(s_t, a_t))$$

3. For each state s_t in the episode, improve the policy (e.g. ϵ -greedy):

$$\pi(s_t, a) = \begin{cases} 1 - \epsilon & \text{if } a = \operatorname{argmax} Q(s, a) \\ \frac{\epsilon}{|\mathcal{A}(s_t)| - 1} & \text{otherwise.} \end{cases}$$

3 - Off-policy Monte Carlo control


Off-policy Monte Carlo control

- Another option to ensure exploration is to generate the sample episodes using a **behavior policy** $b(s, a)$ different from the **learned policy** $\pi(s, a)$ of the agent.
- The **behavior policy** $b(s, a)$ used to generate the episodes is only required to select at least occasionally the same actions as the **learned policy** $\pi(s, a)$ (coverage assumption).

$$\pi(s, a) > 0 \Rightarrow b(s, a) > 0$$

- There are mostly two choices regarding the behavior policy:
 1. An ϵ -soft behavior policy over the **Q-values** as in on-policy MC is often enough, while a deterministic (greedy) policy can be learned implicitly.
 2. The behavior policy could also come from **expert knowledge**, i.e. known episodes from the MDP generated by somebody else (human demonstrator, classical algorithm).

Offline RL: process control



Source: <https://deepmind.com/blog/deepmind-ai-reduces-google-data-centre-cooling-bill-40/>

- 40% reduction of energy consumption when using deep RL to control the cooling of Google's datacenters.
- The RL algorithm learned passively from the **behavior policy** (expert decisions) what the optimal policy should be.
- Learning from data (a.k.a **learning from demonstrations**) is often referred to as **offline RL**.

Importance sampling

- But are we mathematically allowed to do this?
- We search for the optimal policy that maximizes in expectation the return of each **trajectory** (episode) possible under the learned policy π :


$$\mathcal{J}(\pi) = \mathbb{E}_{\tau \sim \rho_\pi} [R(\tau)]$$

- ρ_π denotes the probability distribution of trajectories achievable using the policy π .
- If we generate the trajectories from the behavior policy $b(s, a)$, we end up maximizing something else:

$$\mathcal{J}'(\pi) = \mathbb{E}_{\tau \sim \rho_b} [R(\tau)]$$

- The policy that maximizes $\mathcal{J}'(\pi)$ is **not** the optimal policy of the MDP.

Importance sampling



- If you try to estimate a parameter of a random distribution π using samples of another distribution b , the sample average will have a strong **bias**.
- We need to **correct** the samples from b in order to be able to estimate the parameters of π correctly:
 - **importance sampling** (IS).

Importance sampling


- We want to estimate the expected return of the trajectories generated by the policy π :

$$\mathcal{J}(\pi) = \mathbb{E}_{\tau \sim \rho_{\pi}} [R(\tau)]$$

- We start by using the definition of the mathematical expectation:

$$\mathcal{J}(\pi) = \int_{\tau} \rho_{\pi}(\tau) R(\tau) d\tau$$

- The expectation is the integral over all possible trajectories of their return $R(\tau)$, weighted by the likelihood $\rho_{\pi}(\tau)$ that a trajectory τ is generated by the policy π .



Importance sampling

- The trick is to introduce the behavior policy b in what we want to estimate:

$$\mathcal{J}(\pi) = \int_{\tau} \frac{\rho_b(\tau)}{\rho_b(\tau)} \rho_{\pi}(\tau) R(\tau) d\tau$$

- $\rho_b(\tau)$ is the likelihood that a trajectory τ is generated by the behavior policy b .
- We shuffle a bit the terms:

$$\mathcal{J}(\pi) = \int_{\tau} \rho_b(\tau) \frac{\rho_{\pi}(\tau)}{\rho_b(\tau)} R(\tau) d\tau$$


and notice that it has the form of an expectation over trajectories generated by b :

$$\mathcal{J}(\pi) = \mathbb{E}_{\tau \sim \rho_b} \left[\frac{\rho_{\pi}(\tau)}{\rho_b(\tau)} R(\tau) \right]$$

- This means that we can sample trajectories from b , but we need to **correct** the observed return by the **importance sampling weight** $\frac{\rho_{\pi}(\tau)}{\rho_b(\tau)}$.

Importance sampling


- The importance sampling weight corrects the mismatch between π and b .




- If the two distributions are the same (on-policy), the IS weight is 1, no need to correct the return.
- If a sample is likely under b but not under π , we should not care about its return: $\frac{\rho_{\pi}(\tau)}{\rho_b(\tau)} \ll 1$
- If a sample is likely under π but not much under b , we increase its importance in estimating the return: $\frac{\rho_{\pi}(\tau)}{\rho_b(\tau)} \gg 1$
- The sampling average of the corrected samples will be closer from the true estimate (unbiased).

Importance sampling

- Great, but how do we compute these probability distributions $\rho_\pi(\tau)$ and $\rho_b(\tau)$ for a trajectory τ ?



- A trajectory τ is a sequence of state-action transitions $(s_0, a_0, s_1, a_1, \dots, s_T)$ whose probability depends on:
 - the probability of choosing an action a_t in state s_t : the **policy** $\pi(s, a)$.
 - the probability of arriving in the state s_{t+1} from the state s_t with the action a_t : the **transition probability** $p(s_{t+1} | s_t, a_t)$.



Importance sampling

- The **likelihood** of a trajectory $\tau = (s_0, a_0, s_1, a_1, \dots, s_T)$ under a policy π depends on the policy and the transition probabilities (Markov property):

$$\rho_{\pi}(\tau) = p_{\pi}(s_0, a_0, s_1, a_1, \dots, s_T) = p(s_0) \prod_{t=0}^{T-1} \pi_{\theta}(s_t, a_t) p(s_{t+1} | s_t, a_t)$$

- $p(s_0)$ is the probability of starting an episode in s_0 , we do not have control over it.
- What is interesting is that the transition probabilities disappear when calculating the **importance sampling weight**:

$$\rho_{0:T-1} = \frac{\rho_{\pi}(\tau)}{\rho_b(\tau)} = \frac{p_0(s_0) \prod_{t=0}^{T-1} \pi(s_t, a_t) p(s_{t+1} | s_t, a_t)}{p_0(s_0) \prod_{t=0}^T b(s_t, a_t) p(s_{t+1} | s_t, a_t)} = \frac{\prod_{t=0}^{T-1} \pi(s_t, a_t)}{\prod_{t=0}^T b(s_t, a_t)} = \prod_{t=0}^{T-1} \frac{\pi(s_t, a_t)}{b(s_t, a_t)}$$

- The importance sampling weight is simply the product over the length of the episode of the ratio between $\pi(s_t, a_t)$ and $b(s_t, a_t)$.

Off-policy Monte Carlo control

- In **off-policy MC control**, we generate episodes using the behavior policy b and update **greedily** the learned policy π .
- For the state s_t , the obtained returns just need to be weighted by the relative probability of occurrence of the **rest of the episode** following the policies π and b :

$$\rho_{t:T-1} = \prod_{k=t}^{T-1} \frac{\pi(s_k, a_k)}{b(s_k, a_k)}$$

$$V^\pi(s_t) = \mathbb{E}_{\tau \sim \rho_b} [\rho_{t:T-1} R_t]$$

- This gives us the updates:

$$V(s_t) \leftarrow V(s_t) + \alpha \rho_{t:T-1} (R_t - V(s_t))$$

and:

$$Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha \rho_{t:T-1} (R_t - Q(s_t, a_t))$$

- Unlikely episodes under π are barely used for learning, likely ones are used a lot.

Off-policy Monte Carlo control

- **while** True:

1. Generate an episode $\tau = (s_0, a_0, r_1, \dots, s_T)$ using the **behavior** policy b .
2. For each state-action pair (s_t, a_t) in the episode, update the estimated Q-value:

$$\rho_{t:T-1} = \prod_{k=t}^{T-1} \frac{\pi(s_k, a_k)}{b(s_k, a_k)}$$

$$Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha \rho_{t:T-1} (R_t - Q(s_t, a_t))$$

3. For each state s_t in the episode, update the **learned** deterministic policy (greedy):

$$\pi(s_t, a) = \begin{cases} 1 & \text{if } a = \operatorname{argmax} Q(s_t, a) \\ 0 & \text{otherwise.} \end{cases}$$

Off-policy Monte Carlo control

- **Problem 1:** if the learned policy is greedy, the IS weight becomes quickly 0 for a non-greedy action a_t :

$$\pi(s_t, a_t) = 0 \rightarrow \rho_{0:T-1} = \prod_{k=0}^{T-1} \frac{\pi(s_k, a_k)}{b(s_k, a_k)} = 0$$

Off-policy MC control only learns from the last greedy actions, what is slow at the beginning.

Solution: π and b should not be very different. Usually π is greedy and b is a softmax (or ϵ -greedy) over it.

- **Problem 2:** if the learned policy is stochastic, the IS weights can quickly **vanish** to 0 or **explode** to infinity:

$$\rho_{t:T-1} = \prod_{k=t}^{T-1} \frac{\pi(s_k, a_k)}{b(s_k, a_k)}$$

If $\frac{\pi(s_k, a_k)}{b(s_k, a_k)}$ is smaller than 1, the products go to 0. If it is bigger than 1, it grows to infinity.

Solution: one can normalize the IS weight between different episodes (see Sutton and Barto) or **clip** it (e.g. restrict it to $[0.9, 1.1]$, see PPO later in this course).

Advantages of off-policy methods

- The main advantage of **off-policy** strategies is that you can learn from other's actions, you don't have to rely on your initially wrong policies to discover the solution by chance.
 - Example: learning to play chess by studying thousands/millions of plays by chess masters.
- In a given state, only a subset of the possible actions are actually executed by experts: the others may be too obviously wrong.
- The exploration is then guided by this expert knowledge, not randomly among all possible actions.
- Off-policy methods greatly reduce the number of transitions needed to learn a policy: very stupid actions are not even considered, but the estimation policy learns an optimal strategy from the "classical" moves.
- Drawback: if a good move is not explored by the behavior policy, the learned policy will never try it.

Properties of Monte Carlo methods

- Monte Carlo evaluation estimates value functions via **sampling** of entire episodes.
- MC for action values is **model-free**: you do not need to know $p(s'|s, a)$ to learn the optimal policy, you just sample transitions (trial and error).
- MC only applies to **episodic tasks**: as you learn at the end of an episode, it is not possible to learn continuing tasks.
- MC suffers from the **exploration-exploitation** problem:
 - **on-policy** MC learns a stochastic policy (ϵ -greedy, softmax) to ensure exploration.
 - **off-policy** MC learns a greedy policy, but explores via a behavior policy (importance sampling).
- Monte Carlo methods have:
 - a **small bias**: with enough sampled episodes, the estimated values converge to the true values.
 - a **huge variance**: the slightest change of the policy can completely change the episode and its return. You will need a lot of samples to form correct estimates: **sample complexity**.

